

The syslog-ng 3.0 Administrator Guide

Ninth Edition

Published November 6, 2009

This manual is the primary documentation of the syslog-ng 3.0 product line, including syslog-ng Open Source Edition (syslog-ng OSE), syslog-ng Premium Edition (syslog-ng PE), and the syslog-ng Agent for Windows (which is a part of syslog-ng PE).



The syslog-ng 3.0 Administrator Guide

Product Marketing and Documentation Department

Revision History

Ninth Revision	November 6, 2009
Corrected the minimal requirements of the System i agent, as well as the XML configuration of the Windows agent.	
Eighth Revision	October 30, 2009
Several minor clarifications and typo corrections.	
Seventh Revision	August 19, 2009
Updated to syslog-ng Agent for Windows version 3.0.4, including the description of the new XML-based configuration file of the agent. Notes on escaping special characters in pattern database rule pattern were also added to Section 8.6.2.3, “Creating pattern databases” (p. 220). The <code>log_disk_fifo_size()</code> parameter was removed from the list of global syslog-ng options, because it must be set separately for every destination.	
Sixth Revision	July 31, 2009
Updated to syslog-ng Open Source Edition version 3.0.3, and documented the pattern database (<code>db_parser()</code>) feature: see Section 2.12, “Classifying log messages” (p. 15), Section 3.9, “Classifying messages” (p. 63), and Section 8.6.2, “Pattern databases” (p. 218) for details. Also many smaller corrections and clarifications.	
Fifth Revision	February 6, 2009
Updated to syslog-ng Premium Edition version 3.0.2	
Fourth Revision	December 12, 2008
Several typo corrections, clarifications, and other cleanups.	
Third Revision	October 31, 2008
Updated Chapter 5, <i>Collecting logs from Windows hosts</i> (p. 93) to cover version 3.0.3 of the syslog-ng Agent for Windows application.	
Second Revision	October 8, 2008
Added documentation on installing syslog-ng from .run binaries, and various updates and corrections.	
First Revision	October 2, 2008
Initial release for review	

Copyright © 2006-2009 BalaBit IT Security Ltd.

This guide is published under the Creative Commons Attribution-Noncommercial-No Derivative Works (by-nc-nd) 3.0 license. See Appendix 4, *Creative Commons Attribution Non-commercial No Derivatives (by-nc-nd) License* (p. 256) for details. The latest version is always available at <http://www.balabit.com/support/documentation>.

This product includes software developed by the OpenSSL Project for use in the OpenSSL Toolkit (<http://www.openssl.org/>). This product includes cryptographic software written by Eric Young (ey@cryptsoft.com)

This documentation and the product it describes are considered protected by copyright according to the applicable laws.

The syslog-ngTM name and the syslog-ngTM logo are registered trademarks of BalaBit.

The BalaBitTM name and the BalaBitTM logo are registered trademarks of BalaBit.

LinuxTM is a registered trademark of Linus Torvalds.

DebianTM is a registered trademark of Software in the Public Interest Inc.

WindowsTM XP, 2003 Server, Vista, and 2008 Server are registered trademarks of Microsoft Corporation.

MySQLTM is a registered trademark of MySQL AB in the United States, the European Union and other countries.

OracleTM, JD EdwardsTM, PeopleSoftTM, and SiebelTM are registered trademarks of Oracle Corporation and/or its affiliates.

Red HatTM, Inc., Red HatTM Enterprise LinuxTM and Red HatTM LinuxTM are trademarks of Red Hat, Inc.

SUSETM is a trademark of SUSE AG, a Novell business.

SolarisTM is a registered trademark of Sun Microsystems, Inc.

AIXTM, AIX 5LTM, AS/400TM, BladeCenterTM, eServerTM, IBMTM, the IBMTM logo, IBM System iTM, IBM System i5TM, IBM System xTM, iSeriesTM, i5/OSTM, NetfinityTM, NetServerTM, OpenPowerTM, OS/400TM, PartnerWorldTM, POWERTM, ServerGuideTM, ServerProvenTM, and xSeriesTM are trademarks or registered trademarks of International Business Machines.

Alliance Log Agent for System iTM is a registered trademark of Patrick Townsend & Associates, Inc.

All other product names mentioned herein are the trademarks of their respective owners.

Some rights reserved.

DISCLAIMER



BalaBit is not responsible for any third-party Web sites mentioned in this document. BalaBit does not endorse and is not responsible or liable for any content, advertising, products, or other material on or available from such sites or resources. BalaBit will not be responsible or liable for any damage or loss caused or alleged to be caused by or in connection with use of or reliance on any such content, goods, or services that are available on or through any such sites or resources.



Table of Contents

Preface	xiv
1. Summary of contents	xiv
2. Target audience and prerequisites	xiv
3. Products covered in this guide	xv
4. Typographical conventions	xv
5. Contact and support information	xvi
5.1. Sales contact	xvi
5.2. Support contact	xvii
5.3. Training	xvii
6. About this document	xvii
6.1. What is new in this main edition of The syslog-ng Administrator Guide?	xvii
6.2. Feedback	xviii
6.3. Acknowledgments	xviii
1. Introduction to syslog-ng	1
1.1. What syslog-ng is	1
1.2. What syslog-ng is not	2
1.3. Why is syslog-ng needed?	2
1.4. What is new in syslog-ng 3.0?	2
1.5. Who uses syslog-ng?	3
1.6. Supported platforms	4
2. The concepts of syslog-ng	6
2.1. The philosophy of syslog-ng	6
2.2. Logging with syslog-ng	6
2.2.1. Embedded log statements	8
2.3. Modes of operation	10
2.3.1. Client mode	10
2.3.2. Relay mode	10
2.3.3. Server mode	11
2.4. Global objects	11
2.5. Timezone handling	12
2.6. Daylight saving changes	13
2.7. Secure logging using TLS	13
2.8. Secure storage of log messages	14
2.9. Formatting messages, filenames, directories, and tablenamees	14
2.10. Segmenting messages	14
2.11. Modifying messages	15
2.12. Classifying log messages	15
2.12.1. The structure of the pattern database	16
2.12.2. How pattern matching works	17
2.12.3. Artificial ignorance	18
2.13. Managing incoming and outgoing messages with flow-control	18
2.13.1. Flow-control and multiple destinations	20



2.14. Using disk-based buffering	20
2.15. Licensing	21
2.16. High availability support	22
2.17. Possible causes of losing log messages	22
2.18. The structure of a log message	24
2.18.1. BSD-syslog or legacy-syslog messages	24
2.18.2. IETF-syslog messages	26
3. Configuring syslog-ng	30
3.1. The syslog-ng configuration file	30
3.1.1. Including configuration files	31
3.1.2. Logging configuration changes	31
3.2. Defining global objects	32
3.2.1. Notes about the configuration syntax	33
3.3. Sources and source drivers	33
3.3.1. Collecting internal messages	36
3.3.2. Collecting messages from text files	37
3.3.3. Collecting messages from named pipes	38
3.3.4. Collecting messages on Sun Solaris	39
3.3.5. Collecting messages using the IETF syslog protocol	39
3.3.6. Collecting messages from remote hosts using the BSD syslog protocol	40
3.3.7. Collecting messages from UNIX domain sockets	41
3.4. Destinations and destination drivers	42
3.4.1. Storing messages in plain-text files	43
3.4.2. Storing messages in encrypted files	44
3.4.3. Sending messages to named pipes	46
3.4.4. Sending messages to external applications	46
3.4.5. Storing messages in an SQL database	47
3.4.6. Sending messages to a remote logserver using the IETF-syslog protocol	51
3.4.7. Sending messages to a remote logserver using the legacy BSD-syslog protocol	52
3.4.8. Sending messages to UNIX domain sockets	53
3.4.9. usrtty()	207
3.5. Log paths	53
3.5.1. Using embedded log statements	55
3.5.2. Configuring flow-control	56
3.6. Filters	57
3.6.1. Optimizing regular expressions in filters	60
3.7. Templates and macros	60
3.8. Parsing messages	62
3.9. Classifying messages	63
3.9.1. Using parser results in filters and templates	64
3.10. Rewriting messages	66
3.11. Configuring global syslog-ng options	67
3.12. Enabling disk-based buffering	67
3.13. Encrypting log messages with TLS	68
3.14. Mutual authentication using TLS	70
3.15. Configuring syslog-ng clients	73
3.16. Configuring syslog-ng relays	73



3.17. Configuring syslog-ng servers	74
3.18. Installing and upgrading the license	75
3.19. Troubleshooting syslog-ng	76
3.19.1. Creating syslog-ng core files	76
3.19.2. Running a failure script	77
3.19.3. Stopping syslog-ng	77
4. Installing syslog-ng	78
4.1. Installing syslog-ng using the .run installer	78
4.1.1. Installing syslog-ng in client or relay mode	79
4.1.2. Installing syslog-ng in server mode	81
4.1.3. Installing syslog-ng without user-interaction	84
4.2. Installing syslog-ng on RPM-based platforms (Red Hat, SUSE, AIX)	85
4.3. Installing syslog-ng on Debian-based platforms	86
4.4. Compiling syslog-ng from source	86
4.5. Uninstalling syslog-ng	89
4.6. Configuring Microsoft SQL Server to accept logs from syslog-ng	89
5. Collecting logs from Windows hosts	93
5.1. Installing the syslog-ng agent	93
5.1.1. Installing the syslog-ng agent in standalone mode	94
5.1.2. Installing the syslog-ng agent on the hosts of a domain	95
5.1.3. Upgrading syslog-ng Agent for Windows 2.x to 3.0.x	96
5.1.4. Upgrading syslog-ng Agent for Windows 3.0.1 to version 3.0.2	96
5.1.5. Upgrading syslog-ng Agent for Windows 3.0.2 to version 3.0.3	97
5.1.6. Upgrading syslog-ng Agent for Windows to version 3.0.4	97
5.2. Configuring destinations	97
5.2.1. Limiting the rate of messages	99
5.3. Configuring message sources	99
5.3.1. Eventlog sources	99
5.3.2. File sources and logrotation	101
5.3.3. Global settings of the syslog-ng agent	102
5.4. Using SSL-encrypted connections with the syslog-ng agent	103
5.4.1. Using mutual authentication with syslog-ng agent	104
5.4.2. Importing certificates with the Microsoft Management Console	105
5.5. Filtering messages	106
5.6. Customizing the message format	107
5.6.1. Customizing the timestamp used by the syslog-ng Agent	109
5.6.2. Macros available in the syslog-ng Agent	109
5.7. Using an XML-based configuration file	113
5.7.1. A sample configuration file for the syslog-ng Agent	115
5.8. Controlling the syslog-ng agent services	116
5.9. Domain versus local settings	117
5.10. Troubleshooting syslog-ng Agent for Windows	117
5.10.1. Sending messages and CPU load	118
5.10.2. Creating core and memory dumps	119
5.11. Configuring the auditing policy on Windows	119
5.11.1. Turning on security logging on Windows XP	119



5.11.2. Turning on security logging for domain controllers	120
5.11.3. Turning on auditing on Windows 2003 Server	120
6. Collecting logs from IBM System i	121
6.1. Supported sources	121
6.2. Supported output formats	121
6.3. Filtering log entries	122
6.4. Installing the syslog-ng Agent for IBM System i	122
6.4.1. Installing from an Internet download	122
6.4.2. Installing from a product CD	122
6.4.3. Upgrading the syslog-ng Agent for IBM System i	122
6.5. Configuring System i security auditing	123
6.5.1. Enabling security auditing manually	123
6.5.2. Enabling user auditing	124
6.5.3. Enabling object auditing	124
6.5.4. Configuring syslog-ng Agent for IBM System i	124
6.5.5. Configuring Alliance Syslog for System i	125
6.5.6. Configuring communication between the syslog-ng Agent and the server	126
6.5.7. Work with security types	128
6.6. Controlling the syslog-ng Agent for IBM System i	129
6.6.1. Starting the Alliance subsystem	129
6.6.2. Automating the start of the Alliance subsystem ALLSYL100	130
6.7. Application maintenance	131
6.8. View application logs	131
6.9. Configuring IBM System i Servers	132
6.9.1. Configuring Apache server logs	132
6.9.2. OpenSSH server logs	133
6.9.3. Other server logs	133
6.10. Troubleshooting the syslog-ng Agent for IBM System i	133
6.10.1. System operator messages	134
6.10.2. Application logging	134
6.10.3. Cannot install the product from CD	134
6.10.4. Logs are not being transferred to my log server	134
6.10.5. I get a license error when trying to use configuration options	134
6.10.6. The product no longer works after a system upgrade	134
6.10.7. Security events are not being captured	135
6.10.8. I am not capturing information about our security administrators	135
6.10.9. I am not capturing information about programs and files	135
6.10.10. I am not capturing QSYSOPR messages	135
6.10.11. I would like to turn off some audit journal events	135
6.10.12. Where do I find error messages?	135
7. Best practices and examples	136
7.1. General recommendations	136
7.2. Handling lots of parallel connections	136
7.3. Handling large message load	137
7.4. Using name resolution in syslog-ng	137
7.4.1. Resolving hostnames locally	138



7.5. Collecting logs from chroot	138
7.6. Replacing klogd on Linux	139
7.7. A note on timezones and timestamps	140
7.8. Dropping messages	140
8. Reference	141
8.1. Source drivers	141
8.1.1. internal()	141
8.1.2. file()	141
8.1.3. pipe()	146
8.1.4. program()	150
8.1.5. sun-streams() driver	154
8.1.6. syslog()	160
8.1.7. tcp(), tcp6(), udp() and udp6()	165
8.1.8. unix-stream() and unix-dgram()	173
8.2. Destination drivers	179
8.2.1. file()	179
8.2.2. logstore()	183
8.2.3. pipe()	186
8.2.4. program()	189
8.2.5. sql()	191
8.2.6. syslog()	195
8.2.7. tcp(), tcp6(), udp(), and udp6(),	200
8.2.8. unix-stream() & unix-dgram()	204
8.2.9. usertty()	207
8.3. Log path flags	207
8.4. Filter functions	208
8.5. Macros	210
8.6. Message parsers	214
8.6.1. CSV parsers	214
8.6.2. Pattern databases	218
8.7. Rewriting messages	225
8.8. Regular expressions	226
8.9. Global options	227
8.10. TLS options	231
1. The syslog-ng manual pages	234
2. BalaBit syslog-ng Premium Edition License contract	243
2.1. SUBJECT OF THE License CONTRACT	243
2.2. DEFINITIONS	243
2.3. Words and expressions	243
2.4. LICENSE GRANTS AND RESTRICTIONS	244
2.5. SUBSIDIARIES	245
2.6. INTELLECTUAL PROPERTY RIGHTS	245
2.7. TRADE MARKS	246
2.8. NEGLIGENT INFRINGEMENT	246
2.9. INTELLECTUAL PROPERTY INDEMNIFICATION	246



2.10. LICENSE FEE	247
2.11. WARRANTIES	247
2.12. <i>DISCLAIMER OF WARRANTIES</i>	247
2.13. LIMITATION OF LIABILITY	247
2.14. DURATION AND TERMINATION	248
2.15. AMENDMENTS	248
2.16. WAIVER	248
2.17. SEVERABILITY	248
2.18. NOTICES	248
2.19. MISCELLANEOUS	249
3. GNU General Public License	250
3.1. Preamble	250
3.2. TERMS AND CONDITIONS FOR COPYING, DISTRIBUTION AND MODIFICATION	251
3.2.1. Section 0	251
3.2.2. Section 1	251
3.2.3. Section 2	251
3.2.4. Section 3	252
3.2.5. Section 4	253
3.2.6. Section 5	253
3.2.7. Section 6	253
3.2.8. Section 7	253
3.2.9. Section 8	254
3.2.10. Section 9	254
3.2.11. Section 10	254
3.2.12. NO WARRANTY Section 11	254
3.2.13. Section 12	254
3.3. How to Apply These Terms to Your New Programs	255
4. Creative Commons Attribution Non-commercial No Derivatives (by-nc-nd) License	256
Glossary	261
Index	265



List of Examples

2.1. Counting log source hosts	22
3.1. A simple configuration file	30
3.2. Using required and optional parameters	32
3.3. A simple source statement	34
3.4. A source statement using two source drivers	34
3.5. Setting default priority and facility	34
3.6. Source statement on a Linux based operating system	35
3.7. Using the internal() driver	141
3.8. Using the file() driver	146
3.9. Using wildcards in the filename	146
3.10. Monitoring multiple directories	146
3.11. Using the pipe() driver	150
3.12. Using the sun-streams() driver	160
3.13. Using the syslog() driver	165
3.14. Using the udp() and tcp() drivers	173
3.15. Using the unix-stream() and unix-dgram() drivers	179
3.16. A simple destination statement	43
3.17. Using the file() driver	182
3.18. Using the file() driver with macros in the file name and a template for the message	182
3.19. Using the logstore() driver	186
3.20. Using the pipe() driver	189
3.21. Using the program() destination driver	191
3.22. Using the sql() driver	193
3.23. Using the sql() driver with an Oracle database	194
3.24. Using the sql() driver with an MSSQL database	194
3.25. Using the syslog() driver	199
3.26. Using the tcp() driver	203
3.27. Using the unix-stream() driver	206
3.28. Using the usertty() driver	53
3.29. A simple log statement	54
3.30. Using log path flags	208
3.31. Using embedded log paths	56
3.32. Sizing parameters for flow-control	57
3.33. A simple filter statement	58
3.34. Optimizing regular expressions in filters	60
3.35. Using templates	61
3.36. Segmenting hostnames separated with a dash	216
3.37. Parsing Apache log files	217
3.38. Segmenting a part of a message	217
3.39. Defining pattern databases	64
3.40. Using classification results	64
3.41. Using classification results for filtering messages	219
3.42. Using pattern parsers as macros	220
3.43. Using substitution rules	225
3.44. Setting message fields to a particular value	226



3.45. Using global options	67
3.46. Enabling disk-based buffering	203
3.47. A destination statement using TLS	69
3.48. A source statement using TLS	69
3.49. Disabling mutual authentication	70
3.50. A destination statement using mutual authentication	71
3.51. A source statement using TLS	72
3.52. A simple configuration for clients	73
3.53. A simple configuration for relays	74
3.54. A simple configuration for servers	75
5.1. Collecting the logs of multiple applications from a single folder	102
7.1. Skipping messages	140
8.1. Using the internal() driver	141
8.2. Using the file() driver	146
8.3. Tailing files	146
8.4. Using wildcards in the filename	146
8.5. Monitoring multiple directories	146
8.6. Using the pipe() driver	150
8.7. Using the program() driver	154
8.8. Using the sun-streams() driver	160
8.9. Using the syslog() driver	165
8.10. Using the udp() and tcp() drivers	173
8.11. Using the unix-stream() and unix-dgram() drivers	179
8.12. Using the file() driver	182
8.13. Using the file() driver with macros in the file name and a template for the message	182
8.14. Using the logstore() driver	186
8.15. Using the pipe() driver	189
8.16. Using the program() destination driver	191
8.17. Using the sql() driver	193
8.18. Using the sql() driver with an Oracle database	194
8.19. Using the sql() driver with an MSSQL database	194
8.20. Using SQL NULL values	195
8.21. Using the syslog() driver	199
8.22. Using the tcp() driver	203
8.23. Enabling disk-based buffering	203
8.24. Using the unix-stream() driver	206
8.25. Using the usertty() driver	207
8.26. Using log path flags	208
8.27. Segmenting hostnames separated with a dash	216
8.28. Parsing Apache log files	217
8.29. Segmenting a part of a message	217
8.30. Adding the end of the message to the last column	217
8.31. Pattern parser syntax	218
8.32. Using the STRING and ESTRING parsers	219
8.33. Using classification results for filtering messages	219
8.34. Using pattern parsers as macros	220
8.35. A V1 pattern database containing a single rule	222
8.36. A V2 pattern database containing a single rule	224



8.37. Using substitution rules	225
8.38. Setting message fields to a particular value	226
8.39. Using Posix regular expressions	226
8.40. Using PCRE regular expressions	227



List of Procedures

2.1. The route of a log message in syslog-ng	7
3.1. Configuring TLS on the syslog-ng clients	68
3.2. Configuring TLS on the syslog-ng server	69
3.3. Configuring TLS on the syslog-ng clients	71
3.4. Configuring TLS on the syslog-ng server	72
3.5. Configuring syslog-ng on client hosts	73
3.6. Configuring syslog-ng on relay hosts	74
3.7. Configuring syslog-ng on server hosts	75
3.8. Creating syslog-ng core files	76
4.1. Installing syslog-ng in client or relay mode	79
4.2. Installing syslog-ng in server mode	81
4.3. Installing syslog-ng on RPM-based systems	85
4.4. Installing syslog-ng on Debian-based systems	86
4.5. Compiling syslog-ng from source	87
4.6. Configuring Microsoft SQL Server to accept logs from syslog-ng	89
5.1. Installing the syslog-ng agent in standalone mode	94
5.2. Installing the syslog-ng agent on the hosts of a domain	95
5.3. Upgrading syslog-ng Agent for Windows 3.0.1 to version 3.0.2	97
5.4. Configuring the destination logservers	98
5.5. Limiting the number of messages	99
5.6. Managing eventlog sources	100
5.7. Determining the name of a custom eventlog container	100
5.8. Managing file sources	101
5.9. Configuring global settings	102
5.10. Disabling sources and filters globally	103
5.11. Enabling encrypted connections	103
5.12. Configuring mutual authentication with the syslog-ng Agent for Windows	104
5.13. Importing certificates with MMC	105
5.14. Filtering eventlog messages	106
5.15. Filtering file messages	107
5.16. Customizing messages using templates	108
5.17. Customizing eventlog messages	108
5.18. Customizing file messages	109
5.19. Creating an XML configuration file for the syslog-ng agent	114
5.20. Turning on security logging on Windows XP	119
5.21. Turning on security logging for domain controllers	120
5.22. Turning on auditing on Windows 2003 Server	120
6.1. Upgrading the syslog-ng Agent for IBM System i	122
6.2. Forwarding Apache server logs from System i	132
6.3. Forwarding OpenSSH server logs from System i	133
7.1. Resolving hostnames locally	138
7.2. Collecting logs from chroot	139
7.3. Replacing klogd on Linux	140



Preface

Welcome to the syslog-ng Administrator Guide!

This document describes how to configure and manage syslog-ng. Background information for the technology and concepts used by the product is also discussed.

1. Summary of contents

Chapter 1, *Introduction to syslog-ng* (p. 1) describes the main functionality and purpose of syslog-ng.

Chapter 2, *The concepts of syslog-ng* (p. 6) discusses the technical concepts and philosophies behind syslog-ng.

Chapter 3, *Configuring syslog-ng* (p. 30) provides detailed description on configuring and managing syslog-ng as a client or a server.

Chapter 4, *Installing syslog-ng* (p. 78) describes how to install syslog-ng on various UNIX-based platforms using the precompiled binaries, and how to compile syslog-ng Open Source Edition from source.

Chapter 5, *Collecting logs from Windows hosts* (p. 93) describes how to install and configure the syslog-ng Agent for Windows application.

Chapter 6, *Collecting logs from IBM System i* (p. 121) describes how to install and configure the syslog-ng Agent for IBM System i application.

Chapter 7, *Best practices and examples* (p. 136) gives recommendations to configure special features of syslog-ng.

Chapter 8, *Reference* (p. 141) is a reference guide of syslog-ng, describing all available parameters and options.

Appendix 1, *The syslog-ng manual pages* (p. 234) contains the manual pages of the syslog-ng application.

Appendix 2, *BalaBit syslog-ng Premium Edition License contract* (p. 243) includes the text of the End-User License Agreement applicable to syslog-ng Premium Edition.

Appendix 3, *GNU General Public License* (p. 250) includes the text of the GNU General Public License applicable to syslog-ng Open Source Edition.

Glossary (p. 261) provides definitions of important terms used in this guide.

Index (p. 265) provides cross-references to important terms used in this guide.

2. Target audience and prerequisites

This guide is intended for system administrators and consultants responsible for designing and maintaining logging solutions and log centers. It is also useful for IT decision makers looking for a tool to implement centralized logging in heterogeneous environments.

The following skills and knowledge are necessary for a successful syslog-ng administrator:



- At least basic system administration knowledge.
- An understanding of networks, TCP/IP protocols, and general network terminology.
- Working knowledge of the UNIX or Linux operating system.
- In-depth knowledge of the logging process of various platforms and applications.
- An understanding of the legacy syslog (BSD-syslog) protocol (see RFC 3164, available at <http://www.ietf.org/rfc/rfc3164.txt>) and the new syslog (IETF-syslog) protocol standard (see RFC 5424-5428, available at <http://tools.ietf.org/html/rfc5424>).

3. Products covered in this guide

This guide describes the use of the following syslog-ng versions:

- syslog-ng Open Source Edition (OSE) v3.0.x
- syslog-ng Premium Edition (PE) v3.0.x and later, including syslog-ng Agent for Windows v3.0.x and later
- syslog-ng Agent for IBM System i

Most of the guide applies equally to both the Open Source and the Premium editions of syslog-ng, with the following exceptions:

- The syslog-ng agent for Microsoft Windows is available only as part of the Premium Edition.
- Disk-based buffering (disk-buffer) is available only in the Premium Edition.
- Only the Premium Edition can store messages in encrypted and timestamped log files (so called logstore).
- The Premium Edition automatically detects configuration changes.
- Only the Premium Edition can read messages from file sources that have wildcards in their path or filename.
- Only the Premium Edition can handle directories recursively: that is, monitor a directory and its subdirectory for log files.
- The Open Source Edition does not require a license file.
- The syslog-ng Agent for IBM System i is a commercial product independent from both syslog-ng OSE and PE, and must be licensed separately.

4. Typographical conventions

Before you start using this guide, it is important to understand the terms and typographical conventions used in the documentation. For more information on specialized terms and abbreviations used in the documentation, see the [Glossary](#) at the end of this document.

The following kinds of text formatting and icons identify special information in the document.

**Tip**

Tips provide best practices and recommendations.

**Note**

Notes provide additional information on a topic and emphasize important facts and considerations.

**Warning**

Warnings mark situations where loss of data or misconfiguration of the device is possible if the instructions are not obeyed.

Command

Commands you have to execute.

Emphasis

Reference items, additional readings.

/path/to/file

File names.

Parameters

Parameter and attribute names.

Label

GUI output messages or dialog labels.

Menu

A submenu in the menu bar.

Button

Buttons in dialog windows.

5. Contact and support information

The syslog-ng Premium Edition and syslog-ng Open Source Edition applications are developed and maintained by BalaBit IT Security Ltd. We are located in Budapest, Hungary. Our address is:

BalaBit IT Security Ltd.
 1464 Budapest P.O. BOX 1279
 Hungary
 Tel: +36 1 371-0540
 Fax: +36 1 208-0875
 E-mail: info@balabit.com
 Web: <http://www.balabit.com/>

5.1. Sales contact

You can directly contact us with sales related topics at the e-mail address [<sales@balabit.com>](mailto:sales@balabit.com).



5.2. Support contact

To subscribe to the mailing list of the syslog-ng community, visit <https://lists.balabit.hu/mailman/listinfo/syslog-ng/>.

To report bugs found in syslog-ng, visit <https://bugzilla.balabit.com/>.

Product support, including 7x24 online support is available for both syslog-ng PE and OSE in various packages. For support options, visit the following page: <http://www.balabit.com/support/packages/>

For syslog-ng OSE, precompiled binary packages are available for free for the supported Linux and BSD platforms at <http://www.balabit.com/network-security/syslog-ng/opensource-logging-system/upgrades/>. See the following link for the list of supported platforms: <http://www.balabit.com/network-security/syslog-ng/opensource-logging-system/support/>

You can register your copy of syslog-ng Premium Edition online on the BalaBit website or by sending the filled registration form. Registration is a prerequisite for all support services. Your product can be registered online at the <http://www.balabit.com/support/registration/> website.

E-mail and telephone support is available for registered users, please write or call us for details.

Support e-mail address: <support@balabit.com>.

Support hotline: +36 1 371 0540 (available from 9 AM to 5 PM CET on weekdays)

The BalaBit Online Support System is available at <https://boss.balabit.com/> and offers 24 hours technical support. This system is available only for registered users with a valid support contract and a MyBalaBit account. To sign up for MyBalaBit, visit the following page: <http://www.balabit.com/mybalabit>.

5.3. Training

BalaBit IT Security Ltd. holds courses for advanced GNU/Linux system administrators. Our experienced system engineers give lectures on syslog-ng administration.

6. About this document

This guide is a work-in-progress document with new versions appearing periodically.

The latest version of this document can be downloaded from the BalaBit website at <http://www.balabit.com/support/documentation/>.

For news and update notifications about the syslog-ng documentation, visit the BalaBit Documentation Blog at <http://robert.blogs.balabit.com>.

6.1. What is new in this main edition of The syslog-ng Administrator Guide?

Version 3.0.x of The syslog-ng Administrator Guide contains the following main changes compared to earlier versions:



- The contents of the guide have been updated for syslog-ng 3.0 and syslog-ng Agent for Windows 3.0.1. Since syslog-ng 3.0 contains many new features (see Section 1.4, “What is new in syslog-ng 3.0?” (p. 2) for details), there are several new sections in the following chapters: Chapter 2, *The concepts of syslog-ng* (p. 6), Chapter 3, *Configuring syslog-ng* (p. 30), and Chapter 8, *Reference* (p. 141).
- Chapter 8, *Reference* (p. 141) has become more like a parameter reference, and most of the descriptions and configuration know-how has been moved to Chapter 3, *Configuring syslog-ng* (p. 30). However, configuration examples are included in both chapters for convenience.
- Earlier versions of this guide contained two chapters called *Best practices and examples* and *Troubleshooting and performance tuning*. Most of the material in these chapters have been moved to the relevant parts of Chapter 2, *The concepts of syslog-ng* (p. 6) and Chapter 3, *Configuring syslog-ng* (p. 30). The remaining material is included in Chapter 7, *Best practices and examples* (p. 136).
- Every driver description in Chapter 8, *Reference* (p. 141) contains every available parameter for the driver, sections like *Common destination driver options* have been removed.
- The syslog-ng Administrator Guide is now published under the Creative Commons Attribution-Non-commercial-No Derivative Works (by-nc-nd) 3.0 license, meaning that it can be freely distributed. See Appendix 4, *Creative Commons Attribution Non-commercial No Derivatives (by-nc-nd) License* (p. 256) for details.

6.2. Feedback

Any feedback is greatly appreciated. General comments, errors found in the text, and any suggestions about how to improve the documentation is welcome at <documentation@balabit.com>.

6.3. Acknowledgments

BalaBit would like to express its gratitude to the syslog-ng users and the syslog-ng community for their invaluable help and support.

Special thanks to Nate Campi for organizing and hosting the syslog-ng FAQ (<http://campin.net/syslog-ng/faq.html>) and for his permission to reproduce parts of his work in this guide.



Chapter 1. Introduction to syslog-ng

This chapter introduces the syslog-ng Premium Edition application in a non-technical manner, discussing how and why it is useful, and the benefits it offers to an existing IT infrastructure.

1.1. What syslog-ng is

The syslog-ng application is a flexible and highly scalable system logging application that is ideal for creating centralized and trusted logging solutions. The main features of syslog-ng are summarized below.

- *Reliable log transfer*: The syslog-ng application enables you to send the log messages of your hosts to remote servers using the latest protocol standards. The logs of different servers can be collected and stored centrally on dedicated log servers. Transferring log messages using the TCP protocol ensures that no messages are lost.
- *Secure logging using TLS*: Log messages may contain sensitive information that should not be accessed by third parties. Therefore, syslog-ng uses the Transport Layer Security (TLS) protocol to encrypt the communication. TLS also allows the mutual authentication of the host and the server using X.509 certificates.
- *Disk-based message buffering*: The Premium Edition of syslog-ng stores messages on the local hard disk if the central log server or the network connection becomes unavailable. The syslog-ng application automatically sends the stored messages to the server when the connection is reestablished, in the same order the messages were received. The disk buffer is persistent – no messages are lost even if syslog-ng is restarted.
- *Direct database access*: Storing your log messages in a database allows you to easily search and query the messages and interoperate with log analyzing applications. The Premium Edition of syslog-ng supports the following databases: MSSQL, MySQL, Oracle, PostgreSQL, and SQLite.
- *Encrypted and timestamped log storage*: The Premium Edition of syslog-ng can store log messages securely in encrypted, compressed, and timestamped binary files. Timestamps can be requested from an external Timestamping Authority (TSA).
- *Heterogeneous environments*: The syslog-ng application is the ideal choice to collect logs in massively heterogeneous environments using several different operating systems and hardware platforms, including Linux, Unix, BSD, Sun Solaris, HP-UX, and AIX. An agent is available to transfer logs from Microsoft Windows hosts to the central syslog-ng server.
- *Filter and classify*: The syslog-ng application can sort the incoming log messages based on their content and various parameters like the source host, application, and priority. Directories, files, and database tables can be created dynamically using macros. Complex filtering using regular expressions and boolean operators offers almost unlimited flexibility to forward only the important log messages to the selected destinations.
- *Parse and rewrite*: The syslog-ng application can segment log messages to named fields or columns, and also modify the values of these fields.
- *IPv4 and IPv6 support*: The syslog-ng application can operate in both IPv4 and IPv6 network environments; it can receive and send messages to both types of networks.



1.2. What syslog-ng is not

The syslog-ng application is not log analysis software. It can filter log messages and select only the ones matching certain criteria. It can even convert the messages and restructure them to a predefined format, or parse the messages and segment them into different fields. But syslog-ng cannot interpret and analyze the meaning behind the messages, or recognize patterns in the occurrence of different messages.

1.3. Why is syslog-ng needed?

Log messages contain information about the events happening on the hosts. Monitoring system events is essential for security and system health monitoring reasons.

The original syslog protocol separates messages based on the priority of the message and the facility sending the message. These two parameters alone are often inadequate to consistently classify messages, as many applications might use the same facility — and the facility itself is not even included in the log message. To make things worse, many log messages contain unimportant information. The syslog-ng application helps you to select only the really interesting messages, and forward them to a central server.

Company policies or other regulations often require log messages to be archived. Storing the important messages in a central location greatly simplifies this process.

For details on how can you use syslog-ng to comply with various regulations, see the *Regulatory compliance and system logging* whitepaper available at <http://www.balabit.com/support/documentation/>

1.4. What is new in syslog-ng 3.0?

Version 3.0 of syslog-ng includes the following main features:

- Support for the new IETF syslog protocol standard — see Section 2.18.2, “IETF-syslog messages” (p. 26), Section 3.3.5, “Collecting messages using the IETF syslog protocol” (p. 39) and Section 3.4.6, “Sending messages to a remote logserver using the IETF-syslog protocol” (p. 51).
- Parsing and segmenting log messages — see Section 3.8, “Parsing messages” (p. 62).
- Rewriting log messages — see Section 3.10, “Rewriting messages” (p. 66).
- Storing log messages in encrypted, timestamped logfiles — see Section 2.8, “Secure storage of log messages” (p. 14) and Section 3.4.2, “Storing messages in encrypted files” (p. 44).
- Complex, embedded log paths — see Section 2.2.1, “Embedded log statements” (p. 8) and Section 3.5.1, “Using embedded log statements” (p. 55).
- File sources with wildcards in their filename or path — see Section 3.3.2, “Collecting messages from text files” (p. 37).
- The syslog-ng application can receive messages directly from external applications using the new `program()` source driver that listens for log messages on the standard output (stdout) — see Section 8.1.4, “program()” (p. 150).
- On Linux, the syslog-ng application can support capabilities and run as a non-root user if compiled with the `--enable-linux-caps` option.



- The syslog-ng application automatically generates a unique sequence ID for every new local message (but not for relayed messages). This ID number is included in outgoing messages that use the IETF-syslog format, and can be included in legacy messages using the `$SEQNUM` macro.
- On-demand log statistics can be requested from syslog-ng via a unix-domain socket. See Section 3.3.1.1, “Log statistics” (p. 36).
- Starting with syslog-ng Open Source Edition 3.0.2, the precompiled binary packages are available for free for the supported Linux and BSD platforms at <http://www.balabit.com/network-security/syslog-ng/opensource-logging-system/upgrades/>.

Version 3.0 of syslog-ng includes the following important changes:

- The `tcp`, `tcp6`, `udp`, `udp6`, `unix-stream`, and `unix-dgram` destination drivers support the `keep-alive` option, enabling them to keep connections open during a HUP and saving the output queue between restarts — see Section 8.2.7, “`tcp()`, `tcp6()`, `udp()`, and `udp6()`,” (p. 200) and Section 8.2.8, “`unix-stream()` & `unix-dgram()`” (p. 204).
- The `log-prefix()` option has been deprecated. Use the new `program-override()` and `host-override()` options instead — see Section 8.2.7, “`tcp()`, `tcp6()`, `udp()`, and `udp6()`,” (p. 200) and Section 8.2.8, “`unix-stream()` & `unix-dgram()`” (p. 204).
- The `keep_hostname`, `keep_timestamp`, `use_dns`, and `use_fqdn` options can be set individually for every source.
- Legacy destination drivers like `tcp` and `file` can output log messages in the new IETF-syslog format if the `flags(syslog-protocol)` option is enabled for the destination. Similarly, legacy sources can receive such messages using this option.
- If syslog-ng is compiled with PCRE support, Perl Compatible Regular Expressions can be used using the `type(pcre)` option.
- You can set the part of the message where the `match()` filter searches for the specified string using macros (e.g., `match("example" value(PROGRAM))`).
- The default value of the `follow_freq` option has been changed to 1.
- The default value of the `chain_hostnames` option has been changed to 0 (no).
- The default value of the `template_escape` option has been changed to 0 (no).
- `NL` characters are not removed by default, to remove these characters, use the `flags(no-multi-line)` option of the destination.
- The installation packages for syslog-ng 3.0 PE are `.run` binaries that include every dependency to simplify the installation process.

1.5. Who uses syslog-ng?

The syslog-ng application is used worldwide by companies and institutions who collect and manage the logs of several hosts, and want to store them in a centralized, organized way. Using syslog-ng is particularly advantageous for:

- Internet Service Providers;
- Financial institutions and companies requiring policy compliance;



- Server, web, and application hosting companies;
- Datacenters;
- Wide area network (WAN) operators;
- Server farm administrators.

The following is a list of public references — companies who use syslog-ng in their production environment:

- Allianz Hungary Insurance Co. (<http://www.allianz.hu/>)
- Navisite Inc. (<http://www.navisite.com/>)
- Svenska Handelsbanken AB (<http://www.handelsbanken.com/>)
- Swedish National Debt Office (<http://www.riksgalden.se>)

1.6. Supported platforms

The syslog-ng PE application is officially supported on the following platforms. Note that the following table is for general reference only, and is not always accurate about the supported platforms and options available for specific platforms. The latest version of this table is available at <http://www.balabit.com/network-security/syslog-ng/central-syslog-server/>.

	x86	x86_64	SUN SPARC	ppc32	ppc64	PA-RISC
AIX 5.2 & 5.3	X	X	X	✓	upon request	X
Debian etch	✓	✓	X	X	X	X
FreeBSD 6.1 *	✓	upon request	upon request	X	X	X
HP-UX 11i	X	X	X	X	X	✓
IBM System i	X	X	X	✓	X	X
Red Hat ES 4 / CentOS 4	✓	✓	X	X	X	X
Red Hat ES 5 / CentOS 5	✓	✓	X	X	X	X
SLES 10 / openSUSE 10.0	✓	upon request	X	X	X	X
SLES 10 SP1 / openSUSE 10.1	✓	✓	X	X	X	X
Solaris 8	X	X	✓	X	X	X
Solaris 9	upon request	X	✓	X	X	X
Solaris 10	upon request	✓	✓	X	X	X
Windows	✓	✓	X	X	X	X

*Oracle database access is not supported

Table 1.1. Platforms supported by syslog-ng PE

The central syslog-ng server cannot be installed on Microsoft Windows platforms. The syslog-ng Agent for Windows capable of forwarding eventlog messages to the central server is available on the x86 and x86_64 architecture for Microsoft Windows XP, Microsoft Windows 2003 Server, Microsoft Windows Vista, and Microsoft Windows 2008 Server. The syslog-ng Agent is available only in syslog-ng Premium Edition.



The central syslog-ng server can be installed on the IBM System i platform, but the syslog-ng Agent for IBM System i is needed to collect the native logs of IBM System i (see Chapter 6, *Collecting logs from IBM System i* (p. 121)). The syslog-ng Agent for IBM System i is a commercial product independent from both syslog-ng OSE and PE, and must be licensed separately.

For syslog-ng OSE, precompiled binary packages are available for free for the supported Linux and BSD platforms at <http://www.balabit.com/network-security/syslog-ng/opensource-logging-system/upgrades/>. Precompiled binary packages for HP-UX, IBM AIX, and Solaris are available for an annual fee at the BalaBit webshop at <http://www.balabit.com/shop/>. For the list of available platforms, see <http://www.balabit.com/network-security/syslog-ng/opensource-logging-system/support/>.



Chapter 2. The concepts of syslog-ng

This chapter discusses the technical concepts of syslog-ng.

2.1. The philosophy of syslog-ng

Typically, syslog-ng is used to manage log messages and implement centralized logging, where the aim is to collect the log messages of several devices on a single, central log server. The different devices — called syslog-ng clients — all run syslog-ng, and collect the log messages from the various applications, files, and other *sources*. The clients send all important log messages to the remote syslog-ng server, where the server sorts and stores them.

2.2. Logging with syslog-ng

The syslog-ng application reads incoming messages and forwards them to the selected *destinations*. The syslog-ng application can receive messages from files, remote hosts, and other *sources*.

Log messages enter syslog-ng in one of the defined sources, and are sent to one or more *destinations*.

Sources and destinations are independent objects; *log paths* define what syslog-ng does with a message, connecting the sources to the destinations. A log path consists of one or more sources and one or more destinations; messages arriving to a source are sent to every destination listed in the log path. A log path defined in syslog-ng is called a *log statement*.

Optionally, log paths can include *filters*. Filters are rules that select only certain messages, for example, selecting only messages sent by a specific application. If a log path includes filters, syslog-ng sends only the messages satisfying the filter rules to the destinations set in the log path.

Other optional elements that can appear in log statements are *parsers* and *rewriting rules*. Parsers segment messages into different fields to help processing the messages, while rewrite rules modify the messages by adding, replacing, or removing parts of the messages.

The following procedure illustrates the route of a log message from its source on the syslog-ng client to its final destination on the central syslog-ng server.

Procedure 2.1. The route of a log message in syslog-ng




Figure 2.1. The route of a log message

1. A device or application sends a log message to a source on the syslog-ng client. For example, an Apache web server running on Linux enters a message into the `/var/log/apache` file.
2. The syslog-ng client running on the web server reads the message from its `/var/log/apache` source.
3. The syslog-ng client processes the first log statement that includes the `/var/log/apache` source.
4. The syslog-ng client performs optional operations (message filtering, parsing, and rewriting) on the message; for example, it compares the message to the filters of the log statement (if any). If the message complies with all filter rules, syslog-ng sends the message to the destinations set in the log statement, for example, to the remote syslog-ng server.



Warning

Message filtering, parsing, and rewriting is performed in the order that the operations appear in the log statement.



Note

The syslog-ng client sends a message to *all* matching destinations by default. As a result, a message may be sent to a destination more than once, if the destination is used in multiple log statements. To prevent such situations, use the `final` flag in the destination statements. See Table 8.16, “Log statement flags” (p. 207) for details.



5. The syslog-ng client processes the next log statement that includes the `/var/log/apache` source, repeating Steps 3-4.
6. The message sent by the syslog-ng client arrives to a source set in the syslog-ng server.
7. The syslog-ng server reads the message from its source and processes the first log statement that includes that source.
8. The syslog-ng server performs optional operations (message filtering, parsing, and rewriting) on the message; for example, it compares the message to the filters of the log statement (if any). If the message complies with all filter rules, syslog-ng sends the message to the destinations set in the log statement.

**Warning**

Message filtering, parsing, and rewriting is performed in the order that the operations appear in the log statement.

9. The syslog-ng server processes the next log statement, repeating Steps 7-9.

**Note**

The syslog-ng application can stop reading messages from its sources if the destinations cannot process the sent messages. This feature is called flow-control and is detailed in Section 2.13, “Managing incoming and outgoing messages with flow-control” (p. 18).

2.2.1. Embedded log statements

Starting from version 3.0, syslog-ng can handle embedded log statements (also called log pipes). Embedded log statements are useful for creating complex, multi-level log paths with several destinations and use filters, parsers, and rewrite rules.

For example, if you want to filter your incoming messages based on the facility parameter, and then use further filters to send messages arriving from different hosts to different destinations, you would use embedded log statements.




Figure 2.2. Embedded log statement

Embedded log statements include sources — and usually filters, parsers, rewrite rules, or destinations — and other log statements that can include filters, parsers, rewrite rules, and destinations. The following rules apply to embedded log statements:

- Only the beginning (also called top-level) log statement can include sources.
- Embedded log statements can include multiple log statements on the same level (i.e., a top-level log statement can include two or more log statements).
- Embedded log statements can include several levels of log statements (i.e., a top-level log statement can include a log statement that includes another log statement, and so on).
- Only another log statement can follow an embedded log statement, filters or other rules cannot.
- Embedded log statements that are on the same level receive the same messages from the higher-level log statement. For example, if the top-level log statement includes a filter, the lower-level log statements receive only the messages that pass the filter.

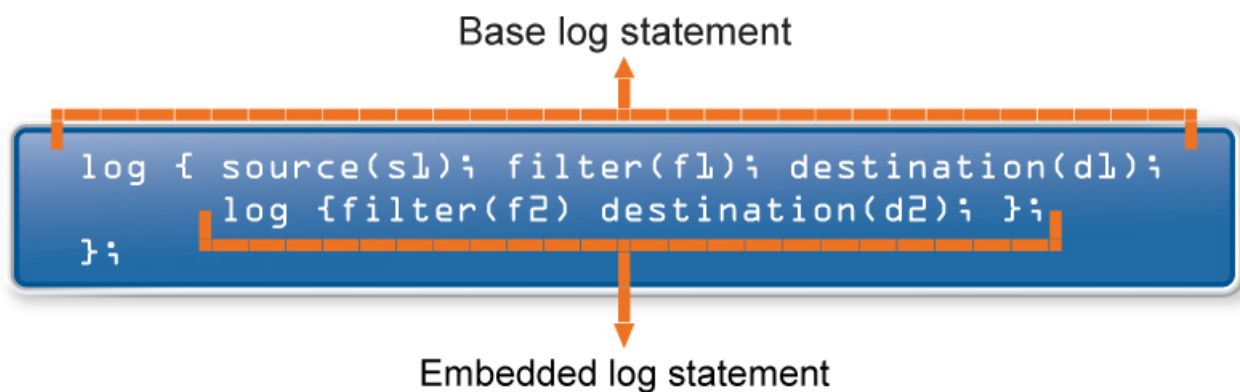


Figure 2.3. Embedded log statements

Embedded log filters can be used to optimize the processing of log messages, for example, to re-use the results of filtering and rewriting operations.



2.3. Modes of operation

The syslog-ng Premium Edition application has three distinct modes of operation: *Client*, *Server*, and *Relay*. The syslog-ng application running on a host determines the mode of operation automatically based on the license and the configuration file.



Note

Microsoft Windows based hosts can run only the syslog-ng agent. The syslog-ng agent operates only in client mode.

2.3.1. Client mode

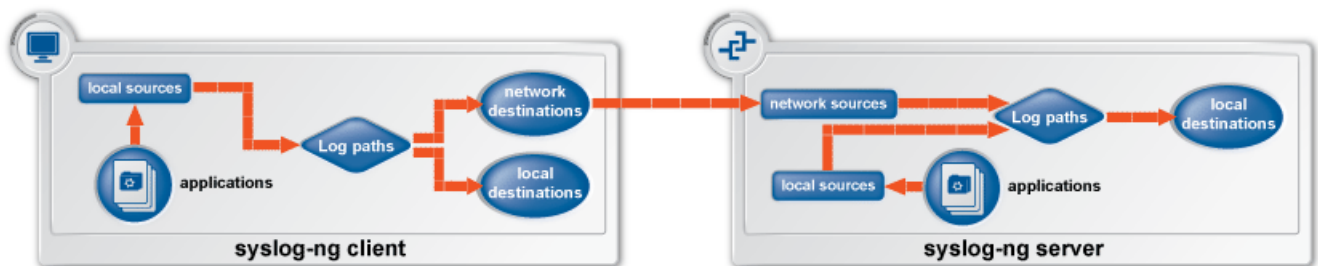


Figure 2.4. Client-mode operation

In client mode, syslog-ng collects the local logs generated by the host and forwards them through a network connection to the central syslog-ng server or to a relay. Clients can also log the messages locally into files.

No license file is required to run syslog-ng in client mode.

2.3.2. Relay mode




Figure 2.5. Relay-mode operation

In relay mode, syslog-ng receives logs through the network from syslog-ng clients and forwards them to the central syslog-ng server using a network connection. Relays can also log the messages from the relay host into a local file, or forward these messages to the central syslog-ng server.

Relays cannot write messages received from the network into local files, only buffer the messages to the hard disk when disk-based buffering is used.



No license file is required to run syslog-ng in relay mode.

2.3.3. Server mode




Figure 2.6. Server-mode operation

In server mode, syslog-ng acts as a central log-collecting server. It receives messages from syslog-ng clients and relays over the network, and stores them locally in files, or passes them to other applications, e.g., log analyzers.

Running syslog-ng Premium Edition in server mode requires a license file. The license determines how many individual hosts can connect to the server.

Running syslog-ng Open Source Edition in server mode does not require a license file.

2.4. Global objects

The syslog-ng application uses the following objects:

- *Source driver*: A communication method used to receive log messages. For example, syslog-ng can receive messages from a remote host via TCP/IP, or read the messages of a local application from a file.
- *Source*: A named collection of configured source drivers.
- *Destination driver*: A communication method used to send log messages. For example, syslog-ng can send messages to a remote host via TCP/IP, or write the messages into a file or database.
- *Destination*: A named collection of configured destination drivers.
- *Filter*: An expression to select messages. For example, a simple filter can select the messages received from a specific host.



- *Macro*: An identifier that refers to a part of the log message. For example, the `$HOST` macro returns the name of the host that sent the message. Macros are often used in templates and filenames.
- *Parser*: A rule that segments messages into separate columns at a predefined separator character (e.g., a comma). Every column has a unique name that can be used as a macro.
- *Rewrite rule*: A rule modifies a part of the message, for example, replaces a string, or sets a field to a specified value.
- *Log paths*: A combination of sources, destinations, and other objects like filters, parsers, and rewrite rules. The syslog-ng application sends messages arriving to the sources of the log paths to the defined destinations, and performs filtering, parsing, and rewriting of the messages. Log paths are also called log statements. Log statements can include other (embedded) log statements to create complex log paths.
- *Template*: A template is a set of macros that can be used to restructure log messages or automatically generate file names. For example, a template can add the hostname and the date to the beginning of every log message.
- *Option*: Options set global parameters of syslog-ng, like the parameters of name resolution and timezone handling.

For details on the above objects, see Section 3.2, “Defining global objects” (p. 32).

2.5. Timezone handling

The syslog-ng application supports messages originating from different timezones. The original syslog protocol does not include timezone information, but syslog-ng provides a solution by extending the syslog protocol to include the timezone in the log messages. The syslog-ng application also enables administrators to supply timezone information for legacy devices which do not support the protocol extension.

Timezone information is associated with messages entering syslog-ng is selected using the following algorithm:

1. The sender application (e.g., the syslog-ng client) or host specifies the timezone of the messages. If the incoming message includes a timezone it is associated with the message. Otherwise, the local timezone is assumed.
2. Specify the `_time_zone()` parameter for the source driver that reads the message. This parameter overrides the original timezone of the message. Each source defaults to the value of the `_recv_time_zone()` global option.
3. Specify the timezone in the destination driver using the `_time_zone()` parameter. Each destination driver might have an associated timezone value; syslog-ng converts message timestamps to this timezone before sending the message to its destination (file or network socket). Each destination defaults to the value of the `_send_time_zone()` global option.



Note

A message can be sent to multiple destination zones. The syslog-ng application converts the timezone information properly for every individual destination zone.

4. If the timezone is not specified, the message is left unchanged.
5. When macro expansions are used in the destination filenames, the local timezone is used.

2.6. Daylight saving changes

The syslog-ng application receives the timezone and daylight saving information from the operating system it is installed on. If the operating system handles daylight saving correctly, so does syslog-ng.

2.7. Secure logging using TLS

The syslog-ng application can send and receive log messages securely over the network using the Transport Layer Security (TLS) protocol. TLS is an encryption protocol over the TCP/IP network protocol, so it can be used only with TCP-based sources and destinations (`tcp()` and `tcp6()`).

TLS uses certificates to authenticate and encrypt the communication, as illustrated on the following figure:



Figure 2.7. Certificate-based authentication

The client authenticates the server by requesting its certificate and public key. Optionally, the server can also request a certificate from the client, thus mutual authentication is also possible.

In order to use TLS encryption in syslog-ng, the following elements are required:

- A certificate on the syslog-ng server that identifies the syslog-ng server.
- The certificate of the Certificate Authority that issued the certificate of the syslog-ng server must be available on the syslog-ng client.

When using mutual authentication to verify the identity of the clients, the following elements are required:

- A certificate must be available on the syslog-ng client. This certificate identifies the syslog-ng client.
- The certificate of the Certificate Authority that issued the certificate of the syslog-ng client must be available on the syslog-ng server.

Mutual authentication ensures that the syslog-ng server accepts log messages only from authorized clients.



See Section 3.13, “Encrypting log messages with TLS” (p. 68) for details on configuring TLS communication in syslog-ng.

2.8. Secure storage of log messages

The Premium Edition of syslog-ng can store log messages securely in encrypted, compressed and timestamped binary files. Timestamps can be requested from an external Timestamping Authority (TSA).

Logstore files consist of individual chunks, every chunk can be encrypted, compressed, and timestamped separately. Chunks contain log message data, chunk size defaults to 128k (about 1MB worth of compressed logs). Chunks are closed when their size reaches the limit set in the *chunk_size* parameter, or when the time limit set in the *chunk_time* parameter expires and a new message arrives. Specifically, when a new message arrives to the logstore, syslog-ng checks if *chunk_time* time has elapsed since the last message has arrived. If it has, then the old chunk is closed and the new message is written into a new chunk.

The syslog-ng PE application generates an SHA-1 hash for every chunk to verify the integrity of the chunk. The hashes of the chunks are chained together to prevent injecting chunks into the logstore file. The syslog-ng application can encrypt the logstore using the *aes128* algorithm in CBC mode; the hashing (HMAC) algorithm is *hmac-sha1*. Currently it is not possible to use other algorithms.

2.9. Formatting messages, filenames, directories, and tablenamees

The syslog-ng application can dynamically create filenames, directories, or names of database tables using macros that help you organize your log messages. Macros refer to a property or a part of the log message, for example, the *\$HOST* macro refers to the name or IP address of the client that sent the log message, while *\$DAY* is the day of the month when syslog-ng has received the message. Using these macros in the path of the destination log files allows you for example to collect the logs of every host into separate files for every day.

A set of macros can be defined as a template object and used in multiple destinations.

Another use of macros and templates is to customize the format of the syslog message, for example to add elements of the message header to the message text. Note that if a message uses the IETF-syslog format, only the text of the message can be customized, the structure of the header is fixed.

For details on using templates and macros, see Section 3.7, “Templates and macros” (p. 60) and Section 8.5, “Macros” (p. 210).

2.10. Segmenting messages

The filters and default macros of syslog-ng work well on the headers and metainformation of the log messages, but are rather limited when processing the content of the messages. Parsers can segment the content of the messages into name-value pairs, and these names can be used as user-defined macros. Subsequent filtering or other type of processing of the message can use these custom macros to refer to parts of the message.

Parsers are global objects most often used together with filters and rewrite rules. For details on using parsers, see Section 3.8, “Parsing messages” (p. 62) and Section 8.6, “Message parsers” (p. 214).



2.11. Modifying messages

The syslog-ng application can rewrite parts of the messages using rewrite rules. Rewrite rules are global objects similar to parsers and filters and can be used in log paths. The syslog-ng application has two methods to rewrite parts of the log messages: replacing (setting) a part of the message to a fix value, and a general search-and-replace mode.

Substitution completely replaces a specific part of the message that is referenced using a built-in or user-defined macro.

General rewriting searches for a string in the entire message (or only a part of the message specified by a macro) and replaces it with another string. Optionally, this replacement string can be a template that contains macros.

For details on using rewrite rules, see Section 3.10, “Rewriting messages” (p. 66) and Section 8.7, “Rewriting messages” (p. 225).

2.12. Classifying log messages

The syslog-ng application can compare the contents of the received log messages to predefined message patterns. By comparing the messages to the known patterns, syslog-ng is able to identify the exact type of the messages, and sort them into message classes. The message classes can be used to classify the type of the event described in the log message. The message classes can be customized, and for example can label the messages as user login, application crash, file transfer, etc. events.

To find the pattern that matches a particular message, syslog-ng uses a method called longest prefix match radix tree. This means that syslog-ng creates a tree structure of the available patterns, where the different characters available in the patterns for a given position are the branches of the tree. This is also illustrated on the following figure:

To classify a message, syslog-ng selects the first character of the message (the text of message, not the header), and selects the patterns starting with this character, other patterns are ignored for the rest of the process. After that, the second character of the message is compared to the second character of the selected patterns. Again, matching patterns are selected, and the others discarded. This process is repeated until a single pattern completely matches the message, or no match is found. In the latter case, the message is classified as unknown, otherwise the class of the matching pattern is assigned to the message.

To make the message classification more flexible and robust, the patterns can contain pattern parsers: elements that match on a set of characters. For example, the NUMBER parser matches on any integer numbers (e.g., 1, 123, 894054, etc.). Other pattern parsers match on various strings and IP addresses. For the details of available pattern parsers, see Section 2.12.3, “Artificial ignorance” (p. 18).

The functionality of the pattern database is similar to that of the logcheck project, but it is much easier to write and maintain the patterns used by syslog-ng, than the regular expressions used by logcheck. Also, it is much easier to understand syslog-ng patterns than regular expressions.

Pattern matching based on regular expressions is computationally very intensive, especially when the number of patterns increases. The solution used by syslog-ng can be performed real-time, and is independent from the number of patterns, so it scales much better. The following patterns describe the same message: *Accepted password for bazsi from 10.50.0.247 port 42156 ssh2*



A regular expression matching this message from the logcheck project: `Accepted \ (gssapi (-with-mic|-keyex)?|rsa|dsa|password|publickey|keyboard-interactive/pam) \ for [^[:space:]]+ from [^[:space:]]+ port [0-9]+((ssh|ssh2))?`

A syslog-ng database pattern for this message: `Accepted @QSTRING:auth_method: @for@QSTRING:username: @from\ @QSTRING:client_addr: @port @NUMBER:port:@ssh2`

For details on using pattern databases to classify log messages, see Section 3.9, “Classifying messages” (p. 63) and Section 8.6.2, “Pattern databases” (p. 218).

2.12.1. The structure of the pattern database

The pattern database is organized as follows:

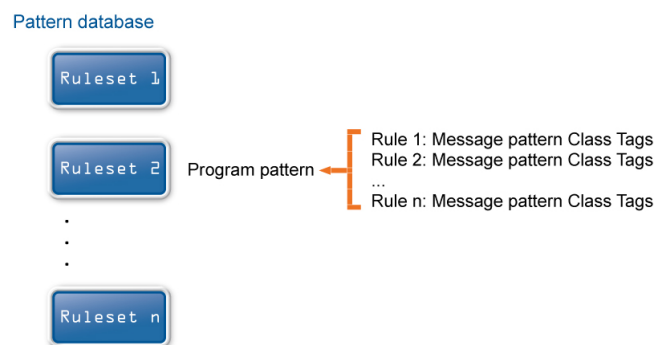


Figure 2.8. The structure of the pattern database

- The pattern database consists of rulesets. A ruleset consists of a Program Pattern and a set of rules: the rules of a ruleset are applied to log messages if the name of the application that sent the message matches the Program Pattern of the ruleset. The name of the application (the content of the \$PROGRAM macro) is compared to the Program Patterns of the available rulesets, and then the rules of the matching rulesets are applied to the message.
- The Program Pattern can be a string that specifies the name of the application or the beginning of its name (e.g., to match for sendmail, the program pattern can be sendmail, or just send), and the Program Pattern can contain pattern parsers. Note that pattern parsers are completely independent from the syslog-ng parsers used to segment messages. Additionally, every rule has a unique identifier: if a message matches a rule, the identifier of the rule is stored together with the message.
- Rules consist of a message pattern and a class. The Message Pattern is similar to the Program Pattern, but is applied to the message part of the log message (the content of the \$MESSAGE macro). If a message pattern matches the message, the class of the rule is assigned to the message (e.g., Security, Violation, etc.).
- Rules can also contain additional information about the matching messages, such as the description of the rule, an URL, or free-form tags.
- Patterns can consist of literals (keywords, or rather, keycharacters) and pattern parsers.

**Note**

If the \$PROGRAM part of a message is empty, rules with an empty Program Pattern are used to classify the message.

If the same Program Pattern is used in multiple rulesets, the rules of these rulesets are merged, and every rule is used to classify the message. Note that message patterns must be unique within the merged rulesets, but the currently only one ruleset is checked for uniqueness.

2.12.2. How pattern matching works

A sample log message:



Figure 2.9. Applying patterns

The followings describe how patterns work. This information applies to program patterns and message patterns alike, even though message patterns are used to illustrate the procedure.

Patterns can consist of literals (keywords, or rather, keycharacters) and pattern parsers. Pattern parsers attempt to parse a sequence of characters according to certain rules.

**Note**

Wildcards and regular expressions cannot be used in patterns. The @ character must be escaped, i.e., to match for this character, you have to write @@ in your pattern. This is required because pattern parsers of syslog-ng are enclosed between @ characters.

When a new message arrives, syslog-ng attempts to classify it using the pattern database. The available patterns are organized alphabetically into a tree, and syslog-ng inspects the message character-by-character, starting from the beginning. This approach ensures that only a small subset of the rules must be evaluated at any given step, resulting in high processing speed. Note that the speed of classifying messages is practically independent from the total number of rules.

For example, if the message begins with the *Apple* string, only patterns beginning with the character *A* are considered. In the next step, syslog-ng selects the patterns that start with *Ap*, and so on, until there is no more specific pattern left.

Note that literal matches take precedence over pattern parser matches: if at a step there is a pattern that matches the next character with a literal, and another pattern that would match it with a parser, the pattern with the literal match is selected. Using the previous example, if at the third step there is the literal pattern *Apport* and a pattern parser *Ap@STRING@*, the *Apport* pattern is matched, even if the pattern parser would result in a better match.

If there are two parsers at the same level (e.g., *Ap@STRING@* and *Ap@QSTRING@*), it is random which pattern is applied (technically, the one that is loaded first). However, if the selected parser cannot parse at least one character of the message, the other parser is used. But having two different parsers at the same level is extremely rare, so the impact of this limitation is much less than it appears.



2.12.3. Artificial ignorance

Artificial ignorance is a method to detect anomalies. When applied to log analysis, it means that you ignore the regular, common log messages - these are the result of the regular behavior of your system, and therefore are not too interesting. However, new messages that have not appeared in the logs before can sign important events, and should be therefore investigated. "By definition, something we have never seen before is anomalous" (Marcus J. Ranum).

The syslog-ng application can classify messages using a pattern database: messages that do not match any pattern are classified as unknown. This provides a way to use artificial ignorance to review your log messages. You can periodically review the unknown messages — syslog-ng can send them to a separate destination - and add patterns for them to the pattern database. By reviewing and manually classifying the unknown messages, you can iteratively classify more and more messages, until the only the really anomalous messages show up as unknown.

Obviously, for this to work, a large number of message patterns are required. The radix-tree matching method used for message classification is very effective, can be performed very fast, and scales very well; basically the time required to perform a pattern matching is independent from the number of patterns in the database.

To simplify the building of pattern databases, BalaBit has released (and will continue to release) sample databases. Currently the following pattern databases are available at the BalaBit Download page <http://www.balabit.com/network-security/syslog-ng/log-server-appliance/>:

- a database for the log messages of Cisco PIX firewalls;
- the database of the Logcheck project (<http://logcheck.org/>) containing message patterns for a large number of open source applications;
- a database for the log messages of the Zorp Application Level Gateway (<http://www.balabit.com/network-security/zorp-gateway/>) (developed by BalaBit IT Security).

2.13. Managing incoming and outgoing messages with flow-control

This section describes the internal message-processing model of syslog-ng, as well as the flow-control feature that can prevent message losses. To use flow-control, the *flow-control* flag must be enabled for the particular log path.

The syslog-ng application monitors (polls) the sources defined in its configuration file, periodically checking each source for messages. When a log message is found in one of the sources, syslog-ng polls every source and reads the available messages. These messages are processed and put into the output buffer of syslog-ng (also called fifo). From the output buffer, the operating system sends the messages to the appropriate destinations.

In large-traffic environments many messages can arrive during a single poll loop, therefore syslog-ng reads only a fixed number of messages from each source. The *log_fetch_limit()* option specifies the number of messages read during a poll loop from a single source.




Figure 2.10. Managing log messages in syslog-ng



Note

The `log_fetch_limit()` parameter can be set as a global option, or for every source individually.

Every destination has its own output buffer. The output buffer is needed because the destination might not be able to accept all messages immediately. The `log_fifo_size()` parameter sets the size of the output buffer. The output buffer must be larger than the `log_fetch_limit()` of the sources, to ensure that every message read during the poll loop fits into the output buffer. If the log path sends messages to a destination from multiple sources, the output buffer must be large enough to store the incoming messages of every source.

TCP and unix-stream sources can receive the logs from several incoming connections (e.g., many different clients or applications). For such sources, syslog-ng reads messages from every connection, thus the `log_fetch_limit()` parameter applies individually to every connection of the source.




Figure 2.11. Managing log messages of TCP sources in syslog-ng

The flow-control of syslog-ng introduces a control window to the source that tracks how many messages can syslog-ng accept from the source. Every message that syslog-ng reads from the source lowers the window size by one; every message that syslog-ng successfully sends from the output buffer increases the window size by one. If the window is full (i.e., its size decreases to zero), syslog-ng stops reading messages from the source. The initial size of the control window is by default 100: the `log_fifo_size()` must be larger than this value in order for flow-control to have any effect. If a source accepts messages from multiple connections, all messages use the same control window.

When flow-control is used, every source has its own control window. As a worst-case situation, the output buffer of the destination must be set to accommodate all messages of every control window, that is, the `log_fifo_size()` of the destination must be greater than `number_of_sources*log_iw_size()`. This applies to every source that sends logs to the particular destination. Thus if two sources having several connec-



tions and heavy traffic send logs to the same destination, the control window of both sources must fit into the output buffer of the destination. Otherwise, syslog-ng does not activate the flow-control, and messages may be lost.

**Note**

Flow-control can be used together with the disk-based buffering feature of syslog-ng PE. See Section 2.14, “Using disk-based buffering” (p. 20) for details.

2.13.1. Flow-control and multiple destinations

Using flow-control on a source has an important side-effect if the messages of the source are sent to multiple destinations. If flow-control is in use and one of the destinations cannot accept the messages, the other destinations do not receive any messages either, because syslog-ng stops reading the source. For example, if messages from a source are sent to a remote server and also stored locally in a file, and the network connection to the server becomes unavailable, neither the remote server nor the local file will receive any messages. This side-effect of the flow-control can be avoided by using the disk-based buffering feature of syslog-ng Premium Edition.

**Note**

Creating separate log paths for the destinations that use the same flow-controlled source does not avoid the problem.

2.14. Using disk-based buffering

The Premium Edition of syslog-ng stores messages on the local hard disk if the central log server or the network connection to the server becomes unavailable. The syslog-ng application automatically sends the stored messages to the server when the connection is reestablished. The disk buffer is used as a queue: when the connection to the server is reestablished, syslog-ng sends the messages to the server in the order they were received.

**Note**

Disk-based buffering can be used in conjunction with flow-control. See Section 2.13, “Managing incoming and outgoing messages with flow-control” (p. 18) for details.

Disk buffers can be used with `tcp()`, `tcp6()`, `syslog()` (when using the `tcp` or `tls` transport methods), and `sql()` destinations. Every such destination uses a separate disk buffer (similarly to the output buffers controlled by `log_fifo_size()`). The hard disk space is not pre-allocated, so ensure that there is always enough free space to store the disk buffers even when the disk buffers are full.

If syslog-ng is restarted (using the `/etc/init.d/syslog-ng restart` command), it automatically saves any unsent messages of the disk buffer and the output queue. After the restart, syslog-ng sends the saved messages to the server. In other words, the disk buffer is persistent.

The syslog-ng application handles outgoing messages the following way:

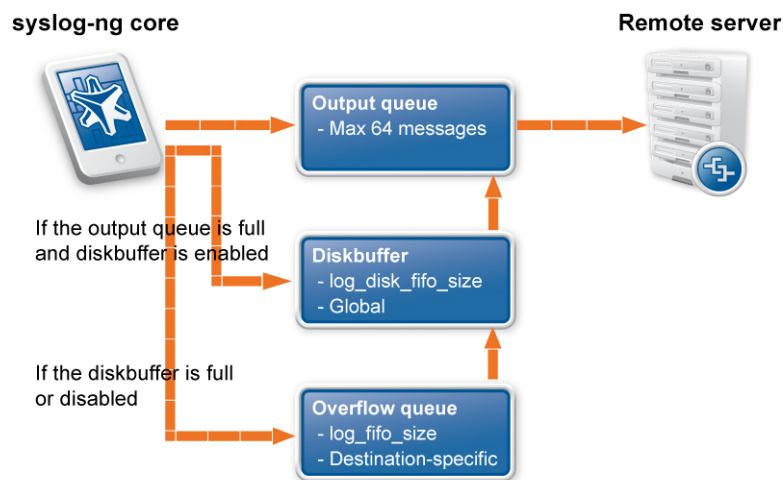


Figure 2.12. Handling outgoing messages in syslog-ng PE

- **Output queue:** Messages from the output queue are sent to the target syslog-ng server. The syslog-ng application puts the outgoing messages directly into the output queue, unless the output queue is full. The output queue can hold 64 messages, this is a fixed value and cannot be modified.
- **Disk buffer:** If the output queue is full and disk-buffering is enabled, syslog-ng puts the outgoing messages into the disk buffer of the destination. The disk buffer is enabled if the `log_disk_fifo_size()` parameter of the destination is larger than 0; the size of the disk buffer is specified in bytes.
- **Overflow queue:** If the output queue is full and the disk buffer is disabled or full, syslog-ng puts the outgoing messages into the overflow queue of the destination. (The overflow queue is identical to the output buffer used by other destinations.) The `log_fifo_size()` parameter specifies the number of messages stored in the overflow queue. See also Section 2.13, “Managing incoming and outgoing messages with flow-control” (p. 18) for details on sizing the `log_fifo_size()` parameter.

2.15. Licensing

The syslog-ng Premium Edition application is licensed on a per-host basis: the syslog-ng server accepts connections only from the number of individual hosts (also called log source hosts) specified in its license file.

A log source host is a host or network device (including syslog-ng clients and relays) that sends logs to the syslog-ng server. Log source hosts can be servers, routers, desktop computers, or other devices capable of sending syslog messages or running syslog-ng. Log source hosts are identified by their IP addresses, so virtual machines and vhosts are separately counted. Licenses are available for 5, 10, 25, 50, 100, 150, 200, 250, 300, 500, 750, 1000, and unlimited number of log source hosts.



Warning

The `chain_hostnames()` option of syslog-ng can interfere with the way syslog-ng counts the log source hosts, causing syslog-ng to think there are more hosts logging to the central server. As `chain_hostnames()` is a deprecated option, disable it on your log sources to avoid any problems related to license counting.



Buying a syslog-ng server license permits you to perform the following:

- Install the syslog-ng application in server mode to a single host. This host acts as the central log server of the network.
- Install the syslog-ng application in relay or client mode on host computers. The total number of hosts permitted to run syslog-ng in relay or client mode is limited by the syslog-ng server license. The client and relay hosts may use any operating system supported by syslog-ng. See Section 1.6, “Supported platforms” (p. 4) for details.
- Download software updates for a year.



Example 2.1. Counting log source hosts

Let's say that you have two facilities (e.g., data centers or server farms), and you have 80 AIX servers and 20 Microsoft Windows host at Facility 1, and 5 HP-UX servers and 40 Debian servers at Facility 2. That is 145 hosts altogether.

- If you want to collect the log messages of these host to a single logserver, then you need a syslog-ng PE license that allows you to accept logs from at least 145 hosts. (In practice this means you have to buy a license for 150 hosts.)
- If you want each facility to have its own logserver, and do not want to have a central server that collects the log messages of both facilities, you need two separate licenses: a license for 100 hosts at Facility 1, and a license for at least 45 hosts at Facility 2 (actually you have to buy license for 50 hosts).
- If you want each facility to have its own local logserver that stores the logs locally, and also want to have a central logserver that collects every log message independently from the two local logserver, you need three licenses: a license for 100 hosts at Facility 1, and a license for at least 45 hosts at Facility 2, and a license for the central logserver. The size of the license on the central logserver should be $100 \text{ (the hosts at Facility 1)} + 45 \text{ (the hosts at Facility 2)} + 2 \text{ (the two local logservers at each facility)} = 147$ — practically thats another 150-host license.



Note

If, for example, the 40 Debian servers at Facility 2 are each running 3 virtual hosts, then the total number of hosts at Facility 2 is 125, and the license sizes should be calculated accordingly.

The syslog-ng Open Source Edition application is distributed under version 2 of the GNU General Public License. See Appendix 3, *GNU General Public License* (p. 250) for details.

2.16. High availability support

Multiple syslog-ng servers can be run in fail-over mode. The syslog-ng application does not include any internal support for this, as clustering support must be implemented on the operating system level. A tool that can be used to create UNIX clusters is Heartbeat (see <http://www.linux-ha.org/> for details).

2.17. Possible causes of losing log messages

During the course of a message from the sending application to the final destination of the message, there are a number of locations where a message may be lost, even though syslog-ng does its best to avoid message loss. Usually losing messages can be avoided with careful planning and proper configuration of syslog-ng and the hosts



running syslog-ng. The following list shows the possible locations where messages may be lost, and provides methods to minimize the risk of losing messages.

**Note**

The following list covers the main possibilities of losing messages, but does not take into account the possible use of flow-control (see Section 2.13, “Managing incoming and outgoing messages with flow-control” (p. 18)). This topic will be addressed in more detail in the future releases of this guide.

- *Between the application and the syslog-ng client:* Make sure to use an appropriate source to receive the logs from the application (e.g., from `/dev/log`). For example, use `unix-stream` instead of `unix-dgram` whenever possible.
- *When syslog-ng is sending messages:* If syslog-ng cannot send messages to the destination and the output buffer gets full, syslog-ng will drop messages. The number of dropped messages is displayed per destination in the log message statistics of syslog-ng (see Section 3.3.1.1, “Log statistics” (p. 36) for details). To prevent such message loss, use the disk buffer of syslog-ng Premium Edition to increase the capacity of your output buffer beyond that would be feasible using only a memory-based buffer.
- *On the network:* When transferring messages using the UDP protocol, messages may be lost without any notice or feedback — such is the nature of the UDP protocol. Always use the TCP protocol to transfer messages over the network whenever possible.
- *In the socket receive buffer:* When transferring messages using the UDP protocol, the UDP datagram (i.e., the message) that reaches the receiving host placed in a memory area called the *socket receive buffer*. If the host receives more messages than it can process, this area overflows, and the kernel drops messages without letting syslog-ng know about it. Using TCP instead of UDP prevents this issue. If you must use the UDP protocol, increase the size of the receive buffer using the `so_rcvbuf()` option.
- *When syslog-ng is receiving messages:* The receiving syslog-ng (e.g., the syslog-ng server or relay) may drop messages if the fifo of the destination file gets full. The number of dropped messages is displayed per destination in the log message statistics of syslog-ng (see Section 3.3.1.1, “Log statistics” (p. 36) for details). To prevent such message loss, adjust the fifo appropriately for the message load and use the disk buffer of syslog-ng Premium Edition. See Section 7.3, “Handling large message load” (p. 137) and Section 2.14, “Using disk-based buffering” (p. 20) for details.
- *When the destination cannot handle large load:* When syslog-ng is sending messages at a high rate into an SQL database, a file, or another destination, it is possible that the destination cannot handle the load, and processes the messages slowly. As a result, the buffers of syslog-ng fill up, syslog-ng cannot process the incoming messages, and starts to lose messages. See the previous entry for details. Use the `throttle` parameter and the disk buffer of syslog-ng Premium Edition (Section 2.14, “Using disk-based buffering” (p. 20)).
- *As a result of an unclean shutdown of the syslog-ng server:* If the host running the syslog-ng server experiences an unclean shutdown, it takes time until the clients realize that the connection to the syslog-ng server is down. Messages that are put into the output TCP buffer of the clients during this period are not sent to the server. Since on Windows the buffer of the TCP stack is 3 MB by default, such a situation can result in significant message loss.



2.18. The structure of a log message

The following sections describe the structure of log messages. Currently there are two standard syslog message formats:

- The old standard described in RFC 3164 (also called the BSD-syslog or the legacy-syslog protocol): see Section 2.18.1, “BSD-syslog or legacy-syslog messages” (p. 24)
- The new standard described in RFC 5424 (also called the IETF-syslog protocol): see Section 2.18.2, “IETF-syslog messages” (p. 26)

2.18.1. BSD-syslog or legacy-syslog messages

This section describes the format of a syslog message, according to the legacy-syslog or BSD-syslog protocol (see RFC 3164 <http://www.ietf.org/rfc/rfc3164.txt>). A syslog message consists of the following parts:

- PRI
- HEADER
- MSG

The total message cannot be longer than 1024 bytes.

The following is a sample syslog message: `<133>Feb 25 14:09:07 webserver syslogd: restart`. The message corresponds to the following format: `<priority>timestamp hostname application: message`. The different parts of the message are explained in the following sections.



Note

The syslog-ng application supports longer messages as well. For details, see the `log_msg_size()` option in Section 8.9, “Global options” (p. 227). However, it is not recommended to enable messages larger than the packet size when using UDP destinations.

2.18.1.1. The PRI message part

The PRI part of the syslog message (known as Priority value) represents the Facility and Severity of the message. Facility represents the part of the system sending the message, while severity marks its importance. The Priority value is calculated by first multiplying the Facility number by 8 and then adding the numerical value of the Severity. The possible facility and severity values are presented below.



Note

Facility codes may slightly vary between different platforms. The syslog-ng application accepts facility codes as numerical values as well.



Numerical Code	Facility
0	kernel messages
1	user-level messages
2	mail system
3	system daemons
4	security/authorization messages
5	messages generated internally by syslogd
6	line printer subsystem
7	network news subsystem
8	UUCP subsystem
9	clock daemon
10	security/authorization messages
11	FTP daemon
12	NTP subsystem
13	log audit
14	log alert
15	clock daemon
16-23	locally used facilities (local0-local7)

Table 2.1. *syslog Message Facilities*

The following table lists the severity values.

Numerical Code	Severity
0	Emergency: system is unusable
1	Alert: action must be taken immediately
2	Critical: critical conditions
3	Error: error conditions
4	Warning: warning conditions
5	Notice: normal but significant condition
6	Informational: informational messages
7	Debug: debug-level messages

Table 2.2. *syslog Message Severities*

2.18.1.2. The HEADER message part

The HEADER part contains a timestamp and the hostname (without the domain name) or the IP address of the device. The timestamp field is the local time in the *Mmm dd hh:mm:ss* format, where:



- *Mmm* is the English abbreviation of the month: Jan, Feb, Mar, Apr, May, Jun, Jul, Aug, Sep, Oct, Nov, Dec.
- *dd* is the day of the month on two digits. If the day of the month is less than 10, the first digit is replaced with a space. (E.g., *Aug 7*.)
- *hh:mm:ss* is the local time. The hour (hh) is represented in a 24-hour format. Valid entries are between 00 and 23, inclusive. The minute (mm) and second (ss) entries are between 00 and 59 inclusive.

**Note**

The syslog-ng application supports other timestamp formats as well, like ISO, or the PIX extended format. For details, see the `ts_format()` option in Section 8.9, “Global options” (p. 227).

2.18.1.3. The MSG message part

The MSG part contains the name of the program or process that generated the message, and the text of the message itself. The MSG part is usually in the following format: *program[pid]: message text*.

2.18.2. IETF-syslog messages

This section describes the format of a syslog message, according to the IETF-syslog protocol (see RFC 5424-5428 <http://tools.ietf.org/html/rfc5424>). A syslog message consists of the following parts:

- HEADER (includes the PRI as well)
- STRUCTURED-DATA
- MSG

The following is a sample syslog message:¹

```
<34>1 2003-10-11T22:14:15.003Z mymachine.example.com su - ID47 - BOM'su
root' failed for lonvick on /dev/pts/8
```

The message corresponds to the following format:

```
<priority>VERSION ISOTIMESTAMP HOSTNAME APPLICATION PID MESSAGEID
STRUCTURED-DATA MSG
```

In this example, the Facility has the value of 4, severity is 2, so PRI is 34. The VERSION is 1. The message was created on 11 October 2003 at 10:14:15pm UTC, 3 milliseconds into the next second. The message originated from a host that identifies itself as "mymachine.example.com". The APP-NAME is "su" and the PROCID is unknown. The MSGID is "ID47". The MSG is "'su root' failed for lonvick...", encoded in UTF-8. The encoding is defined by the BOM. There is no STRUCTURED-DATA present in the message, this is indicated by "-" in the STRUCTURED-DATA field. The MSG is "'su root' failed for lonvick..."

¹Source: <http://tools.ietf.org/html/rfc5424>



The HEADER part of the message must be in plain ASCII format, the parameter values of the STRUCTURED-DATA part must be in UTF-8, while the MSG part should be in UTF-8. The different parts of the message are explained in the following sections.

2.18.2.1. The PRI message part

The PRI part of the syslog message (known as Priority value) represents the Facility and Severity of the message. Facility represents the part of the system sending the message, while severity marks its importance. The Priority value is calculated by first multiplying the Facility number by 8 and then adding the numerical value of the Severity. The possible facility and severity values are presented below.



Note

Facility codes may slightly vary between different platforms. The syslog-ng application accepts facility codes as numerical values as well.

Numerical Code	Facility
0	kernel messages
1	user-level messages
2	mail system
3	system daemons
4	security/authorization messages
5	messages generated internally by syslogd
6	line printer subsystem
7	network news subsystem
8	UUCP subsystem
9	clock daemon
10	security/authorization messages
11	FTP daemon
12	NTP subsystem
13	log audit
14	log alert
15	clock daemon
16-23	locally used facilities (local0-local7)

Table 2.3. syslog Message Facilities

The following table lists the severity values.

Numerical Code	Severity
0	Emergency: system is unusable
1	Alert: action must be taken immediately
2	Critical: critical conditions
3	Error: error conditions
4	Warning: warning conditions
5	Notice: normal but significant condition
6	Informational: informational messages
7	Debug: debug-level messages

Table 2.4. *syslog Message Severities*

2.18.2.2. The HEADER message part

The HEADER part contains the following elements:

- **VERSION**: Version number of the syslog protocol standard. Currently this can only be *1*.
- **ISOTIMESTAMP**: The time when the message was generated in the ISO 8601 compatible standard timestamp format (yyyy-mm-ddThh:mm:ss+-ZONE), e.g.: *2006-06-13T15:58:00.123+01:00*.
- **HOSTNAME**: The machine that originally sent the message.
- **APPLICATION**: The device or application that generated the message
- **PID**: The process name or process ID of the syslog application that sent the message. It is not necessarily the process ID of the application that generated the message.
- **MESSAGEID**: The ID number of the message.



Note

The syslog-ng application supports other timestamp formats as well, like ISO, or the PIX extended format. The timestamp used in the IETF-syslog protocol is derived from RFC3339, which is based on ISO8601. For details, see the `ts_format()` option in Section 8.9, “Global options” (p. 227).

2.18.2.3. The STRUCTURED-DATA message part

The STRUCTURED-DATA message part may contain meta- information about the syslog message, or application-specific information such as traffic counters or IP addresses. STRUCTURED-DATA consists of data blocks enclosed in brackets (`[]`). Every block include the ID of the block, and one or more *name=value* pairs. The syslog-ng application automatically parses the STRUCTURED-DATA part of syslog messages, which can be referenced in macros (see Section 8.5, “Macros” (p. 210) for details). An example STRUCTURED-DATA block looks like:

```
[exampleSDID@0 iut="3" eventSource="Application"
eventID="1011"] [examplePriority@0 class="high"]
```



2.18.2.4. The MSG message part

The MSG part contains the text of the message itself. The encoding of the text must be UTF-8 if the BOM character is present in the message. If the message does not contain the BOM character, the encoding is treated as unknown. Usually messages arriving from legacy sources do not include the BOM character.



Chapter 3. Configuring syslog-ng

This chapter describes how to configure syslog-ng.

3.1. The syslog-ng configuration file

The syslog-ng application is configured by editing the `syslog-ng.conf` file. Use any regular text editor application to modify the file. The precompiled syslog-ng packages include sample configuration files as well.

Every syslog-ng configuration file must begin with a line containing the version information of syslog-ng. For syslog-ng version 3.0, this line looks like:

```
@version:3.0
```

If the configuration file does not contain the version information, syslog-ng assumes that the file is for syslog-ng version 2.x. In this case it interprets the configuration and sends warnings about the parts of the configuration that should be updated. Version 3.0 of syslog-ng will correctly operate with configuration files of version 2.x, but the default values of certain parameters are different in 3.0.

All identifiers, option names and attributes, and any other strings used in the syslog-ng configuration file are case sensitive. Objects must be defined before they are referenced in another statement.



Example 3.1. A simple configuration file

The following is a very simple configuration file for syslog-ng: it collects the internal messages of syslog-ng and the messages from `/dev/log` into the `/var/log/messages_syslog-ng.log` file.

```
@version:3.0

source s_local { unix-stream("/dev/log"); internal(); };

destination d_file_normal { file("/var/log/messages_syslog-ng.log"); };

log { source(s_local); destination(d_file); };
```



Tip

Before activating a new configuration, check that your configuration file is syntactically correct using the **syslog-ng --syntax** command.

To activate the configuration, reload the configuration of syslog-ng using the **/etc/init.d/syslog-ng reload** command.

The `syslog-ng.conf` and `license.txt` files are located under the `/opt/syslog-ng/etc/` directory.



Note

Earlier versions of syslog-ng PE stored the configuration and license files under different directories, depending on the platform; typically under `/etc/syslog-ng/`.



On Microsoft Windows platforms the syslog-ng agent stores its configuration in the system registry, and can be configured from a graphical interface. See Chapter 5, *Collecting logs from Windows hosts* (p. 93) for details.

3.1.1. Including configuration files

The syslog-ng application supports including external files in its configuration file, so parts of its configuration can be managed separately. To include the contents of a file in the syslog-ng configuration, use the following syntax

```
include "filename";
```

This imports the entire file into the configuration of syslog-ng, at the location of the include statement. If you specify a directory, syslog-ng will try to include every file in alphabetic order. When including configuration files, consider the following points:

- If an object is defined twice (e.g., the original syslog-ng configuration file and the file imported into this configuration file both define the same option, source, or other object), then the object that is defined later in the configuration file will be effective. For example, if you set a global option at the beginning of the configuration file, and later include a file that defines the same option with a different value, then the option defined in the imported file will be used.
- Files can be embedded into each other: the included files can contain include statements as well, up to a maximum depth of 15 levels.
- Include statements can only be used at top level of the configuration file. For example, the following is correct:

```
@version:3.0
include "example.conf";
```

But the following is not:

```
source s_example {
    include "example.conf"
};
```



Warning

The syslog-ng application will not start if it cannot find a file that is to be included in its configuration. Always double-check the filenames, paths, and access rights when including configuration files, and use the **--syntax-only** command-line option to check your configuration.

3.1.2. Logging configuration changes

Every time syslog-ng is started, or its configuration is reloaded, it automatically logs the SHA-1 fingerprint of its configuration file using the *internal* message source. That way any modification of the configuration of your syslog-ng clients is visible in the central logs. Note that the log message does not contain the exact change, nor can the configuration file be retrieved from the fingerprint. Only the fact of the configuration change can be detected.



The fingerprint can be examined with the **logchksign** command-line application, which detects that the fingerprint was indeed generated by a syslog-ng application. Just paste the hashes from the log message after the **logchksign** command like in the following example: **bin/logchksign "cfg-fingerprint='832ef664ff79df8afc66cd955c0c8aaa3c343f31', cfg-signature='785223cfa19ad52b855550be141b00306347b0a9' " cfg-nonce-ndx='0',**

3.2. Defining global objects

Global objects (e.g., sources, destinations, log paths, or filters) are defined in the syslog-ng configuration file. Object definitions consist of the following elements:

- *Type of the object:* One of *source*, *destination*, *log*, *filter*, *parser*, *rewrite rule*, or *template*.
- *Identifier of the object:* A unique name identifying the object. When using a reserved word as an identifier, enclose the identifier in quotation marks.



Tip

Use identifiers that refer to the type of the object they identify. For example, prefix source objects with *s_*, destinations with *d_*, and so on.

- *Parameters:* The parameters of the object, enclosed in braces *{parameters}*.
- *Semicolon:* Object definitions end with a semicolon (;).

The syntax is summarized as follows:

```
type identifier { parameters };
```

Objects have parameters; some of them are required, others are optional. Required parameters are positional, meaning that they must be specified in a defined order. Optional arguments can be specified in any order using the *option(value)* format. If a parameter (optional or required) is not specified, its default value is used. The parameters and their default values are listed in the reference section of the particular object. See Chapter 8, *Reference* (p. 141) for details.



Example 3.2. Using required and optional parameters

The *unix-stream()* source driver has a single required argument: the name of the socket to listen on. Optional parameters follow the socket name in any order, so the following source definitions have the same effect:

```
source s_demo_stream1 {
    unix-stream("/dev/log" max-connections(10) group(log)); };
source s_demo_stream2 {
    unix-stream("/dev/log" group(log) max-connections(10)); };
```

To add comments to the configuration file, start a line with *#* and write your comments. These lines are ignored by syslog-ng.



```
# Comment: This is a stream source
source s_demo_stream {
    unix-stream("/dev/log" max-connections(10) group(log)); };
```

3.2.1. Notes about the configuration syntax

When you are editing the syslog-ng configuration file, note the following points:

- When writing the names of options and parameters (or other reserved words), the hyphen (-) and underscore () characters are equivalent, e.g., *max-connections(10)* and *max_connections(10)* are both correct.
- Number can be prefixed with + or - to indicate positive or negative values. Numbers beginning with zero (0) or 0x are treated as hexadecimal or octal numbers, respectively.
- You can use commas (,) to separate options or other parameters for readability; syslog-ng completely ignores them. The following declarations are equivalent:

```
source s_demo_stream {
    unix-stream("/dev/log" max-connections(10)
group(log)); };
```

```
source s_demo_stream {
    unix-stream("/dev/log", max-connections(10),
group(log)); };
```

- Strings between single quotes ('*string*') are treated literally, you do not have to escape special characters. This makes writing and reading regular expressions much more simple: it is recommended to use single quotes when writing regular expressions.
- When enclosing strings between double-quotes ("*string*"), you have to escape special characters: e.g., when enclosing a regular expression that uses the \ character to escape a special character, you have to add an extra \ (e.g., "\ \n"). It is recommended to use single quotes instead.
- Enclosing normal strings between double-quotes ("*string*") is not necessary, you can just omit the double-quotes. E.g., when writing filters, *match("sometext")* and *match(sometext)* will both match for the *sometext* string.
- When enclosing object IDs (e.g., the name of a destination) between double-quotes ("*mydestination*"), the ID can include whitespace as well, e.g.:

```
source "s demo stream" {
    unix-stream("/dev/log" max-connections(10)
group(log)); };
```

3.3. Sources and source drivers

A source is where syslog-ng receives log messages. Sources consist of one or more drivers, each defining where and how messages are received.



To define a source, add a source statement to the syslog-ng configuration file using the following syntax:

```
source <identifier> { source-driver(params); source-driver(params); ... };
```



Example 3.3. A simple source statement

The following source statement receives messages on the TCP port 1999 of the interface having the 10.1.2.3 IP address.

```
source s_demo_tcp { tcp(ip(10.1.2.3) port(1999)); };
```



Example 3.4. A source statement using two source drivers

The following source statement receives messages on the 1999 TCP port and the 1999 UDP port of the interface having the 10.1.2.3 IP address.

```
source s_demo_two_drivers {
    tcp(ip(10.1.2.3) port(1999));
    udp(ip(10.1.2.3) port(1999)); };
```



Example 3.5. Setting default priority and facility

If the message received by the source does not have a proper syslog header, you can use the *default-facility()* and *default-priority()* options to set the facility and priority of the messages. Note that these values are applied only to messages that do not set these parameters in their header.

```
source headerless_messages { udp(default-facility(syslog) default-priority(emerg));
};
```

Define a source only once. The same source can be used in several log paths. Duplicating sources causes syslog-ng to open the source (TCP/IP port, file, etc.) more than once, which might cause problems. For example, include the */dev/log* file source only in one source statement, and use this statement in more than one log path if needed.

To collect log messages on a specific platform, it is important to know how the native *syslogd* communicates on that platform. The following table summarizes the operation methods of *syslogd* on some of the tested platforms:



Platform	Method
Linux	A <i>SOCK_STREAM</i> unix socket named <i>/dev/log</i> ; some of the distributions switched over to using <i>SOCK_DGRAM</i> , though applications still work with either method.
BSD flavors	A <i>SOCK_DGRAM</i> unix socket named <i>/var/run/log</i> .
Solaris (2.5 or below)	An SVR4 style <i>STREAMS</i> device named <i>/dev/log</i> .
Solaris (2.6 or above)	In addition to the <i>STREAMS</i> device used in earlier versions, 2.6 uses a new multithreaded IPC method called door. By default the door used by <i>syslogd</i> is <i>/etc/.syslog_door</i> .
HP-UX 11 or later	HP-UX uses a named pipe called <i>/dev/log</i> that is padded to 2048 bytes, e.g., <i>source s_hp-ux {pipe ("/dev/log" pad_size(2048))}</i> .
AIX 5.2 and 5.3	A <i>SOCK_STREAM</i> or <i>SOCK_DGRAM</i> unix socket called <i>/dev/log</i> .

Table 3.1. Communication methods used between the applications and *syslogd*

Each possible communication mechanism has a corresponding source driver in *syslog-ng*. For example, to open a unix socket with *SOCK_DGRAM* style communication use the driver *unix-dgram*. The same socket using the *SOCK_STREAM* style — as used under Linux — is called *unix-stream*.



Example 3.6. Source statement on a Linux based operating system

The following source statement collects the following log messages:

- *internal()*: Messages generated by *syslog-ng*.
- *udp(ip(0.0.0.0) port(514))*: Messages arriving to the 514/UDP port of any interface of the host.
- *unix-stream("/dev/log")*: Messages arriving to the */dev/log* socket.

```
source s_demo {
    internal();
    udp(ip(0.0.0.0) port(514));
    unix-stream("/dev/log"); };
```

The following table lists the source drivers available in *syslog-ng*.



Name	Description
<code>internal()</code>	Messages generated internally in syslog-ng.
<code>file()</code>	Opens the specified file and reads messages.
<code>pipe()</code> , <code>fifo</code>	Opens the specified named pipe and reads messages.
<code>program()</code>	Opens the specified application and reads messages from its standard output.
<code>sun-stream()</code> , <code>sun-streams()</code>	Opens the specified <i>STREAMS</i> device on Solaris systems and reads incoming messages.
<code>syslog()</code>	Listens for incoming messages using the new <u>IETF-standard syslog protocol</u> .
<code>tcp()</code> , <code>tcp6()</code>	Listens on the specified TCP port for incoming messages using the <u>BSD-syslog protocol</u> over IPv4 and IPv6 networks, respectively.
<code>udp()</code> , <code>udp6()</code>	Listens on the specified UDP port for incoming messages using the <u>BSD-syslog protocol</u> over IPv4 and IPv6 networks, respectively.
<code>unix-dgram()</code>	Opens the specified unix socket in <i>SOCK_DGRAM</i> mode and listens for incoming messages.
<code>unix-stream()</code>	Opens the specified unix socket in <i>SOCK_STREAM</i> mode and listens for incoming messages.

Table 3.2. Source drivers available in syslog-ng

For a complete description of the parameters of the above drivers, see Section 8.1, “Source drivers” (p. 141).

3.3.1. Collecting internal messages

All messages generated internally by syslog-ng use this special source. To collect warnings, errors and notices from syslog-ng itself, include this source in one of your source statements.

```
internal()
```

The syslog-ng application will issue a warning upon startup if none of the defined log paths reference this driver.



Example 3.7. Using the `internal()` driver

```
source s_local { internal(); };
```

3.3.1.1. Log statistics

Periodically, syslog-ng sends a message containing statistics about the received messages, and about any lost messages since the last such message. It includes a *processed* entry for every source and destination, listing the number of messages received or sent, and a *dropped* entry including the IP address of the server for every destination where syslog-ng has lost messages. The *center (received)* entry shows the total number of messages received from every configured sources.



The following is a sample log statistics message for a configuration that has a single source (*s_local*) and a network and a local file destination (*d_network* and *d_local*, respectively). All incoming messages are sent to both destinations.

```
Log statistics;
    dropped='tcp(AF_INET(192.168.10.1:514))=6439',
    processed='center(received)=234413',
    processed='destination(d_tcp)=234413',
    processed='destination(d_local)=234413',
    processed='source(s_local)=234413'
```

Log statistics can be also retrieved on-demand using the **echo STATS | nc -U var/run/syslog-ng.ctl** command. This returns a list of source groups and destinations, as well as the number of processed messages for each. The verbosity of the statistics can be set using the *stats_level()* option. See Section 8.9, “Global options” (p. 227) for details.



Note

To query the statistics, you need the OpenBSD-style **netcat** application. The **netcat** included in most Linux distributions is a GNU-style version that is not suitable to query the statistics of syslog-ng. An alternative is to use the **socat** application:
echo STATS | socat -vv UNIX-CONNECT:/opt/syslog-ng/var/run/syslog-ng.ctl -.

3.3.2. Collecting messages from text files

Collects log messages from plain-text files, e.g., from the logfiles of an Apache webserver.

The syslog-ng application notices if a file is renamed or replaced with a new file, so it can correctly follow the file even if logrotation is used. When syslog-ng is restarted, it records the position of the last sent log message, and continues to send messages from this position after the restart.

The file driver has a single required parameter specifying the file to open. For the list of available optional parameters, see Section 8.1.2, “file()” (p. 141).

Declaration:

```
file(filename);
```

In syslog-ng PE, the filename (but not the pathname) may include wildcard characters (e.g., *). Note that when using wildcards in filenames, always set how often syslog-ng should check the file for new messages using the *follow_freq()* parameter.

When using wildcards, syslog-ng PE monitors every matching file, and can receive new log messages from any of the files. However, monitoring (polling) many files (i.e., more than ten) has a significant overhead and may affect performance. On Linux this overhead is not so significant, because syslog-ng PE uses the inotify feature of the kernel.

**Example 3.8. Using the file() driver**

```
source s_file { file("/var/log/messages");
```

**Example 3.9. Using wildcards in the filename**

The following example monitors every file with the `.log` extension in the `/var/application` directory for log messages. Note that only syslog-ng PE supports wildcards in the file and pathnames.

```
source s_file { file("/var/application/*.log" follow_freq(1));
```

**Example 3.10. Monitoring multiple directories**

The following example reads files having the `.log` extension from the `/var/application/` directory and its subdirectories. Note that only syslog-ng PE supports recursive directory handling and wildcards in the file and pathnames.

```
source s_file_subdirectories { file("/var/application/*.log"
    recursive(yes)
    follow_freq(1)
    log_fetch_limit(100)
);};
```

The kernel usually sends log messages to a special file (`/dev/kmsg` on BSDs, `/proc/kmsg` on Linux). The `file()` driver reads log messages from such files. The syslog-ng application can periodically check the file for new log messages if the `follow_freq()` option is set.

**Note**

On Linux, the `klogd` daemon can be used in addition to syslog-ng to read kernel messages and forward them to syslog-ng. `klogd` used to preprocess kernel messages to resolve symbols etc., but as this is deprecated by `ksymoops` there is really no point in running both `klogd` and syslog-ng in parallel. Also note that running two processes reading `/proc/kmsg` at the same time might result in dead-locks.

When using syslog-ng to read messages from the `/proc/kmsg` file, syslog-ng automatically disables the `follow_freq()` parameter to avoid blocking the file.

3.3.3. Collecting messages from named pipes

The pipe driver opens a named pipe with the specified name and listens for messages. It is used as the native message delivery protocol on HP-UX.

The pipe driver has a single required parameter, specifying the filename of the pipe to open. For the list of available optional parameters, see Section 8.1.3, “pipe()” (p. 146).

Declaration:

```
pipe(filename);
```

**Note**

As of syslog-ng Open Source Edition 3.0.2, pipes are created automatically. In earlier versions, you had to create the pipe using the `mkfifo(1)` command.



Pipe is very similar to the *file()* driver, but there are a few differences, for example *pipe()* opens its argument in read-write mode, therefore it is not recommended to be used on special files like */proc/kmsg*.



Warning

It is not recommended to use *pipe()* on anything else than real pipes.



Example 3.11. Using the pipe() driver

```
source s_pipe { pipe("/dev/pipe" pad_size(2048)); };
```

3.3.4. Collecting messages on Sun Solaris

Solaris uses its *STREAMS* framework to send messages to the *syslogd* process. Solaris 2.5.1 and above uses an IPC called *door* in addition to *STREAMS*, to confirm the delivery of a message. The syslog-ng application supports the IPC mechanism via the *door()* option (see below).



Note

The *sun-streams()* driver must be enabled when the syslog-ng application is compiled (see *./configure --help*). The pre-compiled binaries of syslog-ng Open Source Edition purchased from BalaBit and syslog-ng Premium Edition have this option enabled.

The *sun-streams()* driver has a single required argument specifying the *STREAMS* device to open, and the *door()* option. For the list of available optional parameters, see Section 8.1.5, “*sun-streams()* driver” (p. 154).

Declaration:

```
sun-streams(name_of_the_streams_device door(filename_of_the_door));
```



Example 3.12. Using the sun-streams() driver

```
source s_stream { sun-streams("/dev/log" door("/etc/.syslog_door")); };
```

3.3.5. Collecting messages using the IETF syslog protocol

The *syslog()* driver enables to receive messages from the network using the new standard syslog protocol and message format (also called IETF-syslog protocol; described in RFC 5424-28, see Section 2.18.2, “IETF-syslog messages” (p. 26)). UDP, TCP, and TLS-encrypted TCP can all be used to transport the messages.

For the list of available optional parameters, see Section 8.1.6, “*syslog()*” (p. 160).



Declaration:

```
syslog(ip() port() transport() options());
```



Example 3.13. Using the `syslog()` driver

TCP source listening on the localhost on port 1999.

```
source s_syslog { syslog(ip(127.0.0.1) port(1999) transport("tcp")); };
```

UDP source with defaults.

```
source s_udp { syslog( transport("udp")); };
```

Encrypted source where the client is also authenticated. See Section 8.10, “TLS options” (p. 231) for details on the encryption settings.

```
source s_syslog_tls{ syslog(
    ip(10.100.20.40)
    transport("tls")
    tls(
        peer-verify(required-trusted)
        ca_dir('/opt/syslog-ng/etc/syslog-ng/keys/ca.d/')
        key_file('/opt/syslog-ng/etc/syslog-ng/keys/server_privatekey.pem')
        cert_file('/opt/syslog-ng/etc/syslog-ng/keys/server_certificate.pem')
    )
);};
```

3.3.6. Collecting messages from remote hosts using the BSD syslog protocol

The `tcp()`, `tcp6()`, `udp()`, `udp6()` drivers can receive messages from the network using the TCP and UDP networking protocols. The `tcp6()` and `udp6()` drivers use the IPv6 network protocol, while `tcp()` and `udp()` use IPv4.

UDP is a simple datagram oriented protocol, which provides "best effort service" to transfer messages between hosts. It may lose messages, and no attempt is made at the protocol level to retransmit such lost messages. The *BSD-syslog* protocol traditionally uses UDP.

TCP provides connection-oriented service, which basically means that the path of the messages is flow-controlled. Along this path, each message is acknowledged, and retransmission is done for lost packets. Generally it is safer to use TCP, because lost connections can be detected, and no messages get lost, assuming that the TCP connection does not break. When a TCP connection is broken the 'in-transit' messages that were sent by syslog-ng but not yet received on the other side are lost. (Basically these messages are still sitting in the socket buffer of the sending host and syslog-ng has no information about the fate of these messages).

The `tcp()` and `udp()` drivers do not have any required parameters. By default they bind to the `0.0.0.0:514` address, which means that syslog-ng will listen on all available interfaces, port 514. To limit accepted connections to only one interface, use the `localip()` parameter as described below. For the list of available optional parameters, see Section 8.1.7, “tcp(), tcp6(), udp() and udp6()” (p. 165).

Declaration:

```
tcp([options]);
udp([options]);
```

**Note**

The tcp port 514 is reserved for use with **rshell**, so select a different port if syslog-ng and **rshell** is used at the same time.

If you specify a multicast bind address to `udp()` and `udp6()`, syslog-ng will automatically join the necessary multicast group. TCP does not support multicasting.

The syslog-ng Premium Edition application supports TLS (Transport Layer Security, also known as SSL) for the `tcp()` and `tcp6()` drivers. See the TLS-specific options below and Section 3.13, “Encrypting log messages with TLS” (p. 68) for details. For the list of available optional parameters, see Section 8.1.5, “sun-streams() driver” (p. 154).

**Example 3.14. Using the `udp()` and `tcp()` drivers**

A simple `udp()` source with default settings.

```
source s_udp { udp(); };# An UDP source with default settings.
```

A TCP source listening on the localhost interface, with a limited number of connections allowed.

```
source s_tcp { tcp(ip(127.0.0.1) port(1999) max-connections(10)); };
```

A TCP source listening on a TLS-encrypted channel.

```
source s_tcp { tcp(ip(127.0.0.1) port(1999)
    tls(peer-verify('required-trusted')
        key_file('/opt/syslog-ng/etc/syslog-ng/syslog-ng.key')
        cert_file('/opt/syslog-ng/etc/syslog-ng/syslog-ng.crt'));
};
```

A TCP source listening for messages using the IETF-syslog message format:

```
source s_tcp_syslog { tcp(ip(127.0.0.1) port(1999) flags(syslog-protocol)); };
```

3.3.7. Collecting messages from UNIX domain sockets

The `unix-stream()` and `unix-dgram()` drivers open an `AF_UNIX` socket and start listening on it for messages. The `unix-stream()` driver is primarily used on Linux and uses `SOCK_STREAM` semantics (connection oriented, no messages are lost); while `unix-dgram()` is used on BSDs and uses `SOCK_DGRAM` semantics: this may result in lost local messages if the system is overloaded.

To avoid denial of service attacks when using connection-oriented protocols, the number of simultaneously accepted connections should be limited. This can be achieved using the `max-connections()` parameter. The default value of this parameter is quite strict, you might have to increase it on a busy system.

Both `unix-stream` and `unix-dgram` have a single required argument that specifies the filename of the socket to create. For the list of available optional parameters, see Section 8.1.8, “`unix-stream()` and `unix-dgram()`” (p. 173)

Declaration:

```
unix-stream(filename [options]);
unix-dgram(filename [options]);
```

**Note**

syslogd on Linux originally used *SOCK_STREAM* sockets, but some distributions switched to *SOCK_DGRAM* around 1999 to fix a possible DoS problem. On Linux you can choose to use whichever driver you like as syslog clients automatically detect the socket type being used.

The difference between the *unix-stream* and *unix-dgram* drivers is similar to the difference between the TCP and UDP network protocols. Use the following guidelines to select which driver to use in a particular situation:

Choose *unix-stream* if you would choose TCP (stream) instead of UDP (datagram). The *unix-stream* driver offers the following features:

- Increased reliability
- Ordered delivery of messages
- Client-side notification of failures

Choose *unix-dgram* if you would choose TCP (stream) over UDP (datagram). The *unix-dgram* driver offers the following features:

- Decreased possibility of Dos by opening too many connections (a local vulnerability)
- Less overhead

However, the client does not notice if a message is lost when using the *unix-dgram* driver.

**Example 3.15. Using the *unix-stream()* and *unix-dgram()* drivers**

```
source s_stream { unix-stream("/dev/log" max-connections(10)); };
source s_dgram { unix-dgram("/var/run/log"); };
```

3.4. Destinations and destination drivers

A destination is where a log message is sent if the filtering rules match. Similarly to sources, destinations consist of one or more drivers, each defining where and how messages are sent.

**Tip**

If no drivers are defined for a destination, all messages sent to the destination are discarded. This is equivalent to omitting the destination from the log statement.

To define a destination, add a destination statement to the *syslog-ng* configuration file using the following syntax:

```
destination <identifier> {
    destination-driver(params); destination-driver(params);
... };
```

**Example 3.16. A simple destination statement**

The following destination statement sends messages to the TCP port 1999 of the 10.1.2.3 host.

```
destination d_demo_tcp { tcp("10.1.2.3" port(1999)); };
```

If name resolution is configured, the hostname of the target server can be used as well.

```
destination d_tcp { tcp("target_host" port(1999); localport(999)); };
```

The following table lists the destination drivers available in syslog-ng.

Name	Description
<code>file()</code>	Writes messages to the specified file.
<code>logstore()</code> *	Writes messages to the specified binary logstore file. *Available only in syslog-ng Premium Edition.
<code>fifo()</code> , <code>pipe()</code>	Writes messages to the specified named pipe.
<code>program()</code>	Forks and launches the specified program, and sends messages to its standard input.
<code>sql()</code>	Sends messages into an SQL database. In addition to the standard syslog-ng packages, the <code>sql()</code> destination requires database-specific packages to be installed. Refer to the section appropriate for your platform in Chapter 4, <i>Installing syslog-ng</i> (p. 78).
<code>syslog()</code>	Sends messages to the specified remote host using the IETF-syslog protocol . The IETF standard supports message transport using the UDP, TCP, and TLS networking protocols.
<code>tcp()</code> and <code>tcp6()</code>	Sends messages to the specified TCP port of a remote host using the BSD-syslog protocol over IPv4 and IPv6, respectively.
<code>udp()</code> and <code>udp6()</code>	Sends messages to the specified UDP port of a remote host using the BSD-syslog protocol over IPv4 and IPv6, respectively.
<code>unix-dgram()</code>	Sends messages to the specified unix socket in <code>SOCK_DGRAM</code> style (BSD).
<code>unix-stream()</code>	Sends messages to the specified unix socket in <code>SOCK_STREAM</code> style (Linux).
<code>usertty()</code>	Sends messages to the terminal of the specified user, if the user is logged in.

Table 3.3. Destination drivers available in syslog-ng

For detailed list of driver parameters, see Section 8.2, “Destination drivers” (p. 179).

3.4.1. Storing messages in plain-text files

The file driver is one of the most important destination drivers in syslog-ng. It allows to output messages to the specified text file, or to a set of files.

The destination filename may include macros which get expanded when the message is written, thus a simple `file()` driver may create several files. For more information on available macros see Section 8.5, “Macros” (p. 210).

If the expanded filename refers to a directory which does not exist, it will be created depending on the `create_dirs()` setting (both global and a per destination option).



The *file()* has a single required parameter that specifies the filename that stores the log messages. For the list of available optional parameters, see Section 8.2.1, “file()” (p. 179).

Declaration:

```
file(filename options());
```



Example 3.17. Using the file() driver

```
destination d_file { file("/var/log/messages" ); };
```



Example 3.18. Using the file() driver with macros in the file name and a template for the message

```
destination d_file {
    file("/var/log/$YEAR.$MONTH.$DAY/messages"
        template("$HOUR:$MIN:$SEC $TZ $HOST [$LEVEL] $MSG $MSG\n")
        template_escape(no));
};
```



Note

When using the *file()* destination, update the configuration of your log rotation program to rotate these files. Otherwise, the log files can become very large.



Warning

Since the state of each created file must be tracked by syslog-ng, it consumes some memory for each file. If no new messages are written to a file within 60 seconds (controlled by the *time_reap()* global option), it is closed, and its state is freed.

Exploiting this, a DoS attack can be mounted against the system. If the number of possible destination files and its needed memory is more than the amount available on the syslog-ng server.

The most suspicious macro is *\$PROGRAM*, where the number of possible variations is rather high. Do not use the *\$PROGRAM* macro in insecure environments.

3.4.2. Storing messages in encrypted files

The *logstore()* driver stores log messages in binary files that can be encrypted, compressed, checked for integrity, and timestamped by an external Timestamping Authority (TSA). Otherwise, it is very similar to the *file()* destination.

Logstore files consist of individual chunks, every chunk can be encrypted, compressed, and timestamped separately. Chunks contain log message data, chunk size defaults to 128k (about 1MB worth of compressed logs).

To display the contents of a logstore file, use the **logcat** command supplied with syslog-ng, e.g., **logcat /var/log/messages.lgs**. To display the contents of encrypted log files, specify the private key of the certificate used to encrypt the file, e.g., **logcat -k private.key /var/log/messages.lgs**. The contents of the file are sent to the standard output, so it is possible to use **grep** and other tools to find particular log messages, e.g., **logcat /var/log/messages.lgs | grep 192.168.1.1**.



Every record that is stored in the logstore has a unique record ID. The **logcat** application can quickly jump to a specified record using the **-- seek** option.

For files that are in use by syslog-ng, the last chunk that is open cannot be read. Chunks are closed when their size reaches the limit set in the *chunk_size* parameter, or when the time limit set in the *chunk_time* parameter expires and no new message arrives.

The syslog-ng PE application generates an SHA-1 hash for every chunk to verify the integrity of the chunk. The hashes of the chunks are chained together to prevent injecting chunks into the logstore file. The syslog-ng application can encrypt the logstore using the *aes128* algorithm in CBC mode; the hashing (HMAC) algorithm is *hmac-sha1*. Currently it is not possible to use other algorithms.



Warning

If the syslog-ng Premium Edition application or the computer crashes, an unclosed chunk remains at the end of the file. This chunk is marked as broken, its data stays there but is not shown by **logcat**.

The destination filename may include macros which get expanded when the message is written, thus a simple *logstore()* driver may create several files. For more information on available macros see Section 8.5, “Macros” (p. 210).

If the expanded filename refers to a directory which does not exist, it will be created depending on the *create_dirs()* setting (both global and a per destination option).

The *logstore()* has a single required parameter that specifies the filename that stores the log messages. For the list of available optional parameters, see Section 8.2.2, “logstore()” (p. 183).

Declaration:

```
logstore(filename options());
```



Example 3.19. Using the logstore() driver

A simple example saving and compressing log messages.

```
destination d_logstore { file("/var/log/messages.lgs" compress(5) ); };
```

A more detailed example that encrypts messages, modifies the parameters for closing chunks, and sets file privileges.

```
destination d_logstore { logstore("/var/log/messages-logstore.lgs"
encrypt_certificate("/opt/syslog-ng/etc/syslog-ng/keys/10-100-20-40/public-certificate-of-the-server.pem")

    chunk_size(100)
    chunk_time(5)
    owner("balabit")
    group("balabit")
    perm(0777)
); };
```



Note

When using the *logstore()* destination, update the configuration of your log rotation program to rotate these files. Otherwise, the log files can become very large.

**Warning**

Since the state of each created file must be tracked by syslog-ng, it consumes some memory for each file. If no new messages are written to a file within 60 seconds (controlled by the `time_reap()` global option), it is closed, and its state is freed.

Exploiting this, a DoS attack can be mounted against the system. If the number of possible destination files and its needed memory is more than the amount available on the syslog-ng server.

The most suspicious macro is `$PROGRAM`, where the number of possible variations is rather high. Do not use the `$PROGRAM` macro in insecure environments.

3.4.3. Sending messages to named pipes

The `pipe()` driver sends messages to a named pipe like `/dev/xconsole`.

The pipe driver has a single required parameter, specifying the filename of the pipe to open. The filename can include macros. For the list of available optional parameters, see Section 8.2.3, “pipe()” (p. 186).

Declaration:

```
pipe(filename);
```

**Warning**

As of syslog-ng Open Source Edition 3.0.2, pipes are created automatically. In earlier versions, you had to create the pipe using the `mkfifo(1)` command.

**Example 3.20. Using the pipe() driver**

```
destination d_pipe { pipe("/dev/xconsole"); };
```

3.4.4. Sending messages to external applications

The `program()` driver starts an external application or script and sends the log messages to its standard input (`stdin`).

The `program()` driver has a single required parameter, specifying a program name to start. The program is executed with the help of the current shell, so the command may include both file patterns and I/O redirections. For the list of available optional parameters, see Section 8.2.4, “program()” (p. 189).

Declaration:

```
program(command_to_run);
```

**Note**

The syslog-ng application automatically restarts the external program if it exits for reliability reasons. However it is not recommended to launch programs for single messages, because if the message rate is high, launching several instances of an application might overload the system, resulting in Denial of Service.



Note that the message format does not include the priority and facility values by default. To add these values, specify a template for the program destination, as shown in the following example.



Example 3.21. Using the `program()` destination driver

```
destination d_prog { program("/bin/script" template("<$PRI>$DATE $HOST $MSG\n");  
};
```

3.4.5. Storing messages in an SQL database

The `sql()` driver sends messages into an SQL database. Currently the Microsoft SQL (MSSQL), MySQL, Oracle, PostgreSQL, and SQLite databases are supported.



Note

In order to use the `sql()` destination, syslog-ng Premium Edition must run in server mode. Typically, only the central syslog-ng Premium Edition server uses the `sql()` destination.

The `sql()` driver has the following required parameters:



Name	Type	Default	Description
type	mssql, mysql, oracle, postgres, or sqlite3	n/a	Specifies the type of the database, i.e., the DBI database driver to use. Use the <i>mssql</i> option to send logs to an MSSQL database. See the examples of the databases on the following sections for details.
database	string	n/a	Name of the database that stores the logs.
table	string	n/a	Name of the database table to use (can include macros). When using macros, note that some databases limit the length of table names.
columns	string list	"date", "facility", "level", "host", "program", "pid", "message"	Name of the columns storing the data in <i>fieldname [dbtype]</i> format. The <i>[dbtype]</i> parameter is optional, and specifies the type of the field. By default, syslog-ng creates <i>text</i> columns. Note that not every database engine can index text fields.
values	string list	"\${R_YEAR}-\${R_MONTH}-\${R_DAY} \${R_HOUR}:\${R_MIN}:\${R_SEC}", "\$FACILITY", "\$LEVEL", "\$HOST", "\$PROGRAM", "\$PID", "\$MSGONLY"	The parts of the message to store in the fields specified in the <i>columns</i> parameter.

Table 3.4. Required parameters of the *sql()* driver

For the list of available optional parameters, see Section 8.2.5, “*sql()*” (p. 191).

Declaration:

```
sql(database_type host_parameters database_parameters [options]);
```



Warning

The syslog-ng application requires read and write access to the SQL table, otherwise it cannot verify that the destination table exists.

Currently the syslog-ng application has default schemas for the different databases and uses these defaults if the database schema (e.g., columns and column types) is not defined in the configuration file. However, these schemas will be deprecated and specifying the exact database schema will be required in later versions of syslog-ng.



Note

In addition to the standard syslog-ng packages, the *sql()* destination requires database-specific packages to be installed. These packages are automatically installed by the binary syslog-ng installer.

The *sql()* driver is currently not available for every platform that is supported by syslog-ng. For a list of platforms that support the *sql()* driver, visit <http://www.balabit.com/network-security/syslog-ng/central-syslog-server/>.

The *table* and *value* parameters can include macros to create tables and columns dynamically (see Section 8.5, “Macros” (p. 210) for details).

**Warning**

When using macros in table names, note that some databases limit the maximum allowed length of table names. Consult the documentation of the database for details.

Inserting the records into the database is performed by a separate thread. The syslog-ng application automatically performs the escaping required to insert the messages into the database.

**Example 3.22. Using the sql() driver**

The following example sends the log messages into a PostgreSQL database running on the *logserver* host. The messages are inserted into the *logs* database, the name of the table includes the exact date and the name of the host sending the messages. The syslog-ng application automatically creates the required tables and columns, if the user account used to connect to the database has the required privileges.

```
destination d_sql {
    sql(type(pgsql)
        host("logserver") username("syslog-ng") password("password")
        database("logs")
        table("messages_${HOST}_${R_YEAR}${R_MONTH}${R_DAY}")
        columns("datetime", "host", "program", "pid", "message")
        values("${R_DATE}", "${HOST}", "${PROGRAM}", "${PID}", "$MSGONLY")
        indexes("datetime", "host", "program", "pid", "message"));
};
```

The following example specifies the type of the database columns as well:

```
destination d_sql {
    sql(type(pgsql)
        host("logserver") username("syslog-ng") password("password")
        database("logs")
        table("messages_${HOST}_${R_YEAR}${R_MONTH}${R_DAY}")
        columns("datetime varchar(16)", "host varchar(32)", "program varchar(20)",
            "pid varchar(8)", "message varchar(200)")
        values("${R_DATE}", "${HOST}", "${PROGRAM}", "${PID}", "$MSGONLY")
        indexes("datetime", "host", "program", "pid", "message"));
};
```

3.4.5.1. Using the sql() driver with an Oracle database

The Oracle sql destination has some special aspects that are important to note.

- The hostname of the database server is set in the `tnsnames.ora` file, not in the `host` parameter of the `sql()` destination.
Make sure to set the Oracle-related environment variables properly, so syslog-ng and the Oracle client will find the file. The following variables must be set: `ORACLE_BASE`, `ORACLE_HOME`, and `ORACLE_SID`. See the documentation of the Oracle Instant Client for details.
- As certain database versions limit the maximum length of table names, macros in the table names should be used with care.
- In the current version of syslog-ng PE, the types of database columns must be explicitly set for the Oracle destination. The column used to store the text part of the syslog messages should be able to store messages as long as the longest message permitted by syslog-ng, therefore it is usually recommended to use the



varchar2 or *clob* column type. (The maximum length of the messages can be set using the *log_msg_size()* option.) See the following example for details.



Example 3.23. Using the *sql()* driver with an Oracle database

The following example sends the log messages into an Oracle database running on the *logserver* host, which must be set in the */etc/tnsnames.ora* file. The messages are inserted into the *LOGS* database, the name of the table includes the exact date when the messages were sent. The *syslog-ng* application automatically creates the required tables and columns, if the user account used to connect to the database has the required privileges.

```
destination d_sql {
    sql(type(oracle)
        username("syslog-ng") password("password")
        database("LOGS")
        table("msgs_${R_YEAR}${R_MONTH}${R_DAY}")
        columns("datetime varchar(16)", "host varchar(32)", "program varchar(32)", "pid
        varchar(8)", "message varchar2")
        values("${R_DATE}", "${HOST}", "${PROGRAM}", "${PID}", "$MSGONLY")
        indexes("datetime", "host", "program", "pid", "message"));
};
```

The Oracle Instant Client retrieves the address of the database server from the */etc/tnsnames.ora* file. Edit or create this file as needed for your configuration. A sample is provided below.

```
LOGS =
(DESCRIPTION =
  (ADDRESS_LIST =
    (ADDRESS = (PROTOCOL = TCP)
      (HOST = logserver)
      (PORT = 1521))
  )
  (CONNECT_DATA =
    (SERVICE_NAME = EXAMPLE.SERVICE)
  )
)
```

3.4.5.2. Using the *sql()* driver with a Microsoft SQL database

The *mssql* database driver can access Microsoft SQL (MSSQL) destinations. This driver has some special aspects that are important to note.

- The date format used by the MSSQL database must be explicitly set in the */etc/locales.conf* file of the *syslog-ng* server. See the following example for details.
- As certain database versions limit the maximum length of table names, macros in the table names should be used with care.
- In the current version of *syslog-ng* PE, the types of database columns must be explicitly set for the MSSQL destination. The column used to store the text part of the *syslog* messages should be able to store messages as long as the longest message permitted by *syslog-ng*. The *varchar* column type can store maximum 4096 bytes-long messages. The maximum length of the messages can be set using the *log_msg_size()* option. See the following example for details.
- Remote access for SQL users must be explicitly enabled on the Microsoft Windows host running the Microsoft SQL Server. See Section 4.6, “Configuring Microsoft SQL Server to accept logs from *syslog-ng*” (p. 89) for details.

**Example 3.24. Using the sql() driver with an MSSQL database**

The following example sends the log messages into an MSSQL database running on the *logserver* host. The messages are inserted into the *syslogng* database, the name of the table includes the exact date when the messages were sent. The syslog-ng application automatically creates the required tables and columns, if the user account used to connect to the database has the required privileges.

```
destination d_mssql {
sql(type(mssql) host("logserver") port("1433")
  username("syslogng") password("syslogng") database("syslogng")
  table("msgs_${R_YEAR}${R_MONTH}${R_DAY}") columns("datetime varchar(16)", "host
  varchar(32)",
  "program varchar(32)", "pid varchar(8)", "message varchar(4096)")
  values("${R_DATE}", "${HOST}", "${PROGRAM}", "${PID}", "${MSGONLY}")
  indexes("datetime", "host", "program", "pid", "message"));
};
```

The date format used by the MSSQL database must be explicitly set in the */etc/locales.conf* file of the syslog-ng server. Edit or create this file as needed for your configuration. A sample is provided below.

```
[default]
date = "%Y-%m-%d %H:%M:%S"
```

3.4.6. Sending messages to a remote logserver using the IETF-syslog protocol

The *syslog()* driver sends messages to a remote host (e.g., a syslog-ng server or relay) on the local intranet or internet using the new standard syslog protocol developed by IETF (see Section 2.18.2, “IETF-syslog messages” (p. 26) for details about the new protocol). The protocol supports sending messages using the UDP, TCP, or the encrypted TLS networking protocols.

The required arguments of the driver are the address of the destination host (where messages should be sent). The transport method (networking protocol) is optional, syslog-ng uses the TCP protocol by default. For the list of available optional parameters, see Section 8.2.6, “syslog()” (p. 195).

Declaration:

```
syslog(host transport [options]);
```

**Note**

Note that the *syslog* destination driver has required parameters, while the source driver defaults to the local bind address, and every parameter is optional.

The *udp* transport method automatically sends multicast packets if a multicast destination address is specified. The *tcp* and *tls* methods do not support multicasting.

**Note**

The default ports for the different transport protocols are as follows: UDP — 514; TLS — 6514.

**Example 3.25. Using the syslog() driver**

```
destination d_tcp { syslog(ip"10.1.2.3" transport("tcp") port(1999);
localport(999)); };
```

If name resolution is configured, the hostname of the target server can be used as well.

```
destination d_tcp { syslog(ip"target_host" transport("tcp") port(1999);
localport(999)); };
```

Send the log messages using TLS encryption and use mutual authentication. See Section 8.10, “TLS options” (p. 231) for details on the encryption and authentication options.

```
destination d_syslog_tls{
    syslog("10.100.20.40"
    transport("tls")
    port(6514)
    tls(peer-verify(required-trusted)
    ca_dir('/opt/syslog-ng/etc/syslog-ng/keys/ca.d/')
    key_file('/opt/syslog-ng/etc/syslog-ng/keys/client_key.pem')
cert_file('/opt/syslog-ng/etc/syslog-ng/keys/client_certificate.pem'))
};;
```

3.4.7. Sending messages to a remote logserver using the legacy BSD-syslog protocol

The `tcp()`, `tcp6()`, `udp()`, and `udp6()` drivers send messages to another host (e.g., a syslog-ng server or relay) on the local intranet or internet using the UDP or TCP protocol. The `tcp6()` and `udp6()` drivers use the IPv6 network protocol.

All four drivers have a single required parameter specifying the destination host address, where messages should be sent. For the list of available optional parameters, see Section 8.2.7, “`tcp()`, `tcp6()`, `udp()`, and `udp6()`,” (p. 200).

The `udp()` and `udp6()` drivers automatically send multicast packets if a multicast destination address is specified. The `tcp()` and `tcp6()` drivers do not support multicasting.

Declaration:

```
tcp(host [options]);
udp(host [options]);
tcp6(host [options]);
udp6(host [options]);
```

**Example 3.26. Using the tcp() driver**

```
destination d_tcp { tcp("10.1.2.3" port(1999); localport(999)); };
```

If name resolution is configured, the hostname of the target server can be used as well.

```
destination d_tcp { tcp("target_host" port(1999); localport(999)); };
```

To send messages using the IETF-syslog message format, enable the `syslog-protocol` flag:

```
destination d_tcp { tcp("10.1.2.3" port(1999); flags(syslog-protocol) );
```



3.4.8. Sending messages to UNIX domain sockets

The *unix-stream()* and *unix-dgram()* drivers send messages to a UNIX domain socket in either *SOCK_STREAM* or *SOCK_DGRAM* mode.

Both drivers have a single required argument specifying the name of the socket to connect to. For the list of available optional parameters, see Section 8.2.8, “*unix-stream()* & *unix-dgram()*” (p. 204).

Declaration:

```
unix-stream(filename [options]);
unix-dgram(filename [options]);
```



Example 3.27. Using the *unix-stream()* driver

```
destination d_unix_stream { unix-stream("/var/run/logs"); };
```

3.4.9. *usertty()*

This driver writes messages to the terminal of a logged-in user.

The *usertty()* driver has a single required argument, specifying a username who should receive a copy of matching messages.

Declaration:

```
usertty(username);
```

The *usertty()* does not have any further options nor does it support templates.



Example 3.28. Using the *usertty()* driver

```
destination d_usertty { usertty("root"); };
```

3.5. Log paths

Log paths determine what happens with the incoming log messages. Messages coming from the sources listed in the log statement and matching all the filters are sent to the listed destinations.

To define a log path, add a log statement to the *syslog-ng* configuration file using the following syntax:

```
log {
    source(s1); source(s2); ...
    optional_element(filter1|parser1|rewrite1);
    optional_element(filter2|parser2|rewrite2);...
    destination(d1); destination(d2); ...
}
```



```
flags(flag1[, flag2...]);
};
```

**Warning**

Log statements are processed in the order they appear in the configuration file, thus the order of log paths may influence what happens to a message, especially when using filters and log flags.

**Example 3.29. A simple log statement**

The following log statement sends all messages arriving to the localhost to a remote server.

```
source s_localhost { tcp(ip(127.0.0.1) port(1999) ); };
destination d_tcp { tcp("10.1.2.3" port(1999); localport(999)); };
log { source(s_localhost); destination(d_tcp); };
```

All matching log statements are processed by default, and the messages are sent to *every* matching destination by default. So a single log message might be sent to the same destination several times, provided the destination is listed in several log statements, and it can be also sent to several different destinations.

This default behavior can be changed using the *flags()* parameter. Flags apply to individual log paths; they are not global options. The following flags available in syslog-ng:

- *final*: Do not send the messages processed by this log path to any further destination.
- *fallback*: Process messages that were not processed by other log paths.
- *catchall*: Process every message, regardless of its source or if it was already processed by other log paths.
- *flow-control*: Stop reading messages from the source if the destination cannot accept them. See Section 2.13, “Managing incoming and outgoing messages with flow-control” (p. 18).

**Warning**

The *final*, *fallback*, and *catchall* flags apply only for the top-level log paths, they have no effect on embedded log paths.

**Example 3.30. Using log path flags**

Let's suppose that you have two hosts (*myhost_A* and *myhost_B*) that run two applications each (*application_A* and *application_B*), and you collect the log messages to a central syslog-ng server. On the server, you create two log paths:

- one that processes only the messages sent by *myhost_A*; and
- one that processes only the messages sent by *application_A*.

This means that messages sent by *application_A* running on *myhost_A* will be processed by both log paths, and the messages of *application_B* running on *myhost_B* will not be processed at all.

- If you add the *final* flag to the first log path, then only this log path will process the messages of *myhost_A*, so the second log path will receive only the messages of *application_A* running on *myhost_B*.
- If you create a third log path that includes the *fallback* flag, it will process the messages not processed by the first two log paths, in this case, the messages of *application_B* running on *myhost_B*.
- Adding a fourth log path with the *catchall* flag would process every message received by the syslog-ng server.



```
log { source(s_localhost); destination(d_file); flags(catchall); };
```

For details on the individual flags, see Section 8.3, “Log path flags” (p. 207). The effect and use of the *flow-control* flag is detailed in Section 2.13, “Managing incoming and outgoing messages with flow-control” (p. 18).

3.5.1. Using embedded log statements

Embedded log statements (see Section 2.2.1, “Embedded log statements” (p. 8)) re-use the results of processing messages (e.g., the results of filtering or rewriting) to create complex log paths. Embedded log statements use the same syntax as regular log statements, but they cannot contain additional sources. To define embedded log statements, use the following syntax:

```
log {
    source(s1); source(s2); ...

    optional_element(filter1|parser1|rewrite1);
    optional_element(filter2|parser2|rewrite2);...

    destination(d1); destination(d2); ...

    flags(flag1[, flag2...]);

    #embedded log statement
    log
    {
        optional_element(filter1|parser1|rewrite1);
        optional_element(filter2|parser2|rewrite2);...
        destination(d1); destination(d2); ...
        #another embedded log statement
        log
        {
            optional_element(filter1|parser1|rewrite1);
            optional_element(filter2|parser2|rewrite2);...
            destination(d1); destination(d2); ...};
        };
    };
};
```



Warning

The *final*, *fallback*, and *catchall* flags apply only for the top-level log paths, they have no effect on embedded log paths.

**Example 3.31. Using embedded log paths**

The following log path sends every message to the `d_file1` and the `d_file2` destinations.

```
log { source(s_localhost); destination(d_file1); destination(d_file2); };
```

The next example is equivalent with the one above, but uses an embedded log statement.

```
log { source(s_localhost); destination(d_file1);
      log {destination(d_file2); };
};
```

The following example sends every message coming from the host `192.168.1.1` into the `d_file1` destination, and sends every message coming from the host `192.168.1.1` and containing the string `example` into the `d_file2` destination.

```
log { source(s_localhost); host(192.168.1.); destination(d_file1);
      log {message("example"); destination(d_file2); };
};
```

The following example collects logs from multiple source groups and uses the `source()` filter in the embedded log statement to select messages of the `s_network` source group.

```
log { source(s_localhost); source(s_network); destination(d_file1);
      log {source(s_network); destination(d_file2); };
};
```

3.5.2. Configuring flow-control

For details on how flow-control works, see Section 2.13, “Managing incoming and outgoing messages with flow-control” (p. 18). The summary of the main points is as follows:

- The syslog-ng application normally reads a maximum of `log_fetch_limit()` number of messages from a source.
- From TCP and unix-stream sources, syslog-ng reads a maximum of `log_fetch_limit()` from every connection of the source. The number of connections to the source is set using the `max_connections()` parameter.
- Every destination has an output buffer (`log_fifo_size()`).
- Flow-control uses a control window to determine if there is free space in the output buffer for new messages. Every source has its own control window; `log_iw_size()` parameter sets the size of the control window.
- When a source accepts multiple connections, the messages of every connection use the same control window.
- The output buffer must be larger than the control window of every source that logs to the destination.
- If the control window is full, syslog-ng stops reading messages from the source until some messages are successfully sent to the destination.
- If the output buffer becomes full, and neither disk-buffering nor flow-control is used, messages may be lost.

**Note**

If you modify the `max_connections()` or the `log_fetch_limit()` parameter, do not forget to adjust the `log_iw_size()` and `log_fifo_size()` parameters accordingly.



Example 3.32. Sizing parameters for flow-control

Suppose that syslog-ng has a source that must accept up to 300 parallel connections. Such situation can arise when a network source receives connections from many clients, or if many applications log to the same socket. Therefore, set the `max_connections()` parameter of the source to 300. However, the `log_fetch_limit()` (default value: 10) parameter applies to every connection of the source individually, while the `log_iw_size()` (default value: 100) parameter applies to the source. In a worst-case scenario, the destination does not accept any messages, while all 300 connections send at least `log_fetch_limit()` number of messages to the source during every poll loop. Therefore, the control window must accommodate at least `max_connections()*log_fetch_limit()` messages to be able to read every incoming message of a poll loop. In the current example this means that (`log_iw_size()` should be greater than $300 * 10 = 3000$). If the control window is smaller than this value, the control window might fill up with messages from the first connections — causing syslog-ng to read only one message of the last connections in every poll loop.

The output buffer of the destination must accommodate at least `log_iw_size()` messages, but use a greater value: in the current example $3000 * 10 = 30000$ messages. That way all incoming messages of ten poll loops fit in the output buffer. If the output buffer is full, syslog-ng does not read any messages from the source until some messages are successfully sent to the destination.

```
source s_localhost {
    tcp(ip(127.0.0.1) port(1999) max-connections(300)); };
    destination d_tcp {
        tcp("10.1.2.3" port(1999); localport(999)); log_fifo_size(30000); };

    log { source(s_localhost); destination(d_tcp); flags(flow-control);
};
```

If other sources send messages to this destination, then the output buffer must be further increased. For example, if a network host with maximum 100 connections also logs into the destination, then increase the `log_fifo_size()` by 10000.

```
source s_localhost {
    tcp(ip(127.0.0.1) port(1999) max-connections(300)); };
    source s_tcp {
        tcp(ip(192.168.1.5) port(1999) max-connections(100)); };
    destination d_tcp {
        tcp("10.1.2.3" port(1999); localport(999)); log_fifo_size(40000); };

    log { source(s_localhost); destination(d_tcp); flags(flow-control);
};
```

See also Section 7.2, “Handling lots of parallel connections” (p. 136).

3.6. Filters

Filters perform log routing within syslog-ng: a message passes the filter if the filter expression is true for the particular message. If a log statement includes filters, the messages are sent to the destinations only if they pass all filters of the log path. For example, a filter can select only the messages originating from a particular host. Complex filters can be created using filter functions and logical boolean expressions.

To define a filter, add a filter statement to the syslog-ng configuration file using the following syntax:

```
filter <identifier> { expression; };
```

The expression may contain the following elements:

- The functions listed in Table 8.17, “Filter functions in syslog-ng” (p. 209). Some of the functions accept extended regular expressions as parameters.
- The boolean operators *and*, *or*, *not*.
- Parentheses to mark the precedence of the operators when using complex filters.



Example 3.33. A simple filter statement

The following filter statement selects the messages that contain the word *deny* and come from the host *example*.

```
filter demo_filter { host("example") and match("deny"); };
```

For the filter to have effect, include it in a log statement:

```
log demo_filteredlog{
    source(s1); source(s2);
    filter(demo_filter);
    destination(d1); destination(d2); };
```

The *host()*, *match()*, and *program()* filter functions accept regular expressions as parameters.

```
filter demo_regexp_filter { host("system.*1") and match("deny"); };
```

To limit the scope of a filter function, add the *value()* option and set its scope with a macro. For example, to limit the scope of the *match()* filter to the text part of the message, use:

```
match("keyword" value("MESSAGE"))
```

The *value()* parameter accepts both built-in macros and user-defined ones created with a parser. Do not prefix the macros with the \$ sign. For details on macros and parsers, see Section 3.7, “Templates and macros” (p. 60) and Section 3.8, “Parsing messages” (p. 62).



Note

When a log statement includes multiple filter statements, syslog-ng sends a message to the destination only if all filters are true for the message. In other words, the filters are connected with the logical *AND* operator. In the following example, no message arrives to the destination, because the filters are exclusive (the hostname of a client cannot be *example1* and *example2* at the same time):

```
filter demo_filter1 { host("example1"); };
filter demo_filter2 { host("example2"); };

log demo_filteredlog{
    source(s1); source(s2);
    filter(demo_filter1); filter(demo_filter2);
    destination(d1); destination(d2); };
```

To select the messages that come from either host *example1* or *example2*, use a single filter expression:

```
filter demo_filter { host("example1") or host("example2"); };

log demo_filteredlog{
    source(s1); source(s2);
    filter(demo_filter);
    destination(d1); destination(d2); };
```

>Use the
not

operator to invert filters, for example, to select the messages that were not sent by host *example1*:

```
filter demo_filter { not host("example1"); };
```

However, to select the messages that were not sent by host *example1* or *example2*, you have to use the *and* operator (that's how boolean logic works):

```
filter demo_filter { not host("example1") and not host("example2"); };
```

Alternatively, you can use parentheses to avoid this confusion:



```
filter demo_filter { not (host("example1") or host("example2")); };
```

In the extended regular expressions, the characters `()[].*?+^$` are used as special symbols. Therefore, these characters have to be preceded with a backslash (`\`) if they are meant literally. For example, the `\$40` expression matches the `$40` string. Backslashes have to be escaped as well if they are meant literally. For example, the `\\d` expression matches the `\d` string.



Tip

If you use single quotes in, you do not need to escape the backslash, e.g. `match("\.\.")` is equivalent to `match('\.')`.

By default, all regular expressions are case sensitive. To disable the case sensitivity of the expression, add the `flags(ignore-case)` option to the regular expression.

```
filter demo_regexp_insensitive { host("system" flags(ignore-case)); };
```



Note

In regular expressions, the asterisk (`*`) character means 0, 1 or any number of the previous expression. For example, in the `f*ilter` expression the asterisk means 0 or more `f` letters. This expression matches for the following strings: `ilter`, `filter`, `ffilter`, etc. To achieve the wildcard functionality commonly represented by the asterisk character in other applications, use `.*` in your expressions, e.g. `f.*ilter`.

The `level()` filter can select messages corresponding to a single importance level, or a level-range. To select messages of a specific level, use the name of the level as a filter parameter, e.g., use the following to select warning messages:

```
level(warning)
```

To select a range of levels, include the beginning and the ending level in the filter, separated with two dots (`..`). For example, to select every message of error or higher level, use the following filter:

```
level(err..emerg)
```

Similarly, messages sent by a range of facilities can also be selected. Note that this is only possible when using the name of the facilities. It is not possible to select ranges the numerical codes of the facilities.

```
facility(local0..local5)
```

For a complete list of the available levels and facilities, see Section 8.4, “Filter functions” (p. 208).

For a complete description on the above functions, see Section 8.4, “Filter functions” (p. 208).



3.6.1. Optimizing regular expressions in filters

Some filter functions accept regular expressions as parameters. But evaluating general regular expressions puts a high load on the CPU, which can cause problems when the message traffic is very high. Often the regular expression can be replaced with simple filter functions and logical operators. Using simple filters and logical operators, the same effect can be achieved at a much lower CPU load.



Example 3.34. Optimizing regular expressions in filters

Suppose you need a filter that matches the following error message logged by the *xntpd* NTP daemon:

```
xntpd[1567]: time error -1159.777379 is too large (set clock manually);
```

The following filter uses regular expressions and matches every instance and variant of this message.

```
filter f_demo_regexp {
    program("demo_program") and
    match("time error .* is too large .* set clock manually"); };
```

Segmenting the *match()* part of this filter into separate *match()* functions greatly improves the performance of the filter.

```
filter f_demo_optimized_regexp {
    program("demo_program") and
    match("time error") and
    match("is too large") and
    match("set clock manually"); };
```

3.7. Templates and macros

The syslog-ng application allows you to define message templates, and reference them from every object that can use a template. Templates can be used to create standard message formats or filenames. Templates can reference one or more macros (e.g., date, the hostname, etc.). See Section 8.5, “Macros” (p. 210) for a list of macros available in the Linux/Unix versions of syslog-ng, and Section 5.6, “Customizing the message format” (p. 107) for the macros of the syslog-ng Agent for Windows application. Fields from the structured data (SD) part of messages using the new IETF-syslog standard can also be used as macros.

Template objects have a single option called *template_escape*, which is disabled by default (*template_escape(no)*). This behavior is useful when the messages are passed to an application that cannot handle escaped characters properly. Enabling template escaping (*template_escape(yes)*) causes syslog-ng to escape the ' and " characters from the messages.



Note

In versions 2.1 and earlier, the *template_escape()* option was enabled by default.

Macros can be included by prefixing the macro name with a \$ sign, just like in Bourne compatible shells. Regarding braces around macro names, the following two formats are equivalent *"\$MSG"* and *"\${MSG}"*.

Default values for macros can also be specified by appending the *: -* characters and the default value to the macro, e.g.,



```
${HOST:-default_hostname}
```



Note

See Section 5.6, “Customizing the message format” (p. 107) for the macros available in the syslog-ng Agent for Windows application.

The macros related to the date of the message (e.g.: *ISODATE*, *HOURL*, etc.) have two further versions each: one with the *S_* and one with the *R_* prefix (e.g.: *S_DATE* and *R_DATE*). The *S_DATE* macro represents the date found in the log message, i.e. when the message was sent by the original application. *R_DATE* is the date when syslog has received the message.

DATE equals either *S_DATE* or *R_DATE*, depending on the global option set in the now deprecated *use_time_recvd()* parameter (see Section 8.9, “Global options” (p. 227)).



Warning

The hostname-related macros (*FULLHOST*, *FULLHOST_FROM*, *HOST*, and *HOST_FROM*) do not have any effect if the *keep_hostname()* option is disabled.

By default, syslog-ng sends messages using the following template: *\$ISODATE \$HOST \$MSGHDR\$MSG*\n. (The *\$MSGHDR\$MSG* part is written together because the *\$MSGHDR* macro includes a trailing whitespace.)



Note

Earlier versions of syslog-ng used templates and scripts to send log messages into SQL databases. Starting from version 2.1, syslog-ng natively supports direct database access using the *sql()* destination. See Section 8.2.5, “sql()” (p. 191) for details.



Example 3.35. Using templates

The following template (*t_demo_filetemplate*) adds the date of the message and the name of the host sending the message to the beginning of the message text. The template is then used in a file destination: messages sent to this destination (*d_file*) will use the message format defined in the template.

```
template t_demo_filetemplate {
    template("$ISODATE $HOST $MSG\n"); template_escape(no); };

destination d_file {
    file("/var/log/messages" template(t_demo_filetemplate)); };
```

Templates can also be used inline, if they are used only at a single location. The following destination is equivalent with the previous example:

```
destination d_file {
    file ("/var/log/messages"
        template("$ISODATE $HOST $MSG\n") template_escape(no) );

    };
```



3.8. Parsing messages

The syslog-ng application can separate parts of log messages (i.e., the contents of the `$MSG` macro) to named fields (columns). These fields act as user-defined macros that can be referenced in message templates, file- and tablenames, etc.

Parsers are similar to filters: they must be defined in the syslog-ng configuration file and used in the log statement.



Note

The order of filters, rewriting rules, and parsers in the log statement is important, as they are processed sequentially.

To create a parser, define the columns of the message, the delimiter or separator characters, and optionally the characters that are used to escape the delimiter characters (quote-pairs). For the list of parser parameters, see Section 8.6, “Message parsers” (p. 214).

Declaration:

```
parser parser_name {
    csv-parser(column1, column2, ...)
    delimiters()
    quote-pairs()
};
```

Column names work like macros. Always use a prefix to identify the columns of the parsers, e.g., `MYPARSER1.COLUMN1`, `MYPARSER2.COLUMN2`, etc. Column names starting with a dot (e.g., `.HOST`) are reserved for use by syslog-ng.



Example 3.36. Segmenting hostnames separated with a dash

The following example separates hostnames like `example-1` and `example-2` into two parts.

```
parser p_hostname_segmentation {
    csv-parser(columns("HOSTNAME.NAME", "HOSTNAME.ID")
    delimiters("-")
    flags(escape-none)
    template("${HOST}");
};
destination d_file { file("/var/log/messages-${HOSTNAME.NAME:-examplehost}"); };
log { source(s_local); parser(p_hostname_segmentation); destination(d_file);};
```



Example 3.37. Parsing Apache log files

The following parser processes the log of Apache web servers and separates them into different fields. Apache log messages can be formatted like:

```
"%h %l %u %t \"%r\" %>s %b \"%{Referer}i\" \"%{User-Agent}i\" %T %v"
```

Here is a sample message:

```
192.168.1.1 - - [31/Dec/2007:00:17:10 +0100] "GET /cgi-bin/example.cgi HTTP/1.1"
200 2708 "-" "curl/7.15.5 (i4 86-pc-linux-gnu) libcurl/7.15.5 OpenSSL/0.9.8c
zlib/1.2.3 libidn/0.6.5" 2 example.balabit
```




To parse such logs, the delimiter character is set to a single whitespace (*delimiters(" ")*). Whitespaces between quotes and brackets are ignored (*quote-pairs('"'[]')*).

```
parser p_apache {
    csv-parser(columns("APACHE.CLIENT_IP", "APACHE.IDENT_NAME", "APACHE.USER_NAME",
        "APACHE.TIMESTAMP", "APACHE.REQUEST_URL", "APACHE.REQUEST_STATUS",
        "APACHE.CONTENT_LENGTH", "APACHE.REFERER", "APACHE.USER_AGENT",
        "APACHE.PROCESS_TIME", "APACHE.SERVER_NAME")
        flags(escape-double-char,strip-whitespace)
        delimiters(" ")
        quote-pairs('"'[]'))
};
```

The results can be used for example to separate log messages into different files based on the `APACHE.USER_NAME` field. If the field is empty, the *nouser* name is assigned.

```
log { source(s_local);
    parser(p_apache); destination(d_file); };
};
destination d_file { file("/var/log/messages-${APACHE.USER_NAME:-nouser}"); };
```

Multiple parsers can be used to split a part of an already parsed message into further segments.



Example 3.38. Segmenting a part of a message

The following example splits the timestamp of a parsed Apache log message into separate fields.

```
parser p_apache_timestamp {
    csv-parser(columns("APACHE.TIMESTAMP.DAY", "APACHE.TIMESTAMP.MONTH",
        "APACHE.TIMESTAMP.YEAR", "APACHE.TIMESTAMP.HOUR", "APACHE.TIMESTAMP.MIN",
        "APACHE.TIMESTAMP.MIN", "APACHE.TIMESTAMP.ZONE")
        delimiters("/: ")
        flags(escape-none)
        template("${APACHE.TIMESTAMP}"));
};
log { source(s_local);
    log { parser(p_apache); parser(p_apache_timestamp); destination(d_file); };
};
```

3.9. Classifying messages

To classify messages using a pattern database, include a *db_parser()* statement in your syslog-ng configuration file using the following syntax:

```
Declaration:
    parser <identifier> {db_parser(file("<database_filename>"))};
```

Note that using the parser in a log statement only performs the classification, but does not automatically do anything with the results of the classification.

**Example 3.39. Defining pattern databases**

The following statement uses the database located at `/opt/syslog-ng/var/db/patterndb.xml`.

```
parser pattern_db {
    db_parser (
        file("/opt/syslog-ng/var/db/patterndb.xml")
    );
};
```

To apply the patterns on the incoming messages, include the parser in a log statement:

```
log {
    source(s_all);
    parser(pattern_db);
    destination(di_messages_class);
};
```

**Note**

The default location of the pattern database file is `/opt/syslog-ng/var/run/patterndb.xml`. The `file` option of the `db-parser` statement can be used to specify a different file, thus different db-parser statements can use different pattern databases. Later versions of syslog-ng will be able to dynamically generate a main database from separate pattern database files.

**Example 3.40. Using classification results**

The following destination separates the log messages into different files based on the class assigned to the pattern that matches the message (e.g., Violation and Security type messages are stored in a separate file), and also adds the ID of the matching rule to the message:

```
destination di_messages_class {
    file("/var/log/messages-${.classifier.class}")
template("${.classifier.rule_id};${S_UNIXTIME};${SOURCEIP};${HOST};${PROGRAM};${PID};${MSG}\n")

    template_escape(no)
};
```

Sample pattern databases are available at the BalaBit Download page <http://www.balabit.com/network-security/syslog-ng/log-server-appliance/>. However, these are not directly usable in syslog-ng 3.0.x, because they are formatted according to the second version (V2) of the pattern database format, which is supported only by the syslog-ng Store Box (SSB) appliance version 1.0.x. The syslog-ng 3.0.x OSE and PE applications only support the first version (V1) of the pattern database; support for the V2 and V3 pattern database formats will be available in syslog-ng 3.1. In the meantime, you can create your own pattern database: see Section 8.6.2.3, “Creating pattern databases” (p. 220) for details.

3.9.1. Using parser results in filters and templates

3.9.1.1. Filtering messages based on classification

The results of message classification and parsing can be used in custom filters and file and database templates as well. There are two built-in macros in syslog-ng that allow you to use the results of the classification: the `.classifier.class` macro contains the class assigned to the message (e.g., violation, security, or unknown),



while the `.classifier.rule_id` macro contains the identifier of the message pattern that matched the message.



Example 3.41. Using classification results for filtering messages

To filter on a specific message class, create a filter that checks the `.classifier.class` macro, and use this filter in a log statement.

```
filter fi_class_violation {
    match("violation"
        value(".classifier.class")
        type("string")
    );
};

log {
    source(s_all);
    parser(pattern_db);
    filter(fi_class_violation);
    destination(di_class_violation);
};
```

Filtering on the *unknown* class selects messages that did not match any rule of the pattern database. Routing these messages into a separate file allows you to periodically review new or unknown messages.

To filter on messages matching a specific classification rule, create a filter that checks the `.classifier.rule_id` macro. The unique identifier of the rule (e.g., `e1e9c0d8-13bb-11de-8293-000c2922ed0a`) is the *id* attribute of the rule in the XML database.

```
filter fi_class_rule {
    match("e1e9c0d8-13bb-11de-8293-000c2922ed0a"
        value(".classifier.rule_id")
        type("string")
    );
};
```

The message-segments parsed by the pattern parsers can also be used as macros as well. To accomplish this, you have to add a name to the parser, and then you can use this name as a macro that refers to the parsed value of the message.



Example 3.42. Using pattern parsers as macros

For example, you want to parse messages of an application that look like `"Transaction: <type>."`, where `<type>` is a string that has different values (e.g., refused, accepted, incomplete, etc.). To parse these messages, you can use the following pattern:

```
'Transaction: @ESTRING:..@'
```

Here the `@ESTRING@` parser parses the message until the next full stop character. To use the results in a filter or a filename template, include a name in the parser of the pattern, e.g.:

```
'Transaction: @ESTRING:TRANSACTIONTYPE:..@'
```

After that, add a custom template to the logpath that uses this template. For example, to select every *accepted* transaction, use the following custom filter in the log path:

```
match("accepted" value("TRANSACTIONTYPE"));
```

**Note**

The above macros can be used in database columns and filename templates as well, if you create custom templates for the destination or logspace.

Use a consistent naming scheme for your macros, for example, `APPLICATIONNAME_MACRONAME`.

3.10. Rewriting messages

The syslog-ng application can rewrite parts of log messages: it can search and replace text, and also set a specific field to a specified value. Rewriting messages is often used in conjunction with message parsing Section 3.8, “Parsing messages” (p. 62).

Rewrite rules are similar to filters: they must be defined in the syslog-ng configuration file and used in the log statement.

**Note**

The order of filters, rewriting rules, and parsers in the log statement is important, as they are processed sequentially.

To create replace a part of the log message, define the string or regular expression to replace, the string to replace the original text (macros can be used as well), and the field of the message that the rewrite rule should process. Substitution rules can operate on any value available via macros, e.g., `HOST`, `MESSAGE`, `PROGRAM`, or any user-defined macros created using parsers (see Section 8.6, “Message parsers” (p. 214) for details.). Substitution rules use the following syntax:

Declaration:

```
rewrite <name_of_the_rule>
    {subst("<string or regular expression to find>",
"<replacement string>", value(<field name>), flags());};
```

A single substitution rule can include multiple substitutions that are applied sequentially to the message. Note that rewriting rules must be included in the log statement to have any effect.

**Tip**

For case-insensitive searches, add the `flags(ignore-case)` option; to replace every occurrence of the string, add `flags(global)` option.

**Example 3.43. Using substitution rules**

The following example replaces the first occurrence of the string `IP` in the text of the message with the string `IP-Address`.

```
rewrite r_rewrite_subst{subst("IP", "IP-Address", value("MESSAGE"))};
```

To replace every occurrence, use:

```
rewrite r_rewrite_subst{subst("IP", "IP-Address", value("MESSAGE"),
flags("global"))};
```



Multiple substitution rules are applied sequentially; the following rules replace the first occurrence of the string *IP* with the string *IP-Addresses*.

```
rewrite r_rewrite_subst{subst("IP", "IP-Address", value("MESSAGE"));  
subst("Address", "Addresses", value("MESSAGE"))};;
```

To set a field of the message to a specific value, define the string to include in the message, and the field where it should be included. Setting a field can operate on any value available via macros, e.g., *HOST*, *MESSAGE*, *PROGRAM*, or any user-defined macros created using parsers (see Section 8.6, “Message parsers” (p. 214) for details.). Note that this operation completely replaces any previous value of that field. Use the following syntax:

Declaration:

```
rewrite <name_of_the_rule>  
{set("<string to include>", value(<field name>))};;
```



Example 3.44. Setting message fields to a particular value

The following example sets the *HOST* field of the message to *myhost*.

```
rewrite r_rewrite_set{set("myhost", value("HOST"))};;
```

3.11. Configuring global syslog-ng options

The syslog-ng application has a number of global options governing DNS usage, the timestamp format used, and other general points. Each option may have parameters, similarly to driver specifications. To set global options, add an option statement to the syslog-ng configuration file using the following syntax:

```
options { option1(params); option2(params); ... };
```



Example 3.45. Using global options

To disable domain name resolving, add the following line to the syslog-ng configuration file:

```
options { use_dns(no); };
```

For a detailed list of the available options, see Section 8.9, “Global options” (p. 227). See Chapter 7, *Best practices and examples* (p. 136) for important global options and recommendations on their use.

3.12. Enabling disk-based buffering

To enable disk-based buffering, use the *log_disk_fifo_size()* parameter to set the size of the disk buffer in bytes. Note that this value applies to every destination separately; every destination will have its own diskbuffer file, even if the parameter is set as a global option. For details on how disk-based buffering works, see Section 2.14, “Using disk-based buffering” (p. 20). Disk buffers can be used with *tcp()*, *tcp6()*, *syslog()* (when using the *tcp* or *tls* transport methods), and *sql()* destinations. The number of messages that the disk buffer can store depends on the size (length) of the actual messages. The maximum length of a message is limited by the *log_msg_size()* parameter, which is 8192 bytes by default.



The disk buffer is located under `/opt/syslog-ng/var/` on every platform.



Example 3.46. Enabling disk-based buffering

The following example turns on disk-based buffering for the destination. The size of the disk buffer is 4 194 304 bytes (4 megabytes). In a worst-case situation, using the default value of the `log_msg_size()` parameter (8192 bytes), this disk buffer can store at least 512 messages. Typical log messages are about 300-500 bytes long, so a disk buffer of 4 megabytes can store over 8000 messages. Set the size of the disk buffer based on the average size and number of messages, and the longest estimated downtime of the server.

```
destination d_tcp {
    tcp("10.1.2.3" port(1999) log_disk_fifo_size(4194304)); };
```

3.13. Encrypting log messages with TLS

This section describes how to configure TLS encryption in syslog-ng. For the concepts of using TLS in syslog-ng, see Section 2.7, “Secure logging using TLS” (p. 13).

Create an X.509 certificate for the syslog-ng server.



Note

The `subject_alt_name` parameter (or the `Common Name` parameter if the `subject_alt_name` parameter is empty) of the server's certificate must contain the hostname or the IP address (as resolved from the syslog-ng clients and relays) of the server (e.g., `syslog-ng.example.com`).

Alternatively, the `Common Name` or the `subject_alt_name` parameter can contain a generic hostname, e.g., `*.example.com`.

Note that if the `Common Name` of the certificate contains a generic hostname, do not specify a specific hostname or an IP address in the `subject_alt_name` parameter.

Complete the following steps on every syslog-ng client host. Examples are provided using both the legacy BSD-syslog protocol (using the `tcp()` driver) and the new IETF-syslog protocol standard (using the `syslog()` driver):

Procedure 3.1. Configuring TLS on the syslog-ng clients

1. Copy the CA certificate (e.g., `cacert.pem`) of the Certificate Authority that issued the certificate of the syslog-ng server to the syslog-ng client hosts, for example into the `/opt/syslog-ng/etc/syslog-ng/ca.d` directory.

Issue the following command on the certificate: **`openssl x509 -noout -hash -in cacert.pem`** The result is a hash (e.g., `6d2962a8`), a series of alphanumeric characters based on the Distinguished Name of the certificate.

Issue the following command to create a symbolic link to the certificate that uses the hash returned by the previous command and the `.0` suffix.

`ln -s cacert.pem 6d2962a8.0`

2. Add a destination statement to the syslog-ng configuration file that uses the `tls(ca_dir(path_to_ca_directory))` option and specify the directory using the CA certificate. The destination must use the `tcp()` or `tcpv6()` destination driver, and the IP address and port parameters of the driver must point to the syslog-ng server.



Example 3.47. A destination statement using TLS

The following destination encrypts the log messages using TLS and sends them to the *6514/TCP* port of the syslog-ng server having the *10.1.2.3* IP address.

```
destination demo_tls_destination {
    tcp("10.1.2.3" port(6514)
    tls( ca_dir("/opt/syslog-ng/etc/syslog-ng/ca.d")) ); };
```

A similar statement using the IETF-syslog protocol and thus the *syslog()* driver:

```
destination demo_tls_syslog_destination { syslog("10.1.2.3" port(6514)
    transport("tls")
    port(3214)
    tls(ca_dir("/opt/syslog-ng/etc/syslog-ng/ca.d")) ); };
```

3. Include the destination created in Step 2 in a log statement.



Warning

The encrypted connection between the server and the client fails if the *Common Name* or the *subject_alt_name* parameter of the server certificate does not contain the hostname or the IP address (as resolved from the syslog-ng clients and relays) of the server.

Do not forget to update the certificate files when they expire.

Complete the following steps on the syslog-ng server:

Procedure 3.2. Configuring TLS on the syslog-ng server

1. Copy the certificate (e.g., *syslog-ng.cert*) of the syslog-ng server to the syslog-ng server host, for example into the */opt/syslog-ng/etc/syslog-ng/cert.d* directory. The certificate must be a valid X.509 certificate in PEM format.
2. Copy the private key (e.g., *syslog-ng.key*) matching the certificate of the syslog-ng server to the syslog-ng server host, for example into the */opt/syslog-ng/etc/syslog-ng/key.d* directory. The key must be in PEM format, and must not be password-protected.
3. Add a source statement to the syslog-ng configuration file that uses the *tls(key_file(key_file_fullpathname) cert_file(cert_file_fullpathname))* option and specify the key and certificate files. The source must use the source driver (*tcp()* or *tcpv6()*) matching the destination driver used by the syslog-ng client.



Example 3.48. A source statement using TLS

The following source receives log messages encrypted using TLS, arriving to the *1999/TCP* port of any interface of the syslog-ng server.

```
source demo_tls_source {
    tcp(ip(0.0.0.0) port(1999)
    tls(
    key_file("/opt/syslog-ng/etc/syslog-ng/key.d/syslog-ng.key")
    cert_file("/opt/syslog-ng/etc/syslog-ng/cert.d/syslog-ng.cert")) ); };
```

A similar source for receiving messages using the IETF-syslog protocol:



```
source demo_tls_syslog_source {
    syslog(ip(0.0.0.0) port(1999)
    transport("tls")
    tls(
key_file("/opt/syslog-ng/etc/syslog-ng/key.d/syslog-ng.key")
cert_file("/opt/syslog-ng/etc/syslog-ng/cert.d/syslog-ng.cert")) ); };
```

4. Disable mutual authentication for the source by setting the following TLS option in the source statement:
`tls(peer_verify(optional-untrusted) ;`
 To configure mutual authentication, see Section 3.14, “Mutual authentication using TLS” (p. 70).



Example 3.49. Disabling mutual authentication

The following source receives log messages encrypted using TLS, arriving to the 1999/TCP port of any interface of the syslog-ng server. The identity of the syslog-ng client is not verified.

```
source demo_tls_source {
    tcp(ip(0.0.0.0) port(1999)
    tls(
key_file("/opt/syslog-ng/etc/syslog-ng/key.d/syslog-ng.key")
cert_file("/opt/syslog-ng/etc/syslog-ng/cert.d/syslog-ng.cert")
    peer_verify(optional-untrusted)) ); };
```

A similar source for receiving messages using the IETF-syslog protocol:

```
source demo_tls_syslog_source {
    syslog(ip(0.0.0.0) port(1999)
    transport("tls")
    tls(
key_file("/opt/syslog-ng/etc/syslog-ng/key.d/syslog-ng.key")
cert_file("/opt/syslog-ng/etc/syslog-ng/cert.d/syslog-ng.cert")
    peer_verify(optional-untrusted)) ); };
```



Warning

Do not forget to update the certificate and key files when they expire.

For the details of the available `tls()` options, see Section 8.10, “TLS options” (p. 231).

3.14. Mutual authentication using TLS

This section describes how to configure mutual authentication between the syslog-ng server and the client. Configuring mutual authentication is similar to configuring TLS (see Section 3.13, “Encrypting log messages with TLS” (p. 68)), but the server verifies the identity of the client as well. Therefore, each client must have a certificate, and the server must have the certificate of the CA that issued the certificate of the clients. For the concepts of using TLS in syslog-ng, see Section 2.7, “Secure logging using TLS” (p. 13).



Complete the following steps on every syslog-ng client host. Examples are provided using both the legacy BSD-syslog protocol (using the `tcp()` driver) and the new IETF-syslog protocol standard (using the `syslog()` driver):

Procedure 3.3. Configuring TLS on the syslog-ng clients

1. Create an X.509 certificate for the syslog-ng client.
2. Copy the certificate (e.g., `client_cert.pem`) and the matching private key (e.g., `client.key`) to the syslog-ng client host, for example into the `/opt/syslog-ng/etc/syslog-ng/cert.d` directory. The certificate must be a valid X.509 certificate in PEM format and must not be password-protected.
3. Copy the CA certificate of the Certificate Authority (e.g., `cacert.pem`) that issued the certificate of the syslog-ng server to the syslog-ng client hosts, for example into the `/opt/syslog-ng/etc/syslog-ng/ca.d` directory.

Issue the following command on the certificate: **`openssl x509 -noout -hash -in cacert.pem`** The result is a hash (e.g., `6d2962a8`), a series of alphanumeric characters based on the Distinguished Name of the certificate.

Issue the following command to create a symbolic link to the certificate that uses the hash returned by the previous command and the `.0` suffix.

`ln -s cacert.pem 6d2962a8.0`

4. Add a destination statement to the syslog-ng configuration file that uses the `tls(ca_dir(path_to_ca_directory))` option and specify the directory using the CA certificate. The destination must use the `tcp()` or `tcpv6()` destination driver, and the IP address and port parameters of the driver must point to the syslog-ng server. Include the client's certificate and private key in the `tls()` options.



Example 3.50. A destination statement using mutual authentication

The following destination encrypts the log messages using TLS and sends them to the `1999/TCP` port of the syslog-ng server having the `10.1.2.3` IP address. The private key and the certificate file authenticating the client is also specified.

```
destination demo_tls_destination {
    tcp("10.1.2.3" port(1999)
        tls( ca_dir("/opt/syslog-ng/etc/syslog-ng/ca.d")
            key_file("/opt/syslog-ng/etc/syslog-ng/key.d/client.key")
            cert_file("/opt/syslog-ng/etc/syslog-ng/cert.d/client_cert.pem"))
    ); };

destination demo_tls_syslog_destination {
    syslog("10.1.2.3" port(1999)
        transport("tls")
        tls( ca_dir("/opt/syslog-ng/etc/syslog-ng/ca.d")
            key_file("/opt/syslog-ng/etc/syslog-ng/key.d/client.key")
            cert_file("/opt/syslog-ng/etc/syslog-ng/cert.d/client_cert.pem")) ); };
```

5. Include the destination created in Step 2 in a log statement.

**Warning**

The encrypted connection between the server and the client fails if the *Common Name* or the *subject_alt_name* parameter of the server certificate does not the hostname or the IP address (as resolved from the syslog-ng clients and relays) of the server.

Do not forget to update the certificate files when they expire.

Complete the following steps on the syslog-ng server:

Procedure 3.4. Configuring TLS on the syslog-ng server

1. Copy the certificate (e.g., `syslog-ng.cert`) of the syslog-ng server to the syslog-ng server host, for example into the `/opt/syslog-ng/etc/syslog-ng/cert.d` directory. The certificate must be a valid X.509 certificate in PEM format.
2. Copy the CA certificate (e.g., `cacert.pem`) of the Certificate Authority that issued the certificate of the syslog-ng clients to the syslog-ng server, for example into the `/opt/syslog-ng/etc/syslog-ng/ca.d` directory.

Issue the following command on the certificate: **`openssl x509 -noout -hash -in cacert.pem`** The result is a hash (e.g., `6d2962a8`), a series of alphanumeric characters based on the Distinguished Name of the certificate.

Issue the following command to create a symbolic link to the certificate that uses the hash returned by the previous command and the `.0` suffix.

`ln -s cacert.pem 6d2962a8.0`

3. Copy the private key (e.g., `syslog-ng.key`) matching the certificate of the syslog-ng server to the syslog-ng server host, for example into the `/opt/syslog-ng/etc/syslog-ng/key.d` directory. The key must be in PEM format, and must not be password-protected.
4. Add a source statement to the syslog-ng configuration file that uses the `tls(key_file(key_file_fullpathname) cert_file(cert_file_fullpathname))` option and specify the key and certificate files. The source must use the source driver (`tcp()` or `tcpv6()`) matching the destination driver used by the syslog-ng client. Also specify the directory storing the certificate of the CA that issued the client's certificate.

**Example 3.51. A source statement using TLS**

The following source receives log messages encrypted using TLS, arriving to the `1999/TCP` port of any interface of the syslog-ng server.

```
source demo_tls_source {
    tcp(ip(0.0.0.0) port(1999)
    tls(
key_file("/opt/syslog-ng/etc/syslog-ng/key.d/syslog-ng.key")
cert_file("/opt/syslog-ng/etc/syslog-ng/cert.d/syslog-ng.cert")
    ca_dir("/opt/syslog-ng/etc/syslog-ng/ca.d")) ); }
```

A similar source for receiving messages using the IETF-syslog protocol:

```
source demo_tls_syslog_source {
    syslog(ip(0.0.0.0) port(1999)
    transport("tls")
    tls(
key_file("/opt/syslog-ng/etc/syslog-ng/key.d/syslog-ng.key")
```



```
cert_file("/opt/syslog-ng/etc/syslog-ng/cert.d/syslog-ng.cert")
      ca_dir("/opt/syslog-ng/etc/syslog-ng/ca.d") ) ; };
```

**Warning**

Do not forget to update the certificate and key files when they expire.

For the details of the available `tls()` options, see Section 8.10, “TLS options” (p. 231).

3.15. Configuring syslog-ng clients

To configure syslog-ng on a client host, complete the following steps:

Procedure 3.5. Configuring syslog-ng on client hosts

1. Install the syslog-ng application on the host. See Chapter 4, *Installing syslog-ng* (p. 78) for details installing syslog-ng on specific operating systems.
2. Configure the local sources that collect the log messages of the host.
3. Create a network destination that points directly to the syslog-ng server, or to a local relay.
4. Create a log statement connecting the local sources to the syslog-ng server or relay.
5. If the logs will also be stored locally on the host, create local file destinations.
6. Create a log statement connecting the local sources to the file destination.
7. Set filters and options (e.g., TLS encryption) as necessary.

**Example 3.52. A simple configuration for clients**

The following is a simple configuration file that collects local log messages and forwards them to a logserver using the IETF-syslog protocol.

```
@version:3.0

options {
    mark_freq(30);
};

source s_local { unix-stream("/dev/log"); internal(); };

destination d_syslog_tcp {
    syslog("192.168.1.1" transport("tcp") port(2010));
};

log { source(s_local); destination(d_syslog_tcp); };
```

3.16. Configuring syslog-ng relays

To configure syslog-ng on a relay host, complete the following steps:



Procedure 3.6. Configuring syslog-ng on relay hosts

1. Install the syslog-ng application on the host. See Chapter 4, *Installing syslog-ng* (p. 78) for details installing syslog-ng on specific operating systems.
2. Configure the network sources that collect the log messages sent by the clients.
3. Create a network destination that points to the syslog-ng server.
4. Create a log statement connecting the network sources to the syslog-ng server.
5. Configure the local sources that collect the log messages of the relay host.
6. Create a log statement connecting the local sources to the syslog-ng server.
7. Set filters and options (e.g., TLS encryption) as necessary.



Note

By default, the syslog-ng server will treat the relayed messages as if they were created by the relay host, not the host that originally sent them to the relay. In order to use the original hostname on the syslog-ng server, use the `keep_hostname (yes)` option both on the syslog-ng relay and the syslog-ng relay. This option can be set individually for every source if needed.

In relay mode, syslog-ng cannot write messages received from network sources into files; the `file()` destination is disabled. The following sources are network sources: `syslog()`, `tcp()`, `tcp6()`, `udp()`, `udp6()`.



Example 3.53. A simple configuration for relays

The following is a simple configuration file that collects local and incoming log messages and forwards them to a logserver using the IETF-syslog protocol.

```
@version:3.0

options {
    mark_freq(30);
    keep_hostname(yes);
};

source s_local { unix-stream("/dev/log"); internal(); };
source s_network { syslog(transport(tcp)); };

destination d_syslog_tcp {
    syslog("192.168.1.5" transport("tcp") port(2010)
);
};

log { source(s_local); source(s_network);
    destination(d_syslog_tcp); };
```

3.17. Configuring syslog-ng servers

To configure syslog-ng on a server host, complete the following steps:



Procedure 3.7. Configuring syslog-ng on server hosts

1. Install the syslog-ng application on the host. See Chapter 4, *Installing syslog-ng* (p. 78) for details installing syslog-ng on specific operating systems.
2. Configure the network sources that collect the log messages sent by the clients and relays.
3. Create local destinations that will store the log messages, e.g., files or programs.
4. Create a log statement connecting the network sources to the local destinations.
5. Configure the local sources that collect the log messages of the syslog-ng server.
6. Create a log statement connecting the local sources to the local destinations.
7. Set filters, options (e.g., TLS encryption) and other advanced features as necessary.



Note

By default, the syslog-ng server will treat the relayed messages as if they were created by the relay host, not the host that originally sent them to the relay. In order to use the original hostname on the syslog-ng server, use the `keep_hostname (yes)` option both on the syslog-ng relay and the syslog-ng relay. This option can be set individually for every source if needed.



Example 3.54. A simple configuration for servers

The following is a simple configuration file for syslog-ng Premium Edition that collects incoming log messages and stores them in a logstore file.

```
@version:3.0
options {
    time_reap(30);
    mark_freq(10);
    keep_hostname(yes);
};

source s_local { unix-stream("/dev/log"); internal(); };
source s_network { syslog(transport(tcp)); };

destination d_logstore {
    logstore(
        "/var/log/syslog-ng-pe/out/logstore.lgs"
    );
};

encrypt_certificate("/opt/syslog-ng/etc/kulcsok/public-certificate.pem")
    chunk_size(10000)
    chunk_time(1)
    compress(2)
    owner("root")
    group("root")
    perm(0777)
};

log { source(s_local); source(s_network); destination(d_logstore); };
```

3.18. Installing and upgrading the license

The syslog-ng Premium Edition server operates only if a valid license file is present on the host. The license file is called `license.txt`, and is located in the same directory as the syslog-ng configuration file.



Warning

The `license.txt` file must be readable to the user running the syslog-ng process.

To install a license file, copy it to the directory where the configuration file is stored. See Section 3.1, “The syslog-ng configuration file” (p. 30) for the location of the license file.

To upgrade a license file, simply overwrite the old license file with the new one.



Note

The license file is needed only when running syslog-ng Premium Edition in server mode.

3.19. Troubleshooting syslog-ng

This section provides tips and guidelines about troubleshooting problems related to syslog-ng. Troubleshooting the syslog-ng Agent for Windows application is discussed in Section 5.10, “Troubleshooting syslog-ng Agent for Windows” (p. 117).



Tip

As a general rule, first try to get logging the messages to a local file. Once this is working, you know that syslog-ng is running correctly and receiving messages, and you can proceed to forwarding the messages to the server.

If the syslog-ng server does not receive the messages, use `tcpdump` or a similar packet sniffer tool on the client to verify that the messages are sent correctly, and on the server to verify that it receives the messages.

If syslog-ng is closing the connections for no apparent reason, be sure to check the log messages of syslog-ng. You might also want to run syslog-ng with the `--verbose` or `--debug` command-line options for more-detailed log messages.

Similarly, build up encrypted connections step-by-step: first create a working unencrypted (e.g., TCP) connection, then add TLS encryption, and finally client authentication if needed.

3.19.1. Creating syslog-ng core files

When syslog-ng crashes for some reason, it can create a core file that contains important troubleshooting information. To enable core files, complete the following procedure:

Procedure 3.8. Creating syslog-ng core files

1. Core files are produced only if the *maximum core file size* ulimit is set to a high value in the init script of syslog-ng. Add the following line to the init script of syslog-ng:

```
ulimit -c unlimited
```

2. Verify that syslog-ng has permissions to write the directory it is started from, e.g., `/opt/syslog-ng/sbin/`.
3. If syslog-ng crashes, it will create a core file in the directory syslog-ng was started from.



4. To test that syslog-ng can create a core file, you can create a crash manually. For this, determine the PID of syslog-ng (e.g., using the **ps -All | grep syslog-ng** command), then issue the following command: **kill -ABRT <syslog-ng pid>**
This should create a core file in the current working directory.

3.19.2. Running a failure script

When syslog-ng is abnormally terminated, it can execute a user-created failure script. This can be used for example to send an automatic e-mail notification. The script must be located at `/opt/syslog-ng/sbin/syslog-ng-failure`.

3.19.3. Stopping syslog-ng

To avoid problems, always use the init scripts to stop syslog-ng (`/etc/init.d/syslog-ng stop`), instead of using the **kill** command. This is especially true on Solaris and HP-UX systems, here use `/etc/init.d/syslog stop`.



Chapter 4. Installing syslog-ng

This chapter explains how to install syslog-ng on the supported platforms using the precompiled binary files.

Version 3.0 of syslog-ng features a unified installer package with identical look on every supported platform (excluding Microsoft Windows and IBM System i — see the respective chapters of this guide for details on installing syslog-ng on [Microsoft Windows](#) and [IBM System i](#)).

**Note**

For instructions on compiling syslog-ng Open Source Edition from the source code, see Section 4.4, “Compiling syslog-ng from source” (p. 86).

As of syslog-ng Open Source Edition 3.0.2, binary installation packages of syslog-ng OSE are available for free for the supported Linux and BSD platforms.

The syslog-ng binaries include all required libraries and dependencies of syslog-ng. The components are installed into the `/opt/syslog-ng` directory. It can automatically re-use existing configuration and license files, and also generate a simple configuration automatically into the `/opt/syslog-ng/etc/syslog-ng.conf` file.

**Note**

There are two versions of every binary release. The one with the *client* suffix does not include the libraries required to log into SQL databases. If you are installing syslog-ng in client or relay mode, or you do not use the *sql()* destination, use these binaries. That way no unnecessary components are installed to your system.

The syslog-ng application can be installed interactively following the on-screen instructions as described in Section 4.1, “Installing syslog-ng using the .run installer” (p. 78), and also without user interaction using the silent installation option — see Section 4.1.3, “Installing syslog-ng without user-interaction” (p. 84).

4.1. Installing syslog-ng using the .run installer

This section describes how to install the syslog-ng application interactively using the binary installer. The installer has a simple interface: use the TAB or the arrow keys of your keyboard to navigate between the options, and Enter to select an option.

- To install syslog-ng on clients or relays, complete Section 4.1.1, “Installing syslog-ng in client or relay mode” (p. 79).
- To install syslog-ng on your central logserver, complete Section 4.1.2, “Installing syslog-ng in server mode” (p. 81).
- To install syslog-ng without any user-interaction, complete Section 4.1.3, “Installing syslog-ng without user-interaction” (p. 84).

**Note**

The installer stops the running syslogd application if it is running, but its components are not removed. The `/etc/init.d/sysklogd` init script is automatically renamed to `/etc/init.d/sysklogd.backup`. Rename this file to its original name if you want to remove or restart the syslogd package.

4.1.1. Installing syslog-ng in client or relay mode

Complete the following steps to install syslog-ng Premium Edition on clients or relays. See Section 2.3, “Modes of operation” (p. 10) for details on the different operation modes of syslog-ng.

Procedure 4.1. Installing syslog-ng in client or relay mode

1. Login to your MyBalabit account (<http://www.balabit.com/mybalabit>) and download the syslog-ng installer package.
2. Start the installer as root using the `./syslog-ng-premium-edition-<version>-<OS>-<platform>.run` command. (Note that the exact name of the file depends on the operating system and platform.) Wait until the package is uncompressed and the welcome screen appears, then select **Continue**.




Figure 4.1. The welcome screen

3. *Accepting the EULA:* You can install syslog-ng only if you understand and accept the terms of the End-User License Agreement (EULA). The full text of the EULA can be displayed during installation by selecting the **Show EULA** option, and is also available in this guide for convenience at Appendix 2, *BalaBit syslog-ng Premium Edition License contract* (p. 243). Select **Accept** to accept the EULA and continue the installation. If you do not accept the terms of the EULA for some reason, select **Reject** to cancel installing syslog-ng.
4. *Detecting platform and operating system:* The installer attempts to automatically detect your operating system and platform. If the displayed information is correct, select **Yes**. Otherwise select **Exit** to abort the installation, and verify that your platform is supported. See Section 1.6, “Supported platforms” (p. 4) for a list of supported platforms. If your platform is supported but not detected correctly, contact your local distributor, reseller, or the BalaBit Support Team. See Section 5, “Contact and support information” (p. xvi) for contact details.

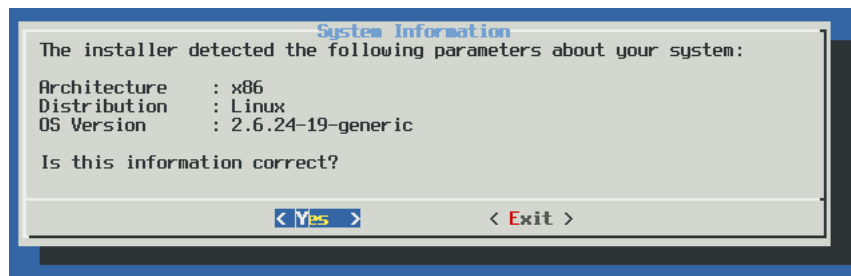


Figure 4.2. Platform detection

5. *Locating the license:* Since you are installing syslog-ng in client or relay mode, simply select **OK**. See Section 2.3, “Modes of operation” (p. 10) for details on the different operation modes of syslog-ng.
6. *Upgrading:* The syslog-ng installer can automatically detect if you have previously installed a version of syslog-ng on your system. To use the configuration file of this previous installation, select **Yes**. To ignore the old configuration file and create a new one, select **No**.
Note that if you decide to use your existing configuration file, the installer automatically checks it for syntax error and displays a list of warnings and errors if it finds any problems.




Figure 4.3. Upgrading syslog-ng

7. *Generating a new configuration file:* The installer displays some questions to generate a new configuration file.
 - a. *Remote sources:* Select **Yes** to accept log messages from the network. TCP, UDP, and SYSLOG messages on every interface will be automatically accepted.

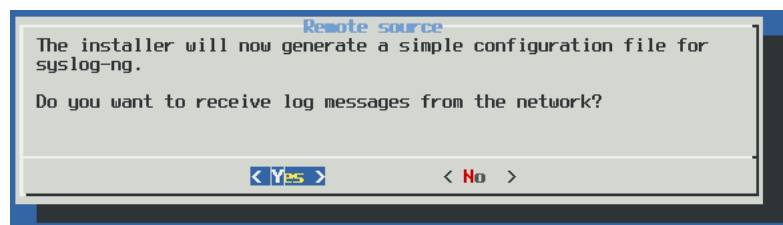


Figure 4.4. Accepting remote messages

- b. *Remote destinations:* Enter the IP address or hostname of your logserver or relay and select **OK**.

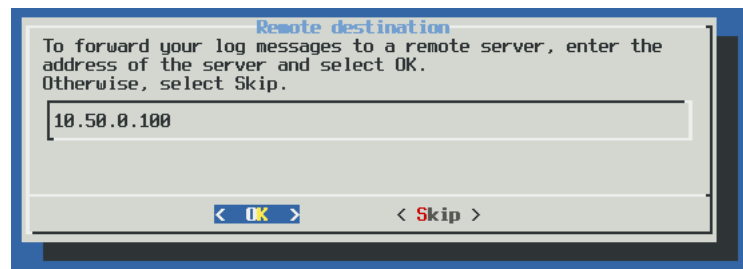


Figure 4.5. Forwarding messages to the logserver



Note

Accepting remote messages and forwarding them to a logserver means that syslog-ng will start in relay mode.

8. After the installation is finished, add the `/opt/syslog-ng/bin` and `/opt/syslog-ng/sbin` directories to your search `PATH` environment variable. That way you can use syslog-ng and its related tools without having to specify the full pathname. Add the following line to your shell profile:

```
PATH=/opt/syslog-ng/bin:$PATH
```



Note

The native logrotation tools do not send a `SIGHUP` to syslog-ng after rotating the log files, causing syslog-ng to write into files already rotated. To solve this problem, the syslog-ng init script links the `/var/run/syslog.pid` file to syslog-ng's pid. Also, on Linux, the `install.sh` script symlinks the initscript of the original syslog daemon to syslog-ng's initscript.

4.1.2. Installing syslog-ng in server mode

Complete the following steps to install syslog-ng on logservers. See Section 2.3, “Modes of operation” (p. 10) for details on the different operation modes of syslog-ng.

Procedure 4.2. Installing syslog-ng in server mode

1. Login to your MyBalabit account (<http://www.balabit.com/mybalabit>) and download the syslog-ng installer package and your syslog-ng Premium Edition license. The license will be required to run syslog-ng in server mode (see Section 2.3.3, “Server mode” (p. 11)) and is needed when you are installing syslog-ng on your central logserver.
2. Start the installer as root using the `./syslog-ng-premium-edition-<version>-<OS>-<platform>.run` command. (Note that the exact name of the file depends on the operating system and platform.) Wait until the package is uncompressed and the welcome screen appears, then select **Continue**.




Figure 4.6. The welcome screen

3. *Accepting the EULA:* You can install syslog-ng only if you understand and accept the terms of the End-User License Agreement (EULA). The full text of the EULA can be displayed during installation by selecting the **Show EULA** option, and is also available in this guide for convenience at Appendix 2, *BalaBit syslog-ng Premium Edition License contract* (p. 243). Select **Accept** to accept the EULA and continue the installation. If you do not accept the terms of the EULA for some reason, select **Reject** to cancel installing syslog-ng.
4. *Detecting platform and operating system:* The installer attempts to automatically detect your operating system and platform. If the displayed information is correct, select **Yes**. Otherwise select **Exit** to abort the installation, and verify that your platform is supported. See Section 1.6, “Supported platforms” (p. 4) for a list of supported platforms. If your platform is supported but not detected correctly, contact your local distributor, reseller, or the BalaBit Support Team. See Section 5, “Contact and support information” (p. xvi) for contact details.

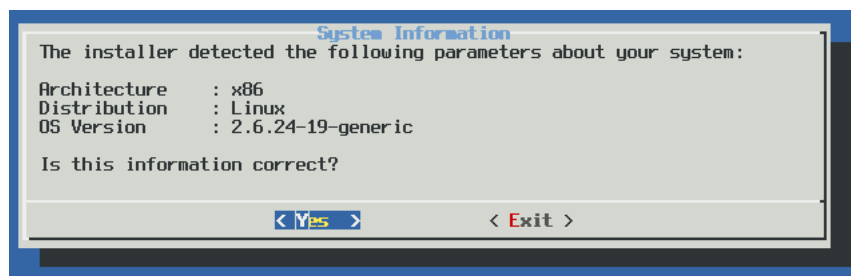


Figure 4.7. Platform detection

5. *Locating the license:* Enter the path to your license file and select **OK**. Typically this is required only for your central logserver. If you are upgrading an existing configuration that already has a license file, the installer automatically detects it.

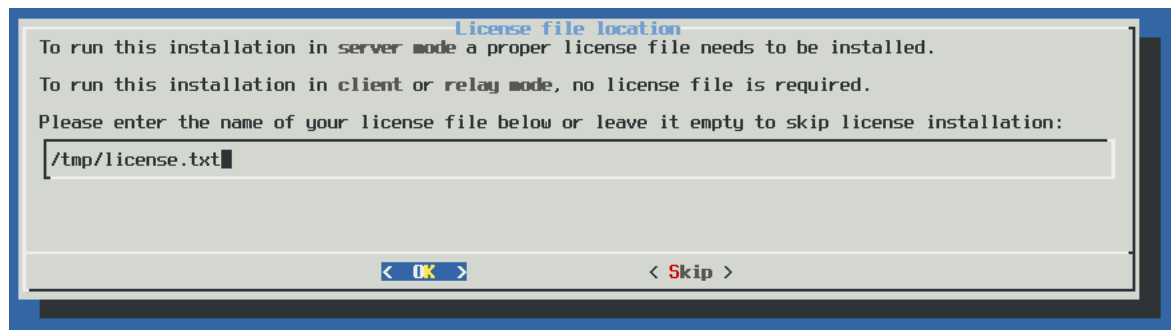


Figure 4.8. Platform detection

6. *Upgrading:* The syslog-ng installer can automatically detect if you have previously installed a version of syslog-ng on your system. To use the configuration file of this previous installation, select **Yes**. To ignore the old configuration file and create a new one, select **No**.
Note that if you decide to use your existing configuration file, the installer automatically checks it for syntax error and displays a list of warnings and errors if it finds any problems.




Figure 4.9. Upgrading syslog-ng

7. *Generating a new configuration file:* The installer displays some questions to generate a new configuration file.
 - a. *Remote sources:* Select **Yes** to accept log messages from the network. TCP, UDP, and SYSLOG messages on every interface will be automatically accepted.

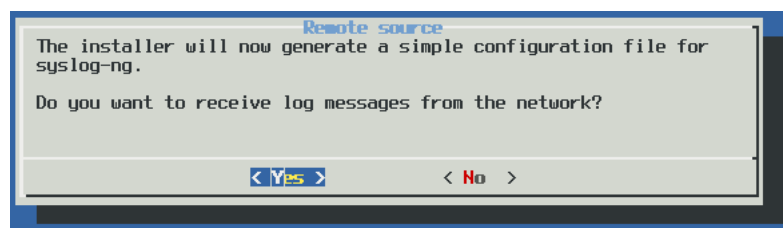


Figure 4.10. Accepting remote messages

- b. *Remote destinations:* Enter the IP address or hostname of your logserver or relay and select **OK**.




Figure 4.11. Forwarding messages to the logserver

**Note**

Accepting remote messages and forwarding them to a logserver means that syslog-ng will start in relay mode.

8. After the installation is finished, add the `/opt/syslog-ng/bin` and `/opt/syslog-ng/sbin` directories to your search `PATH` environment variable. That way you can use syslog-ng and its related tools without having to specify the full pathname. Add the following line to your shell profile:

```
PATH=/opt/syslog-ng/bin:$PATH
```

**Note**

The native logrotation tools do not send a `SIGHUP` to syslog-ng after rotating the log files, causing syslog-ng to write into files already rotated. To solve this problem, the syslog-ng init script links the `/var/run/syslog.pid` file to syslog-ng's pid. Also, on Linux, the `install.sh` script symlinks the initscript of the original syslog daemon to syslog-ng's initscript.

4.1.3. Installing syslog-ng without user-interaction

The syslog-ng application can be installed in silent mode without any user-interaction by specifying the required parameters from the command line. Answers to every question of the installer can be set in advance using command-line parameters.

```
./syslog-ng-premium-edition-<version>.run -- [options]
```

**Warning**

The `--` characters between the executable and the parameters are mandatory, like in the following example:
`./syslog-ng-premium-edition-3.0.1b-solaris-10-sparc-client.run -- --accept-eula -l /var/tmp/license.txt`

To display the list of parameters, execute the `./syslog-ng-premium-edition-<version>.run -- --h` command. Currently the following options are available:

- `--accept-eula` or `-a`: Accept the EULA.



- `--license-file <file> or -l <file>`: Path to the license file.
- `--upgrade | -u`: Perform automatic upgrade — use the configuration file and license file from an existing installation.
- `--remote <destination host>`: Send logs to the specified remote server. Not available when performing an upgrade.
- `--network`: Accept messages from the network. Not available when performing an upgrade.
- `--configuration <file>`: Use the specified configuration file.

4.2. Installing syslog-ng on RPM-based platforms (Red Hat, SUSE, AIX)

To install syslog-ng on operating systems that use the Red Hat Package Manager (RPM), complete the following steps. Installing syslog-ng automatically replaces the original syslog service. The following supported operating systems use RPM:

- AIX 5.2 and 5.3
- CentOS 4 and 5
- openSUSE Linux Enterprise Server 10.0 and 10.1
- Red Hat Enterprise Server 4 and 5
- SUSE Linux Enterprise Server 10 and 10 SP1

Procedure 4.3. Installing syslog-ng on RPM-based systems

1. Login to your MyBalabit account (<http://www.balabit.com/mybalabit>) and download the syslog-ng RPM package for your system from <http://www.balabit.com/network-security/syslog-ng/central-syslog-server/upgrades/>.
2.
 - If the host already uses syslog-ng for logging, execute the following command as root. Otherwise, skip this step.

```
rpm -U syslog-ng-premium-edition-<version>-<OS>-<arch>.rpm
```

The syslog-ng Premium Edition application and all its dependencies will be installed, and the configuration of the existing syslog-ng installation will be used.



Note

If you are upgrading from syslog-ng version 2.1, note that the location of the configuration file has been moved to `/opt/syslog-ng/etc/syslog-ng.conf`

- Execute the following command as root:

```
rpm -i syslog-ng-premium-edition-<version>-<OS>-<arch>.rpm
```

The syslog-ng Premium Edition application and all its dependencies will be installed.



3. Answer the configuration questions of syslog-ng. These are described in detail in Section 4.1, “Installing syslog-ng using the .run installer” (p. 78).



Warning

When performing an upgrade, the package manager might automatically execute the post-uninstall script of the upgraded package, stopping syslog-ng and starting syslogd. If this happens, stop syslogd and start syslog-ng by issuing the following commands:

```
/etc/init.d/syslogd stop
/etc/init.d/syslog-ng start
```

This behavior has been detected on CentOS 4 systems, but may occur on other rpm-based platforms as well.

5. *Optional step for AIX systems:* To redirect the messages of the AIX Error log into syslog, create a file (e.g., /tmp/syslog-ng.add) with the following contents:

```
errnotify:
en_name = "syslog1"
en_persistenceflg = 1
en_method = "logger Msg from Error Log: `errpt -l $1 | grep -v 'ERROR_ID
TIMESTAMP'`"
```

Then execute the following command as root: **odmadd /tmp/syslog-ng.add**.

4.3. Installing syslog-ng on Debian-based platforms

To install syslog-ng on operating systems that use the Debian Software Package (deb) format, complete the following steps. The following supported operating systems use this format:

- Debian etch

Procedure 4.4. Installing syslog-ng on Debian-based systems

1. Login to your MyBalabit account (<http://www.balabit.com/mybalabit>) and download the syslog-ng DEB package for your system from <http://www.balabit.com/network-security/syslog-ng/central-syslog-server/upgrades/>.
2. Issue the following command as root:
dpkg -i syslog-ng-premium-edition-<version>-<OS>-<arch>.deb
3. Answer the configuration questions of syslog-ng. These are described in detail in Section 4.1, “Installing syslog-ng using the .run installer” (p. 78).

4.4. Compiling syslog-ng from source

To compile syslog-ng Open Source Edition (OSE) from the source code, complete the following steps. Alternatively, you can use the precompiled binary packages. Precompiled binary packages are available for free for the supported Linux and BSD platforms at <http://www.balabit.com/network-security/syslog-ng/opensource-logging-system/upgrades/>. Precompiled binary packages for HP-UX, IBM AIX, and Solaris are available for an annual fee at the BalaBit webshop at



<http://www.balabit.com/shop/>. When you buy a binary package, you automatically receive the latest version of syslog-ng OSE for your platform, and all updates for a year.

Procedure 4.5. Compiling syslog-ng from source

1. Download the latest version of syslog-ng OSE from <https://www.balabit.com/downloads/files/syslog-ng/sources/stable/>. The source code is available as a tar.gz archive file.
2. Download the latest version of the EventLog library available at <https://www.balabit.com/downloads/files/eventlog/0.2/>.
3. Install the following packages that are required to compile syslog-ng. These packages are available for most UNIX/Linux systems. Alternatively, you can also download the sources and compile them.
 - the *gcc* C compiler (at least version 2.7.2),
 - the *GNU flex* lexical analyser generator, available at <http://flex.sourceforge.net/>;
 - the *bison* parser generator, available at <http://ftp.gnu.org/gnu/bison/>;
 - and the development files of the *glib* library, available at <http://freshmeat.net/projects/glib/>.
4. If you want to use the spoof-source function of syslog-ng, install the development files of the *libnet* library, available at <http://libnet.sourceforge.net>.
5. If you want to use the */etc/hosts.deny* and */etc/hosts.allow* for TCP access, install the development files of the *libwrap* (also called TCP-wrappers) library, available at <ftp://ftp.porcupine.org/pub/security/index.html> [<ftp://ftp.porcupine.org/pub/security/index.html>].
6. Uncompress the eventlog archive using the

```
$ tar xvfz eventlog-x.x.x.x.tar.gz
```

or the

```
$ gunzip -c eventlog-x.x.x.x.tar.gz | tar xvf -
```

command. A new directory containing the source code of eventlog will be created.

7. By default, eventlog creates a file used by the syslog-ng configure script in the */usr/local/lib/pkgconfig* directory. Issue the following command to add this directory to your `PKG_CONFIG_PATH`:

```
PKG_CONFIG_PATH=/usr/local/lib/pkgconfig:$PKG_CONFIG_PATH
```

8. Enter the new directory and issue the following commands:

```
$ ./configure
$ make
$ make install
```

9. Uncompress the syslog-ng archive using the

```
tar xvfz syslog-ng-x.xx.tar.gz
```



or the

```
unzip -c syslog-ng-x.xx.tar.gz | tar xvf -
```

command. A new directory containing the source code of syslog-ng will be created.

10. Enter the new directory and issue the following commands:

```
$ ./configure
$ make
$ make install
```

These commands will build syslog-ng using its default options.

11. If needed, use the following options to change how syslog-ng is compiled using the following command syntax:

```
$ ./configure --compile-time-option-name
```



Note

You can also use *--disable options*, to explicitly disable a feature and override autodetection. For example, to disable the TCP-wrapper support, use the *--disable-tcp-wrapper* option.



Note

Note that the pre-compiled binary packages of syslog-ng Open Source Edition (OSE) and the syslog-ng Premium Edition packages (both available from the BalaBit webshop at <http://www.balabit.com/shop/>) are compiled with all options enabled.

Execute **syslog-ng --version** to display the list of enabled build parameters of a syslog-ng binary.



Warning

Starting with syslog-ng Open Source Edition 3.0.2, default linking mode of syslog-ng is *dynamic*. This means that syslog-ng might not be able to start up if the */usr* directory is on NFS. On platforms where syslog-ng is used as a system logger, the *--enable-mixed-linking* is preferred.

- *--enable-debug* Include debug information.
- *--enable-dynamic-linking* Compile syslog-ng as a completely dynamic binary. If not specified syslog-ng uses mixed linking (*--enable-mixed-linking*): it links dynamically to system libraries and statically to everything else.
- *--enable-ipv6* Enable IPv6 support.
- *--enable-linux-caps* Enable support for capabilities on Linux.
- *--enable-pcre* Enable using PCRE-type regular expressions. Requires the *libpcre* library package.
- *--enable-spoof-source* Enable spoof_source feature (disabled by default).
- *--enable-static-linking* Compile syslog-ng as a static binary.
- *--enable-sun-door* Enable Sun door support even if not detected (autodetected by default).



- `--enable-sun-streams` Enable Sun STREAMS support even if not detected (autodetected by default).
- `--enable-tcp-wrapper` Enable using `/etc/hosts.deny` and `/etc/hosts.allow` for TCP access (enabled automatically if the `libwrap` libraries are detected).
- `--with-timezone-dir` Specifies the directory where syslog-ng looks for the timezone files to resolve the `time_zone()` and `local_time_zone()` options. If not specified, the `/opt/syslog-ng/share/zoneinfo/` and `/usr/share/zoneinfo/` directories are checked, respectively. Note that HP-UX uses a unique file format (`tztab`) to describe the timezone information; that format is currently not supported in syslog-ng. As a workaround, copy the zoneinfo files from another, non-HP-UX system to the `/opt/syslog-ng/share/zoneinfo/` directory of your HP-UX system.

For information on configuring syslog-ng, see the Chapter 3, *Configuring syslog-ng* (p. 30).

4.5. Uninstalling syslog-ng

If you need to uninstall syslog-ng for some reason, you have the following options:

- *If you have installed syslog-ng using the .run installer:* Execute the **uninstall.sh** script located at `/opt/syslog-ng/bin/uninstall.sh`. The uninstall script will automatically restore the syslog daemon used before installing syslog-ng. To completely remove syslog-ng, including the configuration files, use the **uninstall.sh --purge** command.
- *If you have installed syslog-ng from a .deb package:* Execute the **dpkg -r syslog-ng-premium-edition** command to remove syslog-ng; or the **dpkg -P syslog-ng-premium-edition** command to remove syslog-ng and the configuration files as well. Note that removing syslog-ng does not restore the syslog daemon used before syslog-ng.
- *If you have installed syslog-ng from an .rpm package:* Execute the **rpm -e syslog-ng-premium-edition** command to remove syslog-ng. Note that removing syslog-ng does not restore the syslog daemon used before syslog-ng.

4.6. Configuring Microsoft SQL Server to accept logs from syslog-ng

Complete the following steps to configure your Microsoft SQL Server to enable remote logins and accept log messages from syslog-ng.

Procedure 4.6. Configuring Microsoft SQL Server to accept logs from syslog-ng

1. Start the SQL Server Management Studio application. Select **Start > Programs > Microsoft SQL Server 2005 > SQL Server Management Studio**.
2. Create a new database.



a.




Figure 4.12. Creating a new MSSQL database 1.

In the Object Explorer, right-click on the **Databases** entry and select **New Database**.

b.




Figure 4.13. Creating a new MSSQL database 2.

Enter the name of the new database (e.g., *syslogng*) into the **Database name** field and click **OK**.

3. Create a new database user and associate it with the new database.

a.




Figure 4.14. Creating a new MSSQL user 1.

In the Object Explorer, select **Security**, right-click on the **Logins** entry, then select **New Login**.



b.




Figure 4.15. Creating a new MSSQL user 2.

Enter a name (e.g., *syslog-ng*) for the user into the **Login name** field.

- c. Select the **SQL Server Authentication** option and enter a password for the user.
- d. In the **Default database** field, select the database created in Step 2 (e.g., *syslogng*).
- e. In the **Default language** field, select the language of log messages that you want to store in the database, then click **OK**.

**Warning**

Incorrect language settings may result in the database converting the messages to a different character-encoding format. That way the log messages may become unreadable, causing information loss.

- f. In the Object Explorer, select **Security > Logins**, then right-click on the new login created in the previous step, and select **Properties**.



g.




Figure 4.16. Associating database with the new user

Select **User Mapping**. In the **Users mapped to this login** option, check the line corresponding to the new login (e.g., *syslogng*). In the **Database role membership** field, check the **db_owner** and **public** options.

4.

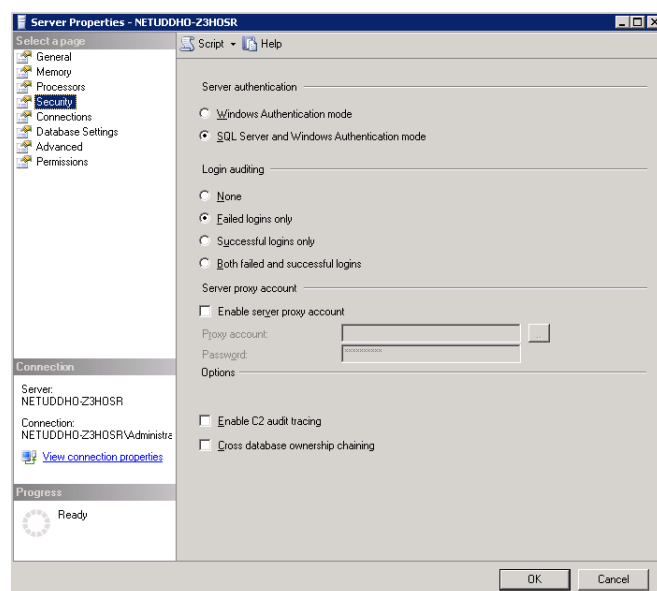


Figure 4.17. Associating database with the new user

Enable remote logins for SQL users.

In the Object Explorer right-click on your database server, and select **Properties > Security**, and set the **Server Authentication** option to **SQL Server and Windows Authentication mode**.



Chapter 5. Collecting logs from Windows hosts

This chapter describes how to install and configure the syslog-ng agent on Microsoft Windows hosts.

The syslog-ng Agent for Windows is a log collector and forwarder application for the Microsoft Windows platform. It collects the log messages of the Windows-based host and forwards them to a syslog-ng server using regular or TLS-encrypted TCP connections.

The features and restrictions of the syslog-ng agent are summarized below:

- Reads messages from eventlog groups and log files.
- Transfers log messages using TCP.
- Supports TLS encryption.
- Authenticates the server using X.509 certificates. Mutual authentication is also supported.
- The format of eventlog messages can be customized using macros.
- Supports multiple destinations both in parallel and fail-over modes.
- Can be managed from a domain controller using group policies.
- Assigns unique message IDs.
- Only basic filtering is supported by the agent, message segmenting, parsing, and classification is not.
- Note that the log messages on Windows come from files — either eventlog containers or custom logfiles — which are already stored on the harddisk, so the agent does not use additional disk buffering.

The syslog-ng agent supports the following operating systems:

- Microsoft Windows Server 2003
- Microsoft Windows XP
- Microsoft Windows Vista
- Microsoft Windows Server 2008



Note

Starting from version 3.0.3, the syslog-ng Agent for Windows application supports the new XML-based eventlog used format on Microsoft Windows Vista and Microsoft Windows Server 2008, and also offers full support for 64-bit operating systems.

5.1. Installing the syslog-ng agent

The syslog-ng Agent for Windows application can be installed in standalone mode on independent hosts. If your hosts are members of a domain, you can install the syslog-ng agent on the domain controller and configure them globally.



- To install the syslog-ng Agent for Windows application in standalone mode, see Section 5.1.1, “Installing the syslog-ng agent in standalone mode” (p. 94).
- To install the syslog-ng Agent for Windows application on the members of a domain, see Section 5.1.2, “Installing the syslog-ng agent on the hosts of a domain” (p. 95).

**Note**

The syslog-ng Agent for Windows application is configured usually using its graphical configuration utility (in standalone mode) or its MMC snap-in (when managed globally from the domain controller). However, it is also possible to use an XML-based configuration file. See Section 5.7, “Using an XML-based configuration file” (p. 113) for details.

5.1.1. Installing the syslog-ng agent in standalone mode

The syslog-ng Agent for Windows application can be installed in standalone mode on independent hosts. If your hosts are members of a domain, install the syslog-ng agent on the domain controller, as described in Section 5.1.2, “Installing the syslog-ng agent on the hosts of a domain” (p. 95). The syslog-ng agent requires about 1 MB hard disk space.

To install the syslog-ng agent in standalone mode, complete the following steps:

**Note**

The syslog-ng Agent for Windows requires the Microsoft .NET Framework version 2.0. This package is usually already installed on most hosts. It can be downloaded at:

<http://www.microsoft.com/downloads/details.aspx?FamilyID=0856EACB-4362-4B0D-8EDD-AAB15C5E04F5&displaylang=en>

Procedure 5.1. Installing the syslog-ng agent in standalone mode

1. Start the installer. Run the `syslog-ng-agent-<versionnumber>-setup.exe` file.

**Note**

Installing the syslog-ng agent requires administrator privileges.

2. Read the End User License Agreement and select **I Agree**.
3. Select the destination folder where you want to install the syslog-ng Agent for Windows application, then select **Next**.
4. Select **Standalone mode**, then click **Next**.
5. Starting from version 3.0.3, the syslog-ng Agent sends only messages that are created after the agent has been installed. If you want to send old log messages to the syslog-ng server, enable the **Send log messages generated before the syslog-ng Agent was installed** option and click **Install**.
6. The installer automatically opens the configuration interface of the syslog-ng agent. As a minimum, you must set the IP address of the destination server, and the agent will automatically start sending eventlog messages to your central logserver from the Application, Security, and System eventlog containers.

**Note**

The installation is completed only after you close the configuration interface.

In standalone mode, to configure an already installed syslog-ng agent, select **Start Menu > Programs > syslog-ng Agent for Windows > Configure syslog-ng Agent**.

5.1.2. Installing the syslog-ng agent on the hosts of a domain

The syslog-ng Agent for Windows application can be installed on the members of a domain from the domain controller, and configured globally using group policies. The syslog-ng agent requires about 1 MB hard disk space.

**Note**

Starting from version 3.0.4, the *.msi* version of the installer does not install the MMC configuration snap-in of the agent, therefore the *.msi* installer does not require the .NET framework. See also Section 5.1.6, “Upgrading syslog-ng Agent for Windows to version 3.0.4” (p. 97) for details.

Procedure 5.2. Installing the syslog-ng agent on the hosts of a domain

**Note**

Starting from version 3.0.3, the syslog-ng Agent sends only messages that are created after the agent has been installed. If you want to send old log messages to the syslog-ng server, download the Orca MSI editor from <http://www.technipages.com/download-orca-msi-editor.html>, open the *.msi* installer of the syslog-ng Agent, select **Property**, and change the value of the **SENDOLDMESSAGES** field to *yes*.

Alternatively, you can also create an XML configuration file for the agent, and configure it to send the old messages. For details on using an XML-based configuration file for the installation, see Section 5.7, “Using an XML-based configuration file” (p. 113).

1. Download both the Microsoft Installer (*.msi*) version and the executable (*.exe*) version of the syslog-ng agent installer to the domain controller host. Make sure to download the executable that includes the MMC snap-in module. Note that separate *.msi* installers are available for 32-bit and 64-bit operating systems.

**Note**

Installing the syslog-ng agent requires administrator privileges, but configuring the related group policies on the domain controller requires domain administrator or higher (e.g., enterprise administrator) privileges.

2. Install the syslog-ng Agent application to your domain controllers using the *.exe* installer.

**Note**

The syslog-ng Agent for Windows requires the Microsoft .NET Framework version 2.0. This package is usually already installed on most hosts. It can be downloaded at:

<http://www.microsoft.com/downloads/details.aspx?FamilyID=0856EACB-4362-4B0D-8EDD-AAB15C5E04F5&displaylang=en>



3. Select **Start > Control Panel > Administrative Tools > Active Directory Users and Computers**, right-click on the Organizational Unit of the domain whose hosts you want to install the syslog-ng agent to, and select **Properties**.
4. Select **Group Policy**, and select the Group Policy object you want to add the syslog-ng agent configuration to. Alternatively, you can create a new group policy object as well.
5. Select **Computer Configuration**, right-click on **Software Settings**, and select **New > Package**.
6. Navigate to the syslog-ng Agent for Windows *.msi* installer and select **Open**.
7. Select **Assigned**, then **OK**.

To configure an already installed syslog-ng agent from the domain controller, install the syslog-ng agent on the domain controller, then select **Start > Control Panel > Administrative Tools > Active Directory Users and Computers**, right-click on the Organizational Unit, select **Properties > Syslog-ng Agent Settings**.

5.1.3. Upgrading syslog-ng Agent for Windows 2.x to 3.0.x

To upgrade the syslog-ng agent application on hosts that are not members of a domain, install the executable (*.exe*) version of the syslog-ng Agent for Windows installer and select **Standalone mode**. The installer automatically receives and converts every setting of version 2.1.x and 2.2beta, and continues to send the log messages to the configured destination. At the end of the installation, the new configuration interface is displayed, where you can start using the new features of the syslog-ng agent.

To upgrade the syslog-ng agent application on hosts that are members of a domain, install the executable (*.exe*) version of the syslog-ng Agent for Windows installer and select **Manage syslog-ng Agent centrally using Group Policy**. After that, the installer asks if you want to use the existing configuration as a Local Policy, or as a Group Policy. (Selecting both options is also possible, although seldom needed.)

If you decide to use it as a Group Policy, enter the unique name for the policy, or select it from the list of available policies. Any local settings are automatically added to the group policy, so these local settings will be applied to every computer that belongs to the selected group policy. Afterwards, the installer converts every setting of version 2.1.x and 2.2beta, and also automatically downloads any group policies that are configured on the domain controller.



Warning

If there are any group policies for the syslog-ng agent configured on the domain controller, downloading the group policies to the clients will overwrite the local settings.



Note

Upgrading from version 2.1 is supported for the 32-bit Windows XP and Server 2003 platforms.

5.1.4. Upgrading syslog-ng Agent for Windows 3.0.1 to version 3.0.2

Due to a minor error in the installer of syslog-ng Agent for Windows 3.0.1, follow the next procedure to upgrade from version 3.0.1 to 3.0.2:



Procedure 5.3. Upgrading syslog-ng Agent for Windows 3.0.1 to version 3.0.2

1. Remove syslog-ng Agent for Windows version 3.0.1 from your group policies, uninstalling the application from your hosts.
2. Verify that syslog-ng Agent for Windows version 3.0.1 was successfully uninstalled from your hosts.
3. Add syslog-ng Agent for Windows version 3.0.2 to your group policies and reboot your hosts. The syslog-ng Agent for Windows version 3.0.2 application will be automatically installed when your hosts are rebooted.

**Note**

Version 3.0.2 of the syslog-ng Agent will retain the existing configuration of the hosts and resume transferring the log messages to the syslog-ng server.

5.1.5. Upgrading syslog-ng Agent for Windows 3.0.2 to version 3.0.3

When upgrading a machine running version 3.0.2 in standalone mode, run the .exe installer and verify that the displayed information is correct.

When upgrading a machine running version 3.0.2 in domain mode, follow the steps described in Section 5.1.2, “Installing the syslog-ng agent on the hosts of a domain” (p. 95) of the documentation. The system will automatically recognize that the new package will update the syslog-ng Agent for Windows version 3.0.2 package.

**Note**

Upgrading to 3.0.3 is not supported on Windows Vista, Server 2008, and 64-bit platforms, as full support for these platforms was introduced only in version 3.0.3.

5.1.6. Upgrading syslog-ng Agent for Windows to version 3.0.4

Starting from version 3.0.4, the .msi version of the installer does not install the MMC configuration snap-in of the agent, therefore the .msi installer does not require the .NET framework.

To upgrade a host that had an earlier version of the agent (one that contained the MMC snap-in) installed, first uninstall the earlier version, then install version 3.0.4, otherwise you may experience erroneous behavior.

5.2. Configuring destinations

The syslog-ng Agent for Windows application can send the log messages of the Windows host to a central log server or relay. It is possible to send the same messages to multiple servers, when each server receives the same messages; and also to configure failover servers, when the agent sends the messages to a primary server, or to a failover server if the primary becomes unavailable. If the agent loses the connection to a destination server and the reconnection fails, it will send an eventlog message. The successful reconnection attempt is also logged. (If the server is unavailable for a long time, the agent sends a log message about the failed connection once in every ten minutes.)



Similarly to the Linux version, the agent now sends MARK messages to the server to indicate that the client host is alive but there are no log messages to send. A MARK message is sent every ten minutes.

To configure a new destination, complete the following steps:

Procedure 5.4. Configuring the destination logservers

1. Start the configuration interface of the syslog-ng Agent for Windows application.
2. Select **syslog-ng Agent Settings > Destinations > Network**, and double-click on **IPv4**.
3. Select **Add**, and enter the hostname or the IP address of the logserver into the **Server Name** field. If your logserver is configured to accept messages on a non-standard port, type the port number into the **Server Port** field.
4. Select the protocol used to transfer log messages. The following protocol templates are available:
 - *Legacy*: Use the legacy BSD-syslog protocol specified in RFC3164. This option uses the following message template: `<${PRI}>${BSDDATE} ${HOST} ${APP_NAME} [${PROCESS_ID}] : ${MSG}`.
 - *Syslog*: Uses the new IETF-syslog protocol specified in RFC 5424-5428 (see <http://www.ietf.org/internet-drafts/draft-ietf-syslog-protocol-23.txt> and <http://www.ietf.org/internet-drafts/draft-ietf-syslog-transport-tls-11.txt>. Starting from version 3.0, syslog-ng also supports the IETF-protocol.
 - *Snare compatible*: Sends log messages in a format compatible with the Snare log monitoring tool, using the following template: `<${PRI}>${BSDDATE} ${HOST} ${MSG}`.



Note

Selecting the **syslog protocol** option is identical to using the `_syslog` driver in syslog-ng 3.0. Similarly, selecting **Legacy syslog** is equivalent to the `_tcp` driver of syslog-ng 3.0.

5. If you have a backup server that can accept log messages if the primary logserver becomes unavailable, select the **Failover Servers** tab, click **Add**, and enter the hostname or the IP address of the backup logserver into the **Server Name** field. Repeat this step if you have more than one backup servers.
6. If you want to send the log messages to more than on server in parallel, so that every server receives every message, repeat Steps 3-4 to add the secondary servers. Secondary servers may have failover servers as well.



Note

The syslog-ng Agent for Windows application considers a message received by the logserver if the primary server of the destination, or one of its failover servers receives it. To modify which server of a destination is the primary server, select **syslog-ng Agent Settings > Destinations > Network > IPv4**, select the server you want to be primary, and select **Edit > Set Primary Server**.

7. Select **Apply**, then **OK**. To activate the changes, restart the syslog-ng Agent service.



5.2.1. Limiting the rate of messages

The syslog-ng Agent can control the rate of messages (message per second) sent to the central server. That way sudden message-bursts can be avoided, and the load of the server is decreased.

To limit the number of messages sent to a destination, complete the following steps:

Procedure 5.5. Limiting the number of messages

1. Start the configuration interface of the syslog-ng Agent for Windows application.
2. Select **syslog-ng Agent Settings > Destinations > Network**, and double-click on **IPv4**.
3. Select the destination server and select **Edit**. To limit the number of messages that the syslog-ng agent sends to the server per second, enter the desired limit into the **Throttle** field. By default (0), the syslog-ng agent does not limit the number of messages sent.



Note

The throttling parameter applies to the total number of messages sent, not to every source independently. The same value applies to the failover servers of the destination.

If you are sending messages to multiple servers, then the speed of the primary server is important: if the primary server cannot accept the messages fast enough, the syslog-ng agent will reduce the number of sent messages to match the speed of the primary server, even if the secondary servers could accept messages faster. If the secondary servers cannot accept messages as fast as the primary server, then the secondary servers will lose messages; the syslog-ng agent will not slow down to wait for them.

4. Select **Apply**, then **OK**. To activate the changes, restart the syslog-ng Agent service.

5.3. Configuring message sources

The syslog-ng Agent for Windows application can read messages from eventlog containers and text files. The following sections explain how to configure these message sources.

- To forward messages from eventlog containers, see Section 5.3.1, “Eventlog sources” (p. 99).
- To forward messages from plain text log files, see Section 5.3.2, “File sources and logrotation” (p. 101).
- Some global settings can apply to both types of sources, these are described in Section 5.3.3, “Global settings of the syslog-ng agent” (p. 102).

5.3.1. Eventlog sources

The syslog-ng Agent for Windows application can collect messages from the standard Windows eventlog containers, as well as from custom containers. The agent automatically forwards the messages from three standard eventlog containers (*Application*, *Security*, *System*). To enable or disable these sources, or to add custom eventlog containers, complete the following steps:

**Note**

The syslog-ng Agent for Windows sends its own log messages into the *Application* eventlog container.

The agent caches in the registry the ID of the last message sent to the destination server, so if the agent is not operating for a time (e.g., it is restarted), then it starts reading messages from the last cached message ID, sending out all the new messages.

**Warning**

If an eventlog container becomes corrupt, the agent will stop processing the event source. A log message (*Eventlog file is corrupt*) is sent directly to the logserver to notify about the error.

Procedure 5.6. Managing eventlog sources

1. Start the configuration interface of the syslog-ng Agent for Windows application.
2. Select **syslog-ng Agent Settings > Eventlog Sources**, and double-click on **Eventlog Containers**.
3.
 - To disable sending messages from an eventlog container, unselect the checkbox before the name of the container.
 - To modify the log facility associated with the messages of the container, select the container, click **Edit**, and select the log facility to use in the **Log Facility** field.
 - To add a custom container, select **Add**, and enter the name of the container into the **Event Container Name** field. If you do not know the name of the container, see Determining the name of a custom eventlog container (p. 100).
4. Select **Apply**, then **OK**. To activate the changes, restart the syslog-ng Agent service.

Procedure 5.7. Determining the name of a custom eventlog container

1. Open the Event Viewer application.
2. Select the custom container you are looking for (e.g., *DNS Server*).
3. Right click on the container and select **Properties**.
4. The name of the container is the name of the file (without the extension) displayed in the **Logname** field (e.g., for *C:\WINDOWS\system32\config\DnsEvent.Evt* it is *DnsEvent*).
5. Use this name as the name of the custom eventlog container during the procedure described in Managing eventlog sources (p. 100).

**Note**

On Windows Vista and Server 2008, some container are not real containers, but show selected messages collected from multiple containers. To forward such messages to the syslog-ng server, you have to find out which real containers are displayed in the container, and add them to the configuration of the syslog-ng Agent.

Some containers have the %4 characters in their names. When adding these to the syslog-ng Agent, replace %4 with the / (slash) character. E.g., write *microsoft-windows-bits-client/analytic* instead of *microsoft-windows-bits-client%4analytic*.

If you are sending old messages to the server as well, the syslog-ng Agent will not send the very first message stored in the container. This is a bug in the Windows API.



5.3.2. File sources and logrotation

The syslog-ng Agent for Windows application can collect log messages from text files, and supports the use of wildcards (*) in filenames and folder names to be able to follow log files that are automatically rotated. To configure file sources, complete the following steps:

Procedure 5.8. Managing file sources

1. Start the configuration interface of the syslog-ng Agent for Windows application.
2. Select **syslog-ng Agent Settings > File Sources**, and double-click on **Sources**.
3. Select **Add > Browse**, and select the log file or the folder containing the log files in the **Base Directory** field. Select or enter the name and extension of the log files in the **File Name Filter** field. Wildcards may be used. The syslog-ng agent will forward log messages from every file that is located in this folder and has a name that matches the filter expression.



Tip

When specifying the Base Directory, you can use the environment variables of Windows, e.g., %WINDIR%, %SYSTEMROOT%, %PROGRAMFILES%, etc.



Warning

Note that when managing members of a domain, the selected path must be available on the domain members, e.g., C:\logs must be available on the client hosts and not on the domain controller.

4.
 - To send messages from the files located in the subfolders of the folder set as Base Directory, select the **Recursive** option.
 - To send messages only from the file that was last modified, select the **Last Modified File Only** option.
If you are forwarding the logs of Internet Information Server (IIS) 5 applications, select the **IIS 5.x Log** option.



Note

If this option is not selected, the syslog-ng agent monitors every matching file in the folder for changes, and sends new log messages from all files.

- To send messages only from the file that was last modified of every subfolder of the Base Directory, select both the **Last Modified File Only** and the **Recursive** options.
- To change the log facility or the log priority associated to the file source, select the desired facility or priority from the **Log Facility** or **Log Priority** fields, respectively.

**Note**

Significant changes to the settings of a file source may cause the syslog-ng Agent to resend the entire contents of the matching files. This means that log messages already sent earlier to the syslog-ng server may be resent and thus duplicated in the server logs. Configuration changes that may result in such behavior are:

- changing the Base Directory,
- changing filter options,
- changing recursivity and Last Modified File Only options.

5. Select **Apply**, then **OK**. To activate the changes, restart the syslog-ng Agent service.

**Note**

If an application writes a message into a log file without ending the line with a new-line character, saves (closes) the file, and later continues to write into the same line, then this is visible in the file as a single line, but the syslog-ng agent interprets them as two separate messages.

**Example 5.1. Collecting the logs of multiple applications from a single folder**

If two applications log into the same folder (e.g., C:\logs), you have to create two file sources. For example, if the name of the log files is *application1-*.log* and *application2-*.log*, respectively, then create two file sources with the C:\logs Base Directory, but with different File Name Filter: *application1-*.log* and *application2-*.log*, respectively.

If other applications log into the C:\logs folder, add a separate expression for each application.

By default, the syslog-ng agent will send every message to the server that arrives into any of the monitored log files. To send only the messages that arrive into the latest file of the source, enable the **Last Modified File Only** option.

5.3.3. Global settings of the syslog-ng agent

The syslog-ng Agent for Windows application has some global settings that can apply to both eventlog and file sources. To configure the global settings, complete the following procedure:

Procedure 5.9. Configuring global settings

1. Start the configuration interface of the syslog-ng Agent for Windows application.
2. Select **syslog-ng Agent Settings** and double-click on **Global Settings**.
3. Set the default log facility associated to the messages.
4. By default, the filters and regular expressions (see Section 5.5, “Filtering messages” (p. 106)) used in the message filters are case-sensitive. To make them case-insensitive, select the **Regular Expressions Ignore Case** or the **Filters Ignore Case** options, or both.

**Note**

The **Regular Expressions Ignore Case** option makes the *Message Contents* filter case-insensitive for both file and eventlog sources. The **Filters Ignore Case** option makes the *Computers*, *Event Sources* and *Categories*, and the *Users* filter case-insensitive.



5. Select **Apply**, then **OK**. To activate the changes, restart the syslog-ng Agent service.

Filters and sources can be disabled globally as well. Disabling filters or sources means that the syslog-ng agent ignores the disabled settings: i.e., if the file sources are disabled, the agent does not send the messages from the files to the server. See the following procedure for details.

Procedure 5.10. Disabling sources and filters globally

1. Start the configuration interface of the syslog-ng Agent for Windows application.
2.
 - To disable file sources, select **syslog-ng Agent Settings**, right-click on **File Sources**, then select **Properties > Disable**.
 - To disable eventlog sources, select **syslog-ng Agent Settings**, right-click on **Eventlog Sources**, then select **Properties > Disable**.
 - To disable file filters, select **syslog-ng Agent Settings > File Sources**, right-click on **Filters**, then select **Properties > Disable**.
 - To disable eventlog filters, select **syslog-ng Agent Settings > Eventlog Sources**, right-click on **Filters**, then select **Properties > Disable**.
3. Select **Apply**, then **OK**. To activate the changes, restart the syslog-ng Agent service.

5.4. Using SSL-encrypted connections with the syslog-ng agent

When connecting to a syslog-ng server using an encrypted connection, the syslog-ng agent verifies the certificate of the server. The connection is established only if the Certificate Authority (CA) that issued the certificate of the server is available in the Certificate Store (**MMC > Certificates > Computer Account > Local Computer > Trusted Root Certificates**) of the Windows-based host.



Note

This certificate (sometimes also called the CACert of the server) is not the certificate of the server: it is the certificate of the CA that signed the certificate of the server. (For details on how certificate-based authentication works, see Section 2.7, “Secure logging using TLS” (p. 13))

To enable SSL-encrypted connections to the server, complete the following steps:

Procedure 5.11. Enabling encrypted connections

1. Start the configuration interface of the syslog-ng Agent for Windows application.
2. Select **syslog-ng Agent Settings > Destinations > Network**, and double-click on **IPv4**.
3. Select the server that accepts encrypted connections and click **Edit**.
4. Select the **Use SSL** option.



Warning

The connection can be established only if the Certificate Authority (CA) that issued the certificate of the server is available in the Certificate Store (**MMC > Certificates > Computer Account > Local Computer > Trusted Root Certificates**) of the Windows-based host. See Section 5.4.2, “Importing certificates with the Microsoft Management Console” (p. 105) for details on importing certificates.

**Note**

The `subject_alt_name` parameter (or the `Common Name` parameter if the `subject_alt_name` parameter is empty) of the server's certificate must contain the hostname or the IP address (as resolved from the syslog-ng clients and relays) of the server (e.g., `syslog-ng.example.com`).

Alternatively, the `Common Name` or the `subject_alt_name` parameter can contain a generic hostname, e.g., `*.example.com`.

Note that if the `Common Name` of the certificate contains a generic hostname, do not specify a specific hostname or an IP address in the `subject_alt_name` parameter.

5. Select **Apply**, then **OK**. To activate the changes, restart the syslog-ng Agent service.

5.4.1. Using mutual authentication with syslog-ng agent

When the syslog-ng server is configured to use mutual authentication, it requests a certificate from the syslog-ng clients. The syslog-ng agent can automatically show the requested certificate to the server when the connection is established if it is available in the **Personal Certificates** store (**MMC > Certificates > Computer Account > Local Computer > Personal Certificates**) of the Local Computer. Use the **Certificate Import Wizard** to import this certificate. See Section 5.4.2, “Importing certificates with the Microsoft Management Console” (p. 105) for details.

**Note**

If a certificate revocation list (CRL) is available in the Local Computer/Personal Certificates store, the syslog-ng agent verifies that the certificate of the syslog-ng server is not on this list.

Procedure 5.12. Configuring mutual authentication with the syslog-ng Agent for Windows

If the syslog-ng server requests authentication from the syslog-ng Agent, complete the following steps.

1. Create certificates for the clients. By default, the syslog-ng agent will look for a certificate that contains the hostname or IP address of the central syslog-ng server in its Common Name. If you use a different Common Name, do not forget to complete Step 3 to set the Common Name of the certificate.
The certificate must contain the private key and must be in PKCS12 format.

**Tip**

To convert a certificate and a key from PEM format to PKCS12 you can use the following command:

```
openssl pkcs12 -export -in agentcertificate.pem -inkey agentprivatekey.pem
-out agentcertificatewithkey.pfx
```

2. Import this certificate into the **Personal Certificate** store of the Local Computer using the Certificate Import Wizard. See Section 5.4.2, “Importing certificates with the Microsoft Management Console” (p. 105) for details.
3. By default, the syslog-ng agent will look for a certificate that contains the hostname or IP address of the central syslog-ng server in its Common Name. (The agent will look for the server name or address set in the **Server Name** field of the destination.) If the certificate of the client has a different Common Name, complete the following steps:



- a. Start the configuration interface of the syslog-ng Agent for Windows application.
- b. Select **syslog-ng Agent Settings > Destinations > Network**, and double-click on **IPv4**.
- c. Select the server that requires mutual authentication and click **Edit**.
- d. Select the **Use SSL** option, click **Browse**, then select the certificate to use.

**Note**

A common way is to use the hostname or the IP address of the agent as the Common Name of the certificate (e.g., *syslog-ng-agent1.example.com*).

4. Select **Apply**, then **OK**. To activate the changes, restart the syslog-ng Agent service.

5.4.2. Importing certificates with the Microsoft Management Console

To import a certificate, complete the following steps.

Procedure 5.13. Importing certificates with MMC

1. Start Microsoft Management Console by executing `mmc.exe` (**Start menu Run application**).

**Note**

Running `mmc.exe` requires administrator privileges.

2. Click on the **Add/Remove snap-in** item of the **File** menu.
3. Click **Add**, select the **Certificates** module, and click **Add**.
4. Select **Computer account** in the displayed window and click **Next**.
5. Select **Local computer** and click **Close**.
6. To import the certificate of the syslog-ng server, navigate to **Console Root > Certificates > Trusted Root Certificate Authorities > Certificates**.
To import a certificate for the syslog-ng agent to perform mutual authentication, navigate to **Console Root > Certificates > Personal > Certificates**.
7. Right-click on the **Certificates** folder and from the appearing menu select **All tasks > Import**. The **Certificate Import Wizard** will be displayed. Click **Next**.
Optional step: Certificates used to authenticate the syslog-ng agent in mutual authentication include the private key. Provide the password for the private key when requested.
8. Windows offers a suitable certificate store by default, so click **Next**.
9. Click **Finish** on the summary window and **Yes** on the window that marks the successful importing of the certificate.



5.5. Filtering messages

The syslog-ng Agent for Windows application can filter log messages in a blacklist-fashion: you can define filters, and any message that matches the filters is ignored by the agent — only messages that do not match the filters are sent to the central server. In other words, the filters are connected to each other with logical OR operations.

Different filters are available for eventlog- and file sources. When the syslog-ng agent processes a message, it checks the relevant filters on-by-one: if it finds a filter that matches the message, the agent stops processing the message without sending it to the server.



Note

By default, all filters are case sensitive. To change this behavior, see Section 5.3.3, “Global settings of the syslog-ng agent” (p. 102).

The following types of filters are available for eventlog sources:

- *Sources and Event ID*: Filter on the source (application) that created the message, and optionally on the identification number of the event. Corresponds with the `EVENT_SOURCE` and `EVENT_ID` macros.
- *Message Contents*: Filter the text of the message, i.e., the contents of the `EVENT_MESSAGE` macro.
- *Sources and Categories*: Filter on the source (application) that created the message, and optionally on the category of the event. Corresponds with the `EVENT_SOURCE` and `EVENT_CATEGORY` macros. Note that leaving the category field empty equals with the *none* category of the Event Viewer.
- *Users*: Filter on the username associated with the event. Corresponds with the `EVENT_USERNAME` macro.
- *Computers*: Filter on the name of the computer (host) that created the event. Corresponds with the `HOST` macro.
- *Event Types*: Filter on the type of the event. Corresponds with the `EVENT_TYPE` macro.

To modify the filters used for eventlog messages, complete the following procedure:

Procedure 5.14. Filtering eventlog messages

1. Start the configuration interface of the syslog-ng Agent for Windows application.
2. Select **syslog-ng Agent Settings > Eventlog Sources > Filters**, and double-click on the type of filter you want to create.
3.
 - To ignore messages sent by a specific application, or messages of the application with a specific event ID, double-click on **Sources and Event ID**, select **Add**, and select the name of the source (application) whose messages you want to ignore from the **Source Name** field. To ignore only specific messages of the application, enter the ID of the event into the **Event ID** field. Select **Add > Apply**.
 - To ignore messages that contain a specific string or text, double-click on **Message Contents**, enter the search term or a regular expression into the **Regular Expression** field, then select **Add > Apply**.
 - To ignore messages sent by a specific application, or messages of the application that fall into a specific category, double-click on **Sources and Categories**, select **Add**, and select the name of



the application whose messages you want to ignore from the **Application Name** field. To ignore only those messages of the application that fall into a specific category, enter the name of the category into the **Category** field. Select **Add > Apply**.

- To ignore messages sent by a specific user, double-click on **Users**, enter the name of the user into the **User** field, then select **Add > Apply**.
- To ignore messages sent by a specific computer (host), double-click on **Computers**, enter the name of the user into the **Computer** field, then select **Add > Apply**.
- *Event Types*: To ignore messages of a specific event-type, double-click on **Event Types**, select the event types to ignore, and select **Ok > Apply**.



Note

Under Windows Vista and Server 2008, Windows labels certain messages as level 3 and the Event Viewer labels such messages as warnings. This is against the official specification: level 3 should not be used; and only level 2 messages are warnings. To filter these events, you have to manually add a new event type to the registry and set its value to 3, e.g.,
`HKEY_LOCAL_MACHINE\SOFTWARE\BalaBit\syslog-ng Agent\Local Settings\EventSources\Filter\Type\Rule0\Type=3`

4. Select **Apply**, then **OK**. To activate the changes, restart the syslog-ng Agent service.

The following types of filters are available for file sources:

- *Message Contents*: Filter the text of the message, i.e., the contents of the FILE_MESSAGE macro.

To modify the filters used for file messages, complete the following procedure:

Procedure 5.15. Filtering file messages

1. Start the configuration interface of the syslog-ng Agent for Windows application.
2. Select **syslog-ng Agent Settings > File Sources > Filters**, and double-click on the type of filter you want to create.
3.
 - To ignore messages that contain a specific string or text, double-click on **Message Contents**, enter the search term or a regular expression into the **Regular Expression** field, then select **Add**.
4. Select **Apply**, then **OK**. To activate the changes, restart the syslog-ng Agent service.

5.6. Customizing the message format

The format of the messages received from the eventlog and the file sources can be customized using templates. You can define separate message format for the eventlog and the file sources. When creating a template to customize the message format, you can use macros, all alphanumeric characters, and the following special characters: `<>, () : ; - + / _`.

To create a template, complete the following procedure:

**Warning**

These macros are available only in the syslog-ng Agent for Windows. To recognize Windows-specific elements of the log message (e.g., eventlog-related macros) on the syslog-ng server, you have to use parsers on the syslog-ng server. The parser must be configured to match the message format set in the syslog-ng Agent. See Section 3.8, “Parsing messages” (p. 62) for details.

Procedure 5.16. Customizing messages using templates

1. Start the configuration interface of the syslog-ng Agent for Windows application.
2. Select **syslog-ng Agent Settings > Destinations > Network**, and double-click on **IPv4**. Select your logserver, and click **Edit**.
3. Type the message format you want to use into the **Template** field. Do not forget to add the \$ character before macros. See Section 5.6.2, “Macros available in the syslog-ng Agent” (p. 109) for a complete list of the available macros.

For example, to send the messages in the *DATE HOSTNAME MESSAGE* format, type *Date:\$DATE Hostname:\$HOST Logmessage:\$MESSAGE*.

Note that the \$MESSAGE macro contains not only the text of the log message, but also additional information received from the message source, such as the name of the eventlog container, or the file, as set in the eventlog-specific and file-specific templates. See Customizing eventlog messages (p. 108) and Customizing eventlog messages (p. 108) for details on modifying the eventlog-specific and file-specific templates.

**Note**

Templates are assigned to a single destination server, so it is possible to use different templates for different servers. However, a server and its failover servers always receive the same message.

**Warning**

If you have more than one destination servers configured (separate servers, not in failover mode), and you want to use the same template for every server, you must manually copy the template into the configuration of each server. Template modifications are not applied automatically to every server.

4. Click **OK**.
5. To activate the changes, restart the syslog-ng Agent service.

To customize the format of eventlog messages, complete the following procedure. This template is applied by the \$MESSAGE macro to format messages received from the eventlog.

Procedure 5.17. Customizing eventlog messages

1. Start the configuration interface of the syslog-ng Agent for Windows application.
2. Select **syslog-ng Agent Settings**, right-click on **Eventlog Sources** and select **Properties**.
3. Type the message format into the **Message Template** field. You can use date- and eventlog-related macros (see Section 5.6.2, “Macros available in the syslog-ng Agent” (p. 109) for a list of macros). The message customized here is included in the server-specific templates using the *MESSAGE* macro.

By default, the following is sent about file messages: *\${EVENT_USERNAME}: \${EVENT_NAME} \${EVENT_SOURCE}: [\${EVENT_TYPE}] \${EVENT_MSG} (EventID \${EVENT_ID})*.



4. Select **Apply**, then **OK**. To activate the changes, restart the syslog-ng Agent service.

To customize the format of file messages, complete the following procedure. This template is applied by the `$MESSAGE` macro to format messages received from the log files.

Procedure 5.18. Customizing file messages

1. Start the configuration interface of the syslog-ng Agent for Windows application.
2. Select **syslog-ng Agent Settings**, right-click on **File Sources** and select **Properties**.
3. Type the message format into the **Message Template** field. You can use date- and file-related macros (see Section 5.6.2, “Macros available in the syslog-ng Agent” (p. 109) for a list of macros). The message customized here is included in the server-specific templates using the `MESSAGE` macro.
By default, the following is sent about file messages: `$FILE_NAME: $FILE_MESSAGE`.
4. Select **Apply**, then **OK**. To activate the changes, restart the syslog-ng Agent service.

5.6.1. Customizing the timestamp used by the syslog-ng Agent

The syslog-ng agent can send the syslog messages using either the ISO or the BSD timestamp format. It is recommended to use the ISO format, because it contains much more information than the BSD format.

Note that in the syslog-ng agent, the macros without prefix (e.g., `DATE`) always refer to the receiving date of the message (e.g., `R_DATE`) when it arrived into the event log container, and are included only for compatibility reasons.



Warning

If a remote host is logging into the event log of the local host that is running syslog-ng Agent for Windows, both hosts should be in the same timezone, because the event log message does not include the timezone information of the sender host. Otherwise, the date of the messages received from the remote host will be incorrect.

5.6.2. Macros available in the syslog-ng Agent

The following tables list the available macros:



Warning

These macros are available only in the syslog-ng Agent for Windows. To recognize Windows-specific elements of the log message (e.g., eventlog-related macros) on the syslog-ng server, you have to use parsers on the syslog-ng server. The parser must be configured to match the message format set in the syslog-ng Agent. See Section 3.8, “Parsing messages” (p. 62) for details.



Note

Note that if you use the Syslog protocol template (meaning that messages are sent using the IETF-syslog protocol), only the message part of the log message can be customized, the structure of the headers and other information is fixed by the protocol.

■ Macros related to protocol headers



- Macros related to the date and time of the message
- Macros related to eventlog sources
- Macros related to file sources

By default, syslog-ng Agent uses the following format: `<${PRI}>${BSDDATE} ${HOST} ${APP_NAME} [${PROCESS_ID}] : ${MESSAGE}`, where `MESSAGE` is `EVENT_USERNAME : EVENT_NAME EVENT_SOURCE : [EVENT_TYPE] EVENT_MSG (EventID EVENT_ID)` for eventlog messages, and `FILE_NAME : FILE_CURRENT_POSITION / FILE_SIZE : FILE_MESSAGE` for file messages.



Macro	Description
HOST	Name of the host sending the message. Hostnames are automatically converted to lowercase.
MESSAGE	The content of the message, including the text of the message and any file- or event-specific macros that are set for the source.
MSG	The content of the message.
PRI	Priority header of the message, storing the facility and the level of the message.

Table 5.1. Protocol-related macros of the syslog-ng agent



Macro	Description
BSDDATE, R_BSDDATE, S_BSDDATE	Date of the message in BSD timestamp format (month/day/hour/minute/second, each expressed in two digits). This is the original syslog time stamp without year information, e.g., <i>Jun 13 15:58:00</i> . If possible, it is recommended to use <i>ISODATE</i> for timestamping.
DATE	An alias of the <i>ISODATE</i> macro.
DAY, R_DAY, S_DAY	The day the message was sent.
FULLDATE, R_FULLDATE, S_FULLDATE	A nonstandard format for the date of the message using the same format as <i>DATE</i> , but including the year as well, e.g.: <i>2006 Jun 13 15:58:00</i> .
HOUR, R_HOUR, S_HOUR	The hour of day the message was sent.
ISODATE, R_ISODATE, S_ISODATE	Date of the message in the ISO 8601 compatible standard timestamp format (yyyy-mm-ddThh:mm:ss+-ZONE), e.g.: <i>2006-06-13T15:58:00.123+01:00</i> . If possible, it is recommended to use <i>ISODATE</i> for timestamping. Note that the syslog-ng agent cannot produce fractions of a second (e.g., milliseconds) in the timestamp.
MIN, R_MIN, S_MIN	The minute the message was sent.
MONTH, R_MONTH, S_MONTH	The month the message was sent as a decimal value, prefixed with a zero if smaller than 10.
MONTHNAME, R_MONTHNAME, S_MONTHNAME	The English name of the month the message was sent, abbreviated to three characters (e.g., Jan, Feb, etc.).
R_DATE	Date when the message was recorded into the eventlog container.
S_DATE	Date when the message was created.
SEC, R_SEC, S_SEC	The second the message was sent.
TZ, R_TZ, S_TZ	The name of the time zone of the host.
TZOFFSET, R_TZOFFSET, S_TZOFFSET	The time-zone as hour offset from GMT; e.g.: <i>-07:00</i> . In syslog-ng 1.6.x this used to be <i>-0700</i> but as <i>ISODATE</i> requires the colon it was added to <i>TZOFFSET</i> as well.
UNIXTIME, R_UNIXTIME, S_UNIXTIME	Standard unix timestamp, represented as the number of seconds since <i>1970-01-01T00:00:00</i> .
YEAR, R_YEAR, S_YEAR	The year the message was sent.
WEEK, R_WEEK, S_WEEK	The week number of the year, prefixed with a zero for the first nine week of the year. (The first Monday in the year marks the first week.)
WEEKDAY, R_WEEKDAY, S_WEEKDAY	The 3-letter name of the day of week the message was sent, e.g. <i>Thu</i> .

Table 5.2. Time-related macros of the syslog-ng agent



Macro	Description
EVENT_CATEGORY	The category of the event.
EVENT_FACILITY	The facility that sent the message.
EVENT_ID	The identification number of the event.
EVENT_LEVEL	Importance level of the message represented as a number: 6 - Success, 5 - Informational, 4- Warning, or 3 - Error).
EVENT_MESSAGE	The content of the message.
EVENT_MESSAGE_XML	Contains the entire message in XML format. Available only on Windows Vista and Server 2008 platforms
EVENT_MSG	The content of the message. This is an alias of the <i>EVENT_MESSAGE</i> .
EVENT_NAME	Name of the Windows event log container (e.g., Application or Security).
EVENT_REC_NUM	The record number of the event in the event log.
EVENT_SID	The security identification number of the event.
EVENT_SID_TYPE	The security identification number resolved into name. One of the following: <i>User</i> , <i>Group</i> , <i>Domain</i> , <i>Alias WellKnownGroup</i> , <i>DeletedAccount</i> , <i>Invalid</i> , <i>Unknown</i> , <i>Computer</i> .
EVENT_SOURCE	The application that created the message.
EVENT_TASK	The task category of the event. Available only on Windows Vista and Server 2008 platforms
EVENT_TYPE	The importance level of the message in text format.
EVENT_USERNAME	The user running the application that created the message.

Table 5.3. Eventlog-related macros of the syslog-ng agent

Macro	Description
FILE_CURRENT_POSITION	The position of the message from the beginning of the file in bytes.
FILE_FACILITY	The facility that sent the message.
FILE_LEVEL	Importance level of the message represented as a number: 6 - Success, 5 - Informational, 4- Warning, or 3 - Error).
FILE_MESSAGE	The content of the message.
FILE_MSG	The content of the message. This is an alias of the <i>FILE_MESSAGE</i> macro.
FILE_NAME	Name of the log file (including its path) from where the syslog-ng Agent received the message.
FILE_SIZE	The current size of the file in bytes.

Table 5.4. File-related macros of the syslog-ng agent

5.7. Using an XML-based configuration file

Starting from syslog-ng Agent for Windows version 3.0.4, it is possible to specify the configuration of the agent in an XML file when installing the agent, and also when starting the agent. The configuration file must be a valid XML



file that complies to the XML schema supplied with the syslog-ng agent.



Note

By default, the XML schema file is called `syslog-ng-agent-conf.xsd` is located in the installation folder of the syslog-ng agent, next to the `syslog-ng-agent.exe` file.

Procedure 5.19. Creating an XML configuration file for the syslog-ng agent

1. Create a new configuration file, or edit the one shown in Section 5.7.1, “A sample configuration file for the syslog-ng Agent” (p. 115). Use a text editor that can validate the file to the XML schema of the configuration file. One such editor is the Microsoft XML Notepad 2007 application, which is available for free at <http://msdn.microsoft.com/en-us/xml/bb190622.aspx>.
2. When creating the configuration file, bear in mind the following points:
 - For details on the format of the XML file, see the sample file at Section 5.7.1, “A sample configuration file for the syslog-ng Agent” (p. 115) and XML schema (.xsd) file installed with the agent.
 - File sources, event sources, servers, and filters must have a unique index, that is, the definition of the first server should start as `<Server0 Enabled="1"`, the second `<Server2 Enabled="1"`, etc.
 - File sources must have a unique identifier (UUID). The agent does not create these identifiers, you must enter them into the configuration file manually.
 - If you do not use throttling, remove the `Throttle` attribute from the destination. Setting the `Throttle` attribute to `0` is not accepted by the agent.
 - If you do not want the agent to send old (already existing) messages to the logserver, use the following in the configuration file:

```
<syslog-ng_Agent SendOldMessages="0">
```

Note that when it starts, the agent automatically removes the `SendOldMessages="0"` attribute from the configuration file, but it will not resend the messages after the agent is restarted.

3. To start the agent and use the configuration file, open a command prompt, and issue the following command: **syslog-ng-agent.exe -c myconfigfile.xml -d**. This command will start the agent in debug mode, and display any errors of the XML configuration file.
4. If there are no errors in the configuration file, start the agent in normal mode: **syslog-ng-agent.exe -c myconfigfile.xml**.

To use the XML file during the installation of the agent, use the same syntax with the installer: **syslog-ng-agent-3.0.4-setup.exe /xmlconfig="fullpath\myconfigfile.xml"**. Note that the XML schema file must be in the same folder as the installer file.



Note

If you want to use the .msi installer with an xml file, use the **syslog-ng-agent-3.0.4-setup.msi SLNGOPTS="/xmlconfig=fullpath\myconfigfile.xml"** command, or edit the installer with the Orca MSI editor (<http://www.technipages.com/download-orca-msi-editor.html>), and add the **SLNGOPTS="/xmlconfig=fullpath\myconfigfile.xml"** to the installation parameters on the **Customization** tab.



5.7.1. A sample configuration file for the syslog-ng Agent

The following is a sample configuration file for the syslog-ng Agent for Windows application.

```
<syslog-ng-agent-configuration>
  <SOFTWARE>
    <BalaBit>
      <syslog-ng_Agent WriteMinidump="1" SendOldMessages="0">
        <Local_Settings Enabled="1" RegExpIgnoreCase="1"
FilterIgnoreCase="0" LogFacility="13">
          <Destinations>
            <Network>
              <IPv4 Enabled="1" PrimaryServer="1">
                <Server Index="1" Enabled="1"
ServerName="10.20.254.158" ServerPort="1999" Throttle="100000" Protocol="0"
ProtocolTemplate="<${PRI}>${BSDDATE} ${HOST} ${MSG}" UseSSL="0"
ClientCertSubject="">
                  <FailoverServers
FailoverServer0="myserver1" FailoverServer1="myserver2"/>
                </Server>
              </IPv4>
            </Network>
          </Destinations>
          <EventSources Enabled="1"
MessageTemplate="${EVENT_USERNAME}: ${EVENT_NAME} ${EVENT_SOURCE}:
[${EVENT_TYPE}] ${EVENT_MSG} (EventID ${EVENT_ID})">
            <Sources Enabled="1">
              <Event Index="1" Enabled="2"
Name="Security"/>
              <Event Index="0" Enabled="1"
Name="Security"/>
            </Sources>
          <Filter>
            <Formatted_Message>
              <Rule Index="0" Regexp="myregexp"/>
              <Rule Index="2" Regexp="myregexp"/>
            </Formatted_Message>
            <Computer>
              <Rule Index="0" Computer="myhost"/>
              <Rule Index="1" Computer="myhost"/>
            </Computer>
            <Type>
              <Rule Index="1" Type="1"/>
              <Rule Index="0" Type="32"/>
            </Type>
            <User>
              <Rule Index="0" Username="myuser"/>
              <Rule Index="1" Username="myuser"/>
          </Filter>
        </Local_Settings>
      </syslog-ng_Agent>
    </BalaBit>
  </SOFTWARE>
</syslog-ng-agent-configuration>
```



```

                                </User>
                                <Source_EventId>
                                    <Rule Index="0" Source="mysource"
EventId="1"/>
                                    <Rule Index="1" Source="mysource"
EventId="2"/>
                                </Source_EventId>
                                <Source_Category>
                                    <Rule Index="0" Source="mysource"
Category="mycategory"/>
                                    <Rule Index="1" Source="mysource"
Category="2"/>
                                </Source_Category>
                            </Filter>
                        </EventSources>
                        <FileSources MessageTemplate="File${File}">
                            <Sources>
                                <File Index="1" Enabled="1"
BaseDirectory="c:\" FileNameFilter="*.log" Recursive="0"
LastModifiedFileOnly="0" id="a455e5ba-d4e9-4b85-8711-e8bf10141028"
PeriodicFileCheck="0" AppName="log"/>
                                <File Index="0" Enabled="1"
BaseDirectory="c:\" FileNameFilter="*.txt" Recursive="0"
LastModifiedFileOnly="0" id="b455e5ba-d4e9-4b85-8711-e8bf10141038"
PeriodicFileCheck="0"/>
                            </Sources>
                        <Filter>
                            <Formatted_Message>
                                <Rule Index="0" Regexp="myregexp"/>
                                <Rule Index="1" Regexp="myregexp"/>
                            </Formatted_Message>
                        </Filter>
                    </FileSources>
                </Local_Settings>
            </syslog-ng_Agent>
        </BalaBit>
    </SOFTWARE>
</syslog-ng-agent-configuration>

```

5.8. Controlling the syslog-ng agent services

During installation, syslog-ng agent registers the *syslog-ng Agent* service that is started automatically when the host boots. To disable the automatic startup of the syslog-ng agent, or manually start or stop the service, use the **Start Menu > Control Panel > Administrative Tools > Services** interface. The service is running with the privileges of the *NT AUTHORITY\SYSTEM* user.



When the syslog-ng Agent service is started or stopped, it sends a syslog message to the central log server and an eventlog message to the Application eventlog container of the host.

**Warning**

If you change the timezone setting of the host while the syslog-ng Agent is running, you have to restart the syslog-ng Agent. Otherwise, it will not receive the updated timezone information and the date of the events will be incorrect.

5.9. Domain versus local settings

Group policies for the syslog-ng Agent can be specified at different levels, e.g., at the domain level, at the organization unit level, at the computer level, or also as a local policy of the computer. When evaluating its configuration settings, the syslog-ng Agent follows the standard policy-inheritance methods of Windows. If the configuration of the syslog-ng Agent is specified at multiple levels (e.g., on the domain level and also at the computer level), then the more specific (or lower level) setting is used (that is, the computer level in the above example). If a setting is not configured at a level, the setting of the next higher level is used (e.g., if something is not configured on the computer level, then the setting of the organization unit — or if it is not specified in the policy of organization unit, then the setting of the domain policy — is used). If a setting is not configured in any group policy, the syslog-ng Agent checks its local policy settings, and uses the local setting if available.

5.10. Troubleshooting syslog-ng Agent for Windows

In case you experience problems with the syslog-ng agent, the following points may be of help.

**Note**

The followings address only problems specific to the syslog-ng agent, and assume that communication between the server and the client is otherwise possible (that is, the server is properly configured to receive messages and is available on the network, and name resolution is properly configured on the client).

- *Configuration changes do not take effect:* Configuration changes take effect only after restarting the syslog-ng service or rebooting the system. Also restart the system after changing the timezone settings of the host, or importing a certificate that you want to use to authenticate the communication between the agent and the server. If the configuration of the agent has changed since the last restart, the syslog-ng agent sends a message of the change, including the hmac-sha-1 hash of the new configuration. Also note that if your clients are managed from a Domain Controller, configuration changes are not instantly downloaded to the client hosts, only at the time of the next group policy update. To update the configuration of a client host earlier, open a command prompt on the client host, and issue the **gpupdate /force** command.

After downloading the configuration from the Domain Controller, the syslog-ng Agent service is automatically restarted if the configuration has changed.

**Note**

Certain domain settings that may affect the syslog-ng Agent are downloaded only when the machine is rebooted. For example, moving the computer from one group policy to another requires a reboot to have effect.

- *The syslog-ng agent does not send messages to the server.* Check the Application eventlog for messages of the syslog-ng agent. In case of connection errors and certificate problems, the syslog-ng agent sends error messages into the eventlog. Ensure that the destination address of the server is correctly set. If you use SSL encryption, verify that the certificate of the Certificate Authority of the server and that the certificate of the client are properly imported. If there are no error messages, check the logs on your logserver: the syslog-ng agent sends a MARK message every ten minutes even if there are no other messages to send.
- *The syslog-ng agent sends only MARK messages to the server.* Verify that you have configured the eventlog and file sources, and that they have not been disabled globally. If these settings are correct but the server still does not send any messages, temporarily disable all filters to see that they are not configured to ignore every message. When using filter, it is also recommended to check the global case-sensitivity settings.
- *Command-line parameters are ignored on Windows Vista and 2008 Server.* Command-line parameters work only for administrators if User Account Control (UAC) is enabled. To execute syslog-ng Agent with command-line parameters, select **Start > Programs > Accessories**, right-click on **Command prompt > Run as administrator**.

If you contact the BalaBit Support Team about a problem with the syslog-ng Agent for Windows, execute the **syslog-ng-agent -V** command from the command line and include every version and platform information it displays in your support request.

- *CPU load is high.* See Section 5.10.1, “Sending messages and CPU load” (p. 118).
- *Losing messages from eventlog containers.* An eventlog container is a special file. The Agent reads this file, formats the messages and sends them to remote log server. Note that the eventlog container can be configured only to a certain size. If the container reaches that size, Windows writes the next message to the beginning of the file. As a result, if the agent is not running (or the destination server is unavailable) so long that the eventlog container is filled up, messages can be lost.

5.10.1. Sending messages and CPU load

The syslog-ng agent application can send messages to the server when the Windows Scheduler provides resources to the syslog-ng agent. When there are many unsent log messages in the log sources, and there is no other significant activity on the host, syslog-ng will start to send the messages to the server, possibly increasing the CPU load to 100%. After all messages have been sent, or if another application requires the resources, the CPU load decreases back to normal.

**Tip**

To avoid the initial large load on the CPU, limit the rate of message sending temporarily. You can remove the limit after the old messages have been sent. See Section 5.2.1, “Limiting the rate of messages” (p. 99) for details.



When relaying the messages from multiple sources, the syslog-ng agent sends one message at a time from each source. That way a single source with a large log traffic does not block other log sources.

5.10.2. Creating core and memory dumps

In certain rare cases, you might have to create core dumps of the syslog-ng Agent to investigate a particular problem. When enabled, the syslog-ng Agent for Windows application creates core dumps automatically when it experiences an unexpected shutdown.

To enable core dumps, set the `HKEY_LOCAL_MACHINE/Software/Balabit/syslog-ng Agent/WriteMinidump` registry key to 1.

Core dumps are written into the installation folder of the syslog-ng Agent under the `syslog-ng-agent.dmp` filename. The size of a core file is typically about 40-50 MB.

5.11. Configuring the auditing policy on Windows

This section describes how to configure the logging and auditing policy on various versions of Microsoft Windows. The syslog-ng agent can transfer log messages only about those events that are actually logged, so the audit policy has to be configured to log the important events.

Microsoft Windows operating systems can record a range of event types, from a system-wide event such as a user logging on, to an attempt by a particular user to read a specific file. Both successful and unsuccessful attempts to perform an action can be recorded. The audit policy specifies the types of events to be audited. When such an event occurs, an entry is added to the log file of the computer.

Following is a brief overview on how to configure the audit policy on various versions of Microsoft Windows. For details, consult the documentation of your operating system, or visit Microsoft TechNet at <http://technet.microsoft.com/>. For details on configuring the auditing and logging of various applications, like the IIS Server or the ISA Server, consult your product documentation.

5.11.1. Turning on security logging on Windows XP

The following procedure describes how to enable security logging on Windows XP Professional hosts.

Procedure 5.20. Turning on security logging on Windows XP

1. Login as an administrator.
2. Click **Start**, click **Run**, and type `mmc /a`.
3. On the **File** menu, click **Add/Remove Snap-in**, and click **Add**.
4. Under **Snap-in**, click **Group Policy**, and click **Add**.
5. In **Select Group Policy Object**, select **Local Computer**, then click **Finish**, click **Close**, and click **OK**.
6. In **Console Root**, select **Local Computer Policy**, then click **Audit Policy**.
7. Right-click the attribute or event you want to audit on the details pane.
8. Set the desired options in the **Properties**.



9. Repeat Steps 7-8 for every other event you want to audit.

**Note**

To remotely enable security logging for workstations, member servers, and domain controllers, see Section 5.11.2, “Turning on security logging for domain controllers” (p. 120).

5.11.2. Turning on security logging for domain controllers

The following procedure describes how to enable security logging on a Windows XP Professional domain controller.

Procedure 5.21. Turning on security logging for domain controllers

1. Login as an administrator.
2. Click **Start**, point to **Programs**, point to **Administrative Tools**, and click **Active Directory Users and Computers**.
3. In the console tree, click **Domain Controllers**.
4. Click **Action**, then click **Properties**.
5. On the **Group Policy** tab, select the policy you want to change, and click **Edit**.
6. In the **Group Policy** window, in the console tree, click **Audit Policy**.
7. Right-click the attribute or event you want to audit on the details pane.
8. Set the desired options in the **Properties**.
9. Repeat Steps 7-8 for every other event you want to audit.

5.11.3. Turning on auditing on Windows 2003 Server

The following procedure describes how to configure auditing on a Windows 2003 Server host.

Procedure 5.22. Turning on auditing on Windows 2003 Server

1. Login as an administrator.
2. Click **Start**, point to **Programs**, point to **Administrative Tools**, and click **Domain Security Policy**.
3. In the console tree, click **Local Policies**, then **Audit Policy**.
4. Double-click on an event and select the **Define these policy settings** option.
5. Select the type of event to log: **Success** or **Failure**.
6. Repeat Steps 4-5 for every other event you want to audit.



Chapter 6. Collecting logs from IBM System i

Patrick Townsend & Associates (<http://www.patowndsend.com>) has partnered with BalaBit IT Security (the developer of syslog-ng) to bring the syslog-ng product to the System i platform. The syslog-ng PE application can be installed and run as a service directly in the Portable Application Solutions Environment (PASE) of the System i platform. Running syslog-ng in PASE allows you to transfer the logs of your server applications that are running in the PASE to a remote syslog-ng server using UDP, TCP, or SSL-encrypted TCP connections (see Section 6.9, “Configuring IBM System i Servers” (p. 132) for details). However, syslog-ng alone cannot access the native logs of the IBM System i, for that you need the syslog-ng Agent for IBM System i application.

The syslog-ng Agent for IBM System i application provides extended support for sending security, operator, server, and user log information to a syslog-ng server, or any syslogd or syslog-ng compatible server. The syslog-ng Agent for IBM System i (also called Alliance LogAgent for System i) application can help you bring your IBM System i into your Security Information Management strategy to meet regulatory compliance requirements and to properly monitor for potential security breaches.

This chapter describes how to use the syslog-ng Agent for IBM System i.

6.1. Supported sources

The syslog-ng Agent for IBM System i application can read logs from the following sources:

- *System audit journal:* The system audit journal QAUDJRN can be configured to capture a large number of security and system management events. The syslog-ng Agent for IBM System i application can capture the entries in the QAUDJRN journal in real time, format them to syslog or CEF format, and send them to a syslog server. The syslog server can be running on the System i or any remote server. The IBM System i Security Reference Manual provides information on how to change system values, user profiles, and objects to capture various system change and security events.
- *System i operator messages:* The system operator message queue QSYSOPR receives application and system messages. The syslog-ng Agent for IBM System i application can capture the messages in the QSYSOPR message queue, format them to the syslog or CEF standard, and send them to a syslog server. The syslog server can be running on the System i or any remote server.
- *User-generated logs:* The syslog-ng Agent for IBM System i application provides command and API interfaces to allow a user program to create syslog or CEF event records. User applications can provide simple text values for messages and specify the priority and facility ID for the message. Any user application can be enabled for syslog or CEF application messages.

6.2. Supported output formats

The syslog-ng Agent for IBM System i application can output messages using the standard syslog (RFC 3164) format and the ArcSight Common Event Format. The ArcSight ESM software product is a security information and event management solution that uses a special log format called the Common Event Format (CEF).

The messages can be sent to remote servers over the network using UDP, TCP, or SSL-encrypted TCP connections.



6.3. Filtering log entries

The syslog-ng Agent for IBM System i application gives you the ability to filter the system audit journal (QAUDJRN) entries that you want to send to a central log server. This can reduce the amount of network traffic and the type of events you transmit to the server. A complete list is available of log system security audit journal event types that the security administrator can edit the events to suppress transmission to the central log server.

6.4. Installing the syslog-ng Agent for IBM System i

The syslog-ng Agent for IBM System i product runs on any version of IBM OS/400 or i5/OS from V5R3 and later.

**Note**

Before you can configure and use the syslog-ng agent, you must enter a temporary or permanent license code during the installation. Please contact your software supplier to receive the license code.

6.4.1. Installing from an Internet download

When you download the Alliance product from the web site, you must unzip the product. Your software provider will supply a pass phrase for extracting the files. See the Readme.txt file in the download for instructions on how to copy the software to the System i using FTP. Once you transfer the Save file to the System i you will restore the library ALLSYL100.

6.4.2. Installing from a product CD

If you received the product on CD you can use the Load and Run (**LODRUN**) command to install the software:

Lodrun dev(opt01) dir('/')

A menu will be displayed that allows you to select the Alliance Syslog product for installation.

6.4.3. Upgrading the syslog-ng Agent for IBM System i

To upgrade the product, complete the following steps:

Procedure 6.1. Upgrading the syslog-ng Agent for IBM System i

1. End the ALLSYL100 subsystem: **Endsbs sbs(allsyl100) option(*immed)**
2. Rename the library: **Rnmobj obj(allsyl100) objtype(*lib) newobj(allsyld)**
3. Install the new version using the Internet download or Alliance product CD instructions above (see Section 6.4, "Installing the syslog-ng Agent for IBM System i" (p. 122)).
4. Use the **Apply Release Upgrade** option on the Installation menu to copy your configuration information from the old library to the new version.



6.4.3.1. Reverting to a previous version after an upgrade

You can revert to a previous version of the product by renaming the new library to a save name, and then renaming the old version to the library ALLSYL100. This will restore the previous version.

6.5. Configuring System i security auditing

The System i can log a wide variety of security events. You may wish to audit all events, or a subset of security events. Please consult the *IBM System i Security Guide* for a description of the events you can log.

To enable security auditing, use the **CHGSECAUD** command. The Change Security Audit (CHGSECAUD) command performs many of the steps to implement security audit through one command step. This command will create the journal receiver, the QAUDJRN journal, and change system values to enable security auditing. The command provides a fast way to implement system I security logging. If you use this command to start security auditing you should review the *IBM System i Security Guide* to determine if there are other security options you would like to enable. You should especially review the Change User Audit (CHGUSRAUD) command and consider logging security administrator user profiles. See Section 6.5.2, “Enabling user auditing” (p. 124) and Section 6.5.3, “Enabling object auditing” (p. 124).



Note

If you will be sending system audit journal information to syslog-ng you may wish to delete older journal receivers before starting the process. The syslog-ng Agent collects journal entries from the beginning of the current chain of journal receivers. The date of all log entries is the date of the actual journal entry, but there may be a lot of historical information that you do not want to process. You should consider making a permanent backup of system audit journals before deleting them. Use the Work With Journal Attributes (WRKJRNA) command to view the journal receivers for the QAUDJRN journal.

6.5.1. Enabling security auditing manually

If you want to manually enable security auditing instead of using the **CHGSECAUD** command, complete the following steps.

1. Create the journal receiver for the security journal. It is recommended that you create a library to contain the journal receiver, and then create the receiver using a 4 digit sequence number in the name. Issue the following commands:

```
CRTLIB LIB(AUDJRN) TEXT('AUDIT JOURNALS')
CRTJRNRCV JRNRCV(AUDJRN/AUDRCV0001)
THRESHOLD(100000) AUT(*EXCLUDE)
TEXT('Auditing Journal Receiver')
```

2. Create the journal QAUDJRN in the system library QSYS and refer to the journal receiver you created above. It is recommended that you allow the system to manage the receivers. Issue the following commands:

```
CRTJRN JRN(QSYS/QAUDJRN) +
JRNRCV(JRNLIB/AUDRCV0001) +
MNGRCV(*SYSTEM)
DLTRCV(*NO) +
```



```
AUT (*EXCLUDE)  
TEXT('Auditing Journal')
```

**Warning**

The system will not automatically delete security audit journals. You will need to periodically backup and delete old journals.

3. Change system values to enable security logging. You can now use the Work With System Values (WRKSYSVAL) command to change the QAUDLVL and QAUDLVL2 settings. These settings are used to select the security audit features and begin security logging. Please see the *IBM System i Security Guide* for a complete description of the audit options.

6.5.2. Enabling user auditing

You can use the Change User Audit (CHGUSRAUD) command to enable the logging of specific user activity. You should consider enabling user auditing for any user with special privileges such as QSECOFR and any user with *SECADM and *AUDIT capabilities.

6.5.3. Enabling object auditing

You may wish to enable specific object auditing using the Change Object Audit (CHGOBJAUD), Change DLO audit (CHGDLOAUD), and Change Audit (CHGAUD) commands. These commands can be used to enable the monitoring of specific objects on your system.

6.5.4. Configuring syslog-ng Agent for IBM System i

The syslog-ng Agent for IBM System i can be configured from a native System i configuration interface. Configuring the syslog-ng Agent application involves configuring the global options for collecting and sending syslog messages, and configuring the communications client application to talk to the syslog server.

Issue the following commands to add the ALLSYL100 library to your library list and display the main menu of the syslog-ng Agent for IBM System i:

```
ADDLIB ALLSYL100  
GO SYMAIN
```

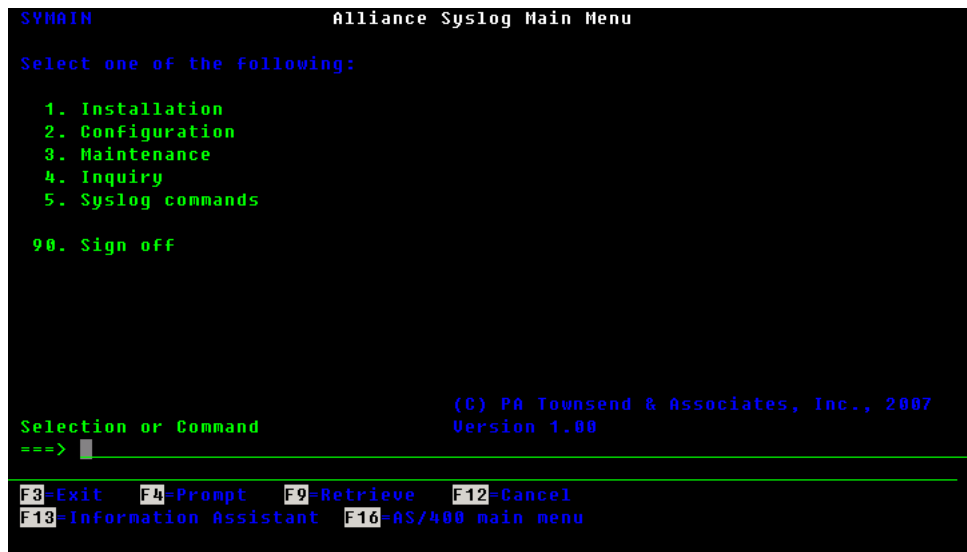


Figure 6.1. The main menu of syslog-ng Agent for IBM System i

6.5.5. Configuring Alliance Syslog for System i

Select the option for Configuration, then select the option to Configure Alliance Syslog. The following panel is displayed:

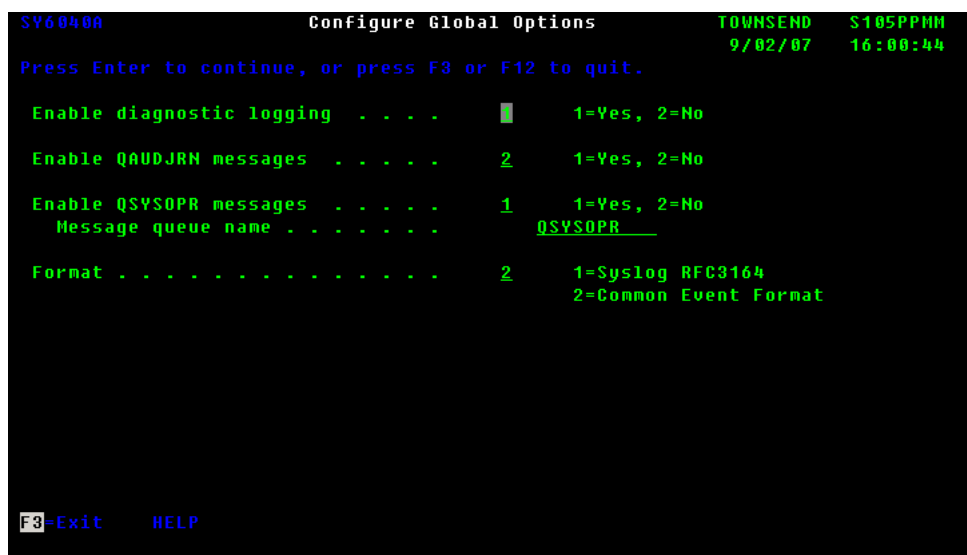


Figure 6.2. Global options of syslog-ng Agent for IBM System i

- *Enable diagnostic logging*: Enter 1 for Yes to enable diagnostic logging. When diagnostic logging is enabled the job descriptions are set for maximum job logs. Enter 2 for No to disable application logging.
- *Enable QAUDJRN messages*: Enter 1 for Yes to enable sending QAUDJRN messages to a syslog server. When enabled the system security audit journal reader job will be started in the Alliance subsystem ALLSYL100. Enter 2 for No to not send audit journal entries to the syslog server.



- *Enable QSYSOPR messages*: Enter 1 for Yes to enable sending QSYSOPR messages to a syslog server. Enter 2 for No to not send QSYSOPR messages to the syslog server.
- *Message queue name*: If you select the option to send QSYSOPR messages to a syslog server enter the name of the message queue. The default is QSYSOPR.
- *Format*: Enter option 1 to create log messages in the Syslog format (RFC 3164). Enter option 2 to create log messages in Common Event Format (CEF).

6.5.6. Configuring communication between the syslog-ng Agent and the server

The syslog-ng Agent for IBM System i can send the log messages to a syslog or syslog-ng server or relay destination. The server can be a remote server, or it can run in the PASE of the System i. To configure the destination server, start the configuration interface of the syslog-ng Agent (**GO SYMAIN**) and select **Configuration > Work With TCP Clients**. The following panel is displayed:

```
SV60400          Configure Global Options          TOWNSEND    S105PPMH
              9/02/07    16:00:44
Press Enter to continue, or press F3 or F12 to quit.

Enable diagnostic logging . . . . . 1 1=Yes, 2=No
Enable QAUDJRN messages . . . . . 2 1=Yes, 2=No
Enable QSYSOPR messages . . . . . 1 1=Yes, 2=No
Message queue name . . . . . QSYSOPR
Format . . . . . 2 1=Syslog RFC3164
                   2=Common Event Format

F3=Exit  HELP
```

Figure 6.3. Configuring communication between the syslog-ng Agent and the server

Three sample configurations are displayed:

- *SYSLOG*: Send log messages to a syslog-ng server using a standard TCP connection.
- *SYSLOGD*: Send log messages to a syslog-ng server using a standard UDP connection.
- *SYSLOGSSL*: Send log messages to a syslog-ng Premium Edition server using an TLS-encrypted connection.



Note
Only TLS encryption is supported, SSL is disabled.

- Use option 2 to change a configuration.



- Use option 3 to copy the configuration to a new definition.
- Use option 4 to delete a configuration.
- Use option 6 to print the configuration details.

When you select option 2 to change the TCP client configuration the following panel is displayed:

```
AL0380-A                               Work With Client Definitions                               9/02/07
                                                                                               16:13:43

Position to . . . .

Type options, press Enter.
  2=Change  3=Copy  4=Delete  6=Print

Opt  Client      Description
---  -
  1  SYSLOG      Syslog-ng TCP client
  2  SYSLOGD     Syslogd UDP client
  3  SYSLOGSSL   Syslog-ng SSL/TLS TCP client

F1=Help  F3=Exit  F5=Refresh  F6=Add  F12=Cancel  F21=Print list
End of file reached.
```

Figure 6.4. Configuring the destination

The following parameters can be configured



Attribute	Description
Client name	The name of this configuration.
Description	Enter a description for this configuration.
Status	Enter 1 for Active or 2 for Inactive. When the status is inactive the TCP client application will not be enabled.
Auto start client	Enter 1 for Yes to automatically start the TCP client communications when the ALLSYL100 subsystem starts. Enter 2 for No to not automatically start the TCP client. Normally you will want to automatically start the TCP client application when the subsystem starts.
Remote host name	Enter the DNS name for the syslog server. Use the IP address field if you do not have a DNS name for the server.
IP address	Enter the IP address of the syslog server if you do not have a DNS name.
Remote port number	Enter the port number for the syslog server. Consult with your network administrator for the port number. This will be the port number for the source syslog TCP service.
Application logging	Enter 1 for Yes to enable application logging. Enter 2 for No to not enable application logging. When this option is enabled detailed log records are written to the file ALL-LOGA. These log entries are not sent to the syslog server.
SSL Application ID	If this client application will use secure TLS communications enter an Application ID. You can use the IBM Digital Certificate Manager to create certificates and associated Application IDs.
SSL certification passthrough	Enter 1 for Yes to enable certificate passthrough. Enter 2 for No to not allow certificate passthrough. Enabling certificate passthrough will disable certificate validity checking, but will not allow un-secure connections.

Table 6.1. Connection parameters of syslog-ng Agent for IBM System i

6.5.7. Work with security types

Use this option to define user-created QAUDJRN journal entries. When a user application sends an entry to the security journal QAUDJRN a user-defined journal entry type is used. This is a two-character value and is different than the journal entry types that are created by i5/OS. In order to report these events you need to define them with this option and provide text and severity values.



```

SV0311-A                                Change Security Type                                9/24/07
                                                                                          07:20:40

Security type . . . . .: AD

Type information, press Enter.

Description . . . . .: Auditing change
Type . . . . .: 1                      1=System, 2=User
Security text . . . . .: Auditing change

Syslog severity . . . . .: 5              Range 0 to 7
Syslog facility . . . . .: 13             Range 0 to 23
CEF severity . . . . .: 7                Range 0 to 10
CEF signature . . . . .: 1001
Send to log server . . . . .: 1           1=Yes, 2=No

F1=Help  F3=Exit  F12=Cancel

```

Figure 6.5. Configuring security types

Attribute	Description
Description	Enter a description for this journal entry type
Type	The type indicates whether the event is a system provided event or a user defined event. This is an output field only.
Security text	Enter the text to be used with the log message. This text should be a brief description of the event type.
Syslog severity	Enter a value for the severity of this event type. The lower the value higher the severity level of the message.
Syslog facility	Enter a facility ID for this event type. See the documentation in RFC 3164 for information on facility IDs. Since the priority of an event is the result of adding the severity by the facility, the lower the facility number the higher the severity of the message.
CEF severity	If you are reporting log events in the Common Event Format enter the CEF severity level. The higher the severity number the more severe the event.
CEF signature	Enter a signature number for this event type. Alliance uses signature values from 1000 to 1999 so you should avoid signature values in this range.
Send to log server	Enter 1 for Yes to send this type of event to a log server. Enter 2 for No to suppress sending this event type to the log server. The default is Yes.

Table 6.2. Parameters of user-created journal entries

6.6. Controlling the syslog-ng Agent for IBM System i

6.6.1. Starting the Alliance subsystem

After configuring the global options and a TCP communications client, you must start the Alliance subsystem ALLSYL100 to start collecting logs. On the configuration menu take the option to **Start Syslog Subsystem**. The



following panel is displayed:

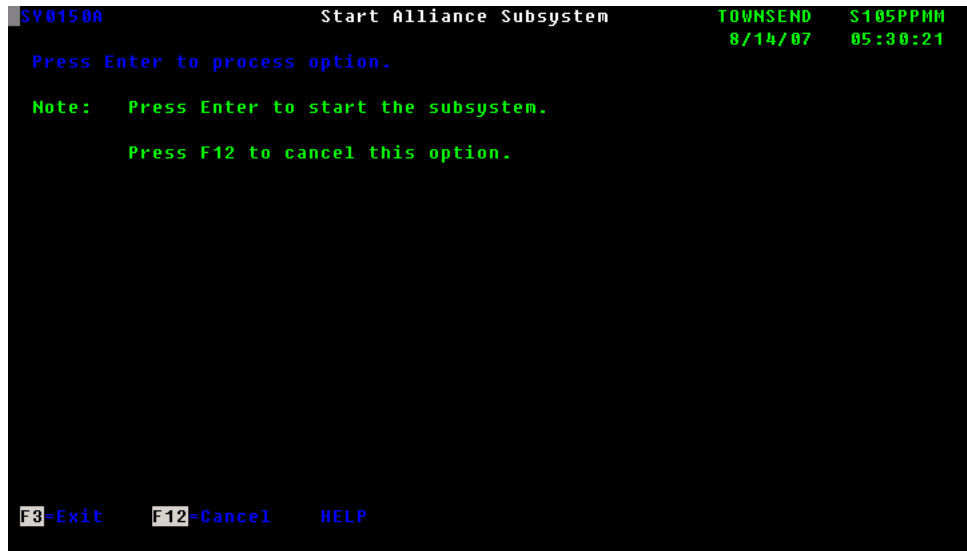


Figure 6.6. Start sending logs to the syslog-ng server

Press **Enter** to start the subsystem. Depending on the configuration options you have selected, the following jobs will appear in the subsystem:

- **OPER_MESG**: extracts messages from QSYSOPR and sends to the internal Syslog queue.
- **QAUDJRN**: extracts audit journal entries and sends to the internal Syslog queue.
- **SYSLOG (2 instances)**: receives Syslog messages from the Syslog queue and uses TCP or SSL/TLS TCP to send to a local or remote instance of Syslog-ng server.

You can use options on the Configuration menu to view active jobs in the Alliance ALLSYL100 subsystem, and to end the subsystem. You can also end the subsystem ALLSYL100 manually using the End Subsystem (**ENDSBS**) command with the ***IMMED** option.



Note

The first time you start the Alliance subsystem the audit journal and operator message queue processes will begin collecting information starting from the earliest message. If there is a substantial amount of history in the journal or message queue it may take time for these messages to be sent to the syslog-ng server.

6.6.2. Automating the start of the Alliance subsystem ALLSYL100

Once you have the configuration the way you want you can automate the start of the ALLSYL100 subsystem by modifying the IPL start up program. The name of the IPL start up program is stored in system value QSTRUPPGM. The program is usually QSTRUP in library QGPL. You can modify this program to add the following statements to start the ALLSYL100 subsystem:

```
QSYS/STRSBS SBSD(ALLSYL100/ALLSYL100)
MONMSG MSGID(CPF0000)
```



You should place these statements after any commands that start the TCP/IP network services.

If you do not have the source for the QSTRUP program you can retrieve the source using the Retrieve CL Source (**RTVCLSRC**) command.

6.7. Application maintenance

The application maintenance option can be used to purge information from the internal diagnostic logs, historical information, and to re-organize any physical files used by the syslog-ng Agent for IBM System i application. From the main menu select **Application Maintenance**, then option 1 to run maintenance. The following panel is displayed:

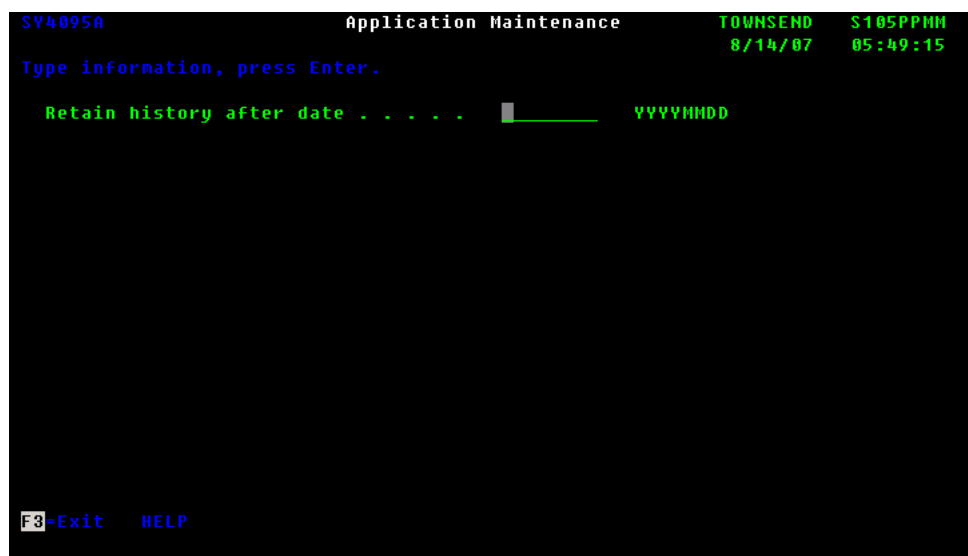


Figure 6.7. Application maintenance

Enter the date in YYYYMMDD format into the **Retain history after date** field to indicate the retention date. Log and historical information more recent than this date will be retained, older data will be deleted. All physical files will be re-organized. Note that you should run this option when the Alliance subsystem ALLSYL100 is not active.

To clear Alliance diagnostic log information manually, issue the following command:

```
clrpfm file(allsyl100/alloga)
```



Note

This command only clears the internal Alliance logs and does not delete any syslog information.

6.8. View application logs

The View Application Logs option can be used to view internal Alliance application logs such as the TCP or TLS TCP application log. From the main menu take the option for Inquiry, then option 1 to view the log. The following panel is displayed:

[illegible]

Figure 6.8. Viewing application logs

Use option 5 to view the log or option 6 to print the log. For large logs it is recommended that you use option 6 to print the log and then use the Work With Spooled Files (**WRKSPLF**) command to view the report



Note

The application logs only contain internal Alliance diagnostic information and do not contain Syslog information that has been collected.

6.9. Configuring IBM System i Servers

This section describes how to enable logging on some applications running on the System i. If you are running syslog-ng in the PASE environment of System i, you can add file sources to transfer the logs of these applications to your central syslog-ng server. For details on configuring file sources, see Section 8.1.2, “file()” (p. 141).

6.9.1. Configuring Apache server logs

To enable logging in the Apache server complete the following steps:

Procedure 6.2. Forwarding Apache server logs from System i

1. Use the Work With Links (**WRKLNK**) command to edit the `/www/(server-name)/conf/httpd_conf` file.
2. Add a `"LogCycle"` directive in order to force the Apache server to create one file:

```
LogCycle Off
CustomLog logs/access log combined
```

Without this directive the log files will have an appended time stamp and the syslog-ng application will not be able to process them.



3. Stop and re-start the Apache web server instance with the Start TCP Server (**STRTCPSVR**) command.
4. Configure a source to read the file in syslog-ng. Apache logs will generally be placed in the /www/ (server-name) /logs directory. See Chapter 3, *Configuring syslog-ng* (p. 30) for details.

6.9.2. OpenSSH server logs

To enable logging in the OpenSSH server, complete the following steps:

Procedure 6.3. Forwarding OpenSSH server logs from System i

1. Use the Work With Links (**WRKLNK**) command to edit the /QOpenSys/QIBM/ProdData/SC1/OpenSSH/openssh-3.5p1/etc/sshd_conf file.
2. Edit the file like this:

```
SyslogFacility AUTH
sLogLevel INFO
```



Note

Consult the documentation on the OpenSSH web site (<http://www.openssh.org> [<http://www.openssh.org/>]) for other syslog options.

3. Create an empty log file. Sign on as QSECOFR, use the **STRQSH** shell, and issue the following commands:

```
mkdir /var/adm
touch /var/adm/sshlog
```

4. Configure a source to read the /var/adm/sshlog file in syslog-ng. See Chapter 3, *Configuring syslog-ng* (p. 30) for details.

6.9.3. Other server logs

A number of other open systems and proprietary applications can be deployed on the IBM System i including MySQL, PHP, Perl, and others. Most of these types of applications can be enabled to collect system logs. Please consult the documentation for these servers on the steps to take to start collecting logs. Once logging is active you can configure a source statement in syslog-ng to capture the logs.

6.10. Troubleshooting the syslog-ng Agent for IBM System i

In the event you have difficulties with an Alliance Syslog application, the following procedures may be helpful.



6.10.1. System operator messages

When Alliance encounters a problem processing a Syslog transaction it may send a message to the system operator message queue. Use the **DSPMSG** command to view these messages. Many of the messages have second level text. You can use F1 or the HELP key to view this text.

6.10.2. Application logging

The Alliance TCP client applications will create extra diagnostic information when the option for application logging is enabled. You should restart the subsystem when changing the logging option. When application logging is enabled there will be additional information written to the job log and to output spooled files in the job.

6.10.3. Cannot install the product from CD

If you have a product CD with the LogAgent product, use the Load and Run (LODRUN) command to install the product: **Lodrun dev(opt01) dir('/')**

Be sure that you are signed on as QSECOFR or similar profile that has authority to restore objects from the optical CD. You can also install the product from an Internet download. Contact your software provider for information on downloading the product.

6.10.4. Logs are not being transferred to my log server

The subsystem ALLSYL100 must be started before logs will transfer. If the subsystem is not started display the main menu SYMAIN and use option 2 (Configuration), then option 10 (Start the subsystem) to start the subsystem. Use option 12 on this menu to view the active jobs in the subsystem.

You must configure at least one communications client in order to send logs to your log server. Use the Configuration menu option to configure TCP clients. Configure one TCP client (syslog, syslog-ng, or syslog-ng with SSL/TLS) and restart the subsystem ALLSYL100 to activate the client. From the Configuration menu select the option to Configure LogAgent. Be sure that you have enabled the option to send security audit journal QAUDJRN messages. After enabling the option to send the security journal messages, restart the ALLSYL100 subsystem.

6.10.5. I get a license error when trying to use configuration options

If error messages appear on your display or in the system operator message queue about a license failure, you should contact your software provider for a temporary or permanent license key. These keys are entered on the Installation menu. If you upgrade your System i operating system or hardware you may need to contact your software supplier for a new license key.

6.10.6. The product no longer works after a system upgrade

The LogAgent license key is tied to the system serial number, model number, processor group, and logical partition number. If you upgrade your System i software or hardware you may need to receive a new license key from your software provider.



6.10.7. Security events are not being captured

The IBM security audit journal (QAUDJRN) is not automatically created by the operating system. You must create the journal receivers and journal manually. For information on creating the QAUDJRN journal please see the IBM iSeries Security Reference manual. This IBM manual provides practical suggestions on creating and managing the journal.

Security events will not be captured even after creating the QAUDJRN journal until you change the security audit system values. There are multiple system values that must be enabled before journal entries are captured. See the LogAgent Reference manual and IBM iSeries Security Reference manual for information on changing the system values for audit collection.

6.10.8. I am not capturing information about our security administrators

You must enable user security journal collection using the Change User Audit (CHGUSRAUD) command in order to capture detailed information about security administrators, or other users. See the IBM iSeries Security Reference manual for information about capturing user information.

6.10.9. I am not capturing information about programs and files

If you have data in sensitive files and you want to capture information about the use of the file, use the Change Object Audit (CHGOBJAUD) command to enable information collection on a file or program. See the IBM iSeries Security Reference manual for more information about this command.

6.10.10. I am not capturing QSYSOPR messages

From the Configuration menu select the option to Configure LogAgent. Be sure that you have enabled the option to send system operator messages (QSYSOPR message queue). After enabling the option to send the operator messages, restart the ALLSYL100 subsystem.

6.10.11. I would like to turn off some audit journal events

You can filter the security audit journal events by changing the LogAgent configuration settings using the Work With Security Types option on the configuration menu. Change the option for Send To Log Server to 2 for No.

6.10.12. Where do I find error messages?

Alliance will write error messages to the system operator message queue QSYSOPR. Use the Display Message (DSPMSG) command to view these messages. For more detailed information about a communications error, you can enable application logging on the TCP client definition. When you restart the TCP client it will write verbose information to the log file ALLOGA. You can use the Inquiry menu to view and print these logs.



Chapter 7. Best practices and examples

This chapter discusses some special examples and recommendations.

7.1. General recommendations

This section provides general tips and recommendations on using syslog-ng. Some of the recommendations are detailed in the subsequent sections.

- Do not base the separation of log messages into different files on the *facility* parameter. As several applications and processes can use the same facility, the facility does not identify the application that sent the message. By default, the *facility* parameter is not even included in the log message itself. In general, sorting the log messages into several different files can make finding specific log messages difficult. If you must create separate log files, use the application name.
- Standard log messages include the local time of the sending host, without any time zone information. It is recommended to replace this timestamp with an ISODATE timestamp, because the ISODATE format includes the year and timezone as well. To convert all timestamps to the ISODATE format, include the following line in the syslog-ng configuration file:

```
options {ts_format(iso)};
```

- Resolving the IP addresses of the clients to domain names can decrease the performance of syslog-ng. See Section 7.4, “Using name resolution in syslog-ng” (p. 137) for details.

7.2. Handling lots of parallel connections

When syslog-ng is receiving messages from a large number of TCP or unix-stream connections, the CPU usage of syslog-ng might increase even if the number of messages is low. By default, syslog-ng processes every message when it is received. To reduce the CPU usage, process the incoming messages in batches. To accomplish this, instruct syslog-ng to wait for a short time before processing a message. During this period additional messages might arrive that can be processed together with the original message. To process log messages in batches, set the *time_sleep()* option (measured in milliseconds) to a non-zero value. Include the following line in your syslog-ng configuration:

```
options { time_sleep(20); };
```

**Note**

It is not recommended to increase the *time_sleep()* parameter above 100ms, as that might distort timestamps, slow down syslog-ng, and cause messages to be dropped.

When modifying the *time_sleep()* option, also adjust the *log_fetch_limit()* and *log_fifo_size()* options accordingly.



The `max_connections()` parameter limits the number of parallel connections for the source.

If adjusting the `time_sleep()` option is not desired for some reason, an alternative solution is to use `unix-stream()`, `udp()` and `unix-dgram()` sources instead of `tcp()` connections.

7.3. Handling large message load

This section provides tips on optimizing the performance of syslog-ng. Optimizing the performance is important for syslog-ng hosts that handle large traffic.

- Disable DNS resolution, or resolve hostnames locally. See Section 7.4, “Using name resolution in syslog-ng” (p. 137) for details.
- Enable flow-control for the TCP sources. See Section 2.13, “Managing incoming and outgoing messages with flow-control” (p. 18) for details.
- Do not use the `usertty()` destination driver. Under heavy load, the users are not be able to read the messages from the console, and it slows down syslog-ng.
- Do not use regular expressions in our filters. Evaluating general regular expressions puts a high load on the CPU. Use simple filter functions and logical operators instead. See Section 3.6.1, “Optimizing regular expressions in filters” (p. 60) for details.
- When receiving lots of messages using the UDP protocol, increase the size of the UDP receive buffer on the syslog-ng hosts. For information about sizing and modifying the UDP buffer, see <http://www.29west.com/docs/THPM/udp-buffer-sizing.html>.

7.4. Using name resolution in syslog-ng

The syslog-ng application can resolve the hostnames of the clients and include them in the log messages. However, the performance of syslog-ng is severely degraded if the domain name server is unaccessible or slow. Therefore, it is not recommended to resolve hostnames in syslog-ng. If you must use name resolution from syslog-ng, consider the following:

- Use DNS caching. Verify that the DNS cache is large enough to store all important hostnames. (By default, the syslog-ng DNS cache stores 1007 entries.)

```
options { dns_cache(2000); };
```

- If the IP addresses of the clients change only rarely, set the expiry of the DNS cache large.

```
options { dns_cache_expire(87600); };
```

- If possible, resolve the hostnames locally. See Section 7.4.1, “Resolving hostnames locally ” (p. 138) for details.



Note

Domain name resolution is important mainly in relay and server mode.



7.4.1. Resolving hostnames locally

Resolving hostnames locally enables you to display hostnames in the log files for frequently used hosts, without having to rely on a DNS server. The known IP address – hostname pairs are stored locally in a file. In the log messages, syslog-ng will replace the IP addresses of known hosts with their hostnames. To configure local name resolution, complete the following steps:

Procedure 7.1. Resolving hostnames locally

1. Add the hostnames and the respective IP addresses to the file used for local name resolution. On Linux and UNIX systems, this is the `/etc/hosts` file. Consult the documentation of your operating system for details.
2. Instruct syslog-ng to resolve hostnames locally. Set the `use_dns()` option of syslog-ng to `persist_only`.
3. Set the `dns_cache_hosts()` option to point to the file storing the hostnames.

```
options {  
    use_dns(persist_only);  
    dns_cache_hosts(/etc/hosts); };
```

7.5. Collecting logs from chroot

To collect logs from a chroot using a syslog-ng client running on the host, complete the following steps:




Figure 7.1. Collecting logs from chroot

Procedure 7.2. Collecting logs from chroot

1. Create a /dev directory within the chroot. The applications running in the chroot send their log messages here.
2. Create a local source in the configuration file of the syslog-ng application running outside the chroot. This source should point to the /dev/log file within the chroot (e.g., to the /chroot/dev/log directory).
3. Include the source in a log statement.



Note

You need to set up timezone information within your chroot as well. This usually means creating a symlink to /etc/localtime.

7.6. Replacing klogd on Linux

The syslog-ng application can replace both the syslogd and klogd daemons on Linux hosts. To replace klogd, complete the following steps:



Procedure 7.3. Replacing klogd on Linux

1. Add a file source pointing to `/proc/kmsg` to the syslog-ng configuration file.

```
source s_kmsg { file("/proc/kmsg"); };
```



Warning

Do not use a pipe source to read `/proc/kmsg`; pipe opens the source in read-write mode and this may cause problems when using SELinux or similar security measures.

2. Include the source defined in Step 1 in a log path.
3. Stop klogd.



Warning

Do not run klogd and syslog-ng simultaneously when using syslog-ng to read `/proc/kmsg`, as it might block syslog-ng.

7.7. A note on timezones and timestamps

If the clients run syslog-ng, then use the ISO timestamp, because it includes timezone information. That way you do not need to adjust the `recv_time_zone()` parameter of syslog-ng.

If you want syslog-ng to output timestamps in Unix (POSIX) time format, use the `S_UNIXTIME` and `R_UNIXTIME` macros. You do not need to change any of the timezone related parameters, because the timestamp information of incoming messages is converted to Unix time internally, and Unix time is a timezone-independent time representation. (Actually, Unix time measures the number of seconds elapsed since midnight of Coordinated Universal Time (UTC) January 1, 1970, but does not count leap seconds.)

7.8. Dropping messages

To skip the processing of a message without sending it to a destination, create a log statement with the appropriate filters, but do not include any destination in the statement, and use the `final` flag.



Example 7.1. Skipping messages

The following log statement drops all *debug* level messages without any further processing.

```
filter demo_debugfilter { level(debug); };
log { source(s_all); filter(demo_debugfilter); flags(final); };
```



Chapter 8. Reference

This chapter documents the drivers and options that can be used in the configuration file. For details on how to use syslog-ng, see Chapter 3, *Configuring syslog-ng* (p. 30).

8.1. Source drivers

8.1.1. `internal()`

All messages generated internally by syslog-ng use this special source. To collect warnings, errors and notices from syslog-ng itself, include this source in one of your source statements.

**Note**

Internal messages always use the local timezone of the host.

```
internal()
```

This driver does not have any parameters.

**Example 8.1. Using the `internal()` driver**

```
source s_local { internal(); };
```

8.1.2. `file()`

Collects log messages from plain-text files. The file driver has a single required parameter specifying the file to open.

Declaration:

```
file(filename);
```

In syslog-ng PE, the filename (but not the pathname) may include wildcard characters (e.g., *). Note that when using wildcards in filenames, always set how often syslog-ng should check the file for new messages using the `follow_freq()` parameter.

When using wildcards, syslog-ng PE monitors every matching file, and can receive new log messages from any of the files. However, monitoring (polling) many files (i.e., more than ten) has a significant overhead and may affect performance. On Linux this overhead is not so significant, because syslog-ng PE uses the inotify feature of the kernel.

**Note**

If the message does not have a proper syslog header, syslog-ng treats messages received from files as sent by the *kern* facility. Use the *default-facility* and *default-priority* options in the source definition to assign a different facility if needed.

The *file()* driver has the following options:

Name	Type	Default	Description
default-facility()	facility string	kern	This parameter assigns a facility value to the messages received from the file source, if the message does not specify one.
default-priority()	priority string		This parameter assigns an emergency level to the messages received from the file source, if the message does not specify one.
file	filename with path		The file to read messages from. Note that only syslog-ng PE supports wildcards in the filename (but not in the pathname). To monitor the subdirectories as well, use the <i>recursive</i> option.
encoding()	string		Specifies the character set (encoding, e.g., <i>UTF-8</i>) of messages using the legacy BSD-syslog protocol. To list the available character sets on a host, execute the iconv -l command.
flags()	empty-lines, kernel, no-multi-line, no-parse, store-legacy-msghdr, syslog-protocol, validate-utf8	empty set	<p>Specifies the log parsing options of the source.</p> <p>Use the <i>empty-lines</i> flag to keep the empty lines of the messages. By default, syslog-ng removes empty lines automatically.</p> <p>The <i>kernel</i> flag makes the source default to the <i>LOG_KERN</i> <i>LOG_CRIT</i> priority if not specified otherwise.</p> <p>The <i>no-multi-line</i> flag disables line-breaking in the messages; the entire message is converted to a single line.</p> <p>By default, syslog-ng parses incoming messages as syslog messages. If a source does not send properly formatted messages, use the <i>no-parse</i> flag to disable message parsing for the source. As a result, syslog-ng will generate a new syslog header and put the entire incoming message into the MSG part of the syslog message.</p> <p>The <i>no-parse</i> flag completely disables syslog message parsing and processes the complete line as the message part of a syslog message. Other information (timestamp, host, etc.) is added automatically. This flag is useful for parsing files not complying to the syslog format.</p> <p>If the <i>store-legacy-msghdr</i> flag is enabled, syslog-ng stores the original incoming header of the log message. This is useful if the original format of a</p>

Name	Type	Default	Description
			<p>non-syslog-compliant message must be retained (syslog-ng automatically corrects minor header errors, e.g., adds a whitespace before <i>msg</i> in the following message: <i>Jan 22 10:06:11 host program:msg</i>). Note that <i>store-legacy-msghdr</i> should be enabled when receiving messages from syslog-ng Agent for Windows clients that use the Snare-compatible mode.</p> <p>The <i>syslog-protocol</i> flag specifies that incoming messages are expected to be formatted according to the new IETF syslog protocol standard. Note that this flag is not needed for the <i>syslog</i> driver.</p> <p>The <i>validate-utf8</i> flag enables encoding-verification for messages formatted according to the new IETF syslog standard (see Section 2.18.2, “IETF-syslog messages” (p. 26) for details). If the BOM character is missing, but the message is otherwise UTF-8 compliant, syslog-ng automatically adds the BOM character to the message.</p>
follow_freq()	number	1	Indicates that the source should be checked periodically instead of being polled. This is useful for files which always indicate readability, even though no new lines were appended. If this value is higher than zero, syslog-ng will not attempt to use <i>poll()</i> on the file, but checks whether the file changed every time the <i>follow_freq()</i> interval (in seconds) has elapsed. Floating-point numbers (e.g., <i>1.5</i>) can be used as well.
keep_timestamp()	yes or no	yes	Specifies whether syslog-ng should accept the timestamp received from the sending application or client. If disabled, the time of reception will be used instead. This option can be specified globally, and per-source as well. The local setting of the source overrides the global option if available.
log_fetch_limit()	number	The value specified by the <code>global log_fetch_limit()</code> option, which defaults to <i>10</i> .	The maximum number of messages fetched from a source during a single poll loop. The destination queues might fill up before flow-control could stop reading if <i>log_fetch_limit()</i> is too high.
log_iw_size()	number	100	The size of the initial window, this value is used during flow control.

Name	Type	Default	Description
log_msg_size()	number	Use the global <i>log_msg_size()</i> option, which defaults to 8192.	Specifies the maximum length of incoming log messages. Uses the value of the <u>global option</u> if not specified.
log_prefix() (DEPRECATED)	string		A string added to the beginning of every log message. It can be used to add an arbitrary string to any log source, though it is most commonly used for adding <i>kernel:</i> to the kernel messages on Linux. NOTE: This option is deprecated. Use <i>program_override()</i> instead.
optional()	yes or no		Instruct syslog-ng to ignore the error if a specific source cannot be initialized. No other attempts to initialize the source will be made until the configuration is reloaded. This option currently applies to the <i>pipe()</i> , <i>unix-dgram</i> , and <i>unix-stream</i> drivers.
pad_size()	number	0	Specifies input padding. Some operating systems (such as HP-UX) pad all 0 messages to block boundary. This option can be used to specify the block size. (HP-UX uses 2048 bytes). Syslog-ng will pad reads from the associated device to the number of bytes set in <i>pad_size()</i> . Mostly used on HP-UX where <i>/dev/log</i> is a named pipe and every write is padded to 2048 bytes.
program_override	string		Replaces the \$PROGRAM part of the message with the parameter string. For example, to mark every message coming from the kernel, include the <i>program_override("kernel")</i> option in the source containing <i>/proc/kmsg</i> . NOTE: This option replaces the deprecated <i>log_prefix()</i> option.
recursive	yes or no	no	When enabled, syslog-ng PE monitors every subdirectory of the directory set in the path of the <i>file</i> parameter, and reads log messages from files with the set filename. The <i>recursive</i> option can be used together with wildcards in the filename.

Name	Type	Default	Description
time_zone()	timezone in the form + / - HH:MM		The default timezone for messages read from the source. Applies only if no timezone is specified within the message itself.

Table 8.1. Options of the file() sources

**Example 8.2. Using the file() driver**

```
source s_file { file("/var/log/messages");
```

**Example 8.3. Tailing files**

The following source checks the access.log file every second for new messages.

```
source s_tail { file("/var/log/apache/access.log"
                    follow_freq(1) flags(no-parse)); };
```

**Example 8.4. Using wildcards in the filename**

The following example monitors every file with the .log extension in the /var/application directory for log messages. Note that only syslog-ng PE supports wildcards in the file and pathnames.

```
source s_file { file("/var/application/*.log" follow_freq(1));
```

**Example 8.5. Monitoring multiple directories**

The following example reads files having the .log extension from the /var/application/ directory and its subdirectories. Note that only syslog-ng PE supports recursive directory handling and wildcards in the file and pathnames.

```
source s_file_subdirectories { file("/var/application/*.log"
                                recursive(yes)
                                follow_freq(1)
                                log_fetch_limit(100)
                                );};
```

8.1.3. pipe()

The pipe driver opens a named pipe with the specified name and listens for messages. It is used as the native message delivery protocol on HP-UX.

The pipe driver has a single required parameter, specifying the filename of the pipe to open.

Declaration:

```
pipe(filename);
```

**Note**

As of syslog-ng Open Source Edition 3.0.2, pipes are created automatically. In earlier versions, you had to create the pipe using the **mkfifo(1)** command.

The *pipe* driver has the following options:

Name	Type	Default	Description
flags()	empty-lines, kernel, no-multi-line, no-parse, store-legacy-msghdr, syslog-protocol, validate-utf8	empty set	<p>Specifies the log parsing options of the source.</p> <p>Use the <i>empty-lines</i> flag to keep the empty lines of the messages. By default, syslog-ng removes empty lines automatically.</p> <p>The <i>kernel</i> flag makes the source default to the <i>LOG_KERN</i> <i>LOG_CRIT</i> priority if not specified otherwise.</p> <p>The <i>no-multi-line</i> flag disables line-breaking in the messages; the entire message is converted to a single line.</p> <p>By default, syslog-ng parses incoming messages as syslog messages. If a source does not send properly formatted messages, use the <i>no-parse</i> flag to disable message parsing for the source. As a result, syslog-ng will generate a new syslog header and put the entire incoming message into the MSG part of the syslog message.</p> <p>The <i>no-parse</i> flag completely disables syslog message parsing and processes the complete line as the message part of a syslog message. Other information (timestamp, host, etc.) is added automatically. This flag is useful for parsing files not complying to the syslog format.</p> <p>If the <i>store-legacy-msghdr</i> flag is enabled, syslog-ng stores the original incoming header of the log message. This is useful if the original format of a non-syslog-compliant message must be retained (syslog-ng automatically corrects minor header errors, e.g., adds a whitespace before <i>msg</i> in the following message: <i>Jan 22 10:06:11 host program:msg</i>). Note that <i>store-legacy-msghdr</i> should be enabled when receiving messages from syslog-ng Agent for Windows clients that use the Snare-compatible mode.</p> <p>The <i>syslog-protocol</i> flag specifies that incoming messages are expected to be formatted according to the new IETF syslog protocol standard. Note that this flag is not needed for the <i>syslog</i> driver.</p> <p>The <i>validate-utf8</i> flag enables encoding-verification for messages formatted according to the new</p>

Name	Type	Default	Description
			IETF syslog standard (see Section 2.18.2, “IETF-syslog messages” (p. 26) for details). If the BOM character is missing, but the message is otherwise UTF-8 compliant, syslog-ng automatically adds the BOM character to the message.
follow_freq()	number	1	Indicates that the source should be checked periodically instead of being polled. This is useful for files which always indicate readability, even though no new lines were appended. If this value is higher than zero, syslog-ng will not attempt to use <code>poll()</code> on the file, but checks whether the file changed every time the <code>follow_freq()</code> interval (in seconds) has elapsed. Floating-point numbers (e.g., <code>1.5</code>) can be used as well.
keep_timestamp()	yes or no	yes	Specifies whether syslog-ng should accept the timestamp received from the sending application or client. If disabled, the time of reception will be used instead. This option can be specified globally, and per-source as well. The local setting of the source overrides the global option if available.
log_fetch_limit()	number	The value specified by the <code>global log_fetch_limit()</code> option, which defaults to <code>10</code> .	The maximum number of messages fetched from a source during a single poll loop. The destination queues might fill up before flow-control could stop reading if <code>log_fetch_limit()</code> is too high.
log_iw_size()	number	100	The size of the initial window, this value is used during flow control.
log_msg_size()	number	Use the global <code>log_msg_size()</code> option, which defaults to <code>8192</code> .	Specifies the maximum length of incoming log messages. Uses the value of the <code>global option</code> if not specified.
log_prefix() (DEPRECATED)	string		A string added to the beginning of every log message. It can be used to add an arbitrary string to any log source, though it is most commonly used for adding <code>kernel:</code> to the kernel messages on Linux. NOTE: This option is deprecated. Use <code>program_override()</code> instead.
optional()	yes or no		Instruct syslog-ng to ignore the error if a specific source cannot be initialized. No other attempts to initialize the source will be made until the configuration is reloaded. This option currently applies to the <code>pipe()</code> , <code>unix-dgram</code> , and <code>unix-stream</code> drivers.

Name	Type	Default	Description
pad_size()	number	0	Specifies input padding. Some operating systems (such as HP-UX) pad all 0 messages to block boundary. This option can be used to specify the block size. (HP-UX uses 2048 bytes). Syslog-ng will pad reads from the associated device to the number of bytes set in <code>pad_size()</code> . Mostly used on HP-UX where <code>/dev/log</code> is a named pipe and every write is padded to 2048 bytes.
pipe	filename with path		The filename of the pipe to read messages from.
program_override	string		Replaces the <code>\$PROGRAM</code> part of the message with the parameter string. For example, to mark every message coming from the kernel, include the <code>program_override("kernel")</code> option in the source containing <code>/proc/kmsg</code> . NOTE: This option replaces the deprecated <code>log_prefix()</code> option.
time_zone()	timezone in the form + / - HH:MM		The default timezone for messages read from the source. Applies only if no timezone is specified within the message itself.

Table 8.2. Options of the pipe() sources

**Example 8.6. Using the pipe() driver**

```
source s_pipe { pipe("/dev/pipe" pad_size(2048)); };
```

8.1.4. program()

The program driver starts an external application and reads messages from the standard output (stdout) of the application. It is mainly useful to receive log messages from daemons that accept incoming messages and convert them to log messages.

The program driver has a single required parameter, specifying the name of the application to start.

Declaration:

```
program(filename);
```

**Note**

The program is restarted automatically if it exits.



The *program* driver has the following options:

Name	Type	Default	Description
flags()	empty-lines, kernel, no-multi-line, no-parse, store-legacy-msghdr, syslog-protocol, validate-utf8	empty set	<p>Specifies the log parsing options of the source.</p> <p>Use the <i>empty-lines</i> flag to keep the empty lines of the messages. By default, syslog-ng removes empty lines automatically.</p> <p>The <i>kernel</i> flag makes the source default to the <i>LOG_KERN</i> <i>LOG_CRIT</i> priority if not specified otherwise.</p> <p>The <i>no-multi-line</i> flag disables line-breaking in the messages; the entire message is converted to a single line.</p> <p>By default, syslog-ng parses incoming messages as syslog messages. If a source does not send properly formatted messages, use the <i>no-parse</i> flag to disable message parsing for the source. As a result, syslog-ng will generate a new syslog header and put the entire incoming message into the MSG part of the syslog message.</p> <p>The <i>no-parse</i> flag completely disables syslog message parsing and processes the complete line as the message part of a syslog message. Other information (timestamp, host, etc.) is added automatically. This flag is useful for parsing files not complying to the syslog format.</p> <p>If the <i>store-legacy-msghdr</i> flag is enabled, syslog-ng stores the original incoming header of the log message. This is useful if the original format of a non-syslog-compliant message must be retained (syslog-ng automatically corrects minor header errors, e.g., adds a whitespace before <i>msg</i> in the following message: <i>Jan 22 10:06:11 host program:msg</i>). Note that <i>store-legacy-msghdr</i> should be enabled when receiving messages from syslog-ng Agent for Windows clients that use the Snare-compatible mode.</p> <p>The <i>syslog-protocol</i> flag specifies that incoming messages are expected to be formatted according to the new IETF syslog protocol standard. Note that this flag is not needed for the <i>syslog</i> driver.</p> <p>The <i>validate-utf8</i> flag enables encoding-verification for messages formatted according to the new</p>

Name	Type	Default	Description
			IETF syslog standard (see Section 2.18.2, “IETF-syslog messages” (p. 26) for details). If the BOM character is missing, but the message is otherwise UTF-8 compliant, syslog-ng automatically adds the BOM character to the message.
follow_freq()	number	1	Indicates that the source should be checked periodically instead of being polled. This is useful for files which always indicate readability, even though no new lines were appended. If this value is higher than zero, syslog-ng will not attempt to use <code>poll()</code> on the file, but checks whether the file changed every time the <code>follow_freq()</code> interval (in seconds) has elapsed. Floating-point numbers (e.g., <code>1.5</code>) can be used as well.
keep_timestamp()	yes or no	yes	Specifies whether syslog-ng should accept the timestamp received from the sending application or client. If disabled, the time of reception will be used instead. This option can be specified globally, and per-source as well. The local setting of the source overrides the global option if available.
log_fetch_limit()	number	The value specified by the <code>global log_fetch_limit()</code> option, which defaults to <code>10</code> .	The maximum number of messages fetched from a source during a single poll loop. The destination queues might fill up before flow-control could stop reading if <code>log_fetch_limit()</code> is too high.
log_iw_size()	number	100	The size of the initial window, this value is used during flow control.
log_msg_size()	number	Use the global <code>log_msg_size()</code> option, which defaults to <code>8192</code> .	Specifies the maximum length of incoming log messages. Uses the value of the <code>global option</code> if not specified.
log_prefix() (DEPRECATED)	string		A string added to the beginning of every log message. It can be used to add an arbitrary string to any log source, though it is most commonly used for adding <code>kernel:</code> to the kernel messages on Linux. NOTE: This option is deprecated. Use <code>program_override()</code> instead.
optional()	yes or no		Instruct syslog-ng to ignore the error if a specific source cannot be initialized. No other attempts to initialize the source will be made until the configuration is reloaded. This option currently applies to the <code>pipe()</code> , <code>unix-dgram</code> , and <code>unix-stream</code> drivers.



Name	Type	Default	Description
pad_size()	number	0	Specifies input padding. Some operating systems (such as HP-UX) pad all 0 messages to block boundary. This option can be used to specify the block size. (HP-UX uses 2048 bytes). Syslog-ng will pad reads from the associated device to the number of bytes set in <code>pad_size()</code> . Mostly used on HP-UX where <code>/dev/log</code> is a named pipe and every write is padded to 2048 bytes.
program	file name with path		The name of the application to start and read messages from.
program_override	string		Replaces the <code>\$PROGRAM</code> part of the message with the parameter string. For example, to mark every message coming from the kernel, include the <code>program_override("kernel")</code> option in the source containing <code>/proc/kmsg</code> . NOTE: This option replaces the deprecated <code>log_prefix()</code> option.
time_zone()	timezone in the form + / - HH:MM		The default timezone for messages read from the source. Applies only if no timezone is specified within the message itself.

Table 8.3. Options of the `program()` source**Example 8.7. Using the `program()` driver**

```
source s_program { program("/etc/init.d/mydaemon"); };
```

8.1.5. sun-streams() driver

Solaris uses its *STREAMS* framework to send messages to the *syslogd* process.

Newer versions of Solaris (2.5.1 and above), use a new IPC in addition to *STREAMS*, called *door* to confirm the delivery of a message. The syslog-ng application supports this new IPC mechanism via the `door()` option (see below).

**Note**

The `sun-streams()` driver must be enabled when the syslog-ng application is compiled (see `./configure --help`).

The `sun-streams()` driver has a single required argument specifying the *STREAMS* device to open, and the `door()` option.



Declaration:

```
sun-streams(name_of_the_streams_device door(filename_of_the_door));
```



Name	Type	Default	Description
door()	string	none	Specifies the filename of a door to open, needed on Solaris above 2.5.1.



Name	Type	Default	Description
flags()	empty-lines, kernel, no-multi-line, no-parse, store-legacy-msghdr, syslog-protocol, validate-utf8	empty set	



Name	Type	Default	Description
			<p>Specifies the log parsing options of the source.</p> <p>Use the <i>empty-lines</i> flag to keep the empty lines of the messages. By default, syslog-ng removes empty lines automatically.</p> <p>The <i>kernel</i> flag makes the source default to the <code>LOG_KERN LOG_CRIT</code> priority if not specified otherwise.</p> <p>The <i>no-multi-line</i> flag disables line-breaking in the messages; the entire message is converted to a single line.</p> <p>By default, syslog-ng parses incoming messages as syslog messages. If a source does not send properly formatted messages, use the <i>no-parse</i> flag to disable message parsing for the source. As a result, syslog-ng will generate a new syslog header and put the entire incoming message into the MSG part of the syslog message.</p> <p>The <i>no-parse</i> flag completely disables syslog message parsing and processes the complete line as the message part of a syslog message. Other information (timestamp, host, etc.) is added automatically. This flag is useful for parsing files not complying to the syslog format.</p> <p>If the <i>store-legacy-msghdr</i> flag is enabled, syslog-ng stores the original incoming header of the log message. This is useful if the original format of a non-syslog-compliant message must be retained (syslog-ng automatically corrects minor header errors, e.g., adds a whitespace before <i>msg</i> in the following message: <i>Jan 22 10:06:11 host program:msg</i>). Note that <i>store-legacy-msghdr</i> should be enabled when receiving messages from syslog-ng Agent for Windows clients that use the Snare-compatible mode.</p> <p>The <i>syslog-protocol</i> flag specifies that incoming messages are expected to be formatted according to the new IETF syslog protocol standard. Note that this flag is not needed for the <i>syslog</i> driver.</p> <p>The <i>validate-utf8</i> flag enables encoding-verification for messages formatted according to the new IETF syslog standard (see Section 2.18.2, “IETF-syslog messages” (p. 26) for details). If the BOM character is</p>

Name	Type	Default	Description
			missing, but the message is otherwise UTF-8 compliant, syslog-ng automatically adds the BOM character to the message.
follow_freq()	number	1	Indicates that the source should be checked periodically instead of being polled. This is useful for files which always indicate readability, even though no new lines were appended. If this value is higher than zero, syslog-ng will not attempt to use <i>poll()</i> on the file, but checks whether the file changed every time the <i>follow_freq()</i> interval (in seconds) has elapsed. Floating-point numbers (e.g., <i>1.5</i>) can be used as well.
keep_timestamp()	yes or no	yes	Specifies whether syslog-ng should accept the timestamp received from the sending application or client. If disabled, the time of reception will be used instead. This option can be specified globally, and per-source as well. The local setting of the source overrides the global option if available.
log_fetch_limit()	number	The value specified by the <code>global log_fetch_limit()</code> option, which defaults to <i>10</i> .	The maximum number of messages fetched from a source during a single poll loop. The destination queues might fill up before flow-control could stop reading if <i>log_fetch_limit()</i> is too high.
log_iw_size()	number	100	The size of the initial window, this value is used during flow control.
log_msg_size()	number	Use the global <i>log_msg_size()</i> option, which defaults to <i>8192</i> .	Specifies the maximum length of incoming log messages. Uses the value of the <u>global option</u> if not specified.
log_prefix() (DEPRECATED)	string		A string added to the beginning of every log message. It can be used to add an arbitrary string to any log source, though it is most commonly used for adding <i>kernel:</i> to the kernel messages on Linux. NOTE: This option is deprecated. Use <i>program_override()</i> instead.
optional()	yes or no		Instruct syslog-ng to ignore the error if a specific source cannot be initialized. No other attempts to initialize the source will be made until the configuration is reloaded. This option currently applies to the <i>pipe()</i> , <i>unix-dgram</i> , and <i>unix-stream</i> drivers.

Name	Type	Default	Description
<code>pad_size()</code>	number	0	Specifies input padding. Some operating systems (such as HP-UX) pad all 0 messages to block boundary. This option can be used to specify the block size. (HP-UX uses 2048 bytes). Syslog-ng will pad reads from the associated device to the number of bytes set in <code>pad_size()</code> . Mostly used on HP-UX where <code>/dev/log</code> is a named pipe and every write is padded to 2048 bytes.
<code>program_override</code>	string		Replaces the <code>\$PROGRAM</code> part of the message with the parameter string. For example, to mark every message coming from the kernel, include the <code>program_override("kernel")</code> option in the source containing <code>/proc/kmsg</code> . NOTE: This option replaces the deprecated <code>log_prefix()</code> option.
<code>time_zone()</code>	timezone in the form +/-HH:MM		The default timezone for messages read from the source. Applies only if no timezone is specified within the message itself.

Table 8.4. Options for *sun-streams***Example 8.8. Using the `sun-streams()` driver**

```
source s_stream { sun-streams("/dev/log" door("/etc/.syslog_door")); };
```

8.1.6. `syslog()`

This driver enables to receive messages from the network using the new standard syslog protocol and message format (see Section 2.18.2, “IETF-syslog messages” (p. 26) for details about the protocol). UDP, TCP, and TLS-encrypted TCP can all be used to transport the messages.

Declaration:

```
syslog(ip() port() transport() options());
```

Name	Type	Default	Description
flags()	<code>empty-lines</code> , <code>kernel</code> , <code>no-multi-line</code> , <code>no-parse</code> , <code>store-legacy-msghdr</code> , <code>syslog-protocol</code> , <code>validate-utf8</code>	empty set	<p>Specifies the log parsing options of the source.</p> <p>Use the <i>empty-lines</i> flag to keep the empty lines of the messages. By default, syslog-ng removes empty lines automatically.</p> <p>The <i>kernel</i> flag makes the source default to the <code>LOG_KERN</code> <code>LOG_CRIT</code> priority if not specified otherwise.</p> <p>The <i>no-multi-line</i> flag disables line-breaking in the messages; the entire message is converted to a single line.</p> <p>By default, syslog-ng parses incoming messages as syslog messages. If a source does not send properly formatted messages, use the <i>no-parse</i> flag to disable message parsing for the source. As a result, syslog-ng will generate a new syslog header and put the entire incoming message into the MSG part of the syslog message.</p> <p>The <i>no-parse</i> flag completely disables syslog message parsing and processes the complete line as the message part of a syslog message. Other information (timestamp, host, etc.) is added automatically. This flag is useful for parsing files not complying to the syslog format.</p> <p>If the <i>store-legacy-msghdr</i> flag is enabled, syslog-ng stores the original incoming header of the log message. This is useful if the original format of a non-syslog-compliant message must be retained (syslog-ng automatically corrects minor header errors, e.g., adds a whitespace before <i>msg</i> in the following message: <i>Jan 22 10:06:11 host program:msg</i>). Note that <i>store-legacy-msghdr</i> should be enabled when receiving messages from syslog-ng Agent for Windows clients that use the Snare-compatible mode.</p> <p>The <i>syslog-protocol</i> flag specifies that incoming messages are expected to be formatted according to the new IETF syslog protocol standard. Note that this flag is not needed for the <i>syslog</i> driver.</p> <p>The <i>validate-utf8</i> flag enables encoding-verification for messages formatted according to the new</p>

Name	Type	Default	Description
			IETF syslog standard (see Section 2.18.2, “IETF-syslog messages” (p. 26) for details). If the BOM character is missing, but the message is otherwise UTF-8 compliant, syslog-ng automatically adds the BOM character to the message.
follow_freq()	number	1	Indicates that the source should be checked periodically instead of being polled. This is useful for files which always indicate readability, even though no new lines were appended. If this value is higher than zero, syslog-ng will not attempt to use <i>poll()</i> on the file, but checks whether the file changed every time the <i>follow_freq()</i> interval (in seconds) has elapsed. Floating-point numbers (e.g., <i>1.5</i>) can be used as well.
host_override()	string		Replaces the \$HOST part of the message with the parameter string.
ip() or localip()	string	0.0.0.0	The IP address to bind to. Note that this is not the address where messages are accepted from.
ip_tos()	number	0	Specifies the Type-of-Service value of outgoing packets.
ip_ttl()	number	0	Specifies the Time-To-Live value of outgoing packets.
keep-alive()	yes or no	yes	Specifies whether connections to sources should be closed when syslog-ng is restarted (upon the receipt of a SIGHUP signal). Note that this applies to the server (source) side of the syslog-ng connections, client-side (destination) connections are always reopened after receiving a HUP signal unless the <i>keep-alive</i> option is enabled for the destination.
keep_hostname()	yes or no	no	Enable or disable hostname rewriting. Enable this option to use hostname-related macros. This option can be specified globally, and per-source as well. The local setting of the source overrides the global option if available. When relaying messages, enable this option on the syslog-ng server and also on every relay, otherwise syslog-ng will treat incoming messages as if they were sent by the last relay.
keep_timestamp()	yes or no	yes	Specifies whether syslog-ng should accept the timestamp received from the sending application or client. If disabled, the time of reception will be used instead. This option can be specified globally, and per-source as well. The local setting of the source overrides the global option if available.

Name	Type	Default	Description
log_fetch_limit()	number	The value specified by the <code>global log_fetch_limit()</code> option, which defaults to 10.	The maximum number of messages fetched from a source during a single poll loop. The destination queues might fill up before flow-control could stop reading if <code>log_fetch_limit()</code> is too high.
log_iw_size()	number	100	The size of the initial window, this value is used during flow control.
log_msg_size()	number	Use the global <code>log_msg_size()</code> option, which defaults to 8192.	Specifies the maximum length of incoming log messages. Uses the value of the <code>global option</code> if not specified.
log_prefix() (DEPRECATED)	string		A string added to the beginning of every log message. It can be used to add an arbitrary string to any log source, though it is most commonly used for adding <code>kernel:</code> to the kernel messages on Linux. NOTE: This option is deprecated. Use <code>program_override()</code> instead.
max-connections()	number	10	Specifies the maximum number of simultaneous connections.
optional()	yes or no		Instruct syslog-ng to ignore the error if a specific source cannot be initialized. No other attempts to initialize the source will be made until the configuration is reloaded. This option currently applies to the <code>pipe()</code> , <code>unix-dgram</code> , and <code>unix-stream</code> drivers.
pad_size()	number	0	Specifies input padding. Some operating systems (such as HP-UX) pad all 0 messages to block boundary. This option can be used to specify the block size. (HP-UX uses 2048 bytes). Syslog-ng will pad reads from the associated device to the number of bytes set in <code>pad_size()</code> . Mostly used on HP-UX where <code>/dev/log</code> is a named pipe and every write is padded to 2048 bytes.
port() or local-port()	number	514	The port number to bind to.
program_override	string		Replaces the <code>\$PROGRAM</code> part of the message with the parameter string. For example, to mark every message coming from the kernel, include the <code>program_override("kernel")</code> option in the source containing <code>/proc/kmsg</code> . NOTE: This option replaces the deprecated <code>log_prefix()</code> option.

Name	Type	Default	Description
<code>so_broadcast()</code>	yes or no	no	This option controls the <code>SO_BROADCAST</code> socket option required to make syslog-ng send messages to a broadcast address. See the socket(7) manual page for details.
<code>so_rcvbuf()</code>	number	0	Specifies the size of the socket receive buffer in bytes. See the socket(7) manual page for details.
<code>so_sndbuf()</code>	number	0	Specifies the size of the socket send buffer in bytes. See the socket(7) manual page for details.
<code>so_broadcast()</code>	yes or no	no	This option controls the <code>SO_BROADCAST</code> socket option required to make syslog-ng send messages to a broadcast address. See the socket(7) manual page for details.
<code>so_keepalive()</code>	yes or no	no	Enables keep-alive messages, keeping the socket open. This only effects TCP and UNIX-stream sockets. See the socket(7) manual page for details.
<code>tcp-keep-alive()</code>	yes or no	no	This is an obsolete alias of the <code>so_keepalive()</code> option.
<code>time_zone()</code>	timezone in the form + / - HH:MM		The default timezone for messages read from the source. Applies only if no timezone is specified within the message itself.
<code>transport</code>	udp, tcp, or tls	tcp	Specifies the protocol used to receive messages from the source.
<code>tls()</code>	tls options	n/a	This option sets various TLS specific options like key/certificate files and trusted CA locations and can only be used with the <code>tcp</code> transport protocols. See Section 8.10, “TLS options” (p. 231) for more information.
<code>use_dns()</code>	yes, no, persist_only	yes	Enable or disable DNS usage. The <code>persist_only</code> option attempts to resolve hostnames locally from file (e.g., from <code>/etc/hosts</code>). syslog-ng blocks on DNS queries, so enabling DNS may lead to a Denial of Service attack. To prevent DoS, protect your syslog-ng network endpoint with firewall rules, and make sure that all hosts which may get to syslog-ng are resolvable. This option can be specified globally, and per-source as well. The local setting of the source overrides the global option if available.



Name	Type	Default	Description
use_fqdn()	yes or no	no	Add Fully Qualified Domain Name instead of short hostname. This option can be specified globally, and per-source as well. The local setting of the source overrides the global option if available.

Table 8.5. Options for syslog() sources

**Example 8.9. Using the syslog() driver**

TCP source listening on the localhost on port 1999.

```
source s_syslog { syslog(ip(127.0.0.1) port(1999) transport("tcp")); };
```

UDP source with defaults.

```
source s_udp { syslog( transport("udp")); };
```

Encrypted source where the client is also authenticated. See Section 8.10, “TLS options” (p. 231) for details on the encryption settings.

```
source s_syslog_tls{ syslog(
    ip(10.100.20.40)
    transport("tls")
    tls(
        peer-verify(required-trusted)
        ca_dir('/opt/syslog-ng/etc/syslog-ng/keys/ca.d/')
key_file('/opt/syslog-ng/etc/syslog-ng/keys/server_privatekey.pem')
cert_file('/opt/syslog-ng/etc/syslog-ng/keys/server_certificate.pem')
    )
);};
```

8.1.7. tcp(), tcp6(), udp() and udp6()

The `tcp()`, `tcp6()`, `udp()`, `udp6()` drivers can receive messages from the network using the TCP and UDP networking protocols. The `tcp6()` and `udp6()` drivers use the IPv6 network protocol, while `tcp()` and `udp()` use IPv4.

The `tcp()` and `udp()` drivers do not have any required parameters. By default they bind to `0.0.0.0:514`, which means that syslog-ng will listen on all available interfaces, port 514. To limit accepted connections to only one interface, use the `localip()` parameter as described below.

**Note**

The tcp port 514 is reserved for use with **rshell**, so select a different port if syslog-ng and **rshell** is used at the same time.

If you specify a multicast bind address to `udp()` and `udp6()`, syslog-ng will automatically join the necessary multicast group. TCP does not support multicasting.



The syslog-ng Premium Edition application supports TLS (Transport Layer Security, also known as SSL) for the `tcp()` and `tcp6()` drivers. See the TLS-specific options below and Section 3.13, “Encrypting log messages with TLS” (p. 68) for details.

Declaration:

```
tcp([options]);  
udp([options]);
```

The following options are valid for `tcp()`, `tcp6()`, `udp()`, and `udp6()` drivers:



Name	Type	Default	Description
encoding()	string		Specifies the character set (encoding, e.g., <i>UTF-8</i>) of messages using the legacy BSD-syslog protocol. To list the available character sets on a host, execute the iconv -I command.



Name	Type	Default	Description
flags()	empty - lines, ker- nel, no- multi-line, no-parse, store-leg- acy-msgh- dr, syslog- protocol, validate- utf8	empty set	

Name	Type	Default	Description
			<p>Specifies the log parsing options of the source.</p> <p>Use the <i>empty-lines</i> flag to keep the empty lines of the messages. By default, syslog-ng removes empty lines automatically.</p> <p>The <i>kernel</i> flag makes the source default to the <i>LOG_KERN</i> <i>LOG_CRIT</i> priority if not specified otherwise.</p> <p>The <i>no-multi-line</i> flag disables line-breaking in the messages; the entire message is converted to a single line.</p> <p>By default, syslog-ng parses incoming messages as syslog messages. If a source does not send properly formatted messages, use the <i>no-parse</i> flag to disable message parsing for the source. As a result, syslog-ng will generate a new syslog header and put the entire incoming message into the MSG part of the syslog message.</p> <p>The <i>no-parse</i> flag completely disables syslog message parsing and processes the complete line as the message part of a syslog message. Other information (timestamp, host, etc.) is added automatically. This flag is useful for parsing files not complying to the syslog format.</p> <p>If the <i>store-legacy-msghdr</i> flag is enabled, syslog-ng stores the original incoming header of the log message. This is useful if the original format of a non-syslog-compliant message must be retained (syslog-ng automatically corrects minor header errors, e.g., adds a whitespace before <i>msg</i> in the following message: <i>Jan 22 10:06:11 host program:msg</i>). Note that <i>store-legacy-msghdr</i> should be enabled when receiving messages from syslog-ng Agent for Windows clients that use the Snare-compatible mode.</p> <p>The <i>syslog-protocol</i> flag specifies that incoming messages are expected to be formatted according to the new IETF syslog protocol standard. Note that this flag is not needed for the <i>syslog</i> driver.</p> <p>The <i>validate-utf8</i> flag enables encoding-verification for messages formatted according to the new</p>

Name	Type	Default	Description
			IETF syslog standard (see Section 2.18.2, “IETF-syslog messages” (p. 26) for details). If the BOM character is missing, but the message is otherwise UTF-8 compliant, syslog-ng automatically adds the BOM character to the message.
follow_freq()	number	1	Indicates that the source should be checked periodically instead of being polled. This is useful for files which always indicate readability, even though no new lines were appended. If this value is higher than zero, syslog-ng will not attempt to use <i>poll()</i> on the file, but checks whether the file changed every time the <i>follow_freq()</i> interval (in seconds) has elapsed. Floating-point numbers (e.g., <i>1.5</i>) can be used as well.
host_override()	string		Replaces the \$HOST part of the message with the parameter string.
ip() or localip()	string	0.0.0.0	The IP address to bind to. Note that this is not the address where messages are accepted from.
ip_tos()	number	0	Specifies the Type-of-Service value of outgoing packets.
ip_ttl()	number	0	Specifies the Time-To-Live value of outgoing packets.
keep-alive()	yes or no	yes	Specifies whether connections to sources should be closed when syslog-ng is restarted (upon the receipt of a SIGHUP signal). Note that this applies to the server (source) side of the syslog-ng connections, client-side (destination) connections are always reopened after receiving a HUP signal unless the <i>keep-alive</i> option is enabled for the destination.
keep_hostname()	yes or no	no	Enable or disable hostname rewriting. Enable this option to use hostname-related macros. This option can be specified globally, and per-source as well. The local setting of the source overrides the global option if available. When relaying messages, enable this option on the syslog-ng server and also on every relay, otherwise syslog-ng will treat incoming messages as if they were sent by the last relay.
keep_timestamp()	yes or no	yes	Specifies whether syslog-ng should accept the timestamp received from the sending application or client. If disabled, the time of reception will be used instead. This option can be specified globally, and per-source as well. The local setting of the source overrides the global option if available.



Name	Type	Default	Description
log_fetch_limit()	number	The value specified by the global <code>log_fetch_limit()</code> option, which defaults to 10.	The maximum number of messages fetched from a source during a single poll loop. The destination queues might fill up before flow-control could stop reading if <code>log_fetch_limit()</code> is too high.
log_iw_size()	number	100	The size of the initial window, this value is used during flow control.
log_msg_size()	number	Use the global <code>log_msg_size()</code> option, which defaults to 8192.	Specifies the maximum length of incoming log messages. Uses the value of the <u>global option</u> if not specified.
log_prefix() (DEPRECATED)	string		A string added to the beginning of every log message. It can be used to add an arbitrary string to any log source, though it is most commonly used for adding <i>kernel:</i> to the kernel messages on Linux. NOTE: This option is deprecated. Use <code>program_override()</code> instead.
max-connections()	number	10	Specifies the maximum number of simultaneous connections.
optional()	yes or no		Instruct syslog-ng to ignore the error if a specific source cannot be initialized. No other attempts to initialize the source will be made until the configuration is reloaded. This option currently applies to the <code>pipe()</code> , <code>unix-dgram</code> , and <code>unix-stream</code> drivers.
pad_size()	number	0	Specifies input padding. Some operating systems (such as HP-UX) pad all 0 messages to block boundary. This option can be used to specify the block size. (HP-UX uses 2048 bytes). Syslog-ng will pad reads from the associated device to the number of bytes set in <code>pad_size()</code> . Mostly used on HP-UX where <code>/dev/log</code> is a named pipe and every write is padded to 2048 bytes.
port() or local-port()	number	514	The port number to bind to.
program_override	string		Replaces the \$PROGRAM part of the message with the parameter string. For example, to mark every message coming from the kernel, include the <code>program_override("kernel")</code> option in the source containing <code>/proc/kmsg</code> . NOTE: This option replaces the deprecated <code>log_prefix()</code> option.

Name	Type	Default	Description
so_broadcast()	yes or no	no	This option controls the <i>SO_BROADCAST</i> socket option required to make syslog-ng send messages to a broadcast address. See the socket(7) manual page for details.
so_keepalive()	yes or no	no	Enables keep-alive messages, keeping the socket open. This only effects TCP and UNIX-stream sockets. See the socket(7) manual page for details.
so_rcvbuf()	number	0	Specifies the size of the socket receive buffer in bytes. See the socket(7) manual page for details.
so_sndbuf()	number	0	Specifies the size of the socket send buffer in bytes. See the socket(7) manual page for details.
tcp-keep-alive()	yes or no	no	This is an obsolete alias of the <i>so_keepalive()</i> option.
time_zone()	timezone in the form + / - HH:MM		The default timezone for messages read from the source. Applies only if no timezone is specified within the message itself.
tls()	tls options	n/a	This option sets various TLS specific options like key/certificate files and trusted CA locations and can only be used with the <i>tcp</i> transport protocols. See Section 8.10, “TLS options” (p. 231) for more information.
use_dns()	yes, no, p e r - sist_only	yes	Enable or disable DNS usage. The <i>persist_only</i> option attempts to resolve hostnames locally from file (e.g., from <i>/etc/hosts</i>). syslog-ng blocks on DNS queries, so enabling DNS may lead to a Denial of Service attack. To prevent DoS, protect your syslog-ng network endpoint with firewall rules, and make sure that all hosts which may get to syslog-ng are resolvable. This option can be specified globally, and per-source as well. The local setting of the source overrides the global option if available.

Name	Type	Default	Description
use_fqdn()	yes or no	no	Add Fully Qualified Domain Name instead of short hostname. This option can be specified globally, and per-source as well. The local setting of the source overrides the global option if available.

Table 8.6. Options for tcp, tcp6, udp, and udp6 drivers

**Example 8.10. Using the udp() and tcp() drivers**

A simple udp() source with default settings.

```
source s_udp { udp(); };# An UDP source with default settings.
```

A TCP source listening on the localhost interface, with a limited number of connections allowed.

```
source s_tcp { tcp(ip(127.0.0.1) port(1999) max-connections(10));};
```

A TCP source listening on a TLS-encrypted channel.

```
source s_tcp { tcp(ip(127.0.0.1) port(1999)
    tls(peer-verify('required-trusted')
        key_file('/opt/syslog-ng/etc/syslog-ng/syslog-ng.key')
        cert_file('/opt/syslog-ng/etc/syslog-ng/syslog-ng.crt')));
};
```

A TCP source listening for messages using the IETF-syslog message format:

```
source s_tcp_syslog { tcp(ip(127.0.0.1) port(1999) flags(syslog-protocol));};
```

8.1.8. unix-stream() and unix-dgram()

These two drivers behave similarly: they open an `AF_UNIX` socket and start listening on it for messages.

Both unix-stream and unix-dgram have a single required argument, specifying the filename of the socket to create.

Declaration:

```
unix-stream(filename [options]);
unix-dgram(filename [options]);
```

The following options can be specified for these drivers:



Name	Type	Default	Description
encoding()	string		Specifies the character set (encoding, e.g., <i>UTF-8</i>) of messages using the legacy BSD-syslog protocol. To list the available character sets on a host, execute the iconv -l command.



Name	Type	Default	Description
flags()	empty - lines, ker- nel, no- multi-line, no-parse, store-leg- acy-msgh- dr, syslog- protocol, validate- utf8	empty set	

Name	Type	Default	Description
			<p>Specifies the log parsing options of the source.</p> <p>Use the <i>empty-lines</i> flag to keep the empty lines of the messages. By default, syslog-ng removes empty lines automatically.</p> <p>The <i>kernel</i> flag makes the source default to the <i>LOG_KERN</i> <i>LOG_CRIT</i> priority if not specified otherwise.</p> <p>The <i>no-multi-line</i> flag disables line-breaking in the messages; the entire message is converted to a single line.</p> <p>By default, syslog-ng parses incoming messages as syslog messages. If a source does not send properly formatted messages, use the <i>no-parse</i> flag to disable message parsing for the source. As a result, syslog-ng will generate a new syslog header and put the entire incoming message into the MSG part of the syslog message.</p> <p>The <i>no-parse</i> flag completely disables syslog message parsing and processes the complete line as the message part of a syslog message. Other information (timestamp, host, etc.) is added automatically. This flag is useful for parsing files not complying to the syslog format.</p> <p>If the <i>store-legacy-msghdr</i> flag is enabled, syslog-ng stores the original incoming header of the log message. This is useful if the original format of a non-syslog-compliant message must be retained (syslog-ng automatically corrects minor header errors, e.g., adds a whitespace before <i>msg</i> in the following message: <i>Jan 22 10:06:11 host program:msg</i>). Note that <i>store-legacy-msghdr</i> should be enabled when receiving messages from syslog-ng Agent for Windows clients that use the Snare-compatible mode.</p> <p>The <i>syslog-protocol</i> flag specifies that incoming messages are expected to be formatted according to the new IETF syslog protocol standard. Note that this flag is not needed for the <i>syslog</i> driver.</p> <p>The <i>validate-utf8</i> flag enables encoding-verification for messages formatted according to the new</p>



Name	Type	Default	Description
			IETF syslog standard (see Section 2.18.2, “IETF-syslog messages” (p. 26) for details). If the BOM character is missing, but the message is otherwise UTF-8 compliant, syslog-ng automatically adds the BOM character to the message.
follow_freq()	number	1	Indicates that the source should be checked periodically instead of being polled. This is useful for files which always indicate readability, even though no new lines were appended. If this value is higher than zero, syslog-ng will not attempt to use <i>poll()</i> on the file, but checks whether the file changed every time the <i>follow_freq()</i> interval (in seconds) has elapsed. Floating-point numbers (e.g., <i>1.5</i>) can be used as well.
group()	string	root	Set the gid of the socket.
host_override()	string		Replaces the \$HOST part of the message with the parameter string.
keep-alive()	yes or no	yes	Selects whether to keep connections open when syslog-ng is restarted; cannot be used with <i>unix-dgram()</i> .
keep_timestamp()	yes or no	yes	Specifies whether syslog-ng should accept the timestamp received from the sending application or client. If disabled, the time of reception will be used instead. This option can be specified globally, and per-source as well. The local setting of the source overrides the global option if available.
log_fetch_limit()	number	The value specified by the <i>global log_fetch_limit()</i> option, which defaults to 10.	The maximum number of messages fetched from a source during a single poll loop. The destination queues might fill up before flow-control could stop reading if <i>log_fetch_limit()</i> is too high.
log_iw_size()	number	100	The size of the initial window, this value is used during flow control.
log_msg_size()	number	Use the global <i>log_msg_size()</i> option, which defaults to 8192.	Specifies the maximum length of incoming log messages. Uses the value of the <u>global option</u> if not specified.
log_prefix() (DEPRECATED)	string		A string added to the beginning of every log message. It can be used to add an arbitrary string to any log source, though it is most commonly used for adding <i>kernel:</i> to the kernel messages on Linux. NOTE: This option is deprecated. Use <i>program_override()</i> instead.



Name	Type	Default	Description
max-connections()	number	256	Limits the number of simultaneously open connections. Cannot be used with <i>unix-dgram()</i> .
optional()	yes or no		Instruct syslog-ng to ignore the error if a specific source cannot be initialized. No other attempts to initialize the source will be made until the configuration is re-loaded. This option currently applies to the <i>pipe()</i> , <i>unix-dgram</i> , and <i>unix-stream</i> drivers.
owner()	string	root	Set the uid of the socket.
pad_size()	number	0	Specifies input padding. Some operating systems (such as HP-UX) pad all 0 messages to block boundary. This option can be used to specify the block size. (HP-UX uses 2048 bytes). Syslog-ng will pad reads from the associated device to the number of bytes set in <i>pad_size()</i> . Mostly used on HP-UX where <i>/dev/log</i> is a named pipe and every write is padded to 2048 bytes.
perm()	number	0666	Set the permission mask. For octal numbers prefix the number with '0', e.g.: use 0755 for <i>rwxr-xr-x</i> .
program_override	string		Replaces the \$PROGRAM part of the message with the parameter string. For example, to mark every message coming from the kernel, include the <i>program_override("kernel")</i> option in the source containing <i>/proc/kmsg</i> . NOTE: This option replaces the deprecated <i>log_prefix()</i> option.
so_broadcast()	yes or no	no	This option controls the <i>SO_BROADCAST</i> socket option required to make syslog-ng send messages to a broadcast address. See the socket(7) manual page for details.
so_keepalive()	yes or no	no	Enables keep-alive messages, keeping the socket open. This only effects TCP and UNIX-stream sockets. See the socket(7) manual page for details.
so_rcvbuf()	number	0	Specifies the size of the socket receive buffer in bytes. See the socket(7) manual page for details.
so_sndbuf()	number	0	Specifies the size of the socket send buffer in bytes. See the socket(7) manual page for details.



Name	Type	Default	Description
time_zone()	timezone in the form + / - HH:MM		The default timezone for messages read from the source. Applies only if no timezone is specified within the message itself.

Table 8.7. Options for `unix-stream()` and `unix-dgram()`



Example 8.11. Using the `unix-stream()` and `unix-dgram()` drivers

```
source s_stream { unix-stream("/dev/log" max-connections(10)); };
source s_dgram { unix-dgram("/var/run/log"); };
```

8.2. Destination drivers

Destination drivers output log messages to somewhere outside syslog-ng e.g., to a file or a network socket.

8.2.1. `file()`

The `file` driver outputs messages to the specified text file, or to a set of files.

The destination filename may include macros which get expanded when the message is written, thus a simple `file()` driver may create several files. For more information on available macros see Section 8.5, “Macros” (p. 210).



Warning

When creating several thousands separate log files, syslog-ng might not be able to open the required number of files. This might happen for example when using the `$HOST` macro in the filename while receiving messages from a large number of hosts. To overcome this problem, adjust the `--fd-limit` command-line parameter of syslog-ng or the global `ulimit` parameter of your host. For setting the `--fd-limit` command-line parameter of syslog-ng see the syslog-ng(8) (p. 235) manual page. For setting the `ulimit` parameter of the host, see the documentation of your operating system.

The `file()` destination has the following options:

Name	Type	Default	Description
create_dirs()	yes or no	no	Enable creating non-existing directories.
dir_group()	string	root	The group of directories created by syslog-ng.
dir_owner()	string	root	The owner of directories created by syslog-ng.
dir_perm()	number	0600	The permission mask of directories created by syslog-ng. Log directories are only created if a file after macro expansion refers to a non-existing directory, and directory creation is enabled (see the <i>create_dirs()</i> option below). For octal numbers prefix the number with 0, e.g., use 0755 for <i>rwxr-xr-x</i> .
flags()	no_multi_line, syslog-protocol	empty set	<p>Flags influence the behavior of the driver.</p> <p>The <i>no-multi-line</i> flag disables line-breaking in the messages; the entire message is converted to a single line.</p> <p>The <i>syslog-protocol</i> flag instructs the driver to format the messages according to the new IETF syslog protocol standard. If this flag is enabled, macros used for the message have effect only for the text of the message, the message header is formatted to the new standard. Note that this flag is not needed for the <i>syslog</i> driver.</p>
flush_lines()	number	Use global setting.	Specifies how many lines are flushed to a destination at a time. Syslog-ng waits for this number of lines to accumulate and sends them off in a single batch. Setting this number high increases throughput as fully filled frames are sent to the network, but also increases message latency. The latency can be limited by the use of the <i>flush_timeout</i> option.
flush_timeout()	time in milliseconds	Use global setting.	Specifies the time syslog-ng waits for lines to accumulate in its output buffer. See the <i>flush_lines</i> option for more information.
frac_digits()	number	0	The syslog-ng application can store fractions of a second in the timestamps according to the ISO8601 format. The <i>frac_digits()</i> parameter specifies the number of digits stored. The digits storing the fractions are padded by zeros if the original timestamp of the message specifies only seconds. Fractions can always be stored for the time the message was received. Note that syslog-ng can add the fractions to non-ISO8601 timestamps as well.

Name	Type	Default	Description
<code>fsync()</code>	yes or no	no	Forces an <i>fsync()</i> call on the destination fd after each write. Note: enabling this option may seriously degrade performance.
<code>group()</code>	string	root	Set the group of the created file to the one specified.
<code>local_time_zone()</code>	name of the timezone or the timezone offset	The local timezone.	Sets the timezone used when expanding filename and tablename templates. The timezone can be specified as using the name of the (e.g., <i>time_zone("Europe/Budapest")</i>), or as the timezone offset (e.g., <i>+01:00</i>). The valid timezone names are listed under the <code>/usr/share/zoneinfo</code> directory.
<code>log_fifo_size()</code>	number	Use global setting.	The number of entries in the output buffer (output fifo).
<code>overwrite_if_older()</code>	number	0	If set to a value higher than 0, syslog-ng checks when the file was last modified before starting to write into the file. If the file is older than the specified amount of time (in seconds), then syslog-ng removes the existing file and opens a new file with the same name. In combination with e.g., the <i>\$WEEKDAY</i> macro, this can be used for simple log rotation, in case not all history has to be kept. (Note that in this weekly log rotation example if its Monday 00:01, then the file from last Monday is not seven days old, because it was probably last modified shortly before 23:59 last Monday, so it is actually not even six days old. So in this case, set the <i>overwrite_if_older()</i> parameter to a-bit-less-than-six-days, for example, to <i>518000</i> seconds.
<code>owner()</code>	string	root	Set the owner of the created file to the one specified.
<code>perm()</code>	number	0600	The permission mask of the file if it is created by syslog-ng. For octal numbers prefix the number with 0, e.g., use <i>0755</i> for <i>rwxr-xr-x</i> .
<code>suppress()</code>	seconds	0 (disabled)	If several identical log messages would be sent to the destination without any other messages between the identical messages (for example, an application repeated an error message ten times), syslog-ng can suppress the repeated messages and send the message only once, followed by the <i>Last message repeated n times.</i> message. The parameter of this option specifies the number of seconds syslog-ng waits for identical messages.

Name	Type	Default	Description
template()	string	A format conforming to the default log file format.	Specifies a template defining the logformat to be used in the destination. Macros are described in Section 8.5, “Macros” (p. 210). Please note that for network destinations it might not be appropriate to change the template as it changes the on-wire format of the syslog protocol which might not be tolerated by stock syslog receivers (like <i>syslogd</i> or <i>syslog-ng</i> itself). For network destinations make sure the receiver can cope with the custom format defined.
template_escape()	yes or no	no	Turns on escaping ' and " in templated output files. This is useful for generating SQL statements and quoting string contents so that parts of the log message are not interpreted as commands to the SQL server.
throttle()	number	0	Sets the maximum number of messages sent to the destination per second. Use this output-rate-limiting functionality only when using disk-buffer as well to avoid the risk of losing messages. Specifying 0 or a lower value sets the output limit to unlimited.
time_zone()	timezone offset in seconds	unspecified	Convert timestamps to the timezone specified by this option. If this option is not set then the original timezone information in the message is used.
ts_format()	rfc3164, bsd, rfc3339, iso	rfc3164	Override the global timestamp format (set in the global <i>ts_format()</i> parameter) for the specific destination. See also Section 7.7, “A note on timezones and timestamps” (p. 140).

Table 8.8. Options for file()

**Example 8.12. Using the file() driver**

```
destination d_file { file("/var/log/messages" ); };
```

**Example 8.13. Using the file() driver with macros in the file name and a template for the message**

```
destination d_file {
    file("/var/log/$YEAR.$MONTH.$DAY/messages"
        template("$HOUR:$MIN:$SEC $TZ $HOST [$LEVEL] $MSG $MSG\n")
        template_escape(no));
};
```


8.2.2. logstore()

The logstore driver stores log messages in binary files that can be encrypted, compressed, checked for integrity, and timestamped by an external Timestamping Authority (TSA). Otherwise, it is very similar to the `file()` destination.

To display the contents of a logstore file, use the **logcat** command supplied with syslog-ng, e.g., **logcat /var/log/messages.lgs**.

The destination filename may include macros which get expanded when the message is written, thus a simple `logstore()` driver may create several files. For more information on available macros see Section 8.5, “Macros” (p. 210).

**Warning**

When creating several thousands separate log files, syslog-ng might not be able to open the required number of files. This might happen for example when using the `$HOST` macro in the filename while receiving messages from a large number of hosts. To overcome this problem, adjust the `--fd-limit` command-line parameter of syslog-ng or the global `ulimit` parameter of your host. For setting the `--fd-limit` command-line parameter of syslog-ng see the syslog-ng(8) (p. 235) manual page. For setting the `ulimit` parameter of the host, see the documentation of your operating system.

The `logstore()` has a single required parameter that specifies the filename that stores the log messages.

Declaration:

```
logstore(filename options());
```

The `logstore()` destination has the following options:

Name	Type	Default	Description
chunk_size()	number	128	Size of a logstore chunk in kilobytes. Note that this size refers to the compressed size of the chunk. Also, the gzip library used for compressing the messages has a 32k long buffer; messages may not appear in the actual logfile until this buffer is not filled. Logstore chunks are closed when they reach the specified size, or when the time limit set in <i>chunk_time</i> expires.
chunk_time()	number	5	Time limit in seconds: syslog-ng PE closes the chunk if no new messages arrive until the time limit expires. Logstore chunks are closed when the time limit expires, or when they reach the size specified in the <i>chunk_size</i> parameter. If the time limit set in the <i>time_reap</i> parameter expires, the entire file is closed.
compress()	number between 0-9	3	Compression level. 0 means uncompressed files, while 1-9 is the compression level used by gzip (9 means the highest but slowest compression, 3 is usually a good compromise).
create_dirs()	yes or no	no	Enable creating non-existing directories.
dir_group()	string	root	The group of the directories created by syslog-ng.
dir_owner()	string	root	The owner of directories created by syslog-ng.
dir_perm()	number	0600	The permission mask of directories created by syslog-ng. Log directories are only created if a file after macro expansion refers to a non-existing directory, and directory creation is enabled (see the <i>create_dirs()</i> option below). For octal numbers prefix the number with 0, e.g., use 0755 for <i>rwxr-xr-x</i> .
encrypt_certificate()	filename	none	Name of a file, that contains an X.509 certificate (and the public key) in PEM format. The syslog-ng application uses this certificate to encrypt the logstore files which can be decrypted using the private key of the certificate.
flags()	no_multi_line, syslog-protocol	empty set	<p>Flags influence the behavior of the driver.</p> <p>The <i>no-multi-line</i> flag disables line-breaking in the messages; the entire message is converted to a single line.</p> <p>The <i>syslog-protocol</i> flag instructs the driver to format the messages according to the new IETF syslog protocol standard. If this flag is enabled, macros used for the message have effect only for the text of the message, the message header is formatted to the new standard. Note that this flag is not needed for the <i>syslog</i> driver.</p>
flush_lines()	number	Use global setting.	Specifies how many lines are flushed to a destination at a time. Syslog-ng waits for this number of lines to accumulate and sends them off in a single batch. Setting this number high increases throughput as fully filled frames are sent to the network, but also increases message latency. The latency can be limited by the use of the <i>flush_timeout</i> option.

Name	Type	Default	Description
flush_timeout()	time in milliseconds	Use global setting.	Specifies the time syslog-ng waits for lines to accumulate in its output buffer. See the <i>flush_lines</i> option for more information.
frac_digits()	number	0	The syslog-ng application can store fractions of a second in the timestamps according to the ISO8601 format.. The <i>frac_digits()</i> parameter specifies the number of digits stored. The digits storing the fractions are padded by zeros if the original timestamp of the message specifies only seconds. Fractions can always be stored for the time the message was received. Note that syslog-ng can add the fractions to non-ISO8601 timestamps as well.
fsync()	yes or no	no	Forces an <i>fsync()</i> call on the destination fd after each write. Note: enabling this option may seriously degrade performance.
group()	string	root	Set the group of the created file to the one specified.
log_fifo_size()	number	Use global setting.	The number of entries in the output buffer (output fifo).
owner()	string	root	Set the owner of the created file to the one specified.
perm()	number	0600	The permission mask of the file if it is created by syslog-ng. For octal numbers prefix the number with 0, e.g., use 0755 for <i>rwxr-xr-x</i> .
suppress()	seconds	0 (disabled)	If several identical log messages would be sent to the destination without any other messages between the identical messages (for example, an application repeated an error message ten times), syslog-ng can suppress the repeated messages and send the message only once, followed by the <i>Last message repeated n times.</i> message. The parameter of this option specifies the number of seconds syslog-ng waits for identical messages.
template()	string	A format conforming to the default log file format.	Specifies a template defining the logformat to be used in the destination. Macros are described in Section 8.5, “Macros” (p. 210). Please note that for network destinations it might not be appropriate to change the template as it changes the on-wire format of the syslog protocol which might not be tolerated by stock syslog receivers (like <i>syslogd</i> or syslog-ng itself). For network destinations make sure the receiver can cope with the custom format defined.
template_escape()	yes or no	no	Turns on escaping ' and " in templated output files. This is useful for generating SQL statements and quoting string contents so that parts of the log message are not interpreted as commands to the SQL server.
throttle()	number	0	Sets the maximum number of messages sent to the destination per second. Use this output-rate-limiting functionality only when using disk-buffer as well to avoid the risk of losing messages. Specifying 0 or a lower value sets the output limit to unlimited.
time_reap()	number	60	The time to wait in seconds before an idle destination file is closed.

Name	Type	Default	Description
timestamp-freq()	number in seconds	Use global setting.	The minimum time that should expire between two timestamping requests. When syslog-ng closes a chunk, it checks how much time has expired since the last timestamping request: if it is higher than the value set in the <i>timestamp-freq</i> parameter, it requests a new timestamp from the authority set in the <i>timestamp-url</i> parameter.
timestamp-url()	string	Use global setting.	The URL of the Timestamping Authority used to request timestamps to sign logstore chunks.
time_zone()	timezone offset in seconds	unspecified	Convert timestamps to the timezone specified by this option. If this option is not set then the original timezone information in the message is used.
ts_format()	rfc3164, bsd, rfc3339, iso	rfc3164	Override the global timestamp format (set in the global <i>ts_format()</i> parameter) for the specific destination. See also Section 7.7, “A note on timezones and timestamps” (p. 140).

Table 8.9. Options for logstore()

**Example 8.14. Using the logstore() driver**

A simple example saving and compressing log messages.

```
destination d_logstore { file("/var/log/messages.lgs" compress(5) ); };
```

A more detailed example that encrypts messages, modifies the parameters for closing chunks, and sets file privileges.

```
destination d_logstore { logstore("/var/log/messages-logstore.lgs"
encrypt_certificate("/opt/syslog-ng/etc/syslog-ng/keys/10-100-20-40/public-certificate-of-the-server.pem")

    chunk_size(100)
    chunk_time(5)
    owner("balabit")
    group("balabit")
    perm(0777)
); };
```

8.2.3. pipe()

This driver sends messages to a named pipe like `/dev/xconsole`.

The pipe driver has a single required parameter, specifying the filename of the pipe to open. The filename can include macros.

Declaration:

```
pipe(filename);
```

**Warning**

As of syslog-ng Open Source Edition 3.0.2, pipes are created automatically. In earlier versions, you had to create the pipe using the **mkfifo(1)** command.

The *pipe()* destination has the following options:

Name	Type	Default	Description
flags()	no_multi_line, syslog-protocol	empty set	<p>Flags influence the behavior of the driver.</p> <p>The <i>no-multi-line</i> flag disables line-breaking in the messages; the entire message is converted to a single line.</p> <p>The <i>syslog-protocol</i> flag instructs the driver to format the messages according to the new IETF syslog protocol standard. If this flag is enabled, macros used for the message have effect only for the text of the message, the message header is formatted to the new standard. Note that this flag is not needed for the <i>syslog</i> driver.</p>
flush_lines()	number	Use global setting.	Specifies how many lines are flushed to a destination at a time. Syslog-ng waits for this number of lines to accumulate and sends them off in a single batch. Setting this number high increases throughput as fully filled frames are sent to the network, but also increases message latency. The latency can be limited by the use of the <i>flush_timeout</i> option.
flush_timeout()	time in milliseconds	Use global setting.	Specifies the time syslog-ng waits for lines to accumulate in its output buffer. See the <i>flush_lines</i> option for more information.
frac_digits()	number	0	The syslog-ng application can store fractions of a second in the timestamps according to the ISO8601 format. The <i>frac_digits()</i> parameter specifies the number of digits stored. The digits storing the fractions are padded by zeros if the original timestamp of the message specifies only seconds. Fractions can always be stored for the time the message was received. Note that syslog-ng can add the fractions to non-ISO8601 timestamps as well.
fsync()	yes or no	no	Forces an <i>fsync()</i> call on the destination fd after each write. Note: enabling this option may seriously degrade performance.
group()	string	root	Set the group of the pipe to the one specified.
log_fifo_size()	number	Use global setting.	The number of entries in the output buffer (output fifo).
owner()	string	root	Set the owner of the pipe to the one specified.
perm()	number	0600	The permission mask of the pipe. For octal numbers prefix the number with '0', e.g.: use 0755 for rwxr-xr-x.
suppress()	seconds	0 (disabled)	If several identical log messages would be sent to the destination without any other messages between the identical messages (for example, an application repeated an error message ten times), syslog-ng can suppress the repeated messages and send the message only once, followed by the <i>Last message repeated n times.</i> message. The parameter of this option specifies the number of seconds syslog-ng waits for identical messages.

Name	Type	Default	Description
template()	string	A format conforming to the default logfile format.	Specifies a template defining the logformat to be used in the destination. Macros are described in Section 8.5, “Macros” (p. 210). Please note that for network destinations it might not be appropriate to change the template as it changes the on-wire format of the syslog protocol which might not be tolerated by stock syslog receivers (like <i>syslogd</i> or <i>syslog-ng</i> itself). For network destinations make sure the receiver can cope with the custom format defined.
template_escape()	yes or no	no	Turns on escaping ' and " in templated output files. This is useful for generating SQL statements and quoting string contents so that parts of the log message are not interpreted as commands to the SQL server.
throttle()	number	0	Sets the maximum number of messages sent to the destination per second. Use this output-rate-limiting functionality only when using disk-buffer as well to avoid the risk of losing messages. Specifying 0 or a lower value sets the output limit to unlimited.
time_zone()	timezone offset in seconds	unspecified	Convert timestamps to the timezone specified by this option. If this option is not set then the original timezone information in the message is used.
ts_format()	rfc3164, bsd, rfc3339, iso	rfc3164	Override the global timestamp format (set in the global <i>ts_format()</i> parameter) for the specific destination. See also Section 7.7, “A note on timezones and timestamps” (p. 140).

Table 8.10. Options for pipe()

**Example 8.15. Using the pipe() driver**

```
destination d_pipe { pipe("/dev/xconsole"); };
```

8.2.4. program()

This driver starts an external application or script and sends the log messages to its standard input (*stdin*).

The *program()* driver has a single required parameter, specifying a program name to start.

Declaration:

```
program(command_to_run);
```

The *program()* destination has the following options:

Name	Type	Default	Description
flags()	no_multi_line, syslog-protocol	empty set	<p>Flags influence the behavior of the driver.</p> <p>The <i>no-multi-line</i> flag disables line-breaking in the messages; the entire message is converted to a single line.</p> <p>The <i>syslog-protocol</i> flag instructs the driver to format the messages according to the new IETF syslog protocol standard. If this flag is enabled, macros used for the message have effect only for the text of the message, the message header is formatted to the new standard. Note that this flag is not needed for the <i>syslog</i> driver.</p>
flush_lines()	number	Use global setting.	Specifies how many lines are flushed to a destination at a time. Syslog-ng waits for this number of lines to accumulate and sends them off in a single batch. Setting this number high increases throughput as fully filled frames are sent to the network, but also increases message latency. The latency can be limited by the use of the <i>flush_timeout</i> option.
flush_timeout()	time in milliseconds	Use global setting.	Specifies the time syslog-ng waits for lines to accumulate in its output buffer. See the <i>flush_lines</i> option for more information.
frac_digits()	number	0	The syslog-ng application can store fractions of a second in the timestamps according to the ISO8601 format. The <i>frac_digits()</i> parameter specifies the number of digits stored. The digits storing the fractions are padded by zeros if the original timestamp of the message specifies only seconds. Fractions can always be stored for the time the message was received. Note that syslog-ng can add the fractions to non-ISO8601 timestamps as well.
fsync()	yes or no	no	Forces an <i>fsync()</i> call on the destination fd after each write. Note: enabling this option may seriously degrade performance.
log_fifo_size()	number	Use global setting.	The number of entries in the output buffer (output fifo).
suppress()	seconds	0 (disabled)	If several identical log messages would be sent to the destination without any other messages between the identical messages (for example, an application repeated an error message ten times), syslog-ng can suppress the repeated messages and send the message only once, followed by the <i>Last message repeated n times.</i> message. The parameter of this option specifies the number of seconds syslog-ng waits for identical messages.

Name	Type	Default	Description
template()	string	A format conforming to the default logfile format.	Specifies a template defining the logformat to be used in the destination. Macros are described in Section 8.5, “Macros” (p. 210). Please note that for network destinations it might not be appropriate to change the template as it changes the on-wire format of the syslog protocol which might not be tolerated by stock syslog receivers (like <i>syslogd</i> or <i>syslog-ng</i> itself). For network destinations make sure the receiver can cope with the custom format defined.
template_escape()	yes or no	no	Turns on escaping ' and " in templated output files. This is useful for generating SQL statements and quoting string contents so that parts of the log message are not interpreted as commands to the SQL server.
throttle()	number	0	Sets the maximum number of messages sent to the destination per second. Use this output-rate-limiting functionality only when using disk-buffer as well to avoid the risk of losing messages. Specifying 0 or a lower value sets the output limit to unlimited.
time_zone()	timezone offset in seconds	unspecified	Convert timestamps to the timezone specified by this option. If this option is not set then the original timezone information in the message is used.
ts_format()	rfc3164, bsd, rfc3339, iso	rfc3164	Override the global timestamp format (set in the global <i>ts_format()</i> parameter) for the specific destination. See also Section 7.7, “A note on timezones and timestamps” (p. 140).

Table 8.11. Options for program()

**Example 8.16. Using the program() destination driver**

```
destination d_prog { program("/bin/script" template("<$PRI>$DATE $HOST $MSG\n");
};
```

8.2.5. sql()

This driver sends messages into an SQL database. The *sql()* driver has the following required parameters: *type*, *database*, *table*, *columns*, *values*.

Declaration:

```
sql(database_type host_parameters database_parameters [options]);
```

The *sql()* destination has the following options:

Name	Type	Default	Description
columns	string list	"date", "facility", "level", "host", "program", "pid", "message"	Name of the columns storing the data in <i>fieldname [dbtype]</i> format. The <i>[dbtype]</i> parameter is optional, and specifies the type of the field. By default, syslog-ng creates <i>text</i> columns. Note that not every database engine can index text fields.
database	string	n/a	Name of the database that stores the logs.
frac_digits()	number	0	The syslog-ng application can store fractions of a second in the timestamps according to the ISO8601 format. The <i>frac_digits()</i> parameter specifies the number of digits stored. The digits storing the fractions are padded by zeros if the original timestamp of the message specifies only seconds. Fractions can always be stored for the time the message was received. Note that syslog-ng can add the fractions to non-ISO8601 timestamps as well.
host	hostname or IP address	n/a	Hostname of the database server. Note that Oracle destinations do not use this parameter, but retrieve the hostname from the <code>/etc/tnsnames.ora</code> file.
indexes	string list	"date", "facility", "host", "program"	The list of columns that are indexed by the database to speed up searching. To disable indexing for the destination, include the empty <i>indexes()</i> parameter in the destination, simply omitting the <i>indexes</i> parameter will cause syslog-ng to request indexing on the default columns.
<i>local_time_zone()</i>	name of the timezone or the timezone offset	The local timezone.	Sets the timezone used when expanding filename and tablename templates. The timezone can be specified as using the name of the (e.g., <i>time_zone("Europe/Budapest")</i>), or as the timezone offset (e.g., <i>+01:00</i>). The valid timezone names are listed under the <code>/usr/share/zoneinfo</code> directory.
log_disk_fifo_size()	number	0	Size of the hard disk space in bytes that is used as disk buffer. Available only in syslog-ng Premium Edition when using the <i>tcp()</i> , <i>tcp6()</i> , <i>syslog()</i> (when using the <i>tcp</i> or <i>tls</i> transport methods), and <i>sql()</i> destinations. Can be also defined as a global option. See Section 2.14, “Using disk-based buffering” (p. 20) for details on using the disk buffer.

Name	Type	Default	Description
log_fifo_size()	number	Use global setting.	The number of entries in the output buffer (output fifo).
null	string		If the content of a column matches the string specified in the <i>null ()</i> parameter, the contents of the column will be replaced with an SQL NULL value. If unset (by default), the option does not match on any string. See the Example 8.20, “Using SQL NULL values” (p. 195) for details.
password	string	n/a	Password of the database user.
table	string	n/a	Name of the database table to use (can include macros). When using macros, note that some databases limit the length of table names.
time_zone()	timezone offset in seconds	unspecified	Convert timestamps to the timezone specified by this option. If this option is not set then the original timezone information in the message is used.
type	m s s q l, m y s q l, o r a c l e, p g s q l, or s q l i t e 3	n/a	Specifies the type of the database, i.e., the DBI database driver to use. Use the <i>mssql</i> option to send logs to an MSSQL database. See the examples of the databases on the following sections for details.
username	string	n/a	Name of the database user.
values	string list	"\${R_YEAR}-\${R_MONTH}-\${R_DAY} \${R_HOUR}:\${R_MIN}:\${R_SEC}", "\$FACILITY", "\$LEVEL", "\$HOST", "\$PROGRAM", "\$PID", "\$MSG-ONLY"	The parts of the message to store in the fields specified in the <i>columns</i> parameter.

Table 8.12. Options for sql()

**Note**

If you specify *host="localhost"*, syslog-ng will use a socket to connect to the local database server. Use *host="127.0.0.1"* to force TCP communication between syslog-ng and the local database server.

To specify the socket to use, set and export the *MYSQL_UNIX_PORT* environment variable, e.g., **MYSQL_UNIX_PORT=/var/lib/mysql/mysql.sock; export MYSQL_UNIX_PORT.**

**Example 8.17. Using the sql() driver**

The following example sends the log messages into a PostgreSQL database running on the *logserver* host. The messages are inserted into the *logs* database, the name of the table includes the exact date and the name of the host sending the messages. The syslog-ng application automatically creates the required tables and columns, if the user account used to connect to the database has the required privileges.

```
destination d_sql {
    sql(type(pgsql))
}
```

```
host("logserver") username("syslog-ng") password("password")
database("logs")
table("messages_${HOST}_${R_YEAR}${R_MONTH}${R_DAY}")
columns("datetime", "host", "program", "pid", "message")
values("${R_DATE}", "${HOST}", "${PROGRAM}", "${PID}", "${MSGONLY}")
indexes("datetime", "host", "program", "pid", "message"));
};
```

The following example specifies the type of the database columns as well:

```
destination d_sql {
    sql(type(pgsql)
    host("logserver") username("syslog-ng") password("password")
    database("logs")
    table("messages_${HOST}_${R_YEAR}${R_MONTH}${R_DAY}")
    columns("datetime varchar(16)", "host varchar(32)", "program varchar(20)",
    "pid varchar(8)", "message varchar(200)")
    values("${R_DATE}", "${HOST}", "${PROGRAM}", "${PID}", "${MSGONLY}")
    indexes("datetime", "host", "program", "pid", "message"));
};
```



Example 8.18. Using the sql() driver with an Oracle database

The following example sends the log messages into an Oracle database running on the *logserver* host, which must be set in the */etc/tnsnames.ora* file. The messages are inserted into the *LOGS* database, the name of the table includes the exact date when the messages were sent. The syslog-ng application automatically creates the required tables and columns, if the user account used to connect to the database has the required privileges.

```
destination d_sql {
    sql(type(oracle)
    username("syslog-ng") password("password")
    database("LOGS")
    table("msgs_${R_YEAR}${R_MONTH}${R_DAY}")
    columns("datetime varchar(16)", "host varchar(32)", "program varchar(32)", "pid
    varchar(8)", "message varchar2")
    values("${R_DATE}", "${HOST}", "${PROGRAM}", "${PID}", "${MSGONLY}")
    indexes("datetime", "host", "program", "pid", "message"));
};
```

The Oracle Instant Client retrieves the address of the database server from the */etc/tnsnames.ora* file. Edit or create this file as needed for your configuration. A sample is provided below.

```
LOGS =
(DESCRIPTION =
(ADDRESS_LIST =
(ADDRESS = (PROTOCOL = TCP)
(HOST = logserver)
(PORT = 1521))
)
(CONNECT_DATA =
(SERVICE_NAME = EXAMPLE.SERVICE)
)
)
```



Example 8.19. Using the sql() driver with an MSSQL database

The following example sends the log messages into an MSSQL database running on the *logserver* host. The messages are inserted into the *syslogng* database, the name of the table includes the exact date when the messages were sent. The syslog-ng application automatically creates the required tables and columns, if the user account used to connect to the database has the required privileges.

```
destination d_mssql {
    sql(type(mssql) host("logserver") port("1433")
    username("syslogng") password("syslogng") database("syslogng")
    table("msgs_${R_YEAR}${R_MONTH}${R_DAY}") columns("datetime varchar(16)", "host
```



```
varchar(32)",
"program varchar(32)", "pid varchar(8)", "message varchar(4096)")
values("$R_DATE", "$HOST", "$PROGRAM", "$PID", "$MSGONLY")
indexes("datetime", "host", "program", "pid", "message"));
};
```

The date format used by the MSSQL database must be explicitly set in the `/etc/locales.conf` file of the syslog-ng server. Edit or create this file as needed for your configuration. A sample is provided below.

```
[default]
date = "%Y-%m-%d %H:%M:%S"
```



Example 8.20. Using SQL NULL values

The `null()` parameter of the SQL driver can be used to replace the contents of a column with a special SQL NULL value. To replace every column that contains an empty string with NULL, use the `null("")` option, e.g.,

```
destination d_sql {
    sql(type(pgsql)
        host("logserver") username("syslog-ng") password("password")

        database("logs")
        table("messages_${HOST}_${R_YEAR}${R_MONTH}${R_DAY}")
        columns("datetime", "host", "program", "pid", "message")
        values("$R_DATE", "$HOST", "$PROGRAM", "$PID", "$MSGONLY")
        indexes("datetime", "host", "program", "pid", "message")
        null(""));
};
```

To replace only a specific column (e.g., `pid`) if it is empty, assign a default value to the column, and use this default value in the `null()` parameter:

```
destination d_sql {
    sql(type(pgsql)
        host("logserver") username("syslog-ng") password("password")

        database("logs")
        table("messages_${HOST}_${R_YEAR}${R_MONTH}${R_DAY}")
        columns("datetime", "host", "program", "pid", "message")
        values("$R_DATE", "$HOST", "$PROGRAM", "${PID:-@NULL@}",
"$MSGONLY")

        indexes("datetime", "host", "program", "pid", "message")
        null("@NULL@"));
};
```

Ensure that the default value you use does not appear in the actual log messages, because other occurrences of this string will be replaced with NULL as well.

8.2.6. syslog()

The `syslog()` driver sends messages to a remote host (e.g., a syslog-ng server or relay) on the local intranet or internet using the new standard syslog protocol developed by IETF (see Section 2.18.2, “IETF-syslog messages” (p. 26) for details about the protocol). The protocol supports sending messages using the UDP, TCP, or the encrypted TLS networking protocols.

The required arguments of the driver are the address of the destination host (where messages should be sent) and the transport method (networking protocol).



The *udp* transport method automatically sends multicast packets if a multicast destination address is specified. The *tcp* and *tls* methods do not support multicasting.

Declaration:

```
syslog(host transport [options]);
```

These destinations have the following options:

Name	Type	Default	Description
flags()	no_multi_line, syslog-protocol	empty set	<p>Flags influence the behavior of the driver.</p> <p>The <i>no-multi-line</i> flag disables line-breaking in the messages; the entire message is converted to a single line.</p> <p>The <i>syslog-protocol</i> flag instructs the driver to format the messages according to the new IETF syslog protocol standard. If this flag is enabled, macros used for the message have effect only for the text of the message, the message header is formatted to the new standard. Note that this flag is not needed for the <i>syslog</i> driver.</p>
flush_lines()	number	Use global setting.	Specifies how many lines are flushed to a destination at a time. Syslog-ng waits for this number of lines to accumulate and sends them off in a single batch. Setting this number high increases throughput as fully filled frames are sent to the network, but also increases message latency. The latency can be limited by the use of the <i>flush_timeout</i> option.
flush_timeout()	time in milliseconds	Use global setting.	Specifies the time syslog-ng waits for lines to accumulate in its output buffer. See the <i>flush_lines</i> option for more information.
frac_digits()	number	0	The syslog-ng application can store fractions of a second in the timestamps according to the ISO8601 format. The <i>frac_digits()</i> parameter specifies the number of digits stored. The digits storing the fractions are padded by zeros if the original timestamp of the message specifies only seconds. Fractions can always be stored for the time the message was received. Note that syslog-ng can add the fractions to non-ISO8601 timestamps as well.
fsync()	yes or no	no	Forces an <i>fsync()</i> call on the destination fd after each write. Note: enabling this option may seriously degrade performance.
ip_tos()	number	0	Specifies the Type-of-Service value of outgoing packets.
ip_ttl()	number	0	Specifies the Time-To-Live value of outgoing packets.
keep-alive()	yes or no	yes	Specifies whether connections to destinations should be closed when syslog-ng is restarted (upon the receipt of a SIGHUP signal). Note that this applies to the client (destination) side of the syslog-ng connections, server-side (source) connections are always re-opened after receiving a HUP signal unless the <i>keep-alive</i> option is enabled for the source. When the <i>keep-alive</i> option is enabled, syslog-ng saves the contents of the output queue of the destination when receiving a HUP signal, reducing the risk of losing messages.
localip()	string	0.0.0.0	The IP address to bind to before connecting to target.
localport()	number	0	The port number to bind to. Messages are sent from this port.

Name	Type	Default	Description
log_disk_fifo_size()	number	0	Size of the hard disk space in bytes that is used as disk buffer. Available only in syslog-ng Premium Edition when using the <i>tcp()</i> , <i>tcp6()</i> , <i>syslog()</i> (when using the <i>tcp</i> or <i>tls</i> transport methods), and <i>sql()</i> destinations. Can be also defined as a global option. See Section 2.14, “Using disk-based buffering” (p. 20) for details on using the disk buffer.
log_fifo_size()	number	Use global setting.	The number of entries in the output buffer (output fifo).
port() or destport()	number	601	The port number to connect to. Note that the default port numbers used by syslog-ng do not comply with the latest RFC which was published after the release of syslog-ng 3.0.2, therefore the default port numbers will change in the future releases.
so_broadcast()	yes or no	no	This option controls the <i>SO_BROADCAST</i> socket option required to make syslog-ng send messages to a broadcast address. See the socket(7) manual page for details.
so_keepalive()	yes or no	no	Enables keep-alive messages, keeping the socket open. This only effects TCP and UNIX-stream sockets. See the socket(7) manual page for details.
so_rcvbuf()	number	0	Specifies the size of the socket receive buffer in bytes. See the socket(7) manual page for details.
so_sndbuf()	number	0	Specifies the size of the socket send buffer in bytes. See the socket(7) manual page for details.
spoof_source()	yes or no	no	Enables source address spoofing. This means that the host running syslog-ng generates UDP packets with the source IP address matching the original sender of the message. It is useful when you want to perform some kind of preprocessing via syslog-ng then forward messages to your central log management solution with the source address of the original sender. This option only works for UDP destinations though the original message can be received by TCP as well. This option is only available if syslog-ng was compiled using the <i>--enable-spoof-source</i> configuration option.
suppress()	seconds	0 (disabled)	If several identical log messages would be sent to the destination without any other messages between the identical messages (for example, an application repeated an error message ten times), syslog-ng can suppress the repeated messages and send the message only once, followed by the <i>Last message repeated n times.</i> message. The parameter of this option specifies the number of seconds syslog-ng waits for identical messages.

Name	Type	Default	Description
template()	string	A format conforming to the default log-file format.	Specifies a template defining the logformat to be used in the destination. Macros are described in Section 8.5, “Macros” (p. 210). Please note that for network destinations it might not be appropriate to change the template as it changes the on-wire format of the syslog protocol which might not be tolerated by stock syslog receivers (like <i>syslogd</i> or <i>syslog-ng</i> itself). For network destinations make sure the receiver can cope with the custom format defined.
template_escape()	yes or no	no	Turns on escaping ' and " in templated output files. This is useful for generating SQL statements and quoting string contents so that parts of the log message are not interpreted as commands to the SQL server.
throttle()	number	0	Sets the maximum number of messages sent to the destination per second. Use this output-rate-limiting functionality only when using disk-buffer as well to avoid the risk of losing messages. Specifying 0 or a lower value sets the output limit to unlimited.
time_zone()	timezone offset in seconds	unspecified	Convert timestamps to the timezone specified by this option. If this option is not set then the original timezone information in the message is used.
tls()	tls options	n/a	This option sets various TLS specific options like key/certificate files and trusted CA locations. TLS can be used only with the <i>tcp</i> transport protocols. See Section 8.10, “TLS options” (p. 231) for more information.
transport	udp, tcp, or tls	tcp	Specifies the protocol used to receive messages from the source.
ts_format()	rfc3164, bsd, rfc3339, iso	rfc3164	Override the global timestamp format (set in the global <i>ts_format()</i> parameter) for the specific destination. See also Section 7.7, “A note on timezones and timestamps” (p. 140).

Table 8.13. Options for syslog() destinations

**Example 8.21. Using the syslog() driver**

```
destination d_tcp { syslog(ip"10.1.2.3" transport("tcp") port(1999);
localport(999)); };
```

If name resolution is configured, the hostname of the target server can be used as well.

```
destination d_tcp { syslog(ip"target_host" transport("tcp") port(1999);
localport(999)); };
```

Send the log messages using TLS encryption and use mutual authentication. See Section 8.10, “TLS options” (p. 231) for details on the encryption and authentication options.

```
destination d_syslog_tls{
    syslog("10.100.20.40"
    transport("tls")
    port(6514)
    tls(peer-verify(required-trusted)
    ca_dir('/opt/syslog-ng/etc/syslog-ng/keys/ca.d/')
    key_file('/opt/syslog-ng/etc/syslog-ng/keys/client_key.pem')
```



```
cert_file('/opt/syslog-ng/etc/syslog-ng/keys/client_certificate.pem'))  
);};
```

8.2.7. tcp(), tcp6(), udp(), and udp6(),

This driver sends messages to another host on the local intranet or internet using the UDP or TCP protocol. The *tcp6()* and *udp6()* drivers use the IPv6 network protocol.

Both drivers have a single required argument specifying the destination host address, where messages should be sent, and several optional parameters. Note that this differs from source drivers, where local bind address is implied, and none of the parameters are required.

The *udp()* and *udp6()* drivers automatically send multicast packets if a multicast destination address is specified. The *tcp()* and *tcp6()* drivers do not support multicasting.

Declaration:

```
tcp(host [options]);  
udp(host [options]);  
tcp6(host [options]);  
udp6(host [options]);
```

These destinations have the following options:

Name	Type	Default	Description
flags()	no_multi_line, syslog-protocol	empty set	<p>Flags influence the behavior of the driver.</p> <p>The <i>no-multi-line</i> flag disables line-breaking in the messages; the entire message is converted to a single line.</p> <p>The <i>syslog-protocol</i> flag instructs the driver to format the messages according to the new IETF syslog protocol standard. If this flag is enabled, macros used for the message have effect only for the text of the message, the message header is formatted to the new standard. Note that this flag is not needed for the <i>syslog</i> driver.</p>
flush_lines()	number	Use global setting.	Specifies how many lines are flushed to a destination at a time. Syslog-ng waits for this number of lines to accumulate and sends them off in a single batch. Setting this number high increases throughput as fully filled frames are sent to the network, but also increases message latency. The latency can be limited by the use of the <i>flush_timeout</i> option.
flush_timeout()	time in milliseconds	Use global setting.	Specifies the time syslog-ng waits for lines to accumulate in its output buffer. See the <i>flush_lines</i> option for more information.
frac_digits()	number	0	The syslog-ng application can store fractions of a second in the timestamps according to the ISO8601 format.. The <i>frac_digits()</i> parameter specifies the number of digits stored. The digits storing the fractions are padded by zeros if the original timestamp of the message specifies only seconds. Fractions can always be stored for the time the message was received. Note that syslog-ng can add the fractions to non-ISO8601 timestamps as well.
fsync()	yes or no	no	Forces an <i>fsync()</i> call on the destination fd after each write. Note: enabling this option may seriously degrade performance.
ip_tos()	number	0	Specifies the Type-of-Service value of outgoing packets.
ip_ttl()	number	0	Specifies the Time-To-Live value of outgoing packets.
keep-alive()	yes or no	yes	Specifies whether connections to destinations should be closed when syslog-ng is restarted (upon the receipt of a SIGHUP signal). Note that this applies to the client (destination) side of the syslog-ng connections, server-side (source) connections are always re-opened after receiving a HUP signal unless the <i>keep-alive</i> option is enabled for the source. When the <i>keep-alive</i> option is enabled, syslog-ng saves the contents of the output queue of the destination when receiving a HUP signal, reducing the risk of losing messages.
localip()	string	0.0.0.0	The IP address to bind to before connecting to target.
localport()	number	0	The port number to bind to. Messages are sent from this port.

Name	Type	Default	Description
log_disk_fifo_size()	number	0	Size of the hard disk space in bytes that is used as disk buffer. Available only in syslog-ng Premium Edition when using the <i>tcp()</i> , <i>tcp6()</i> , <i>syslog()</i> (when using the <i>tcp</i> or <i>tls</i> transport methods), and <i>sql()</i> destinations. Can be also defined as a global option. See Section 2.14, “Using disk-based buffering” (p. 20) for details on using the disk buffer.
log_fifo_size()	number	Use global setting.	The number of entries in the output buffer (output fifo).
port() or destport()	number	514	The port number to connect to. Note that the default port numbers used by syslog-ng do not comply with the latest RFC which was published after the release of syslog-ng 3.0.2, therefore the default port numbers will change in the future releases.
so_broadcast()	yes or no	no	This option controls the <i>SO_BROADCAST</i> socket option required to make syslog-ng send messages to a broadcast address. See the socket(7) manual page for details.
so_keepalive()	yes or no	no	Enables keep-alive messages, keeping the socket open. This only effects TCP and UNIX-stream sockets. See the socket(7) manual page for details.
so_rcvbuf()	number	0	Specifies the size of the socket receive buffer in bytes. See the socket(7) manual page for details.
so_sndbuf()	number	0	Specifies the size of the socket send buffer in bytes. See the socket(7) manual page for details.
spoof_source()	yes or no	no	Enables source address spoofing. This means that the host running syslog-ng generates UDP packets with the source IP address matching the original sender of the message. It is useful when you want to perform some kind of preprocessing via syslog-ng then forward messages to your central log management solution with the source address of the original sender. This option only works for UDP destinations though the original message can be received by TCP as well. This option is only available if syslog-ng was compiled using the <i>--enable-spoof-source</i> configuration option.
suppress()	seconds	0 (disabled)	If several identical log messages would be sent to the destination without any other messages between the identical messages (for example, an application repeated an error message ten times), syslog-ng can suppress the repeated messages and send the message only once, followed by the <i>Last message repeated n times.</i> message. The parameter of this option specifies the number of seconds syslog-ng waits for identical messages.

Name	Type	Default	Description
template()	string	A format conforming to the default log-file format.	Specifies a template defining the logformat to be used in the destination. Macros are described in Section 8.5, “Macros” (p. 210). Please note that for network destinations it might not be appropriate to change the template as it changes the on-wire format of the syslog protocol which might not be tolerated by stock syslog receivers (like <i>syslogd</i> or <i>syslog-ng</i> itself). For network destinations make sure the receiver can cope with the custom format defined.
template_escape()	yes or no	no	Turns on escaping ' and " in templated output files. This is useful for generating SQL statements and quoting string contents so that parts of the log message are not interpreted as commands to the SQL server.
throttle()	number	0	Sets the maximum number of messages sent to the destination per second. Use this output-rate-limiting functionality only when using disk-buffer as well to avoid the risk of losing messages. Specifying 0 or a lower value sets the output limit to unlimited.
time_zone()	timezone offset in seconds	unspecified	Convert timestamps to the timezone specified by this option. If this option is not set then the original timezone information in the message is used.
tls()	tls options	n/a	This option sets various TLS specific options like key/certificate files and trusted CA locations. TLS can be used only with the <i>tcp</i> transport protocols. See Section 8.10, “TLS options” (p. 231) for more information.
ts_format()	rfc3164, bsd, rfc3339, iso	rfc3164	Override the global timestamp format (set in the global <i>ts_format()</i> parameter) for the specific destination. See also Section 7.7, “A note on timezones and timestamps” (p. 140).

Table 8.14. Options for tcp, tcp6, udp, and udp6 destinations

**Example 8.22. Using the tcp() driver**

```
destination d_tcp { tcp("10.1.2.3" port(1999); localport(999)); };
```

If name resolution is configured, the hostname of the target server can be used as well.

```
destination d_tcp { tcp("target_host" port(1999); localport(999)); };
```

To send messages using the IETF-syslog message format, enable the *syslog-protocol* flag:

```
destination d_tcp { tcp("10.1.2.3" port(1999); flags(syslog-protocol) ); }
```

**Example 8.23. Enabling disk-based buffering**

The following example turns on disk-based buffering for the destination. The size of the disk buffer is 4 194 304 bytes (4 megabytes). In a worst-case situation, using the default value of the *log_msg_size()* parameter (8192 bytes), this disk buffer can store at least 512 messages. Typical log messages are about 300-500 bytes long, so a disk buffer of 4 megabytes can store over 8000 messages. Set the size of the disk buffer based on the average size and number of messages, and the longest estimated downtime of the server.



```
destination d_tcp {  
    tcp("10.1.2.3" port(1999) log_disk_fifo_size(4194304)); };
```

8.2.8. unix-stream() & unix-dgram()

These drivers send messages to a unix socket in either *SOCK_STREAM* or *SOCK_DGRAM* mode.

Both drivers have a single required argument specifying the name of the socket to connect to.

Declaration:

```
unix-stream(filename [options]);  
unix-dgram(filename [options]);
```

The *unix-stream()* and *unix-dgram()* destinations have the following options:

Name	Type	Default	Description
flags()	no_multi_line, syslog-protocol	empty set	<p>Flags influence the behavior of the driver.</p> <p>The <i>no-multi-line</i> flag disables line-breaking in the messages; the entire message is converted to a single line.</p> <p>The <i>syslog-protocol</i> flag instructs the driver to format the messages according to the new IETF syslog protocol standard. If this flag is enabled, macros used for the message have effect only for the text of the message, the message header is formatted to the new standard. Note that this flag is not needed for the <i>syslog</i> driver.</p>
flush_lines()	number	Use global setting.	Specifies how many lines are flushed to a destination at a time. Syslog-ng waits for this number of lines to accumulate and sends them off in a single batch. Setting this number high increases throughput as fully filled frames are sent to the network, but also increases message latency. The latency can be limited by the use of the <i>flush_timeout</i> option.
flush_timeout()	time in milliseconds	Use global setting.	Specifies the time syslog-ng waits for lines to accumulate in its output buffer. See the <i>flush_lines</i> option for more information.
frac_digits()	number	0	The syslog-ng application can store fractions of a second in the timestamps according to the ISO8601 format.. The <i>frac_digits()</i> parameter specifies the number of digits stored. The digits storing the fractions are padded by zeros if the original timestamp of the message specifies only seconds. Fractions can always be stored for the time the message was received. Note that syslog-ng can add the fractions to non-ISO8601 timestamps as well.
fsync()	yes or no	no	Forces an <i>fsync()</i> call on the destination fd after each write. Note: enabling this option may seriously degrade performance.
log_fifo_size()	number	Use global setting.	The number of entries in the output buffer (output fifo).
keep-alive()	yes or no	yes	Specifies whether connections to destinations should be closed when syslog-ng is restarted (upon the receipt of a SIGHUP signal). Note that this applies to the client (destination) side of the syslog-ng connections, server-side (source) connections are always reopened after receiving a HUP signal unless the <i>keep-alive</i> option is enabled for the source. When the <i>keep-alive</i> option is enabled, syslog-ng saves the contents of the output queue of the destination when receiving a HUP signal, reducing the risk of losing messages.
so_broadcast()	yes or no	no	This option controls the <i>SO_BROADCAST</i> socket option required to make syslog-ng send messages to a broadcast address. See the socket(7) manual page for details.



Name	Type	Default	Description
so_keepalive()	yes or no	no	Enables keep-alive messages, keeping the socket open. This only effects TCP and UNIX-stream sockets. See the socket(7) manual page for details.
so_rcvbuf()	number	0	Specifies the size of the socket receive buffer in bytes. See the socket(7) manual page for details.
so_sndbuf()	number	0	Specifies the size of the socket send buffer in bytes. See the socket(7) manual page for details.
suppress()	seconds	0 (disabled)	If several identical log messages would be sent to the destination without any other messages between the identical messages (for example, an application repeated an error message ten times), syslog-ng can suppress the repeated messages and send the message only once, followed by the <i>Last message repeated n times.</i> message. The parameter of this option specifies the number of seconds syslog-ng waits for identical messages.
template()	string	A format conforming to the default logfile format.	Specifies a template defining the logformat to be used in the destination. Macros are described in Section 8.5, “Macros” (p. 210). Please note that for network destinations it might not be appropriate to change the template as it changes the on-wire format of the syslog protocol which might not be tolerated by stock syslog receivers (like <i>syslogd</i> or syslog-ng itself). For network destinations make sure the receiver can cope with the custom format defined.
template_escape()	yes or no	no	Turns on escaping ' and " in templated output files. This is useful for generating SQL statements and quoting string contents so that parts of the log message are not interpreted as commands to the SQL server.
throttle()	number	0	Sets the maximum number of messages sent to the destination per second. Use this output-rate-limiting functionality only when using disk-buffer as well to avoid the risk of losing messages. Specifying 0 or a lower value sets the output limit to unlimited.
time_zone()	timezone offset in seconds	unspecified	Convert timestamps to the timezone specified by this option. If this option is not set then the original timezone information in the message is used.
ts_format()	rfc3164, bsd, rfc3339, iso	rfc3164	Override the global timestamp format (set in the global <i>ts_format()</i> parameter) for the specific destination. See also Section 7.7, “A note on timezones and timestamps” (p. 140).

Table 8.15. Options for unix-stream() and unix-dgram()

**Example 8.24. Using the unix-stream() driver**

```
destination d_unix_stream { unix-stream("/var/run/logs"); };
```


8.2.9. usertty()

This driver writes messages to the terminal of a logged-in user.

The `usertty()` driver has a single required argument, specifying a username who should receive a copy of matching messages.

Declaration:

```
usertty(username);
```

The `usertty()` does not have any further options nor does it support templates.



Example 8.25. Using the `usertty()` driver

```
destination d_usertty { usertty("root"); };
```

8.3. Log path flags

Flags influence the behavior of syslog-ng, and the way it processes messages. The following flags may be used in the log paths, as described in Section 3.5, “Log paths” (p. 53).

Flag	Description
catchall	This flag means that the source of the message is ignored, only the filters are taken into account when matching messages. A log statement using the <code>catchall</code> flag processes every message that arrives to any of the defined sources.
fallback	This flag makes a log statement 'fallback'. Fallback log statements process messages that were not processed by other, 'non-fallback' log statements.
final	This flag means that the processing of log messages processed by the log statement ends here, other log statements appearing later in the configuration file will not process the messages processed by the log statement labeled as 'final'. Note that this does not necessarily mean that matching messages will be stored only once, as there can be matching log statements processed prior the current one.
flow-control	Enables flow-control to the log path, meaning that syslog-ng will stop reading messages from the sources of this log statement if the destinations are not able to process the messages at the required speed. If disabled, syslog-ng will drop messages if the destination queues are full. If enabled, syslog-ng will only drop messages if the destination queues/window sizes are improperly sized.

Table 8.16. Log statement flags



Warning

The `final`, `fallback`, and `catchall` flags apply only for the top-level log paths, they have no effect on embedded log paths.

**Example 8.26. Using log path flags**

Let's suppose that you have two hosts (*myhost_A* and *myhost_B*) that run two applications each (*application_A* and *application_B*), and you collect the log messages to a central syslog-ng server. On the server, you create two log paths:

- one that processes only the messages sent by *myhost_A*; and
- one that processes only the messages sent by *application_A*.

This means that messages sent by *application_A* running on *myhost_A* will be processed by both log paths, and the messages of *application_B* running on *myhost_B* will not be processed at all.

- If you add the *final* flag to the first log path, then only this log path will process the messages of *myhost_A*, so the second log path will receive only the messages of *application_A* running on *myhost_B*.
- If you create a third log path that includes the *fallback* flag, it will process the messages not processed by the first two log paths, in this case, the messages of *application_B* running on *myhost_B*.
- Adding a fourth log path with the *catchall* flag would process every message received by the syslog-ng server.

```
log { source(s_localhost); destination(d_file); flags(catchall); };
```

8.4. Filter functions

The following functions may be used in the filter statement, as described in Section 3.6, “Filters” (p. 57).



Name	Synopsis	Description
facility()	facility(facility[,facility])	Match messages having one of the listed facility code. An alternate syntax permits the use of an arbitrary facility codes.
facility()	facility(<numeric facility code>)	An alternate syntax for <i>facility</i> permitting the use of an arbitrary facility code. Facility codes 0-23 are predefined and can be referenced by their usual name. Facility codes above 24 are not defined but can be used by this alternate syntax.
filter()	filter(filtername)	Call another filter rule and evaluate its value.
host()	host(regex)	Match messages by using a regular expression against the hostname field of log messages.
level() or priority()	level(pri[,pri1..pri2[,pri3]])	Match messages based on priority.
match()	match(regex)	Match a regular expression to the headers and the message itself (i.e., the values returned by the <i>MSGHDR</i> and <i>MSG</i> macros). Note that in syslog-ng version 2.1 and earlier, the <i>match()</i> filter was applied only to the text of the message, excluding the headers. This functionality has been moved to the <i>message()</i> filter. To limit the scope of the match to a specific part of the message (identified with a macro), use the <i>match(regex value("MACRO"))</i> syntax. Do not include the \$ sign in the parameter of the <i>value()</i> option.
message()	message(regex)	Match a regular expression to the text of the log message, excluding the headers (i.e., the value returned by the <i>MSG</i> macros). Note that in syslog-ng version 2.1 and earlier, this functionality was performed by the <i>match()</i> filter.
netmask()	netmask(ip/mask)	Select only messages sent by a host whose IP address belongs to the specified IP subnet. Note that this filter checks the IP address of the last-hop relay (the host that actually sent the message to syslog-ng), not the contents of the <i>HOST</i> field of the message.
program()	program(regex)	Match messages by using a regular expression against the program name field of log messages.
source()	string	Select messages of a source statement. This filter can be used in embedded log statements if the parent statement contains multiple source groups — only messages originating from the selected source group are sent to the destination of the embedded log statement.

Table 8.17. Filter functions in syslog-ng

The *host()*, *match()*, and *program()* filter functions accept regular expressions as parameters. See Section 8.8, “Regular expressions” (p. 226) for details.

The *level()* filter accepts the following levels: *emerg*, *alert*, *crit*, *err*, *warning*, *notice*, *info*, *debug*.



The *facility()* filter accepts both the name and the numerical code of the facility or the importance level. The syslog-ng application recognizes the following facilities: (Note that some of these facilities are available only on specific platforms.)

Numerical Code	Facility name	Facility
0	kern	kernel messages
1	user	user-level messages
2	mail	mail system
3	daemon	system daemons
4	auth	security/authorization messages
5	syslog	messages generated internally by syslogd
6	lpr	line printer subsystem
7	news	network news subsystem
8	uucp	UUCP subsystem
9	cron	clock daemon
10	auth	security/authorization messages
11	ftp	FTP daemon
12		NTP subsystem
13		log audit
14		log alert
15	cron	clock daemon
16-23	local0..local7	locally used facilities (local0-local7)

Table 8.18. *syslog Message Facilities recognized by the facility() filter*

8.5. Macros

Certain parts of syslog-ng (e.g., destination filenames and message content templates) can refer to one or more macros, which get expanded as a message is processed. The table below summarizes the macros available in syslog-ng.



Note

See Section 5.6, “Customizing the message format” (p. 107) for the macros available in the syslog-ng Agent for Windows application.

Macros can be included by prefixing the macro name with a *\$* sign, just like in Bourne compatible shells. Regarding braces around macro names, the following two formats are equivalent *"\$MSG"* and *"\${MSG}"*.

Default values for macros can also be specified by appending the *: -* characters and the default value to the macro, e.g.,



```
${HOST:-default_hostname}
```



Name	Description
BSDTAG	Facility/priority information in the format used by the FreeBSD syslogd: a priority number followed by a letter that indicates the facility. The priority number can range from 0 to 7. The facility letter can range from A to Y, where A corresponds to facility number zero (LOG_KERN), B corresponds to facility 1 (LOG_USER), etc.
DATE, R_DATE, S_DATE	Date of the message using the BSD-syslog style timestamp format (month/day/hour/minute/second, each expressed in two digits). This is the original syslog time stamp without year information, e.g.: <i>Jun 13 15:58:00</i> .
DAY, R_DAY, S_DAY	The day the message was sent.
FACILITY	The facility that sent the message.
FULLDATE, R_FULLDATE, S_FULLDATE	A nonstandard format for the date of the message using the same format as <i>DATE</i> , but including the year as well, e.g.: <i>2006 Jun 13 15:58:00</i> .
FULLHOST	The full FQDN of the host name chain (without trimming chained hosts), including the domain name. To use this macro, make sure that the <i>keep_hostname()</i> option is enabled.
FULLHOST_FROM	FQDN of the host that sent the message to syslog-ng as resolved by syslog-ng using DNS. If the message traverses several hosts, this is the last host in the chain. To use this macro, make sure that the <i>keep_hostname()</i> option is enabled.
HOUR, R_HOUR, S_HOUR	The hour of day the message was sent.
HOST	The name of the source host where the message originates from. If the message traverses several hosts and the <i>chain_hostnames()</i> option is on, the first host in the chain is used. To use this macro, make sure that the <i>keep_hostname()</i> option is enabled.
HOST_FROM	Name of the host that sent the message to syslog-ng, as resolved by syslog-ng using DNS. If the message traverses several hosts, this is the last host in the chain. To use this macro, make sure that the <i>keep_hostname()</i> option is enabled.
ISODATE, R_ISODATE, S_ISODATE	Date of the message in the ISO 8601 compatible standard timestamp format (<i>yyyy-mm-ddThh:mm:ss+-ZONE</i>), e.g.: <i>2006-06-13T15:58:00.123+01:00</i> . If possible, it is recommended to use <i>ISODATE</i> for timestamping. Note that syslog-ng can produce fractions of a second (e.g., milliseconds) in the timestamp by using the <i>frac_digits()</i> global or per-destination option.
MIN, R_MIN, S_MIN	The minute the message was sent.
MONTH, R_MONTH, S_MONTH	The month the message was sent as a decimal value, prefixed with a zero if smaller than 10.
MONTH_ABBREV, R_MONTH_ABBREV, S_MONTH_ABBREV	The English abbreviation of the month name (3 letters).



Name	Description
M O N T H _ N A M E , R _ M O N T H _ N A M E , S_MONTH_NAME	The English name of the month name.
M O N T H _ W E E K , R _ M O N T H _ W E E K , S_MONTH_WEEK	The number of the week in the given month (1-5).
MSG or MESSAGE	Text contents of the log message without the program name and pid. Note that this has changed in syslog-ng version 3.0; in earlier versions this macro included the program name and the pid. In syslog-ng 3.0, the <i>MSG</i> macro became equivalent with the <i>MSGONLY</i> macro. The program name and the pid together are available in the <i>MSGHDR</i> macro.
MSGHDR	The name and the pid of the program that sent the log message in <i>PROGRAM:PID</i> format. Includes a trailing whitespace. Note that the macro returns an empty value if both the program and pid fields of the message are empty.
MSGONLY	Message contents without the program name or pid.
PID	The PID of the program sending the message.
PRI	The priority and facility encoded as a 2 or 3 digit decimal number as it is present in syslog messages.
PRIORITY or LEVEL	The priority of the message.
PROGRAM	The name of the program sending the message. Note that the content of the <i>\$PROGRAM</i> variable may not be completely trusted as it is provided by the client program that constructed the message.
SDATA.SDID.SDNAME	The syslog-ng application automatically parses the STRUCTURED-DATA part of IETF-syslog messages, which can be referenced in macros. For example, if a log message contains the following structured data: <i>[exampleSDID@0 iut="3" eventSource="Application" eventId="1011"] [examplePriority@0 class="high"]</i> you can use macros like: <i>\${SDATA.EXAMPLE.EVENTSOURCE}</i> .
SEC, R_SEC, S_SEC	The second the message was sent.
SEQNUM	The sequence number of the message is a unique identifier of the message between the end-points. The syslog-ng client calculates this number when processing a new message from a local source; it is not calculated for relayed messages. The sequence number increases for every message, and is not lost even if syslog-ng is reloaded or restarted. The sequence number is a part of every message that uses the new IETF-syslog protocol (<i>.SDATA.meta.sequenceId</i>), and can be added to BSD-syslog messages using this macro.
SOURCEIP	IP address of the host that sent the message to syslog-ng (I.e. the IP address of the host in the <i>FULLHOST_FROM</i> macro.) Please note that when a message traverses several relays, this macro contains the IP of the last relay.



Name	Description
STAMP, R_STAMP, S_STAMP	A timestamp formatted according to the <code>_ts_format()</code> global or per-destination option.
TAG	The priority and facility encoded as a 2 digit hexadecimal number.
TZ, R_TZ, S_TZ	Equivalent to TZOFFSET, used to mean the time zone name abbreviation in syslog-ng 1.6.x.
TZOFFSET, R_TZOFFSET, S_TZOFFSET	The time-zone as hour offset from GMT; e.g.: <code>-07:00</code> . In syslog-ng 1.6.x this used to be <code>-0700</code> but as <i>ISODATE</i> requires the colon it was added to <i>TZOFFSET</i> as well.
UNIXTIME, R_UNIXTIME, S_UNIXTIME	Standard unix timestamp, represented as the number of seconds since <code>1970-01-01T00:00:00</code> .
YEAR, R_YEAR, S_YEAR	The year the message was sent.
WEEK, R_WEEK, S_WEEK	The week number of the year, prefixed with a zero for the first nine week of the year. (The first Monday in the year marks the first week.)
WEEK_ABBREV, R_WEEK_ABBREV, S_WEEK_ABBREV	The English abbreviation of the name of the day (3 letters).
WEEK_DAY, R_WEEK_DAY, S_WEEK_DAY	The day of the week as a numerical value (1-7).
WEEKDAY, R_WEEKDAY, S_WEEKDAY	The 3-letter name of the day of week the message was sent, e.g. <i>Thu</i> .
WEEK_DAY_NAME, R_WEEK_DAY_NAME, S_WEEK_DAY_NAME	The English name of the day.

Table 8.19. Available macros

8.6. Message parsers

The following sections provide reference for the parsers available in syslog-ng.

- To segment structured messages like comma-separated values, see Section 8.6.1, “CSV parsers” (p. 214).
- To classify messages using a pattern database, see Section 8.6.2, “Pattern databases” (p. 218).

8.6.1. CSV parsers

The syslog-ng application can separate parts of log messages (i.e., the contents of the `$MSG` macro) to named fields (columns). These fields act as user-defined macros that can be referenced in message templates, file- and tablenames, etc.

To create a parser, define the columns of the message, the delimiter or separator characters, and optionally the characters that are used to escape the delimiter characters (quote-pairs).



Declaration:

```
parser parser_name {  
    csv-parser(column1, column2, ...)  
    delimiters()  
    quote-pairs()  
};
```

Column names work like macros. Always use a prefix to identify the columns of the parsers, e.g., *MYPARSER1.COLUMN1*, *MYPARSER2.COLUMN2*, etc. Column names starting with a dot (e.g., *.HOST*) are reserved for use by syslog-ng.



Name	Synopsis	Description
csv-parser	<code>csv-parser(columns("PARSER.COLUMN1", "PARSER.COLUMN2", ...))</code>	Specifies the type of parser to use, and the name of the columns to separate messages to. Currently only the <i>csv-parser</i> is implemented, which can separate columns based on delimiter characters and strings.
delimiters	<code>delimiters("<delimiter_characters>")</code>	The character that separates the columns in the message.
flags()	<code>drop-invalid, escape-none, escape-backslash, escape-double-char, greedy, strip-whitespace</code>	<p>When the <i>drop-invalid</i> option is set, the parser does not process messages that have less columns than defined in the parser. Using this option practically turns the parser into a special filter, that matches messages that have the predefined number of columns (using the specified delimiters).</p> <p>The <i>escape-none</i>, <i>escape-backslash</i>, <i>escape-double-char</i> flags set the escaping rules used by the parser.</p> <p>The <i>greedy</i> option assigns the remainder of the message to the last column, regardless of the delimiter characters set. You can use this option to process messages where the number of columns varies.</p> <p>The <i>strip-whitespace</i> flag removes trailing whitespaces from the beginning and the end of the columns.</p>
quote-pairs()	<code>quote-pairs("<quote_pairs>")</code>	List quote-pairs between single quotes. Delimiter characters enclosed between quote characters are ignored. Note that the beginning and ending quote character does not have to be identical, e.g., [} can also be a quote-pair.
template()	<code>template("\${<macro-name>}")</code>	The macro that contains the part of the message that the parser will process. It can also be a macro created by a previous parser of the log path. By default, this is empty and the parser processes the entire message.

Table 8.20. Parser parameters

**Example 8.27. Segmenting hostnames separated with a dash**

The following example separates hostnames like *example-1* and *example-2* into two parts.

```
parser p_hostname_segmentation {
    csv-parser(columns("HOSTNAME.NAME", "HOSTNAME.ID"))
    delimiters("-")
    flags(escape-none)
    template("${HOST}");
};
destination d_file { file("/var/log/messages-${HOSTNAME.NAME:-examplehost}"); };
log { source(s_local); parser(p_hostname_segmentation); destination(d_file);};
```



Example 8.28. Parsing Apache log files

The following parser processes the log of Apache web servers and separates them into different fields. Apache log messages can be formatted like:

```
"%h %l %u %t \"%r\" %>s %b \"%{Referer}i\" \"%{User-Agent}i\" %T %v"
```

Here is a sample message:

```
192.168.1.1 - - [31/Dec/2007:00:17:10 +0100] "GET /cgi-bin/example.cgi HTTP/1.1"
200 2708 "-" "curl/7.15.5 (i4 86-pc-linux-gnu) libcurl/7.15.5 OpenSSL/0.9.8c
zlib/1.2.3 libidn/0.6.5" 2 example.balabit
```

To parse such logs, the delimiter character is set to a single whitespace (`delimiters(" ")`). Whitespaces between quotes and brackets are ignored (`quote-pairs('\"\"[]')`).

```
parser p_apache {
    csv-parser(columns("APACHE.CLIENT_IP", "APACHE.IDENT_NAME", "APACHE.USER_NAME",
        "APACHE.TIMESTAMP", "APACHE.REQUEST_URL", "APACHE.REQUEST_STATUS",
        "APACHE.CONTENT_LENGTH", "APACHE.REFERER", "APACHE.USER_AGENT",
        "APACHE.PROCESS_TIME", "APACHE.SERVER_NAME")
        flags(escape-double-char,strip-whitespace)
        delimiters(" ")
        quote-pairs('\"\"[]'))
    };
};
```

The results can be used for example to separate log messages into different files based on the `APACHE.USER_NAME` field. If the field is empty, the `nouser` name is assigned.

```
log { source(s_local);
    parser(p_apache); destination(d_file); };
};
destination d_file { file("/var/log/messages-${APACHE.USER_NAME:-nouser}"); };
```



Example 8.29. Segmenting a part of a message

The following example splits the timestamp of a parsed Apache log message into separate fields.

```
parser p_apache_timestamp {
    csv-parser(columns("APACHE.TIMESTAMP.DAY", "APACHE.TIMESTAMP.MONTH",
        "APACHE.TIMESTAMP.YEAR", "APACHE.TIMESTAMP.HOUR", "APACHE.TIMESTAMP.MIN",
        "APACHE.TIMESTAMP.MIN", "APACHE.TIMESTAMP.ZONE")
        delimiters("/: ")
        flags(escape-none)
        template("${APACHE.TIMESTAMP}"));
    };
log { source(s_local);
    log { parser(p_apache); parser(p_apache_timestamp); destination(d_file); };
};
```



Example 8.30. Adding the end of the message to the last column

If the *greedy* option is enabled, the syslog-ng application adds the not-yet-parsed part of the message to the last column, ignoring any delimiter characters that may appear in this part of the message.

For example, you receive the following comma-separated message: *example 1, example2, example3*, and you segment it with the following parser:

```
csv_parser(columns("COLUMN1", "COLUMN2", "COLUMN3") delimiters(",");
```

The `COLUMN1`, `COLUMN2`, and `COLUMN3` variables will contain the strings *example1*, *example2*, and *example3*, respectively. If the message looks like *example 1, example2, example3, some more information*, then any text appearing after the third comma (i.e., *some more information*) is not parsed, and possibly lost if you use only the variables to reconstruct the message (for example, to send it to different columns of an SQL table).



Using the *greedy* flag will assign the remainder of the message to the last column, so that the *COLUMN1*, *COLUMN2*, and *COLUMN3* variables will contain the strings *example1*, *example2*, and *example3, some more information*.

```
csv_parser(columns("COLUMN1", "COLUMN2", "COLUMN3") delimiters(",") flags(greedy));
```

8.6.2. Pattern databases

8.6.2.1. Using pattern parsers

Pattern parsers attempt to parse a part of the message using rules specific to the type of the parser. Parsers are enclosed between @ characters. The syntax of parsers is the following:

- a beginning @ character;
- the type of the parser written in capitals;
- optionally a name;
- parameters of the parser, if any;
- a closing @ character.



Example 8.31. Pattern parser syntax

A simple parser:

```
@STRING@
```

A named parser:

```
@STRING:myparser_name@
```

A named parser with a parameter:

```
@STRING:myparser_name:*@
```

A parser with a parameter, but without a name:

```
@STRING::*@
```

The following parsers are available:

- **@DOUBLE@**: A floating-point number that may contain a dot (.) character.
- **@ESTRING@**: This parser has a required parameter that acts as the stopcharacter: the parser parses everything until it find the stopcharacter. For example to stop by the next " character, use **@ESTRING::"@**.
- **@IPv4@**: Parses an IPv4 IP address (numbers separated with a maximum of 3 dots).
- **@IPv6@**: Parses any valid IPv6 IP address.
- **@IPvANY@**: Parses any IP address.
- **@NUMBER@**: A sequence of decimal (0-9) numbers (e.g., 1, 0687, etc.). Note that if the number starts with the 0x characters, it is parsed as a hexadecimal number, but only if at least one valid character follows 0x.



- **@QSTRING@**: Parse a string between the quote characters specified as parameter. Note that the quote character can be different at the beginning and the end of the quote, e.g.: `@QSTRING::"@` parses everything between two quotation marks ("), while `@QSTRING:<>@` parses from an opening bracket to the closing bracket.
- **@STRING@**: A sequence of alphanumeric characters (0-9, A-z), not including any whitespace. Optionally, other accepted characters can be listed as parameters (e.g., to parse a complete sentence, add the whitespace as parameter, like: `@STRING:: @`). Note that the @ character cannot be a parameter, nor can line-breaks or tabs.

Patterns and literals can be mixed together. For example, to parse a message that begins with the `Host:` string followed by an IP address (e.g., `Host: 192.168.1.1`), the following pattern can be used: `Host:@IPv4@`.



Note

Note that using parsers is a CPU-intensive operation. Use the `ESTRING` and `QSTRING` parsers whenever possible, as these can be processed much faster than the other parsers.



Example 8.32. Using the `STRING` and `ESTRING` parsers

For example, if the message is `user=joe96 group=somegroup, @STRING:mytext:@` parses only to the first non-alphanumeric character (=), parsing only `user`. `@STRING:mytext:=@` parses the equals sign as well, and proceeds to the next non-alphanumeric character (the whitespace), resulting in `user=joe96` being parsed. `@STRING:mytext:=@` will parse the whitespace as well, and proceed to the next non-alphanumeric non-equals sign non-whitespace character, resulting in `user=joe96 group=somegroup`.

Of course, usually it is better to parse the different values separately, like this: `"user=@STRING:user@ group=@STRING:group@"`.

If the username or the group may contain non-alphanumeric characters, you can either include these in the second parameter of the parser (as shown at the beginning of this example), or use an `ESTRING` parser to parse the message till the next whitespace: `"user=@ESTRING:user: @group=@ESTRING:group: @"`.

8.6.2.2. Filtering messages based on classification

The results of message classification and parsing can be used in custom filters and file and database templates as well. There are two built-in macros in `syslog-ng` that allow you to use the results of the classification: the `.classifier.class` macro contains the class assigned to the message (e.g., violation, security, or unknown), while the `.classifier.rule_id` macro contains the identifier of the message pattern that matched the message.



Example 8.33. Using classification results for filtering messages

To filter on a specific message class, create a filter that checks the `.classifier.class` macro, and use this filter in a log statement.

```
filter fi_class_violation {
    match("violation"
        value(".classifier.class")
        type("string")
    );
};

log {
    source(s_all);
    parser(pattern_db);
    filter(fi_class_violation);
}
```



```
destination(di_class_violation);
};
```

Filtering on the *unknown* class selects messages that did not match any rule of the pattern database. Routing these messages into a separate file allows you to periodically review new or unknown messages.

To filter on messages matching a specific classification rule, create a filter that checks the `.classifier_rule_id` macro. The unique identifier of the rule (e.g., `e1e9c0d8-13bb-11de-8293-000c2922ed0a`) is the *id* attribute of the rule in the XML database.

```
filter fi_class_rule {
    match("e1e9c0d8-13bb-11de-8293-000c2922ed0a"
        value(".classifier_rule_id")
        type("string")
    );
};
```

The message-segments parsed by the pattern parsers can also be used as macros as well. To accomplish this, you have to add a name to the parser, and then you can use this name as a macro that refers to the parsed value of the message.



Example 8.34. Using pattern parsers as macros

For example, you want to parse messages of an application that look like `"Transaction: <type>."`, where `<type>` is a string that has different values (e.g., refused, accepted, incomplete, etc.). To parse these messages, you can use the following pattern:

```
'Transaction: @ESTRING:..@'
```

Here the `@ESTRING@` parser parses the message until the next full stop character. To use the results in a filter or a filename template, include a name in the parser of the pattern, e.g.:

```
'Transaction: @ESTRING:TRANSACTIONTYPE:..@'
```

After that, add a custom template to the logpath that uses this template. For example, to select every *accepted* transaction, use the following custom filter in the log path:

```
match("accepted" value("TRANSACTIONTYPE"));
```



Note

The above macros can be used in database columns and filename templates as well, if you create custom templates for the destination or logspace.

Use a consistent naming scheme for your macros, for example, `APPLICATIONNAME_MACRONAME`.

8.6.2.3. Creating pattern databases

Pattern databases are XML files that contain rules describing the message patterns.

The XML schema of the V1 pattern database used in syslog-ng OSE and PE 3.0.X is the following:



Warning

This is an experimental database format that will change in the future releases of syslog-ng. When the new format will be released, an upgrading script will be available to convert the existing databases to the new format. Note that the sample pattern databases available at the BalaBit website already use the new format (dubbed V2).



- **<patterndb>**: The container element of the pattern database. For example:

```
<patterndb version='1' pub_date='2008-08-25'>
```

- *version*: The schema version of the pattern database. The current version is 2.
- *pubdate*: The publication date of the XML file.
- **<program>**: A container element to group log patterns for an application or program. For example:

```
<program name='su' id='480de478-d4a6-4a7f-bea4-0c0245d361e1'>
```

<patterndb> element may contain any number of **<program>** elements.

- *name*: The name of the application. Note that the function of this attribute is to make the database more readable, syslog-ng uses the *<pattern>* element to identify the applications sending log messages.
- *id*: A unique ID of the application, for example, the md5 sum of the *name* attribute.
- **pattern**: The name of the application — syslog-ng matches this value to the \$PROGRAM header of the syslog message to find the rulesets applicable to the syslog message. This element is also called *program pattern*. E.g.,

```
<pattern>su</pattern>
```

- **description**: OPTIONAL — A description of the ruleset or the application.
- **url**: OPTIONAL — An URL referring to further information about the ruleset or the application.
- **<rules>**: A container element for the rules of the ruleset.
 - **<rule>**: An element containing message patterns and how a message that matches these patterns is classified. For example:

```
<rule provider='balabit'
id='f57196aa-75fd-11dd-9bba-001e6806451b' class='violation'>
```



Note

If the following characters appear in the message, they must be escaped in the rule as follows:

- @: Use `&@`, e.g., `user&@example.com`
- <: Use `<`;
- >: Use `>`;
- &: Use `&`;

The **<rules>** element may contain any number of **<rule>** elements.

- *provider*: The provider of the rule. This is used to distinguish between who supplied the rule; i.e., if it has been created by BalaBit, or added to the xml by a local user.
- *id*: The globally unique ID of the rule.



- *class*: The class of the rule — syslog-ng assigns this class to the messages matching a pattern of this rule.
- **<pattern>**: A pattern describing a log message. This element is also called *message pattern*. For example:

```
<pattern>+ ??? root-</pattern>
```



Example 8.35. A V1 pattern database containing a single rule

The following pattern database contains a single rule that matches log messages of the *PF* packet-filtering application. A sample log message looks like:

```
PF: DROP filter/INPUT IN=eth0 OUT= MAC=00:1A:4B:80:90:C9:00:1A:4B:80:90:C6
SRC=192.168.155.11 DST=192.168.155.1 LEN=60 TOS=0x10 PREC=0x00 TTL=64 ID=51939
DF PROTO=TCP SPT=34407 DPT=80 WINDOW=32792 RES=0x00 SYN URGP=0
```

The following is a simple pattern database containing a matching rule.

```
<patterndb version='1' pub_date='2009-04-17'>
  <program name='PF'>
    <pattern>PF</pattern>
    <rule id='1' class='pf'>
      <pattern>@STRING:PF.VERDICT@ @STRING:PF.CHAIN:@
IN=@STRING:PF.IN_IFACE@ OUT= MAC=@STRING:PF.MAC::@ SRC=@IPV4:PF.SRC_IP@
DST=@IPV4:PF.DST_IP@ LEN=@NUMBER:PF.PKT_LEN@ TOS=@STRING:PF.TOS@
PREC=@STRING:PF.PREC@ TTL=@NUMBER:PF.TTL@ ID=@NUMBER:PF.ID@ DF
PROTO=@STRING:PF.PROTO@ SPT=@NUMBER:PF.SRC_PORT@ DPT=@NUMBER:PF.DST_PORT@
WINDOW=@NUMBER:PF.TCP_WINDOW@ RES=@STRING:PF.RES@ SYN
URGP=@NUMBER:PF.TCP_URGP@</pattern>
    </rule>
  </program>
</patterndb>
```

Note that the rule uses macros that refer to parts of the message, for example, you can use the *\$PF.DST_IP* macro refer to the destination IP address of the logged connection+.

The following scheme describes the V2 format of the pattern database. This format is used by the syslog-ng Store Box (SSB) appliance version 1.0.x (see <http://www.balabit.com/network-security/syslog-ng/log-server-appliance/> for details).

For a sample database containing only a single pattern, see Example 8.36, “A V2 pattern database containing a single rule” (p. 224).

- **<patterndb>**: The container element of the pattern database. For example:

```
<patterndb version='2' pub_date='2008-08-25'>
```

- *version*: The schema version of the pattern database. The current version is 2.
- *pubdate*: The publication date of the XML file.
- **<ruleset>**: A container element to group log patterns for an application or program. For example:

```
<ruleset name='su' id='480de478-d4a6-4a7f-bea4-0c0245d361e1'>
```

A **<patterndb>** element may contain any number of **<ruleset>** elements.



- *name*: The name of the application. Note that the function of this attribute is to make the database more readable, syslog-ng uses the `<pattern>` element to identify the applications sending log messages.
- *id*: A unique ID of the application, for example, the md5 sum of the *name* attribute.
- **description**: OPTIONAL — A description of the ruleset or the application.
- **url**: OPTIONAL — An URL referring to further information about the ruleset or the application.
- **pattern**: The name of the application — syslog-ng matches this value to the \$PROGRAM header of the syslog message to find the rulesets applicable to the syslog message. This element is also called *program pattern*. E.g.,

```
<pattern>su</pattern>
```


Note

If the `<pattern>` element of a ruleset is not specified, -ng will use this ruleset as a fallback ruleset: it will apply the ruleset to messages that have an empty PROGRAM header, or if none of the program patterns matched the PROGRAM header of the incoming message.

- **<rules>**: A container element for the rules of the ruleset.
- **<rule>**: An element containing message patterns and how a message that matches these patterns is classified. For example:

```
<rule provider='balabit'
  id='f57196aa-75fd-11dd-9bba-001e6806451b'
  class='violation'>
```


Note

If the following characters appear in the message, they must be escaped in the rule as follows:

- @: Use @@, e.g., `user@@example.com`
- <: Use <
- >: Use >
- &: Use &

The **<rules>** element may contain any number of **<rule>** elements.

- *provider*: The provider of the rule. This is used to distinguish between who supplied the rule; i.e., if it has been created by BalaBit, or added to the xml by a local user.
- *id*: The globally unique ID of the rule.
- *class*: The class of the rule — syslog-ng assigns this class to the messages matching a pattern of this rule.



- **<patterns>**: An element containing the patterns of the rule. If a **<patterns>** element contains multiple **<pattern>** elements, the class of the **<rule>** is assigned to every syslog message matching any of the patterns.
- **<pattern>**: A pattern describing a log message. This element is also called *message pattern*. For example:

```
<pattern>+ ??? root-</pattern>
```

- **description**: OPTIONAL — A description of the pattern or the log message matching the pattern.
- **urls**: OPTIONAL — An element containing one or more URLs referring to further information about the patterns or the matching log messages.
 - **url**: OPTIONAL — An URL referring to further information about the patterns or the matching log messages.
- **tags**: OPTIONAL — An element containing custom keywords (tags) about the rules. The tags can be used to label specific events (e.g., user logons).
- **tag**: OPTIONAL — A keyword or tags applied to messages matching the rule. For example:

```
<tags><tag>UserLogin</tag></tags>
```



Example 8.36. A V2 pattern database containing a single rule

The following pattern database contains a single rule that matches a log message of the *ssh* application. A sample log message looks like:

```
Accepted password for sampleuser from 10.50.0.247 port 42156 ssh2
```

The following is a simple pattern database containing a matching rule.

```
<patterndb version='2' pub_date='2009-04-17'>
  <ruleset name='ssh' id='123456678'>
    <pattern>ssh</pattern>
    <rules>
      <rule provider='me' id='182437592347598' class='system'>
        <patterns>
          <pattern>Accepted @QSTRING:SSH.AUTH_METHOD: @
for@QSTRING:SSH_USERNAME: @from\ @QSTRING:SSH_CLIENT_ADDRESS: @port
@NUMBER:SSH_PORT_NUMBER:@ ssh2</pattern>
        </patterns>
      </rule>
    </rules>
  </ruleset>
</patterndb>
```

Note that the rule uses macros that refer to parts of the message, for example, you can use the `$SSH_USERNAME` macro refer to the username used in the connection.



8.7. Rewriting messages

The syslog-ng application can rewrite parts of log messages: it can search and replace text, and also set a specific field to a specified value. Rewriting messages is often used in conjunction with message parsing Section 8.6, “Message parsers” (p. 214).

To create replace a part of the log message, define the string or regular expression to replace, the string to replace the original text (macros can be used as well), and the field of the message that the rewrite rule should process. Substitution rules can operate on any value available via macros, e.g., HOST, MESSAGE, PROGRAM, or any user-defined macros created using parsers (see Section 8.6, “Message parsers” (p. 214) for details). Substitution rules use the following syntax:

Declaration:

```
rewrite <name_of_the_rule>
{subst("<string or regular expression to find>", "<replacement string>",
value(<field name>) type() flags());};
```

The `type()` and `flags()` options are optional. The `type()` specifies the type of regular expression to use; while the `flags()` are the flags of the regular expressions (see Section 8.8, “Regular expressions” (p. 226) for details):

Name	Description
posix	Use POSIX regular expressions. If the <code>type()</code> parameter is not specified, syslog-ng uses POSIX regular expressions by default.
pcre	Use PCRE regular expressions. This is available only if syslog-ng was compiled with the <code>--enable-pcre</code> option. Execute the syslog-ng -V command to list the options supported by your binary.
string	Match the strings literally, without regular expression support.

Table 8.21. Rewrite rule types



Example 8.37. Using substitution rules

The following example replaces the first occurrence of the string `IP` in the text of the message with the string `IP-Address`.

```
rewrite r_rewrite_subst{subst("IP", "IP-Address", value("MESSAGE"))};
```

To replace every occurrence, use:

```
rewrite r_rewrite_subst{subst("IP", "IP-Address", value("MESSAGE"),
flags("global"))};
```

Multiple substitution rules are applied sequentially; the following rules replace the first occurrence of the string `IP` with the string `IP-Addresses`.

```
rewrite r_rewrite_subst{subst("IP", "IP-Address", value("MESSAGE"));
subst("Address", "Addresses", value("MESSAGE"))};
```

To set a field of the message to a specific value, define the string to include in the message, and the field where it should be included. Setting a field can operate on any value available via macros, e.g., HOST, MESSAGE, PROGRAM, or any user-defined macros created using parsers (see Section 8.6, “Message parsers” (p. 214) for details). Note that this operation completely replaces any previous value of that field. Use the following syntax:



Declaration:

```
rewrite <name_of_the_rule>
{set("<string to include>", value(<field name>) flags());};
```



Example 8.38. Setting message fields to a particular value

The following example sets the HOST field of the message to *myhost*.

```
rewrite r_rewrite_set{set("myhost", value("HOST"));};
```

8.8. Regular expressions

Filters and substitution rewrite rules can use regular expressions. By default, syslog-ng uses POSIX-style regular expressions, but if compiled with the `--enable-pcre` option, Perl Compatible Regular Expressions can be used as well. The regular expressions can use up to 255 regexp matches (`${1}` ... `${255}`), but only from the last filter. For case-insensitive searches, use the `flags("ignore-case")` option.

To use Perl Compatible Regular Expressions (PCRE), add the `type("pcre")` option after the regular expression. Note that PCRE expressions can be used only if syslog-ng was compiled with the `--enable-pcre` option. Execute the **syslog-ng -V** command to list the options supported by your binary.

Posix regular expressions have the following flag options:

Name	Description
global	Usable only in rewrite rules; match for every occurrence of the expression, not only the first one.
ignore-case	Disable case-sensitivity.
newline	
utf8	Use UTF-8 matching.

Table 8.22. Posix options

PCRE regular expressions have the following flag options:



Example 8.39. Using Posix regular expressions

```
filter f_message { message("keyword" flags("utf8" "ignore-case") );
```



Name	Description
global	Usable only in rewrite rules; match for every occurrence of the expression, not only the first one.
ignore-case	Disable case-sensitivity.
newline	Process the string until the end of the string, ignoring newline characters.
nobackref	Do not store back references for the matches — improves performance.
unicode	Use Unicode support for UTF-8 matches.
utf8	Use UTF-8 matching.

Table 8.23. PCRE options**Example 8.40. Using PCRE regular expressions**

```
rewrite r_rewrite_subst
{subst("a*", "?", field("message") type("pcre") flags("utf8"
"global")); }
```

8.9. Global options

The following options can be specified in the options statement, as described in Section 3.11, “Configuring global syslog-ng options” (p. 67).



Name	Accepted values	Default	Description
bad_hostname()	regular expression	no	A regexp containing hostnames which should not be handled as hostnames.
chain_hostnames()	yes or no	no	Enable or disable the chained hostname format.
check_hostname()	yes or no	no	Enable or disable checking whether the hostname contains valid characters.
create_dirs()	yes or no	no	Enable or disable directory creation for destination files.
dir_group()	groupid	root	The default group for newly created directories.
dir_owner()	userid	root	The default owner of newly created directories.
dir_perm()	permission value	0700	The default permission for newly created directories.
dns_cache()	yes or no	yes	Enable or disable DNS cache usage.
dns_cache_expire()	number	3600	Number of seconds while a successful lookup is cached.
dns_cache_expire_failed()	number	60	Number of seconds while a failed lookup is cached.
dns_cache_hosts()	filename	unset	Name of a file in <code>/etc/hosts</code> format that contains static IP->hostname mappings. Use this option to resolve hostnames locally without using a DNS. Note that any change to this file triggers a reload in syslog-ng and is instantaneous.
dns_cache_size()	number	1007	Number of hostnames in the DNS cache.
time_zone()	timezone offset in seconds	unset	Convert timestamps to the timezone specified by this option. If this option is not set then the original timezone information in the message is used.
flush_lines()	number	0	Specifies how many lines are flushed to a destination at a time. Syslog-ng waits for this number of lines to accumulate and sends them off in a single batch. Setting this number high increases throughput as fully filled frames are sent to the network, but also increases message latency. The latency can be limited by the use of the <code>flush_timeout</code> option.
flush_timeout()	time in milliseconds	10000	Specifies the time syslog-ng waits for lines to accumulate in its output buffer. See the <code>flush_lines()</code> option for more information.
group()	groupid	root	The default group of output files. By default, syslog-ng changes the privileges of accessed files (e.g., <code>/dev/null</code>) to <code>root.root 0600</code> . To disable modifying privileges, use this option with the <code>-1</code> value.



Name	Accepted values	Default	Description
keep_hostname()	yes or no	no	Enable or disable hostname rewriting. Enable this option to use hostname-related macros. This option can be specified globally, and per-source as well. The local setting of the source overrides the global option if available. When relaying messages, enable this option on the syslog-ng server and also on every relay, otherwise syslog-ng will treat incoming messages as if they were sent by the last relay.
keep_timestamp()	yes or no	yes	Specifies whether syslog-ng should accept the timestamp received from the sending application or client. If disabled, the time of reception will be used instead. This option can be specified globally, and per-source as well. The local setting of the source overrides the global option if available.
log_fifo_size()	number	100	The number of lines fitting to the output queue
log_msg_size()	number	8192	Maximum length of a message in bytes.
normalize_hostnames()	yes or no	no	Normalize hostnames, which currently translates to converting them to lower case. (requires 1.9.9)
owner()	userid	root	The default owner of output files. By default, syslog-ng changes the privileges of accessed files (e.g., <code>/dev/null</code>) to <code>root.root 0600</code> . To disable modifying privileges, use this option with the <code>-1</code> value.
mark()	number	1200	An alias for the obsolete <code>mark_freq()</code> option, retained for compatibility with syslog-ng version 1.6.x.
mark_freq()	number	1200	The number of seconds between two <code>MARK</code> messages. <code>MARK</code> messages are generated when there was no message traffic to inform the receiver that the connection is still alive. Note that only local messages postpone the sending of the <code>MARK</code> message, relayed messages do not. If set to zero (<code>0</code>), no <code>MARK</code> messages are sent.
perm()	permission value	0600	The default permission for output files. By default, syslog-ng changes the privileges of accessed files (e.g., <code>/dev/null</code>) to <code>root.root 0600</code> . To disable modifying privileges, use this option with the <code>-1</code> value.
recv_time_zone()	time offset (e.g.: <code>+03:00</code>)	local timezone	Specifies the time zone associated with the incoming messages, if not specified otherwise in the message or in the source driver. See also Section 2.5, “Timezone handling” (p. 12) and Section 7.7, “A note on timezones and timestamps” (p. 140) for details.
send_time_zone()	time offset (e.g.: <code>+03:00</code>)	local timezone	Specifies the time zone associated with the messages sent by syslog-ng, if not specified otherwise in the message or in the destination driver. See Section 2.5, “Timezone handling” (p. 12) for details.
stats_freq()	number	600	The period between two <code>STATS</code> messages in seconds. <code>STATS</code> are log messages sent by syslog-ng, containing statistics about dropped log messages. Set to <code>0</code> to disable the <code>STATS</code> messages.



Name	Accepted values	Default	Description
stats_level()	0, 1, or 2	0	Specifies the detail of statistics syslog-ng collects about the processed messages. Level 0 collects only statistics about the sources and destinations; level 1 contains details about the different connections and log files, but has a slight memory overhead; while level 2 can display detailed statistics based on message parameters (e.g., hostname). Note that level 2 increases the memory requirements and CPU load.
sync() or sync_freq() (DEPRECATED)	number	0	Obsolete aliases for <i>flush_lines()</i>
time_reap()	number	60	The time to wait in seconds before an idle destination file is closed.
time_reopen()	number	60	The time to wait in seconds before a dead connection is reestablished.
time_sleep()	number	0	The time to wait in milliseconds between each invocation of the <i>poll()</i> iteration.
ts_format()	rfc3164, bsd, rfc3339, iso	rfc3164	Specifies the timestamp format used when syslog-ng itself formats a timestamp and nothing else specifies a format (e.g.: <i>STAMP</i> macros, internal messages, messages without original timestamps). See also Section 7.7, “A note on timezones and timestamps” (p. 140).
use_dns()	yes, no, persist_only	yes	Enable or disable DNS usage. The <i>persist_only</i> option attempts to resolve hostnames locally from file (e.g., from <i>/etc/hosts</i>). syslog-ng blocks on DNS queries, so enabling DNS may lead to a Denial of Service attack. To prevent DoS, protect your syslog-ng network endpoint with firewall rules, and make sure that all hosts which may get to syslog-ng are resolvable. This option can be specified globally, and per-source as well. The local setting of the source overrides the global option if available.
use_fqdn()	yes or no	no	Add Fully Qualified Domain Name instead of short hostname. This option can be specified globally, and per-source as well. The local setting of the source overrides the global option if available.



Name	Accepted values	Default	Description
use_time_recvd() (DEPRECATED)	yes or no	no	<p>This option controls how the time related macros are expanded in filename and content templates. If set to yes, then the non-prefixed versions of the time related macros (e.g.: <i>HOURL</i> instead of <i>R_HOURL</i> and <i>S_HOURL</i>) refer to the time when the message was received, otherwise it refers to the timestamp which is in the message.</p> <p>NOTE: The timestamps in the messages are generated by the originating host and might not be accurate.</p> <p>This option is deprecated as many users assumed that it controls the timestamp as it is written to logfiles/destinations, which is not the case. To change how messages are formatted, specify a content-template referring to the appropriate prefixed (<i>S_</i> or <i>R_</i>) time macro.</p>

Table 8.24. List of global options supported in syslog-ng

8.10. TLS options

The syslog-ng application is able to encrypt incoming and outgoing syslog message flows using SSL/TLS, if the TCP transport protocol (the *tcp()* or *tcp6()* sources or destination) is used.



Note

The format of the TLS connections used by syslog-ng is similar to using syslog-ng and stunnel, but the source IP information is not lost.

To encrypt connections, use the *tls()* option in the source and destination statements.

The *tls()* option can include the following settings:



Name	Accepted values	Default	Description
<code>ca_dir()</code>	Directory name	none	Name of a directory, that contains a set of trusted CA certificates in PEM format. The CA certificates are named after the 32-bit hash of the subject's name. This naming can be created using the <code>c_rehash</code> utility in the <code>openssl</code> package.
<code>cert_file()</code>	File name	none	Name of a file, that contains an X.509 certificate in PEM format, suitable as a TLS certificate, mat
<code>crl_dir()</code>	Directory name	none	Name of a directory that contains the Certificate Revocation Lists for trusted CAs. Similarly to <code>ca_dir()</code> , the 32-bit hash of the name of the issuing CAs as filenames. The extension of the files must be <code>.crl</code> .
<code>key_file()</code>	File name	none	Name of a file, that contains an unencrypted private key in PEM format, suitable as a TLS key.
<code>peer_verify()</code>	optional-trusted optional-untrusted required-trusted required-untrusted	required-trusted	Verification method of the peer, the four possible values is a combination of two properties of va is required to provide a certificate (required or optional prefix), and whether the certificate provided is trusted or not. For untrusted certificates only the existence of the certificate is checked, but it does not have to be signed by a trusted CA. For required-trusted certificates, the certificate must be signed by a trusted CA and not expired. For required-untrusted certificates, the certificate must be signed by a trusted CA and not expired, but it does not have to be signed by a trusted CA.
<code>trusted_dn()</code>	list of accepted distinguished names	none	To accept connections only from hosts using certain certificates signed by the trusted CAs, list the distinguished names of the accepted certificates in this parameter. E.g., using <code>trusted_dn("*", O=Example Inc, ST=Some-State)</code> will accept only certificates issued for the <i>Example Inc</i> organization in <i>Some-State</i> state.
<code>trusted_keys()</code>	list of accepted SHA-1 fingerprints	none	To accept connections only from hosts using certain certificates having specific SHA-1 fingerprints, list the fingerprints of the accepted certificates in this parameter. E.g., using <code>trusted_keys("SHA1:00:EF:ED:A4:CE:00:D1:14:A4:AB:43:00:EF:00:91:00:00:00:00:00:00", "SHA1:0C:42:00:3E:B2:60:36:64:00:E2:83:F0:80:46:AD:00:A8:9D:00:00:00:00")</code> will accept only certificates having the specified SHA-1 fingerprints.

Table 8.25. List of TLS options



Note

When using the `trusted_keys()` and `trusted_dn()` parameters, note the following:

- First, the `trusted_keys()` parameter is checked. If the fingerprint of the peer is listed, the certificate validation is performed.



- If the fingerprint of the peer is not listed in the `trusted_keys()` parameter, the `trusted_dn()` parameter is checked. If the DN of the peer is not listed in the `trusted_dn()` parameter, the authentication of the peer fails and the connection is closed.



Appendix 1. The syslog-ng manual pages



Name

syslog-ng — syslog-ng system logger application

Synopsis

syslog-ng [options]

Description

NOTE: This manual page covers both editions of syslog-ng: syslog-ng Open Source Edition and the commercial syslog-ng Premium Edition. Features that are only included in the Premium Edition are marked with an asterisk (*). For details, see the official syslog-ng website: <http://www.balabit.com/network-security/syslog-ng/>.

This manual page is only an abstract; for the complete documentation of syslog-ng, see **The syslog-ng Administrator Guide** [<http://www.balabit.com/support/documentation/>].

The syslog-ng application is a flexible and highly scalable system logging application. Typically, syslog-ng is used to manage log messages and implement centralized logging, where the aim is to collect the log messages of several devices on a single, central log server. The different devices - called syslog-ng clients - all run syslog-ng, and collect the log messages from the various applications, files, and other *sources*. The clients send all important log messages to the remote syslog-ng server, where the server sorts and stores them.

Options

--cfgfile <file> or -f <file>	Use the specified configuration file.
--chroot <dir> or -C <dir>	Change root to the specified directory after reading the configuration file. The directory must be set up accordingly. Note that it is not possible to reload the syslog-ng configuration after chrooting.
--debug or -d	Start syslog-ng in debug mode.
--enable-core	Enable syslog-ng to write core files in case of a crash to help support and debugging.
--fd-limit	Set the minimal number of required file descriptors (fd-s); this sets how many files syslog-ng can keep open simultaneously. Default value: 4096. Note that this does not override the global ulimit setting of the host.
--foreground or -F	Do not daemonize, run in the foreground.
--group <group> or -g <group>	Switch to the specified group after initializing the configuration file.
--help or -h	Display a brief help message.
--no-caps	Run syslog-ng as root, without capability-support. This is the default behavior. On Linux, it is possible to run syslog-ng as non-root with capability-support if syslog-ng was compiled with the



	<code>--enable-linux-caps</code> option enabled. (Execute syslog-ng --version to display the list of enabled build parameters.)
<code>--persist-file <persist-file> or -R <persist-file></code>	Set the path and name of the <code>syslog-ng.persist</code> file where the persistent options and data are stored.
<code>--pidfile <pidfile> or -p <pidfile></code>	Set path to the PID file where the pid of the main process is stored.
<code>--process-mode <pidfile></code>	Sets how to run syslog-ng: in the <i>foreground</i> (mainly used for debugging), in the <i>background</i> as a daemon, or in <i>safe-background</i> mode, when syslog-ng restarts itself if it crashes. By default, syslog-ng runs in <i>safe-background</i> mode.
<code>--qdisk-dir <path> or -Q <path></code>	Specify the location of the file used for disk-based buffering. By default, this file is located at <code>/var/lib/syslog-ng/</code> .
<code>--stderr or -e</code>	Log internal messages of syslog-ng to stderr. Mainly used for debugging purposes in conjunction with the <code>--foreground</code> option.
<code>--syntax-only or -s</code>	Verify that the configuration file is syntactically correct and exit.
<code>--user <user> or -u <user></code>	Switch to the specified user after initializing the configuration file (and optionally chrooting). Note that it is not possible to reload the syslog-ng configuration if the specified user has no privilege to create the <code>/dev/log</code> file.
<code>--verbose or -v</code>	Enable verbose logging used to troubleshoot syslog-ng.
<code>--version or -V</code>	Display version number and compilation information.

Files

`/opt/syslog-ng/etc/syslog-ng/`

`/opt/syslog-ng/etc/syslog-ng/syslog-ng.conf`

See also

[syslog-ng.conf\(5\)](#)

The syslog-ng Administrator Guide [<http://www.balabit.com/support/documentation/>]

If you experience any problems or need help with syslog-ng, visit the **syslog-ng mailing list** [<https://lists.balabit.hu/mailman/listinfo/syslog-ng>]

Author

This manual page was written by the BalaBit Documentation Team <documentation@balabit.com>.



Copyright

Copyright © 2007 BalaBit IT Security Ltd. Published under the Creative Commons Attribution-Noncommercial-No Derivative Works (by-nc-nd) 3.0 license. See Appendix 4, *Creative Commons Attribution Non-commercial No Derivatives (by-nc-nd) License* (p. 256) for details. The latest version is always available at <http://www.balabit.com/support/documentation>.



Name

syslog-ng.conf — syslog-ng configuration file

Synopsis

syslog-ng.conf

Description

NOTE: This manual page covers both editions of syslog-ng: syslog-ng Open Source Edition and the commercial syslog-ng Premium Edition. Features that are only included in the Premium Edition are marked with an asterisk (*). For details, see the official syslog-ng website: <http://www.balabit.com/network-security/syslog-ng/>.

This manual page is only an abstract; for the complete documentation of syslog-ng, see **The syslog-ng Administrator Guide** [<http://www.balabit.com/support/documentation/>].

The syslog-ng application is a flexible and highly scalable system logging application. Typically, syslog-ng is used to manage log messages and implement centralized logging, where the aim is to collect the log messages of several devices on a single, central log server. The different devices - called syslog-ng clients - all run syslog-ng, and collect the log messages from the various applications, files, and other *sources*. The clients send all important log messages to the remote syslog-ng server, where the server sorts and stores them.

The syslog-ng application reads incoming messages and forwards them to the selected *destinations*. The syslog-ng application can receive messages from files, remote hosts, and other *sources*.

Log messages enter syslog-ng in one of the defined sources, and are sent to one or more *destinations*.

Sources and destinations are independent objects; *log paths* define what syslog-ng does with a message, connecting the sources to the destinations. A log path consists of one or more sources and one or more destinations; messages arriving to a source are sent to every destination listed in the log path. A log path defined in syslog-ng is called a *log statement*.

Optionally, log paths can include *filters*. Filters are rules that select only certain messages, for example, selecting only messages sent by a specific application. If a log path includes filters, syslog-ng sends only the messages satisfying the filter rules to the destinations set in the log path.

Configuring syslog-ng

Global objects (e.g., sources, destinations, log paths, or filters) are defined in the syslog-ng configuration file. Object definitions consist of the following elements:

- *Type of the object*: One of *source*, *destination*, *log*, *filter*, *parser*, *rewrite* rule, or *template*.
- *Identifier of the object*: A unique name identifying the object. When using a reserved word as an identifier, enclose the identifier in quotation marks.

**Tip**

Use identifiers that refer to the type of the object they identify. For example, prefix source objects with `s_`, destinations with `d_`, and so on.

- *Parameters*: The parameters of the object, enclosed in braces `{parameters}`.
- *Semicolon*: Object definitions end with a semicolon (`;`).

The syntax is summarized as follows:

The syntax of log statements is as follows:

```
log {
    source(s1); source(s2); ...
    optional_element(filter1|parser1|rewrite1);
optional_element(filter2|parser2|rewrite2);...
    destination(d1); destination(d2); ...
    flags(flag1[, flag2...]);
};
```

The following log statement sends all messages arriving to the localhost to a remote server.

```
source s_localhost { tcp(ip(127.0.0.1) port(1999) ); };
destination d_tcp { tcp("10.1.2.3" port(1999); localport(999)); };
log { source(s_localhost); destination(d_tcp); };
```

The syslog-ng application has a number of global options governing DNS usage, the timestamp format used, and other general points. Each option may have parameters, similarly to driver specifications. To set global options, add an option statement to the syslog-ng configuration file using the following syntax:

```
options { option1(params); option2(params); ... };
```

The sources, destinations, and filters available in syslog-ng are listed below. For details, see **The syslog-ng Administrator Guide** [<http://www.balabit.com/support/documentation/>].



Name	Description
<code>internal()</code>	Messages generated internally in syslog-ng.
<code>file()</code>	Opens the specified file and reads messages.
<code>pipe()</code> , <code>fifo</code>	Opens the specified named pipe and reads messages.
<code>program()</code>	Opens the specified application and reads messages from its standard output.
<code>sun-stream()</code> , <code>sun-streams()</code>	Opens the specified <i>STREAMS</i> device on Solaris systems and reads incoming messages.
<code>syslog()</code>	Listens for incoming messages using the new IETF-standard syslog protocol .
<code>tcp()</code> , <code>tcp6()</code>	Listens on the specified TCP port for incoming messages using the BSD-syslog protocol over IPv4 and IPv6 networks, respectively.
<code>udp()</code> , <code>udp6()</code>	Listens on the specified UDP port for incoming messages using the BSD-syslog protocol over IPv4 and IPv6 networks, respectively.
<code>unix-dgram()</code>	Opens the specified unix socket in <i>SOCK_DGRAM</i> mode and listens for incoming messages.
<code>unix-stream()</code>	Opens the specified unix socket in <i>SOCK_STREAM</i> mode and listens for incoming messages.

Table 1.1. Source drivers available in syslog-ng

Name	Description
<code>file()</code>	Writes messages to the specified file.
<code>logstore()*</code>	Writes messages to the specified binary logstore file. *Available only in syslog-ng Premium Edition.
<code>fifo()</code> , <code>pipe()</code>	Writes messages to the specified named pipe.
<code>program()</code>	Forks and launches the specified program, and sends messages to its standard input.
<code>sql()</code>	Sends messages into an SQL database. In addition to the standard syslog-ng packages, the <i>sql()</i> destination requires database-specific packages to be installed. Refer to the section appropriate for your platform in Chapter 4, <i>Installing syslog-ng</i> (p. 78).
<code>syslog()</code>	Sends messages to the specified remote host using the IETF-syslog protocol . The IETF standard supports message transport using the UDP, TCP, and TLS networking protocols.
<code>tcp()</code> and <code>tcp6()</code>	Sends messages to the specified TCP port of a remote host using the BSD-syslog protocol over IPv4 and IPv6, respectively.
<code>udp()</code> and <code>udp6()</code>	Sends messages to the specified UDP port of a remote host using the BSD-syslog protocol over IPv4 and IPv6, respectively.
<code>unix-dgram()</code>	Sends messages to the specified unix socket in <i>SOCK_DGRAM</i> style (BSD).
<code>unix-stream()</code>	Sends messages to the specified unix socket in <i>SOCK_STREAM</i> style (Linux).
<code>usertty()</code>	Sends messages to the terminal of the specified user, if the user is logged in.

Table 1.2. Destination drivers available in syslog-ng



Name	Synopsis	Description
facility()	facility(facility[,facility])	Match messages having one of the listed facility code. An alternate syntax permits the use an arbitrary facility codes.
facility()	facility(<numeric facility code>)	An alternate syntax for <i>facility</i> permitting the use of an arbitrary facility code. Facility codes 0-23 are predefined and can be referenced by their usual name. Facility codes above 24 are not defined but can be used by this alternate syntax.
filter()	filter(filtername)	Call another filter rule and evaluate its value.
host()	host(regex)	Match messages by using a regular expression against the hostname field of log messages.
level() or priority()	level(pri[,pri1..pri2[,pri3]])	Match messages based on priority.
match()	match(regex)	Match a regular expression to the headers and the message itself (i.e., the values returned by the <i>MSGHDR</i> and <i>MSG</i> macros). Note that in syslog-ng version 2.1 and earlier, the <i>match()</i> filter was applied only to the text of the message, excluding the headers. This functionality has been moved to the <i>message()</i> filter. To limit the scope of the match to a specific part of the message (identified with a macro), use the <i>match(regex value("MACRO"))</i> syntax. Do not include the \$ sign in the parameter of the <i>value()</i> option.
message()	message(regex)	Match a regular expression to the text of the log message, excluding the headers (i.e., the value returned by the <i>MSG</i> macros). Note that in syslog-ng version 2.1 and earlier, this functionality was performed by the <i>match()</i> filter.
netmask()	netmask(ip/mask)	Select only messages sent by a host whose IP address belongs to the specified IP subnet. Note that this filter checks the IP address of the last-hop relay (the host that actually sent the message to syslog-ng), not the contents of the <i>HOST</i> field of the message.
program()	program(regex)	Match messages by using a regular expression against the program name field of log messages.
source()	string	Select messages of a source statement. This filter can be used in embedded log statements if the parent statement contains multiple source groups — only messages originating from the selected source group are sent to the destination of the embedded log statement.

Table 1.3. Filter functions in syslog-ng

Files

/opt/syslog-ng/etc/syslog-ng/

/opt/syslog-ng/etc/syslog-ng/syslog-ng.conf



See also

syslog-ng(8)

The syslog-ng Administrator Guide [<http://www.balabit.com/support/documentation/>]

If you experience any problems or need help with syslog-ng, visit the **syslog-ng mailing list** [<https://lists.balabit.hu/mailman/listinfo/syslog-ng>]

Author

This manual page was written by the BalaBit Documentation Team <documentation@balabit.com>.

Copyright

Copyright © 2007 BalaBit IT Security Ltd. Published under the Creative Commons Attribution-Noncommercial-No Derivative Works (by-nc-nd) 3.0 license. See Appendix 4, *Creative Commons Attribution Non-commercial No Derivatives (by-nc-nd) License* (p. 256) for details. The latest version is always available at <http://www.balabit.com/support/documentation>.



Appendix 2. BalaBit syslog-ng Premium Edition License contract

2.1. SUBJECT OF THE License CONTRACT

This License Contract is entered into by and between BalaBit and Licensee and sets out the terms and conditions under which Licensee and/or Licensee's Authorized Subsidiaries may use the BalaBit syslog-ng Premium Edition product.

2.2. DEFINITIONS

In this License Contract, the following words shall have the following meanings:

Company name: *BalaBit IT Security Ltd.*

Registered office: H-1115 Budapest, Bártfai u. 54. Hungary

Company registration number: 01-09-687127

Tax number: HU11996468

2.3. Words and expressions

Annexed Software	Any third party software that is a not a BalaBit Product contained in the install media of the BalaBit Product.
Authorized Subsidiary	Any subsidiary organization: (i) in which Licensee possesses more than fifty percent (50%) of the voting power and (ii) which is located within the Territory.
BalaBit Product	Any software, hardware or service Licensed, sold, or provided by BalaBit including any installation, education, support and warranty services, with the exception of the Annexed Software.
License Contract	The present BalaBit syslog-ng Premium Edition License Contract.
Product Documentation	Any documentation referring to the BalaBit syslog-ng Premium Edition or any module thereof, with special regard to the administration guide, the product description, the installation guide, user guides and manuals.
Number of Log Source Hosts	Number of all host or server computers including virtual machines, active or passive networking devices from which a BalaBit syslog-ng Premium Edition server receives log messages.



Protected Objects	The entire BalaBit syslog-ng Premium Edition software including all of its modules, all the related Product Documentation; the source code, the structure of the databases, all registered information reflecting the structure of the BalaBit syslog-ng Premium Edition and all the adaptation and copies of the Protected Objects that presently exist or that are to be developed in the future, or any product falling under the copyright of BalaBit.
BalaBit syslog-ng Premium Edition	The BalaBit Product designed for aggregate, filter, format, send or receive over network or local connection the log messages and eventlogs as defined by the Product Description.
Warranty Period	The period of twelve (12) months from the date of delivery of the BalaBit syslog-ng Premium Edition to Licensee.
Territory	The countries or areas specified above in respect of which Licensee shall be entitled to install and/or use BalaBit syslog-ng Premium Edition.
End-user Certificate	The document signed by Licensor which contains a) identification data of Licensee; b) configuration of BalaBit syslog-ng Premium Edition, maximum Number of Log Source Hosts and designation of Licensed modules thereof; c) designation of the Territory; d) declaration of the parties on accepting the terms and conditions of this License Contract; and e) declaration of Licensee that is in receipt of the install media and the hardware appliance.

2.4. LICENSE GRANTS AND RESTRICTIONS

For the BalaBit syslog-ng Premium Edition licensed under this License Contract, BalaBit grants to Licensee a non-exclusive, non-transferable, perpetual license to use such BalaBit Product under the terms and conditions of this License Contract and the applicable End-user Certificate.

Licensee shall use the BalaBit syslog-ng Premium Edition in the configuration and in the quantities specified in the End-user Certificate within the Territory.

On the install media (CD-ROM) all modules of the BalaBit syslog-ng Premium Edition will be presented, however, Licensee shall not be entitled to use any module which was not Licensed to it. Access rights to modules and maximum Number of Log Source Hosts are controlled by an “electronic key” accompanying the BalaBit syslog-ng Premium Edition.

Licensee shall be entitled to make one back-up copy of the install media containing the BalaBit syslog-ng Premium Edition.

Licensee shall make available the Protected Objects at its disposal solely to its own employees and those of the Authorized Subsidiaries.

Licensee shall take all reasonable steps to protect BalaBit’s rights with respect to the Protected Objects with special regard and care to protecting it from any unauthorized access.



Licensee shall, in 5 working days, properly answer the queries of BalaBit referring to the actual usage conditions of the BalaBit syslog-ng Premium Edition that may differ or allegedly differs from the License conditions.

Licensee shall not modify the BalaBit syslog-ng Premium Edition in any way, with special regard to the functions inspecting the usage of the software. Licensee shall install the code permitting the usage of the BalaBit syslog-ng Premium Edition according to the provisions defined for it by BalaBit. Licensee may not modify or cancel such codes. Configuration settings of the BalaBit syslog-ng Premium Edition in accordance with the possibilities offered by the system shall not be construed as modification of the software.

Licensee shall only be entitled to analyze the structure of the BalaBit Products (decompilation or reverse-engineering) if concurrent operation with a software developed by a third party is necessary, and upon request to supply the information required for concurrent operation BalaBit does not provide such information within 60 days from the receipt of such a request.

These user actions are limited to parts of the BalaBit Product which are necessary for concurrent operation. *Any information obtained as a result of applying the previous Section (i) cannot be used for purposes other than concurrent operation with the BalaBit Product; (ii) cannot be disclosed to third parties unless it is necessary for concurrent operation with the BalaBit Product; (iii) cannot be used for the development, production or distribution of a different software which is similar to the BalaBit Product in its form of expression, or for any other act violating copyright.*

For any Annexed Software contained by the same install media as the BalaBit Product, the terms and conditions defined by its copyright owner shall be properly applied. BalaBit does not grant any License rights to any Annexed Software.

Any usage of the BalaBit syslog-ng Premium Edition exceeding the limits and restrictions defined in this License Contract shall qualify as material breach of the License Contract.

The Number of Log Source Hosts shall not exceed the amount defined in the End-user Certificate.

Licensee shall have the right to obtain and use content updates only if Licensee concludes a maintenance contract that includes such content updates, or if Licensee has otherwise separately acquired the right to obtain and use such content updates. This License Contract does not otherwise permit Licensee to obtain and use content updates.

2.5. SUBSIDIARIES

Authorized Subsidiaries may also utilize the services of the BalaBit syslog-ng Premium Edition under the terms and conditions of this License Contract. Any Authorized Subsidiary utilizing any service of the BalaBit syslog-ng Premium Edition will be deemed to have accepted the terms and conditions of this License Contract.

2.6. INTELLECTUAL PROPERTY RIGHTS

Licensee agrees that BalaBit owns all rights, titles, and interests related to the BalaBit syslog-ng Premium Edition and all of BalaBit's patents, trademarks, trade names, inventions, copyrights, know-how, and trade secrets relating to the design, manufacture, operation or service of the BalaBit Products.

The use by Licensee of any of these intellectual property rights is authorized only for the purposes set forth herein, and upon termination of this License Contract for any reason, such authorization shall cease.



The BalaBit Products are Licensed only for internal business purposes in every case, under the condition that such License does not convey any license, expressly or by implication, to manufacture, duplicate or otherwise copy or reproduce any of the BalaBit Products. No other rights than expressly stated herein are granted to Licensee.

Licensee will take appropriate steps with its Authorized Subsidiaries, as BalaBit may request, to inform them of and assure compliance with the restrictions contained in the License Contract.

2.7. TRADE MARKS

BalaBit hereby grants to Licensee the non-exclusive right to use the trade marks of the BalaBit Products in the Territory in accordance with the terms and for the duration of this License Contract.

BalaBit makes no representation or warranty as to the validity or enforceability of the trade marks, nor as to whether these infringe any intellectual property rights of third parties in the Territory.

2.8. NEGLIGENT INFRINGEMENT

In case of negligent infringement of BalaBit's rights with respect to the BalaBit syslog-ng Premium Edition, committed by violating the restrictions and limitations defined by this License Contract, Licensee shall pay liquidated damages to BalaBit. The amount of the liquidated damages shall be twice as much as the price of the BalaBit Product concerned, on BalaBit's current Price List.

2.9. INTELLECTUAL PROPERTY INDEMNIFICATION

BalaBit shall pay all damages, costs and reasonable attorney's fees awarded against Licensee in connection with any claim brought against Licensee to the extent that such claim is based on a claim that Licensee's authorized use of the BalaBit Product infringes a patent, copyright, trademark or trade secret. Licensee shall notify BalaBit in writing of any such claim as soon as Licensee learns of it and shall cooperate fully with BalaBit in connection with the defense of that claim. BalaBit shall have sole control of that defense (including without limitation the right to settle the claim).

If Licensee is prohibited from using any BalaBit Product due to an infringement claim, or if BalaBit believes that any BalaBit Product is likely to become the subject of an infringement claim, BalaBit shall at its sole option, either: (i) obtain the right for Licensee to continue to use such BalaBit Product, (ii) replace or modify the BalaBit Product so as to make such BalaBit Product non-infringing and substantially comparable in functionality or (iii) refund to Licensee the amount paid for such infringing BalaBit Product and provide a pro-rated refund of any unused, prepaid maintenance fees paid by Licensee, in exchange for Licensee's return of such BalaBit Product to BalaBit.

Notwithstanding the above, BalaBit will have no liability for any infringement claim to the extent that it is based upon: (i) modification of the BalaBit Product other than by BalaBit, (ii) use of the BalaBit Product in combination with any product not specifically authorized by BalaBit to be combined with the BalaBit Product or (iii) use of the BalaBit Product in an unauthorized manner for which it was not designed.



2.10. LICENSE FEE

The allowed maximum Number of the Log Source Hosts, the configuration and the modules licensed shall serve as the calculation base of the License fee.

Licensee acknowledges that payment of the License fees is a condition of lawful usage.

License fees do not contain any installation or post charges.

2.11. WARRANTIES

BalaBit warrants that during the Warranty Period, the magnetic or optical media upon which the BalaBit Product is recorded will not be defective under normal use. BalaBit will replace any defective media returned to it, accompanied by a dated proof of purchase, within the Warranty Period at no charge to Licensee. Upon receipt of the allegedly defective BalaBit Product, BalaBit will at its option, deliver a replacement BalaBit Product or BalaBit's current equivalent to Licensee at no additional cost. BalaBit will bear the delivery charges to Licensee for the replacement Product.

In case of installation by BalaBit, BalaBit warrants that during the Warranty Period, the BalaBit syslog-ng Premium Edition, under normal use in the operating environment defined by BalaBit, and without unauthorized modification, will perform in substantial compliance with the Product Documentation accompanying the BalaBit Product, when used on that hardware for which it was installed, in compliance with the provisions of the user manuals and the recommendations of BalaBit. The date of the notification sent to BalaBit shall qualify as the date of the failure. Licensee shall do its best to mitigate the consequences of that failure. If, during the Warranty Period, the BalaBit Product fails to comply with this warranty, and such failure is reported by Licensee to BalaBit within the Warranty Period, BalaBit's sole obligation and liability for breach of this warranty is, at BalaBit's sole option, either: (i) to correct such failure, (ii) to replace the defective BalaBit Product or (iii) to refund the license fees paid by Licensee for the applicable BalaBit Product.

2.12. DISCLAIMER OF WARRANTIES

EXCEPT AS SET OUT IN THIS LICENSE CONTRACT, BALABIT MAKES NO WARRANTIES OF ANY KIND WITH RESPECT TO THE BALABIT SYSLOG-NG PREMIUM EDITION. TO THE MAXIMUM EXTENT PERMITTED BY APPLICABLE LAW, BALABIT EXCLUDES ANY OTHER WARRANTIES, INCLUDING BUT NOT LIMITED TO ANY IMPLIED WARRANTIES OF SATISFACTORY QUALITY, MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE AND NON-INFRINGEMENT OF INTELLECTUAL PROPERTY RIGHTS.

2.13. LIMITATION OF LIABILITY

SOME STATES AND COUNTRIES, INCLUDING MEMBER COUNTRIES OF THE EUROPEAN UNION, DO NOT ALLOW THE LIMITATION OR EXCLUSION OF LIABILITY FOR INCIDENTAL OR CONSEQUENTIAL DAMAGES AND, THEREFORE, THE FOLLOWING LIMITATION OR EXCLUSION MAY NOT APPLY TO THIS LICENSE CONTRACT IN THOSE STATES AND COUNTRIES. TO THE MAXIMUM EXTENT PERMITTED BY APPLICABLE LAW AND REGARDLESS OF WHETHER ANY REMEDY SET OUT IN THIS LICENSE CONTRACT FAILS OF ITS ESSENTIAL PURPOSE, IN NO



EVENT SHALL BALABIT BE LIABLE TO LICENSEE FOR ANY SPECIAL, CONSEQUENTIAL, INDIRECT OR SIMILAR DAMAGES OR LOST PROFITS OR LOST DATA ARISING OUT OF THE USE OR INABILITY TO USE THE BALABIT SYSLOG-NG PREMIUM EDITION EVEN IF BALABIT HAS BEEN ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.

IN NO CASE SHALL BALABIT'S TOTAL LIABILITY UNDER THIS LICENSE CONTRACT EXCEED THE FEES PAID BY LICENSEE FOR THE BALABIT SYSLOG-NG PREMIUM EDITION LICENSED UNDER THIS LICENSE CONTRACT.

2.14. DURATION AND TERMINATION

This License Contract shall come into effect on the date of signature of the End-user Certificate by the duly authorized representative of BalaBit.

Licensee may terminate the License Contract at any time by written notice sent to BalaBit and by simultaneously destroying all copies of the Protected Objects licensed under this License Contract.

BalaBit may terminate this License Contract with immediate effect by written notice to Licensee, if Licensee is in material or persistent breach of the License Contract and either that breach is incapable of remedy or Licensee shall have failed to remedy that breach within 30 days after receiving written notice requiring it to remedy that breach.

2.15. AMENDMENTS

Save as expressly provided in this License Contract, no amendment or variation of this License Contract shall be effective unless in writing and signed by a duly authorized representative of the parties to it.

2.16. WAIVER

The failure of a party to exercise or enforce any right under this License Contract shall not be deemed to be a waiver of that right nor operate to bar the exercise or enforcement of it at any time or times thereafter.

2.17. SEVERABILITY

If any part of this License Contract becomes invalid, illegal or unenforceable, the parties shall in such an event negotiate in good faith in order to agree on the terms of a mutually satisfactory provision to be substituted for the invalid, illegal or unenforceable provision which as nearly as possible validly gives effect to their intentions as expressed in this License Contract.

2.18. NOTICES

Any notice required to be given pursuant to this License Contract shall be in writing and shall be given by delivering the notice by hand, or by sending the same by prepaid first class post (airmail if to an address outside the country of posting) to the address of the relevant party set out in this License Contract or such other address as either party notifies to the other from time to time. Any notice given according to the above procedure shall be deemed to have been given at the time of delivery (if delivered by hand) and when received (if sent by post).



2.19. MISCELLANEOUS

Headings are for convenience only and shall be ignored in interpreting this License Contract.

This License Contract and the rights granted in this License Contract may not be assigned, sublicensed or otherwise transferred in whole or in part by Licensee without BalaBit's prior written consent. This consent shall not be unreasonably withheld or delayed.

An independent third party auditor, reasonably acceptable to BalaBit and Licensee, may upon reasonable notice to Licensee and during normal business hours, but not more often than once each year, inspect Licensee's relevant records in order to confirm that usage of the BalaBit syslog-ng Premium Edition complies with the terms and conditions of this License Contract. BalaBit shall bear the costs of such audit. All audits shall be subject to the reasonable safety and security policies and procedures of Licensee.

This License Contract constitutes the entire agreement between the parties with regard to the subject matter hereof.

Any modification of this License Contract must be in writing and signed by both parties.



Appendix 3. GNU General Public License

Version 2, June 1991

Copyright © 1989, 1991 Free Software Foundation, Inc.

Free Software Foundation, Inc.

51 Franklin Street, Fifth Floor,

Boston, MA 02110-1301

USA

Everyone is permitted to copy and distribute verbatim copies of this license document, but changing it is not allowed.

Version 2, June 1991

3.1. Preamble

The licenses for most software are designed to take away your freedom to share and change it. By contrast, the GNU General Public License is intended to guarantee your freedom to share and change free software - to make sure the software is free for all its users. This General Public License applies to most of the Free Software Foundation's software and to any other program whose authors commit to using it. (Some other Free Software Foundation software is covered by the GNU Library General Public License instead.) You can apply it to your programs, too.

When we speak of free software, we are referring to freedom, not price. Our General Public Licenses are designed to make sure that you have the freedom to distribute copies of free software (and charge for this service if you wish), that you receive source code or can get it if you want it, that you can change the software or use pieces of it in new free programs; and that you know you can do these things.

To protect your rights, we need to make restrictions that forbid anyone to deny you these rights or to ask you to surrender the rights. These restrictions translate to certain responsibilities for you if you distribute copies of the software, or if you modify it.

For example, if you distribute copies of such a program, whether gratis or for a fee, you must give the recipients all the rights that you have. You must make sure that they, too, receive or can get the source code. And you must show them these terms so they know their rights.

We protect your rights with two steps:

1. copyright the software, and
2. offer you this license which gives you legal permission to copy, distribute and/or modify the software.

Also, for each author's protection and ours, we want to make certain that everyone understands that there is no warranty for this free software. If the software is modified by someone else and passed on, we want its recipients to know that what they have is not the original, so that any problems introduced by others will not reflect on the original authors' reputations.



Finally, any free program is threatened constantly by software patents. We wish to avoid the danger that redistributors of a free program will individually obtain patent licenses, in effect making the program proprietary. To prevent this, we have made it clear that any patent must be licensed for everyone's free use or not licensed at all.

The precise terms and conditions for copying, distribution and modification follow.

3.2. TERMS AND CONDITIONS FOR COPYING, DISTRIBUTION AND MODIFICATION

3.2.1. Section 0

This License applies to any program or other work which contains a notice placed by the copyright holder saying it may be distributed under the terms of this General Public License. The “Program”, below, refers to any such program or work, and a “work based on the Program” means either the Program or any derivative work under copyright law: that is to say, a work containing the Program or a portion of it, either verbatim or with modifications and/or translated into another language. (Hereinafter, translation is included without limitation in the term “modification”.) Each licensee is addressed as “you”.

Activities other than copying, distribution and modification are not covered by this License; they are outside its scope. The act of running the Program is not restricted, and the output from the Program is covered only if its contents constitute a work based on the Program (independent of having been made by running the Program). Whether that is true depends on what the Program does.

3.2.2. Section 1

You may copy and distribute verbatim copies of the Program's source code as you receive it, in any medium, provided that you conspicuously and appropriately publish on each copy an appropriate copyright notice and disclaimer of warranty; keep intact all the notices that refer to this License and to the absence of any warranty; and give any other recipients of the Program a copy of this License along with the Program.

You may charge a fee for the physical act of transferring a copy, and you may at your option offer warranty protection in exchange for a fee.

3.2.3. Section 2

You may modify your copy or copies of the Program or any portion of it, thus forming a work based on the Program, and copy and distribute such modifications or work under the terms of [Section 1](#) above, provided that you also meet all of these conditions:

- a. You must cause the modified files to carry prominent notices stating that you changed the files and the date of any change.
- b. You must cause any work that you distribute or publish, that in whole or in part contains or is derived from the Program or any part thereof, to be licensed as a whole at no charge to all third parties under the terms of this License.
- c. If the modified program normally reads commands interactively when run, you must cause it, when started running for such interactive use in the most ordinary way, to print or display an announcement including an appropriate copyright notice and a notice that there is no warranty (or else, saying that you



provide a warranty) and that users may redistribute the program under these conditions, and telling the user how to view a copy of this License. (Exception: If the Program itself is interactive but does not normally print such an announcement, your work based on the Program is not required to print an announcement.)

These requirements apply to the modified work as a whole. If identifiable sections of that work are not derived from the Program, and can be reasonably considered independent and separate works in themselves, then this License, and its terms, do not apply to those sections when you distribute them as separate works. But when you distribute the same sections as part of a whole which is a work based on the Program, the distribution of the whole must be on the terms of this License, whose permissions for other licensees extend to the entire whole, and thus to each and every part regardless of who wrote it.

Thus, it is not the intent of this section to claim rights or contest your rights to work written entirely by you; rather, the intent is to exercise the right to control the distribution of derivative or collective works based on the Program.

In addition, mere aggregation of another work not based on the Program with the Program (or with a work based on the Program) on a volume of a storage or distribution medium does not bring the other work under the scope of this License.

3.2.4. Section 3

You may copy and distribute the Program (or a work based on it, under [Section 2](#) in object code or executable form under the terms of [Sections 1](#) and [2](#) above provided that you also do one of the following:

- a. Accompany it with the complete corresponding machine-readable source code, which must be distributed under the terms of Sections 1 and 2 above on a medium customarily used for software interchange; or,
- b. Accompany it with a written offer, valid for at least three years, to give any third party, for a charge no more than your cost of physically performing source distribution, a complete machine-readable copy of the corresponding source code, to be distributed under the terms of Sections 1 and 2 above on a medium customarily used for software interchange; or,
- c. Accompany it with the information you received as to the offer to distribute corresponding source code. (This alternative is allowed only for noncommercial distribution and only if you received the program in object code or executable form with such an offer, in accord with Subsection b above.)

The source code for a work means the preferred form of the work for making modifications to it. For an executable work, complete source code means all the source code for all modules it contains, plus any associated interface definition files, plus the scripts used to control compilation and installation of the executable. However, as a special exception, the source code distributed need not include anything that is normally distributed (in either source or binary form) with the major components (compiler, kernel, and so on) of the operating system on which the executable runs, unless that component itself accompanies the executable.

If distribution of executable or object code is made by offering access to copy from a designated place, then offering equivalent access to copy the source code from the same place counts as distribution of the source code, even though third parties are not compelled to copy the source along with the object code.



3.2.5. Section 4

You may not copy, modify, sublicense, or distribute the Program except as expressly provided under this License. Any attempt otherwise to copy, modify, sublicense or distribute the Program is void, and will automatically terminate your rights under this License. However, parties who have received copies, or rights, from you under this License will not have their licenses terminated so long as such parties remain in full compliance.

3.2.6. Section 5

You are not required to accept this License, since you have not signed it. However, nothing else grants you permission to modify or distribute the Program or its derivative works. These actions are prohibited by law if you do not accept this License. Therefore, by modifying or distributing the Program (or any work based on the Program), you indicate your acceptance of this License to do so, and all its terms and conditions for copying, distributing or modifying the Program or works based on it.

3.2.7. Section 6

Each time you redistribute the Program (or any work based on the Program), the recipient automatically receives a license from the original licensor to copy, distribute or modify the Program subject to these terms and conditions. You may not impose any further restrictions on the recipients' exercise of the rights granted herein. You are not responsible for enforcing compliance by third parties to this License.

3.2.8. Section 7

If, as a consequence of a court judgment or allegation of patent infringement or for any other reason (not limited to patent issues), conditions are imposed on you (whether by court order, agreement or otherwise) that contradict the conditions of this License, they do not excuse you from the conditions of this License. If you cannot distribute so as to satisfy simultaneously your obligations under this License and any other pertinent obligations, then as a consequence you may not distribute the Program at all. For example, if a patent license would not permit royalty-free redistribution of the Program by all those who receive copies directly or indirectly through you, then the only way you could satisfy both it and this License would be to refrain entirely from distribution of the Program.

If any portion of this section is held invalid or unenforceable under any particular circumstance, the balance of the section is intended to apply and the section as a whole is intended to apply in other circumstances.

It is not the purpose of this section to induce you to infringe any patents or other property right claims or to contest validity of any such claims; this section has the sole purpose of protecting the integrity of the free software distribution system, which is implemented by public license practices. Many people have made generous contributions to the wide range of software distributed through that system in reliance on consistent application of that system; it is up to the author/donor to decide if he or she is willing to distribute software through any other system and a licensee cannot impose that choice.

This section is intended to make thoroughly clear what is believed to be a consequence of the rest of this License.



3.2.9. Section 8

If the distribution and/or use of the Program is restricted in certain countries either by patents or by copyrighted interfaces, the original copyright holder who places the Program under this License may add an explicit geographical distribution limitation excluding those countries, so that distribution is permitted only in or among countries not thus excluded. In such case, this License incorporates the limitation as if written in the body of this License.

3.2.10. Section 9

The Free Software Foundation may publish revised and/or new versions of the General Public License from time to time. Such new versions will be similar in spirit to the present version, but may differ in detail to address new problems or concerns.

Each version is given a distinguishing version number. If the Program specifies a version number of this License which applies to it and “any later version”, you have the option of following the terms and conditions either of that version or of any later version published by the Free Software Foundation. If the Program does not specify a version number of this License, you may choose any version ever published by the Free Software Foundation.

3.2.11. Section 10

If you wish to incorporate parts of the Program into other free programs whose distribution conditions are different, write to the author to ask for permission. For software which is copyrighted by the Free Software Foundation, write to the Free Software Foundation; we sometimes make exceptions for this. Our decision will be guided by the two goals of preserving the free status of all derivatives of our free software and of promoting the sharing and reuse of software generally.

3.2.12. NO WARRANTY Section 11

BECAUSE THE PROGRAM IS LICENSED FREE OF CHARGE, THERE IS NO WARRANTY FOR THE PROGRAM, TO THE EXTENT PERMITTED BY APPLICABLE LAW. EXCEPT WHEN OTHERWISE STATED IN WRITING THE COPYRIGHT HOLDERS AND/OR OTHER PARTIES PROVIDE THE PROGRAM “AS IS” WITHOUT WARRANTY OF ANY KIND, EITHER EXPRESSED OR IMPLIED, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE. THE ENTIRE RISK AS TO THE QUALITY AND PERFORMANCE OF THE PROGRAM IS WITH YOU. SHOULD THE PROGRAM PROVE DEFECTIVE, YOU ASSUME THE COST OF ALL NECESSARY SERVICING, REPAIR OR CORRECTION.

3.2.13. Section 12

IN NO EVENT UNLESS REQUIRED BY APPLICABLE LAW OR AGREED TO IN WRITING WILL ANY COPYRIGHT HOLDER, OR ANY OTHER PARTY WHO MAY MODIFY AND/OR REDISTRIBUTE THE PROGRAM AS PERMITTED ABOVE, BE LIABLE TO YOU FOR DAMAGES, INCLUDING ANY GENERAL, SPECIAL, INCIDENTAL OR CONSEQUENTIAL DAMAGES ARISING OUT OF THE USE OR INABILITY TO USE THE PROGRAM (INCLUDING BUT NOT LIMITED TO LOSS OF DATA OR DATA BEING RENDERED INACCURATE OR LOSSES SUSTAINED BY YOU OR THIRD PARTIES OR A



FAILURE OF THE PROGRAM TO OPERATE WITH ANY OTHER PROGRAMS), EVEN IF SUCH HOLDER OR OTHER PARTY HAS BEEN ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.

END OF TERMS AND CONDITIONS

3.3. How to Apply These Terms to Your New Programs

If you develop a new program, and you want it to be of the greatest possible use to the public, the best way to achieve this is to make it free software which everyone can redistribute and change under these terms.

To do so, attach the following notices to the program. It is safest to attach them to the start of each source file to most effectively convey the exclusion of warranty; and each file should have at least the “copyright” line and a pointer to where the full notice is found.

<one line to give the program's name and a brief idea of what it does.> Copyright (C) <year> <name of author>

This program is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 51 Franklin Street, Fifth Floor, Boston, MA 02110-1301 USA

Also add information on how to contact you by electronic and paper mail.

If the program is interactive, make it output a short notice like this when it starts in an interactive mode:

Gnomovision version 69, Copyright (C) year name of author Gnomovision comes with ABSOLUTELY NO WARRANTY; for details type “show w”. This is free software, and you are welcome to redistribute it under certain conditions; type “show c” for details.

The hypothetical commands “show w” and “show c” should show the appropriate parts of the General Public License. Of course, the commands you use may be called something other than “show w” and “show c”; they could even be mouse-clicks or menu items--whatever suits your program.

You should also get your employer (if you work as a programmer) or your school, if any, to sign a “copyright disclaimer” for the program, if necessary. Here is a sample; alter the names:

Yoyodyne, Inc., hereby disclaims all copyright interest in the program “Gnomovision” (which makes passes at compilers) written by James Hacker.

<signature of Ty Coon>, 1 April 1989 Ty Coon, President of Vice

This General Public License does not permit incorporating your program into proprietary programs. If your program is a subroutine library, you may consider it more useful to permit linking proprietary applications with the library. If this is what you want to do, use the GNU Library General Public License instead of this License.



Appendix 4. Creative Commons Attribution Non-commercial No Derivatives (by-nc-nd) License

THE WORK (AS DEFINED BELOW) IS PROVIDED UNDER THE TERMS OF THIS CREATIVE COMMONS PUBLIC LICENSE ("CCPL" OR "LICENSE"). THE WORK IS PROTECTED BY COPYRIGHT AND/OR OTHER APPLICABLE LAW. ANY USE OF THE WORK OTHER THAN AS AUTHORIZED UNDER THIS LICENSE OR COPYRIGHT LAW IS PROHIBITED. BY EXERCISING ANY RIGHTS TO THE WORK PROVIDED HERE, YOU ACCEPT AND AGREE TO BE BOUND BY THE TERMS OF THIS LICENSE. TO THE EXTENT THIS LICENSE MAY BE CONSIDERED TO BE A CONTRACT, THE LICENSOR GRANTS YOU THE RIGHTS CONTAINED HERE IN CONSIDERATION OF YOUR ACCEPTANCE OF SUCH TERMS AND CONDITIONS.

1. *Definitions*

- a. "Adaptation" means a work based upon the Work, or upon the Work and other pre-existing works, such as a translation, adaptation, derivative work, arrangement of music or other alterations of a literary or artistic work, or phonogram or performance and includes cinematographic adaptations or any other form in which the Work may be recast, transformed, or adapted including in any form recognizably derived from the original, except that a work that constitutes a Collection will not be considered an Adaptation for the purpose of this License. For the avoidance of doubt, where the Work is a musical work, performance or phonogram, the synchronization of the Work in timed-relation with a moving image ("synching") will be considered an Adaptation for the purpose of this License.
- b. "Collection" means a collection of literary or artistic works, such as encyclopedias and anthologies, or performances, phonograms or broadcasts, or other works or subject matter other than works listed in Section 1(f) below, which, by reason of the selection and arrangement of their contents, constitute intellectual creations, in which the Work is included in its entirety in unmodified form along with one or more other contributions, each constituting separate and independent works in themselves, which together are assembled into a collective whole. A work that constitutes a Collection will not be considered an Adaptation (as defined above) for the purposes of this License.
- c. "Distribute" means to make available to the public the original and copies of the Work through sale or other transfer of ownership.
- d. "Licensor" means the individual, individuals, entity or entities that offer(s) the Work under the terms of this License.
- e. "Original Author" means, in the case of a literary or artistic work, the individual, individuals, entity or entities who created the Work or if no individual or entity can be identified, the publisher; and in addition (i) in the case of a performance the actors, singers, musicians, dancers, and other persons who act, sing, deliver, declaim, play in, interpret or otherwise perform literary or artistic works or expressions of folklore; (ii) in the case of a phonogram the producer being the person or legal entity who first fixes the sounds of a performance or other sounds; and, (iii) in the case of broadcasts, the organization that transmits the broadcast.
- f. "Work" means the literary and/or artistic work offered under the terms of this License including without limitation any production in the literary, scientific and artistic domain, whatever may be the mode or form of its expression including digital form, such as a book, pamphlet and other writing;



a lecture, address, sermon or other work of the same nature; a dramatic or dramatico-musical work; a choreographic work or entertainment in dumb show; a musical composition with or without words; a cinematographic work to which are assimilated works expressed by a process analogous to cinematography; a work of drawing, painting, architecture, sculpture, engraving or lithography; a photographic work to which are assimilated works expressed by a process analogous to photography; a work of applied art; an illustration, map, plan, sketch or three-dimensional work relative to geography, topography, architecture or science; a performance; a broadcast; a phonogram; a compilation of data to the extent it is protected as a copyrightable work; or a work performed by a variety or circus performer to the extent it is not otherwise considered a literary or artistic work.

- g. "You" means an individual or entity exercising rights under this License who has not previously violated the terms of this License with respect to the Work, or who has received express permission from the Licensor to exercise rights under this License despite a previous violation.
 - h. "Publicly Perform" means to perform public recitations of the Work and to communicate to the public those public recitations, by any means or process, including by wire or wireless means or public digital performances; to make available to the public Works in such a way that members of the public may access these Works from a place and at a place individually chosen by them; to perform the Work to the public by any means or process and the communication to the public of the performances of the Work, including by public digital performance; to broadcast and rebroadcast the Work by any means including signs, sounds or images.
 - i. "Reproduce" means to make copies of the Work by any means including without limitation by sound or visual recordings and the right of fixation and reproducing fixations of the Work, including storage of a protected performance or phonogram in digital form or other electronic medium.
2. *Fair Dealing Rights.* Nothing in this License is intended to reduce, limit, or restrict any uses free from copyright or rights arising from limitations or exceptions that are provided for in connection with the copyright protection under copyright law or other applicable laws.
3. *License Grant.* Subject to the terms and conditions of this License, Licensor hereby grants You a worldwide, royalty-free, non-exclusive, perpetual (for the duration of the applicable copyright) license to exercise the rights in the Work as stated below:
- a. to Reproduce the Work, to incorporate the Work into one or more Collections, and to Reproduce the Work as incorporated in the Collections; and,
 - b. to Distribute and Publicly Perform the Work including as incorporated in Collections.
- The above rights may be exercised in all media and formats whether now known or hereafter devised. The above rights include the right to make such modifications as are technically necessary to exercise the rights in other media and formats, but otherwise you have no rights to make Adaptations. Subject to 8(f), all rights not expressly granted by Licensor are hereby reserved, including but not limited to the rights set forth in Section 4(d).
4. *Restrictions.* The license granted in Section 3 above is expressly made subject to and limited by the following restrictions:
- a. You may Distribute or Publicly Perform the Work only under the terms of this License. You must include a copy of, or the Uniform Resource Identifier (URI) for, this License with every copy of the Work You Distribute or Publicly Perform. You may not offer or impose any terms on the Work that restrict the terms of this License or the ability of the recipient of the Work to exercise the rights granted to that recipient under the terms of the License. You may not sublicense the Work. You must keep intact all notices that refer to this License and to the disclaimer of warranties with every



copy of the Work You Distribute or Publicly Perform. When You Distribute or Publicly Perform the Work, You may not impose any effective technological measures on the Work that restrict the ability of a recipient of the Work from You to exercise the rights granted to that recipient under the terms of the License. This Section 4(a) applies to the Work as incorporated in a Collection, but this does not require the Collection apart from the Work itself to be made subject to the terms of this License. If You create a Collection, upon notice from any Licensor You must, to the extent practicable, remove from the Collection any credit as required by Section 4(c), as requested.

- b. You may not exercise any of the rights granted to You in Section 3 above in any manner that is primarily intended for or directed toward commercial advantage or private monetary compensation. The exchange of the Work for other copyrighted works by means of digital file-sharing or otherwise shall not be considered to be intended for or directed toward commercial advantage or private monetary compensation, provided there is no payment of any monetary compensation in connection with the exchange of copyrighted works.
- c. If You Distribute, or Publicly Perform the Work or Collections, You must, unless a request has been made pursuant to Section 4(a), keep intact all copyright notices for the Work and provide, reasonable to the medium or means You are utilizing: (i) the name of the Original Author (or pseudonym, if applicable) if supplied, and/or if the Original Author and/or Licensor designate another party or parties (e.g., a sponsor institute, publishing entity, journal) for attribution ("Attribution Parties") in Licensor's copyright notice, terms of service or by other reasonable means, the name of such party or parties; (ii) the title of the Work if supplied; (iii) to the extent reasonably practicable, the URI, if any, that Licensor specifies to be associated with the Work, unless such URI does not refer to the copyright notice or licensing information for the Work. The credit required by this Section 4(c) may be implemented in any reasonable manner; provided, however, that in the case of a Collection, at a minimum such credit will appear, if a credit for all contributing authors of Collection appears, then as part of these credits and in a manner at least as prominent as the credits for the other contributing authors. For the avoidance of doubt, You may only use the credit required by this Section for the purpose of attribution in the manner set out above and, by exercising Your rights under this License, You may not implicitly or explicitly assert or imply any connection with, sponsorship or endorsement by the Original Author, Licensor and/or Attribution Parties, as appropriate, of You or Your use of the Work, without the separate, express prior written permission of the Original Author, Licensor and/or Attribution Parties.
- d. For the avoidance of doubt:
 - i. Non-waivable Compulsory License Schemes. In those jurisdictions in which the right to collect royalties through any statutory or compulsory licensing scheme cannot be waived, the Licensor reserves the exclusive right to collect such royalties for any exercise by You of the rights granted under this License;
 - ii. Waivable Compulsory License Schemes. In those jurisdictions in which the right to collect royalties through any statutory or compulsory licensing scheme can be waived, the Licensor reserves the exclusive right to collect such royalties for any exercise by You of the rights granted under this License if Your exercise of such rights is for a purpose or use which is otherwise than noncommercial as permitted under Section 4(b) and otherwise waives the right to collect royalties through any statutory or compulsory licensing scheme; and,
 - iii. Voluntary License Schemes. The Licensor reserves the right to collect royalties, whether individually or, in the event that the Licensor is a member of a collecting society that administers voluntary licensing schemes, via that society, from any exercise by You of the rights granted under this Li-



cense that is for a purpose or use which is otherwise than noncommercial as permitted under Section 4(b).

- e. Except as otherwise agreed in writing by the Licensor or as may be otherwise permitted by applicable law, if You Reproduce, Distribute or Publicly Perform the Work either by itself or as part of any Collections, You must not distort, mutilate, modify or take other derogatory action in relation to the Work which would be prejudicial to the Original Author's honor or reputation.
5. *Representations, Warranties and Disclaimer* UNLESS OTHERWISE MUTUALLY AGREED BY THE PARTIES IN WRITING, LICENSOR OFFERS THE WORK AS-IS AND MAKES NO REPRESENTATIONS OR WARRANTIES OF ANY KIND CONCERNING THE WORK, EXPRESS, IMPLIED, STATUTORY OR OTHERWISE, INCLUDING, WITHOUT LIMITATION, WARRANTIES OF TITLE, MERCHANTIBILITY, FITNESS FOR A PARTICULAR PURPOSE, NONINFRINGEMENT, OR THE ABSENCE OF LATENT OR OTHER DEFECTS, ACCURACY, OR THE PRESENCE OF ABSENCE OF ERRORS, WHETHER OR NOT DISCOVERABLE. SOME JURISDICTIONS DO NOT ALLOW THE EXCLUSION OF IMPLIED WARRANTIES, SO SUCH EXCLUSION MAY NOT APPLY TO YOU.
6. *Limitation on Liability*. EXCEPT TO THE EXTENT REQUIRED BY APPLICABLE LAW, IN NO EVENT WILL LICENSOR BE LIABLE TO YOU ON ANY LEGAL THEORY FOR ANY SPECIAL, INCIDENTAL, CONSEQUENTIAL, PUNITIVE OR EXEMPLARY DAMAGES ARISING OUT OF THIS LICENSE OR THE USE OF THE WORK, EVEN IF LICENSOR HAS BEEN ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.
7. *Termination*
 - a. This License and the rights granted hereunder will terminate automatically upon any breach by You of the terms of this License. Individuals or entities who have received Collections from You under this License, however, will not have their licenses terminated provided such individuals or entities remain in full compliance with those licenses. Sections 1, 2, 5, 6, 7, and 8 will survive any termination of this License.
 - b. Subject to the above terms and conditions, the license granted here is perpetual (for the duration of the applicable copyright in the Work). Notwithstanding the above, Licensor reserves the right to release the Work under different license terms or to stop distributing the Work at any time; provided, however that any such election will not serve to withdraw this License (or any other license that has been, or is required to be, granted under the terms of this License), and this License will continue in full force and effect unless terminated as stated above.
8. *Miscellaneous*
 - a. Each time You Distribute or Publicly Perform the Work or a Collection, the Licensor offers to the recipient a license to the Work on the same terms and conditions as the license granted to You under this License.
 - b. If any provision of this License is invalid or unenforceable under applicable law, it shall not affect the validity or enforceability of the remainder of the terms of this License, and without further action by the parties to this agreement, such provision shall be reformed to the minimum extent necessary to make such provision valid and enforceable.
 - c. No term or provision of this License shall be deemed waived and no breach consented to unless such waiver or consent shall be in writing and signed by the party to be charged with such waiver or consent.



- d. This License constitutes the entire agreement between the parties with respect to the Work licensed here. There are no understandings, agreements or representations with respect to the Work not specified here. Licensor shall not be bound by any additional provisions that may appear in any communication from You. This License may not be modified without the mutual written agreement of the Licensor and You.
- e. The rights granted under, and the subject matter referenced, in this License were drafted utilizing the terminology of the Berne Convention for the Protection of Literary and Artistic Works (as amended on September 28, 1979), the Rome Convention of 1961, the WIPO Copyright Treaty of 1996, the WIPO Performances and Phonograms Treaty of 1996 and the Universal Copyright Convention (as revised on July 24, 1971). These rights and subject matter take effect in the relevant jurisdiction in which the License terms are sought to be enforced according to the corresponding provisions of the implementation of those treaty provisions in the applicable national law. If the standard suite of rights granted under applicable copyright law includes additional rights not granted under this License, such additional rights are deemed to be included in the License; this License is not intended to restrict the license of any rights under applicable law.



Glossary

alias IP	An additional IP address assigned to an interface that already has an IP address. The normal and alias IP addresses both refer to the same physical interface.
authentication	The process of verifying the authenticity of a user or client before allowing access to a network system or service.
auditing policy	The auditing policy determines which events are logged on host running Microsoft Windows operating systems.
BSD-syslog protocol	The old syslog protocol standard described in RFC 3164 http://www.ietf.org/rfc/rfc3164.txt . Sometimes also referred to as the legacy-syslog protocol.
CA	A Certificate Authority (CA) is an institute that issues certificates.
certificate	A certificate is a file that uniquely identifies its owner. Certificates contains information identifying the owner of the certificate, a public key itself, the expiration date of the certificate, the name of the CA that signed the certificate, and some other data.
client mode	In client mode, syslog-ng collects the local logs generated by the host and forwards them through a network connection to the central syslog-ng server or to a relay.
destination	A named collection of configured destination drivers.
destination driver	A communication method used to send log messages.
destination, network	A destination that sends log messages to a remote host (i.e., a syslog-ng relay or server) using a network connection.
destination, local	A destination that transfers log messages within the host, e.g., writes them to a file, or passes them to a log analyzing application.
disk buffer	The Premium Edition of syslog-ng can store messages on the local hard disk if the central log server or the network connection to the server becomes unavailable.
disk queue	See <i>disk buffer</i> .
domain name	The name of a network, e.g.: <i>balabit.com</i> .
embedded log statement	A log statement that is included in another log statement to create a complex log path.
filter	An expression to select messages.
gateway	A device that connect two or more parts of the network, e.g.: your local intranet and the external network (the Internet). Gateways act as entrances into other networks.



high availability	High availability uses a second syslog-ng server unit to ensure that the logs are received even if the first unit breaks down.
host	A computer connected to the network.
hostname	A name that identifies a host on the network.
IETF-syslog protocol	The syslog-protocol standard developed by the Internet Engineering Task Force (I E T F), described in R F C 5 4 2 4 - 5 4 2 8 http://www.ietf.org/internet-drafts/draft-ietf-syslog-protocol-23.txt .
key pair	A private key and its related public key. The private key is known only to the owner; the public key can be freely distributed. Information encrypted with the private key can only be decrypted using the public key.
license	The syslog-ng license determines the number of distinct hosts (clients and relays) that can connect to the syslog-ng server.
log path	A combination of sources, filters, parsers, rewrite rules, and destinations: syslog-ng examines all messages arriving to the sources of the logpath and sends the messages matching all filters to the defined destinations.
logstore	A binary logfile format that can encrypt, compress, and timestamp log messages.
LSH	See <i>log source host</i> .
log source host	A host or network device (including syslog-ng clients and relays) that sends logs to the syslog-ng server. Log source hosts can be servers, routers, desktop computers, or other devices capable of sending syslog messages or running syslog-ng.
log statement	See <i>log path</i> .
name server	A network computer storing the IP addresses corresponding to domain names.
Oracle Instant Client	The Oracle Instant Client is a small set of libraries, which allow you to connect to an Oracle Database. A subset of the full Oracle Client, it requires minimal installation but has full functionality.
output buffer	A part of the memory of the host where syslog-ng stores outgoing log messages if the destination cannot accept the messages immediately.
output queue	Messages from the output queue are sent to the target syslog-ng server. The syslog-ng application puts the outgoing messages directly into the output queue, unless the output queue is full. The output queue can hold 64 messages, this is a fixed value and cannot be modified.
overflow queue	See <i>output buffer</i> .
parser	A set of rules to segment messages into named fields or columns.



ping	A command that sends a message from a host to another host over a network to test connectivity and packet loss.
port	A number ranging from 1 to 65535 that identifies the destination application of the transmitted data. E.g.: SSH commonly uses port 22, web servers (HTTP) use port 80, etc.
Public-key authentication	An authentication method that uses encryption key pairs to verify the identity of a user or a client.
regular expression	A regular expression is a string that describes or matches a set of strings. The syslog-ng application supports extended regular expressions (also called POSIX modern regular expressions).
relay mode	In relay mode, syslog-ng receives logs through the network from syslog-ng clients and forwards them to the central syslog-ng server using a network connection.
rewrite rule	A set of rules to modify selected elements of a log message.
template	A user-defined structure that can be used to restructure log messages or automatically generate file names.
server mode	In server mode, syslog-ng acts as a central log-collecting server. It receives messages from syslog-ng clients and relays over the network, and stores them locally in files, or passes them to other applications, e.g., log analyzers.
source	A named collection of configured source drivers.
source, network	A source that receives log messages from a remote host using a network connection. The following sources are network sources: <code>tcp()</code> , <code>tcp6()</code> , <code>udp()</code> , <code>udp6()</code> .
source, local	A source that receives log messages from within the host, e.g., from a file.
source driver	A communication method used to receive log messages.
SSL	See <i>TLS</i> .
syslog-ng	The syslog-ng application is a flexible and highly scalable system logging application, typically used to manage log messages and implement centralized logging.
syslog-ng agent	The syslog-ng agent for Windows is a log collector and forwarder application for the Microsoft Windows platform. It collects the log messages of the Windows-based host and forwards them to a syslog-ng server using regular or SSL-encrypted TCP connections.
syslog-ng client	A host running syslog-ng in client mode.
syslog-ng Premium Edition	The syslog-ng Premium Edition is the commercial version of the open-source application. It offers additional features, like encrypted message transfer and an agent for Microsoft Windows platforms.



syslog-ng relay	A host running syslog-ng in relay mode.
syslog-ng server	A host running syslog-ng in server mode.
TLS	Transport Layer Security (TLS) and its predecessor, Secure Sockets Layer (SSL), are cryptographic protocols which provide secure communications on the Internet. The syslog-ng Premium Edition application can encrypt the communication between the clients and the server using TLS to prevent unauthorized access to sensitive log messages.
traceroute	A command that shows all routing steps (the path of a message) between two hosts.
unix domain socket	A Unix domain socket (UDS) or IPC socket (inter-procedure call socket) is a virtual socket, used for inter-process communication.



Index

A

AIX

- installing syslog-ng, 85
- redirecting errorlog to syslog-ng, 86

ALLSYL100, 129

artificial ignorance

- message classification, 218

AS/400, 121

auditing policy, 119

- configuring on Windows 2003 Server, 120
- configuring on Windows XP, 119-120

authentication, 13, 68

- syslog-ng agent, 103

B

batch processing, 136

C

CentOS

- installing syslog-ng, 85

certificate revocation lists

- syslog-ng agent, 104

certificates, 13, 103

- importing on Windows, 105

CHGSECAUD, 123

chroots, 138

classifying messages

- concepts of, 15
- configuration, 63
- creating databases, 220
- filtering, 64, 219
- pattern matching concepts, 17

client authentication

- syslog-ng agent, 104

client mode, 10

client-side failover

- syslog-ng agent, 98

compatibility with Snare, 144, 148, 152, 158, 161, 169, 176

compiling syslog-ng OSE, 86

configuration file

- detecting changes, 31
- including other files, 31

configuring syslog-ng

- on IBM System i, 124
- on Linux/Unix, 30
- on Windows, 93
- syslog-ng agent, 95-96

Coordinated Universal Time, 140

core files, 76

CRL

- syslog-ng agent, 104

CSV parsers, 214

D

daylight saving changes, 13

defining global objects, 32

deleting syslog-ng, 89

destination drivers, 11, 42

- database driver, 47, 191
- file() driver, 43, 179
- list of, 43, 240
- logstore() driver, 44, 183
- pipe() driver, 46, 186
- program() driver, 46, 189
- reference, 179
- sql() driver, 47, 191
- syslog() driver, 51, 195
- tcp() driver, 52, 200
- tcp6() driver, 52, 200
- udp() driver, 52, 200
- udp6() driver, 52, 200
- unix-dgram() driver, 53, 204
- unix-stream() driver, 53, 204
- usertty() driver, 53, 207

destinations, 6, 11, 42

- defining, 34, 42
- FreeTDS configuration, 89
- Microsoft SQL Server configuration, 89
- MSSQL configuration, 89
- sql() configuration, 49-51, 193-195
- syslog-ng agent, 97

disk buffer, 20

- location of, 68
- on Windows, 93

disk queue (see disk buffer)

disk-based buffering, 20

dropping messages, 140

E

embedded log statements, 8



- encrypted log files, 44
- encrypting log files, 14
- encrypting log messages, 13, 68
 - on the hard disk, 14
- error solving, 76
- eventlog
 - remote logging, 109

F

- facilities, 25, 27, 136, 210
- fail-over, 22
- failure script, 77
- fd limit, 179, 183
- file descriptors, 179, 183
- file encryption, 14
- filters, 6, 11, **57**, 60, 137
 - defining, 57
 - facilities, , 210
 - facility and priority (level) ranges, 59
 - priorities, 209
 - reference, 208
 - wildcards, 59
- flags, 54, 207
- flow-control, 18, 56
 - example, 57
 - multiple destinations, 20
- formatting messages, 14

G

- global objects, 11
 - defining, 32
- global options, **67**
 - reference, 227

I

- IBM iSeries, 121
- IBM System i, 121
- importing certificates, 105
- inheriting settings on Windows, 117
- installation path, 78
- installing syslog-ng, **78**
 - from DEB package, 86
 - from RPM package, 85
 - in silent mode, 84
 - on AIX, 85
 - on CentOS, 85
 - on clients and relays, 79

- on logservers, 81
- on Red Hat Enterprise Server, 85
- on SUSE Linux Enterprise Server, 85
- on Windows, 93
 - syslog-ng agent on domain controllers, 93
- installing syslog-ng OSE from source, 86

L

- license, 11, 21
 - installing, 75
- local time, 25, 28
- log messages, structure, 24
 - BSD-syslog protocol, 24
 - IETF-syslog protocol, 26
 - legacy-syslog protocol, 24
 - RFC 3164, 24
 - RFC 5424, 26
- log paths, 6, **53**
 - defining, 53
 - flags, 54, 207
 - flow-control, 18, 56-57
- log pipes (see embedded log statements)
- log statements, 6, 12 (see log paths)
 - embedded, **8**
- log statistics, 36
 - on unix-socket, 37
- logcat, 44
- logchksign, 31
- logging procedure, 7
- logstore, 14, 44
- losing messages, 22
 - from eventlog containers, 118

M

- macros, 12, 14
 - date and time, 112
 - eventlog sources, 113
 - file sources, 113
 - protocol, 111
 - reference, 210
 - syslog-ng agent, 111-113
- message classification, 63-64, 219-220
- message facilities, 25, 27, 210
- message filtering
 - syslog-ng agent, 106
 - using parsers, 64
- message format
 - syslog-ng agent, 107



- message loss, 22
- message parsing, 62-64, 214, 219
- message rate
 - on Windows, 99
- message templates, 14
- Microsoft SQL
 - sql() configuration, 50
- Microsoft SQL Server configuration, 89
- modes of operation, 10
 - client mode, 10
 - relay mode, 10
 - server mode, 11
- MSSQL
 - sql() configuration, 50-51, 194
- mutual authentication, 13, 70
 - syslog-ng agent, 104

N

- name resolution, 136-137
 - local, 138
- number of open files, 179, 183

O

- optimizing syslog-ng performance, 137
 - regular expressions, 60
- options, 12
 - reference, 227
- Oracle
 - sql() configuration, 49-50, 194
- output buffer, 19, 56
- output queue, 21
- overflow queue (see output buffer)
- overriding facility, 34

P

- parallel connections, 136
- parameters
 - log_disk_fifo_size(), 20
 - log_fetch_limit() , 18, 56, 136
 - log_fifo_size() , 18, 56, 136
 - log_iw_size() , 19, 56
 - max_connections() , 19, 56, 136
 - time_sleep(), 136
- parsers, 6, 12, 62-64, 219
 - reference, 214
- parsing messages, 62-64, 214, 218-219
 - concepts of, 14

- filtering parsed messages, 64
- pattern database, 63-64, 219-220, 222, 224
 - creating parsers, 218
 - structure of, 16
 - using the results, 64
- pattern databases
 - concepts of, 15
 - pattern matching precedence, 17
- pattern matching
 - procedure of, 17
- PostgreSQL
 - sql() configuration, 49, 193
- preventing message loss (see flow-control)

Q

- QAUDJRN, 121, 123, 128

R

- reading messages from external applications, 150
- Red Hat Enterprise Server
 - installing syslog-ng, 85
- regular expressions, 57, 60, 137, 226
 - case-insensitive, 59
 - escaping, 59
 - pcre, 226
 - posix, 209
- relay mode, 10
- removing syslog-ng, 89
- replacing message text, 66, 225
- rewrite
 - reference, 225
- rewrite rules, 6, 12, 66
- rewriting messages, 66, 225
 - concepts of, 15

S

- sedding messages, 66, 225
- segmenting messages, 62, 214
- server license, 21
- server mode, 11
- setting facility, 34
- setting message fields, 67, 225
- signing log files, 14
- skipping messages, 140
- snare, 98
- Snare



- receiving Snare-compatible messages, 144, 148, 152, 158, 161, 169, 176
- Snare-compatibility, 144, 148, 152, 158, 161, 169, 176
- source drivers, 11, **33**
 - file() driver, 37, 141
 - internal() driver, 141
 - list of, 36, 240
 - pipe() driver, 38, 146
 - program() driver, 150
 - reference, 141
 - sun-streams() driver, 39, 154
 - syslog() driver, 39, 160
 - tcp() driver, 40, 165
 - tcp6() driver, 40, 165
 - udp() driver, 40, 165
 - udp6() driver, 40, 165
 - unix-dgram() driver, 173
 - unix-stream() driver, 173
- sources, 6, 11, **33**
 - eventlog, 99
 - on different platforms, 35
 - windows log files, 101
- splitting messages, 62, 214
- SQL NULL values, 195
- statistics, 36-37
- supported architectures, 4
- supported operating systems, 4, 93
- SUSE Linux Enterprise Server
 - installing syslog-ng, 85
- syslog-ng
 - troubleshooting, 76
- syslog-ng agent, 93
 - certificate revocation lists, 104
 - client authentication, 104
 - client-side failover, 98
 - configuring the logserver, 97
 - creating core dumps, 119
 - CRL, 104
 - default message format, 110
 - destinations, 97
 - disabling sources and filters, 103
 - eventlog sources, 99
 - failover servers, 98
 - file sources, 101
 - filtering messages, 106
 - importing certificates, 105
 - installing, 93
 - installing the agent from the domain controller, 95
 - installing the agent in standalone mode, 94
 - message format, 107
 - mutual authentication, 104
 - Snare-compatibility, 144, 148, 152, 158, 161, 169, 176
 - throttle, 99
 - timestamp, 109
 - troubleshooting, 117, 119
 - upgrading, **96**
 - upgrading 3.0.1, 96
 - upgrading 3.0.2, 97
 - upgrading to 3.0.4, 97
- syslog-ng Agent
 - configuration file, 113
 - for IBM System i, 121
 - inheriting settings, 117
 - System i configuration, 124
 - timezone, 117
 - XML, 113
- syslog-ng Agent for IBM System i, 121
 - downgrading, 123
 - filtering, 122
 - installation, 122
 - reverting to older version, 123
 - supported log sources, 121
 - supported output formats, 121
 - upgrading, 122
- syslog-ng binaries
 - location of, 78
- syslog-ng clients
 - configuring, 73
- syslog-ng relays
 - configuring, 73
- syslog-ng servers
 - configuring, 74
- syslog-ng.conf, **30**
 - fingerprint, 31
 - includes, 31
- System i
 - ALLSYL100 subsystem, 129
 - clearing logs manually, 131
 - configuring server applications, 132
 - configuring the syslog-ng Agent, 124
 - custom journal entries, 128
 - destination server, 126
 - enable logging in Apache, 132
 - enable logging in OpenSSH, 133
 - enabling security auditing, 123
 - enabling security auditing manually, 123



- enabling user auditing, 124
- QAUDJRN, 128
- security auditing, 123
- security types, 128
- starting the subsystem, 129
- troubleshooting, 133
- user-created journal entries, 128
- viewing logs, 131

System i security audit journal, 121

T

templates, 12, 14, **60**

- defining, 61
- example, 61

throttle

- on Windows, 99

timestamp, 25, 28, 136, 140

- syslog-ng agent, 109

Timestamping Authority, 14

timezone

- in chroots, 139
- Windows, 109, 117

timezones, 12, 140

TLS, 13, 39, 41, 160, 166

- configuring, 68, 70
- reference, 231
- syslog-ng agent, 103

transport layer security (see TLS)

troubleshooting, 76

- core files, 76
- failure scrip, 77
- syslog-ng, 76-77
- syslog-ng agent, 117

TSA, 14

U

ulimit, 179, 183

uninstalling syslog-ng, 89

UTC, 140

W

wildcards

- in file sources, 37, 141

Windows auditing policy, 119

- configuring on Windows 2003 Server, 120
- configuring on Windows XP, 119-120